

Branch-and-Bound with Peer-to-Peer for Large-Scale Grids

Alexandre di Costanzo

Grid Computing and Distributed Systems (GRIDs) Laboratory
Department of Computer Science and Software Engineering

University of Melbourne, Australia

Monday 2 June 2008

Agenda

- About Me / Previous Experience
- PhD Main Contributions
 - Branch-and-Bound Framework for Grids
 - Desktop Grid with Peer-to-Peer
 - Mixing Desktops & Clusters
- Perspectives
- Possible Support & Conclusion

About Me

- June 2004 **Master in Networking and Distributed Systems**
 - INRIA Sophia Antipolis - OASIS research group / University of Nice Sophia Antipolis, France
 - “*P2P Model and Infrastructure for Communicating Applications*”, supervised by Prof. Denis Caromel
- July-September 2004 **Internship at Indiana University, USA**
 - “*Asynchronous Peer-to-Peer Web Services and Firewalls*”, supervised by Prof. Dennis Gannon
- October 2007 **PhD in Computer Science**
 - INRIA Sophia Antipolis - OASIS research group / University of Nice Sophia Antipolis, France
 - “*Branch-and-Bound with Peer-to-Peer on Large-Scale Grids*”, supervised by Denis Caromel

Agenda

- About Me / Previous Experience
- **PhD Main Contributions**
 - Branch-and-Bound Framework for Grids
 - Desktop Grid with Peer-to-Peer
 - Mixing Desktops & Clusters
 - Perspectives
 - Possible Support & Conclusion

The Big Picture

The Big Picture

● Objective

The Big Picture

● Objective

**Solving combinatorial optimization problems
with Grids**

The Big Picture

- Objective

**Solving combinatorial optimization problems
with Grids**

- Approach

The Big Picture

- Objective

**Solving combinatorial optimization problems
with Grids**

- Approach

Parallel Branch-and-Bound and Peer-to-Peer

The Big Picture

- Objective

**Solving combinatorial optimization problems
with Grids**

- Approach

Parallel Branch-and-Bound and Peer-to-Peer

- Contributions

The Big Picture

● Objective

**Solving combinatorial optimization problems
with Grids**

● Approach

Parallel Branch-and-Bound and Peer-to-Peer

● Contributions

I. Branch-and-Bound framework for Grids

The Big Picture

- Objective

**Solving combinatorial optimization problems
with Grids**

- Approach

Parallel Branch-and-Bound and Peer-to-Peer

- Contributions

1. Branch-and-Bound framework for Grids

2. Peer-to-Peer Infrastructure for Grids

The Big Picture

● Objective

**Solving combinatorial optimization problems
with Grids**

● Approach

Parallel Branch-and-Bound and Peer-to-Peer

● Contributions

1. Branch-and-Bound framework for Grids

2. Peer-to-Peer Infrastructure for Grids

3. Large-scale experiments

Context

Context

- Combinatorial Optimization Problems (COPs)
- costly to solve (finding the best solution)

Context

- Combinatorial Optimization Problems (COPs)
 - costly to solve (finding the best solution)
- Branch-and-Bound (BnB)
 - well adapted for solving COPs
 - relatively easy to provide parallel version

Context

- Combinatorial Optimization Problems (COPs)
 - costly to solve (finding the best solution)
- Branch-and-Bound (BnB)
 - well adapted for solving COPs
 - relatively easy to provide parallel version
- Grid Computing
 - large pool of resources
 - large-scale environment

Branch-and-Bound

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

- Feasible solutions are organized as a tree: **search-tree**

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

- Feasible solutions are organized as a tree: **search-tree**
- 3 operations:

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

- Feasible solutions are organized as a tree: **search-tree**
- 3 operations:
- **Branching**: split in sub-problems

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

- Feasible solutions are organized as a tree: **search-tree**
- 3 operations:
 - Branching**: split in sub-problems
 - Bounding**: compute lower/upper bounds (*objective function*)

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

- Feasible solutions are organized as a tree: **search-tree**
- 3 operations:
 - Branching**: split in sub-problems
 - Bounding**: compute lower/upper bounds (*objective function*)
 - Pruning**: eliminate bad branches

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

- Feasible solutions are organized as a tree: **search-tree**
- 3 operations:
 - Branching**: split in sub-problems
 - Bounding**: compute lower/upper bounds (*objective function*)
 - Pruning**: eliminate bad branches
- Well adapted for solving COPs [Papadimitriou 98]

Branch-and-Bound

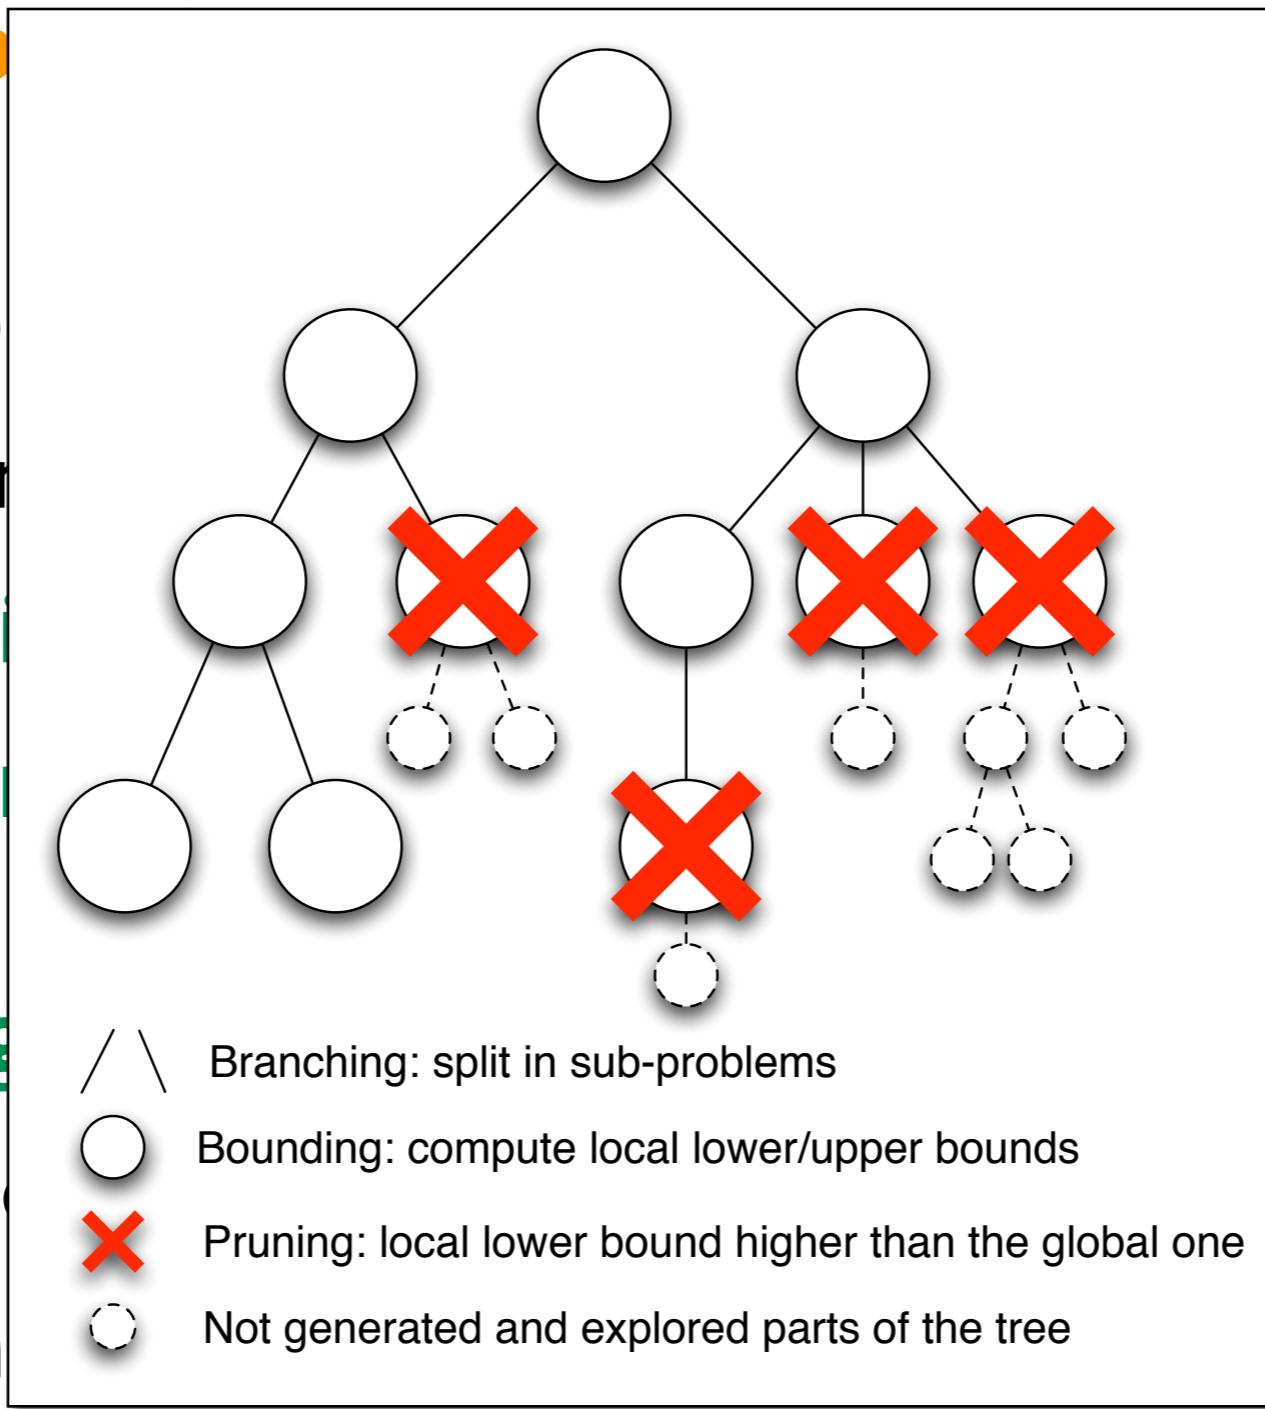
Consists of a partial enumeration of all feasible solutions and returns the guaranteed optimal solution

- Feasible solutions are organized as a tree: **search-tree**
- 3 operations:
 - Branching**: split in sub-problems
 - Bounding**: compute lower/upper bounds (*objective function*)
 - Pruning**: eliminate bad branches
- Well adapted for solving COPs [Papadimitriou 98]
- return the best combination out of all

Branch-and-Bound

Consists of a partial enumeration of all feasible solutions guaranteed

- Feasible solution
- 3 operations
 - Branching**: split in sub-problems
 - Bounding**: compute local lower/upper bounds
 - Pruning**: local lower bound higher than the global one
- Well adapted to discrete optimization problems
- return the best solution found



Grid Computing

Grids involve new concepts for development & execution
of applications

Grid Computing

Grids involve new concepts for development & execution
of applications

Grid Fabric

Schedulers, Networking, OS
Federated Hardware Resources

Grid Computing

Grids involve new concepts for development & execution of applications

Grid Middleware Infrastructure

Super-Schedulers, Resource Trading, Information, Security, etc.

Grid Fabric

Schedulers, Networking, OS
Federated Hardware Resources

Grid Computing

Grids involve new concepts for development & execution of applications

Grid Programming

Models, Tools,
High-Level Access to Middleware

Grid Middleware Infrastructure

Super-Schedulers, Resource Trading,
Information, Security, etc.

Grid Fabric

Schedulers, Networking, OS
Federated Hardware Resources

Grid Computing

Grids involve new concepts for development & execution of applications

Grid Applications & Portals

Grid Programming

Grid Middleware Infrastructure

Grid Fabric

Models, Tools,
High-Level Access to Middleware

Super-Schedulers, Resource Trading,
Information, Security, etc.

Schedulers, Networking, OS
Federated Hardware Resources

Grid Computing

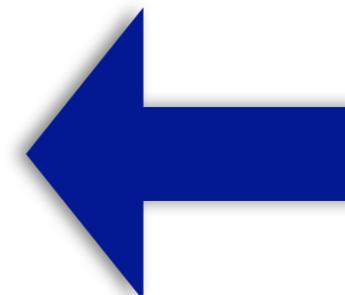
Grids involve new concepts for development & execution of applications

Grid Applications & Portals

Grid Programming

Grid Middleware Infrastructure

Grid Fabric



Branch-and-Bound API

Super-Schedulers, Resource Trading, Information, Security, etc.

Schedulers, Networking, OS
Federated Hardware Resources

Challenges

Combinatorial Optimization Problems
costly to solve

Challenges

Combinatorial Optimization Problems
costly to solve

Branch-and-Bound
adapted for solving COPs
easy to parallelize

Challenges

Combinatorial Optimization Problems

costly to solve

Branch-and-Bound
adapted for solving COPs
easy to parallelize

Grid Computing
huge number of resources

Challenges

Combinatorial Optimization Problems
costly to solve

Branch-and-Bound
adapted for solving COPs
easy to parallelize



Grid Computing
huge number of resources

Use Branch-and-Bound on Grids for solving COPs

Challenges

Combinatorial Optimization Problems
costly to solve

Branch-and-Bound
adapted for solving COPs
easy to parallelize



Grid Computing
huge number of resources

Use Branch-and-Bound on Grids for solving COPs

Efficient communications with Grids is difficult
Problem with sharing bounds

BnB on Grids: Problems and Solutions

BnB on Grids: Problems and Solutions

BnB on Grids: Problems and Solutions

Latency	

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-worker

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-worker
Solution tree size	

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-worker
Solution tree size	Dynamically generated by splitting tasks

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-worker
Solution tree size	Dynamically generated by splitting tasks
Share the best bounds	

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-worker
Solution tree size	Dynamically generated by splitting tasks
Share the best bounds	Efficient parallelism & communication

BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-worker
Solution tree size	Dynamically generated by splitting tasks
Share the best bounds	Efficient parallelism & communication
Faults	

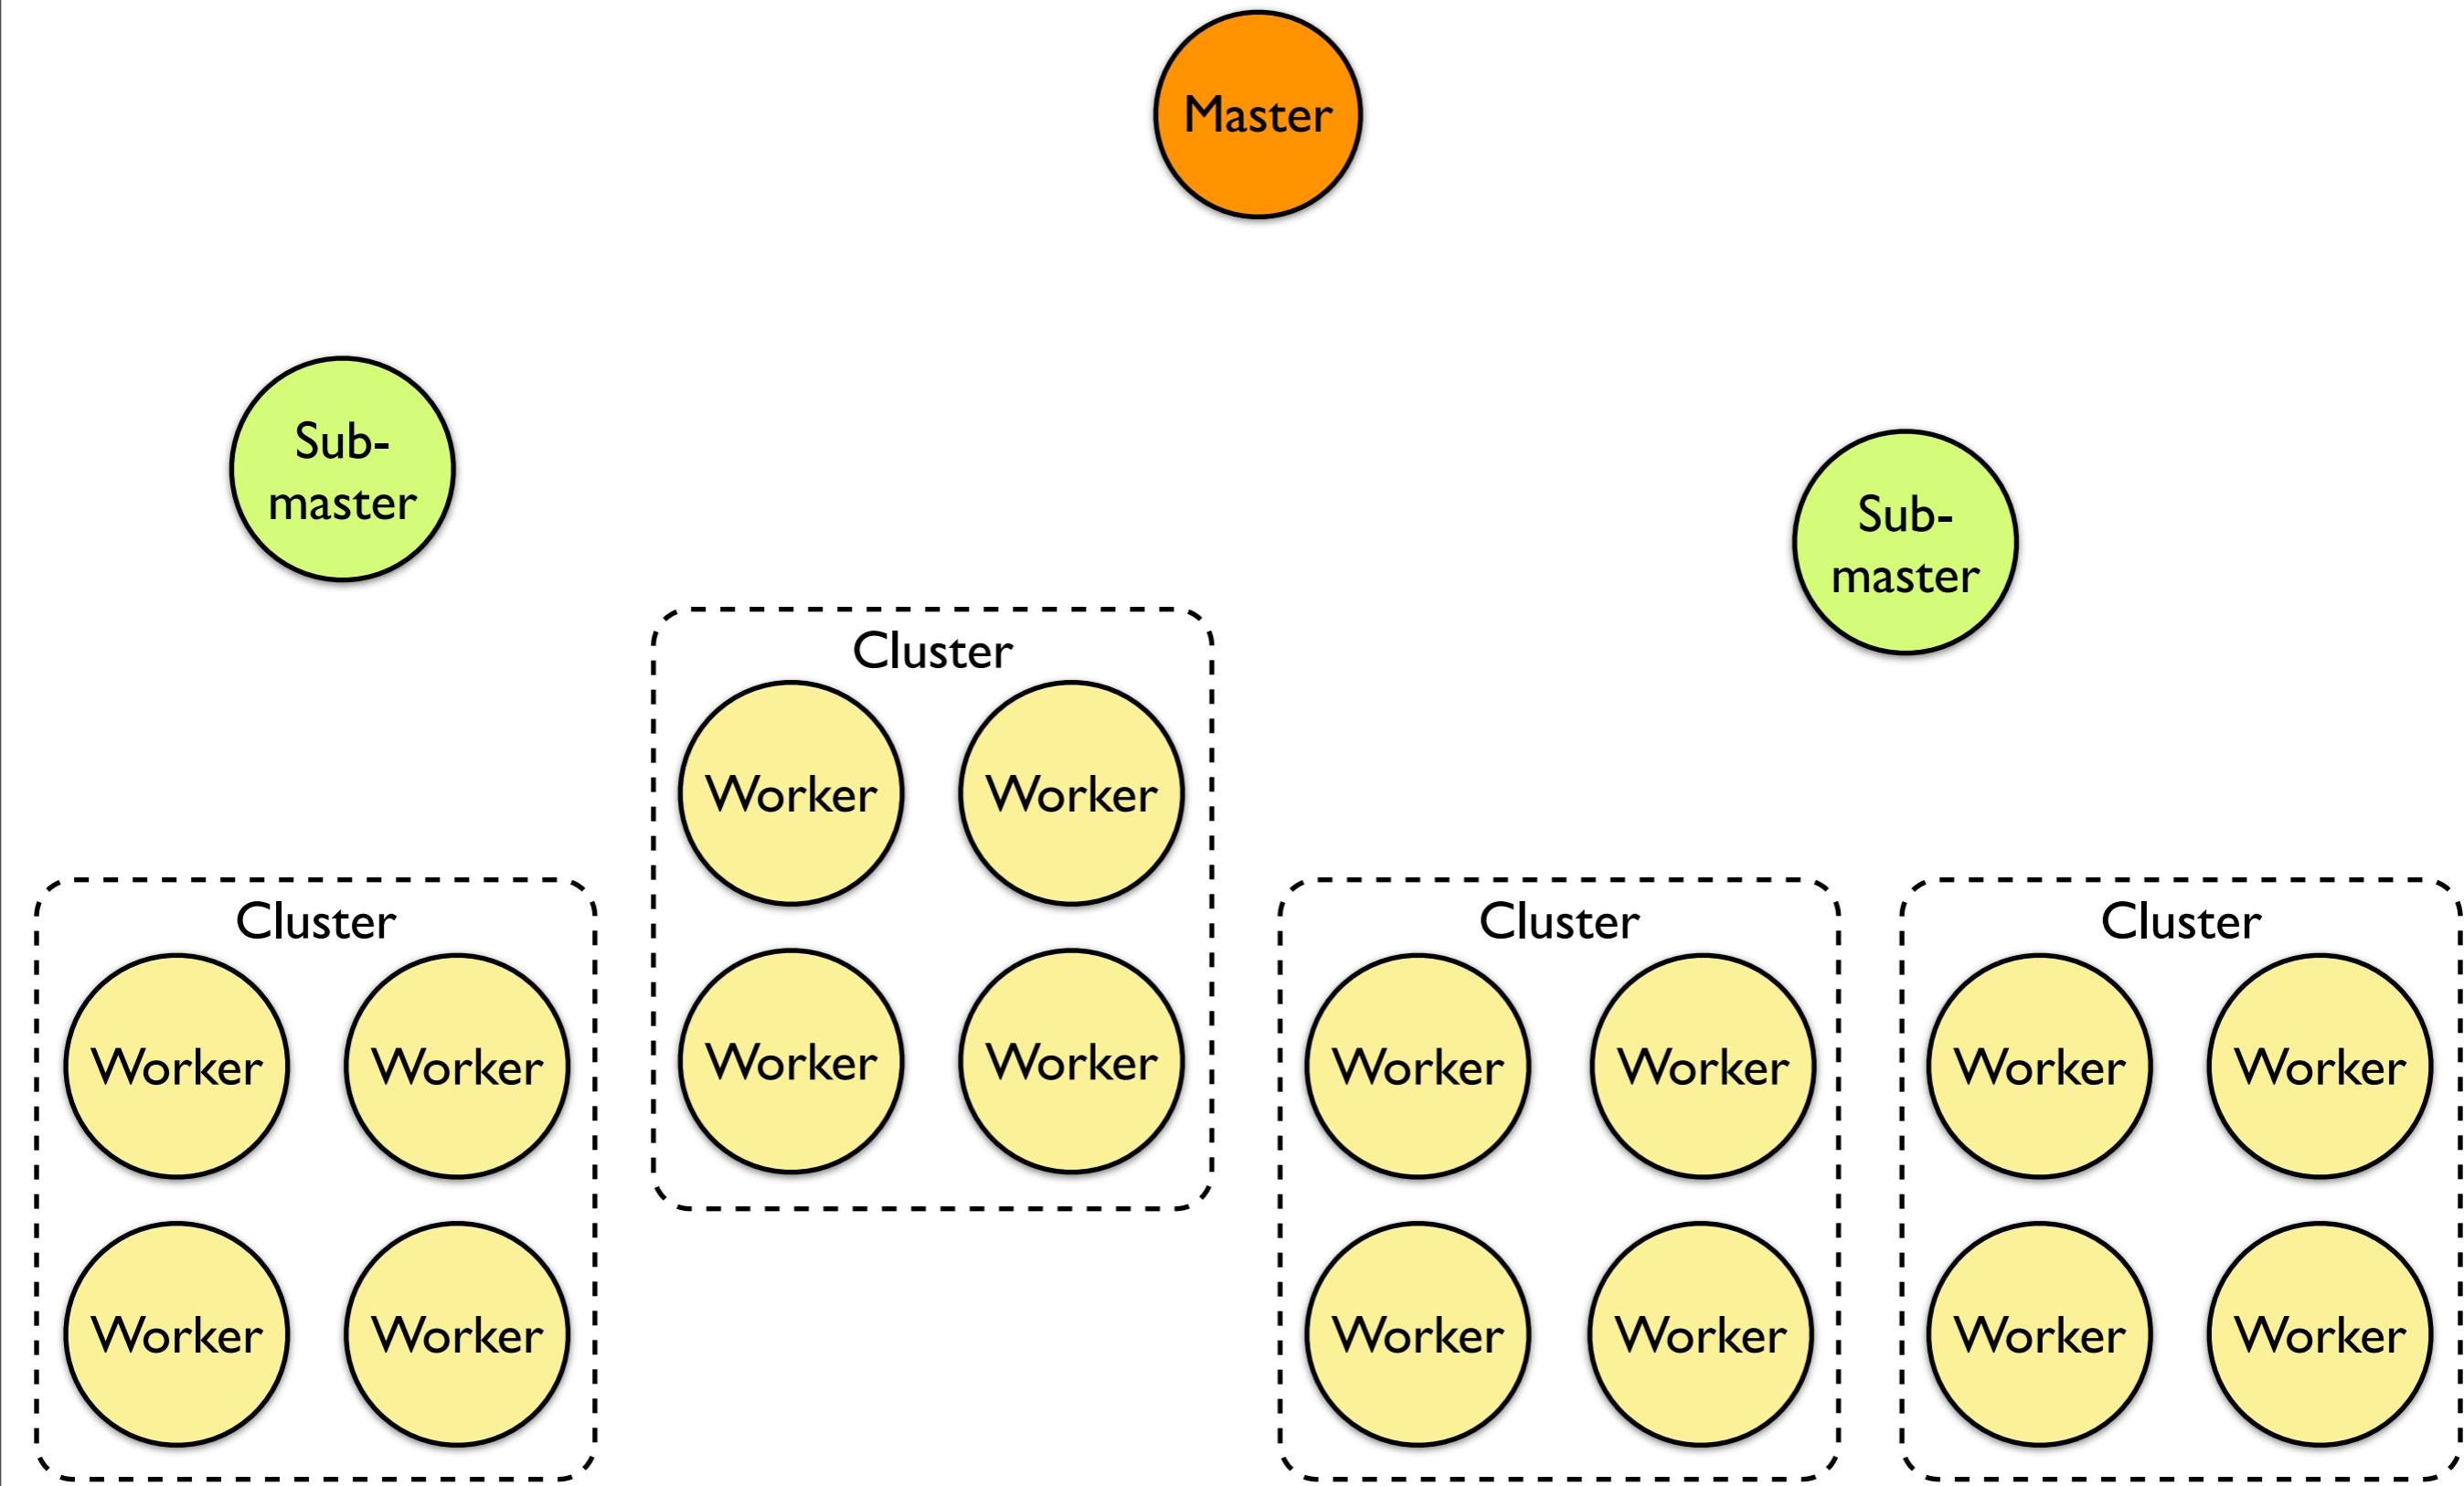
BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-worker
Solution tree size	Dynamically generated by splitting tasks
Share the best bounds	Efficient parallelism & communication
Faults	Fault-tolerance

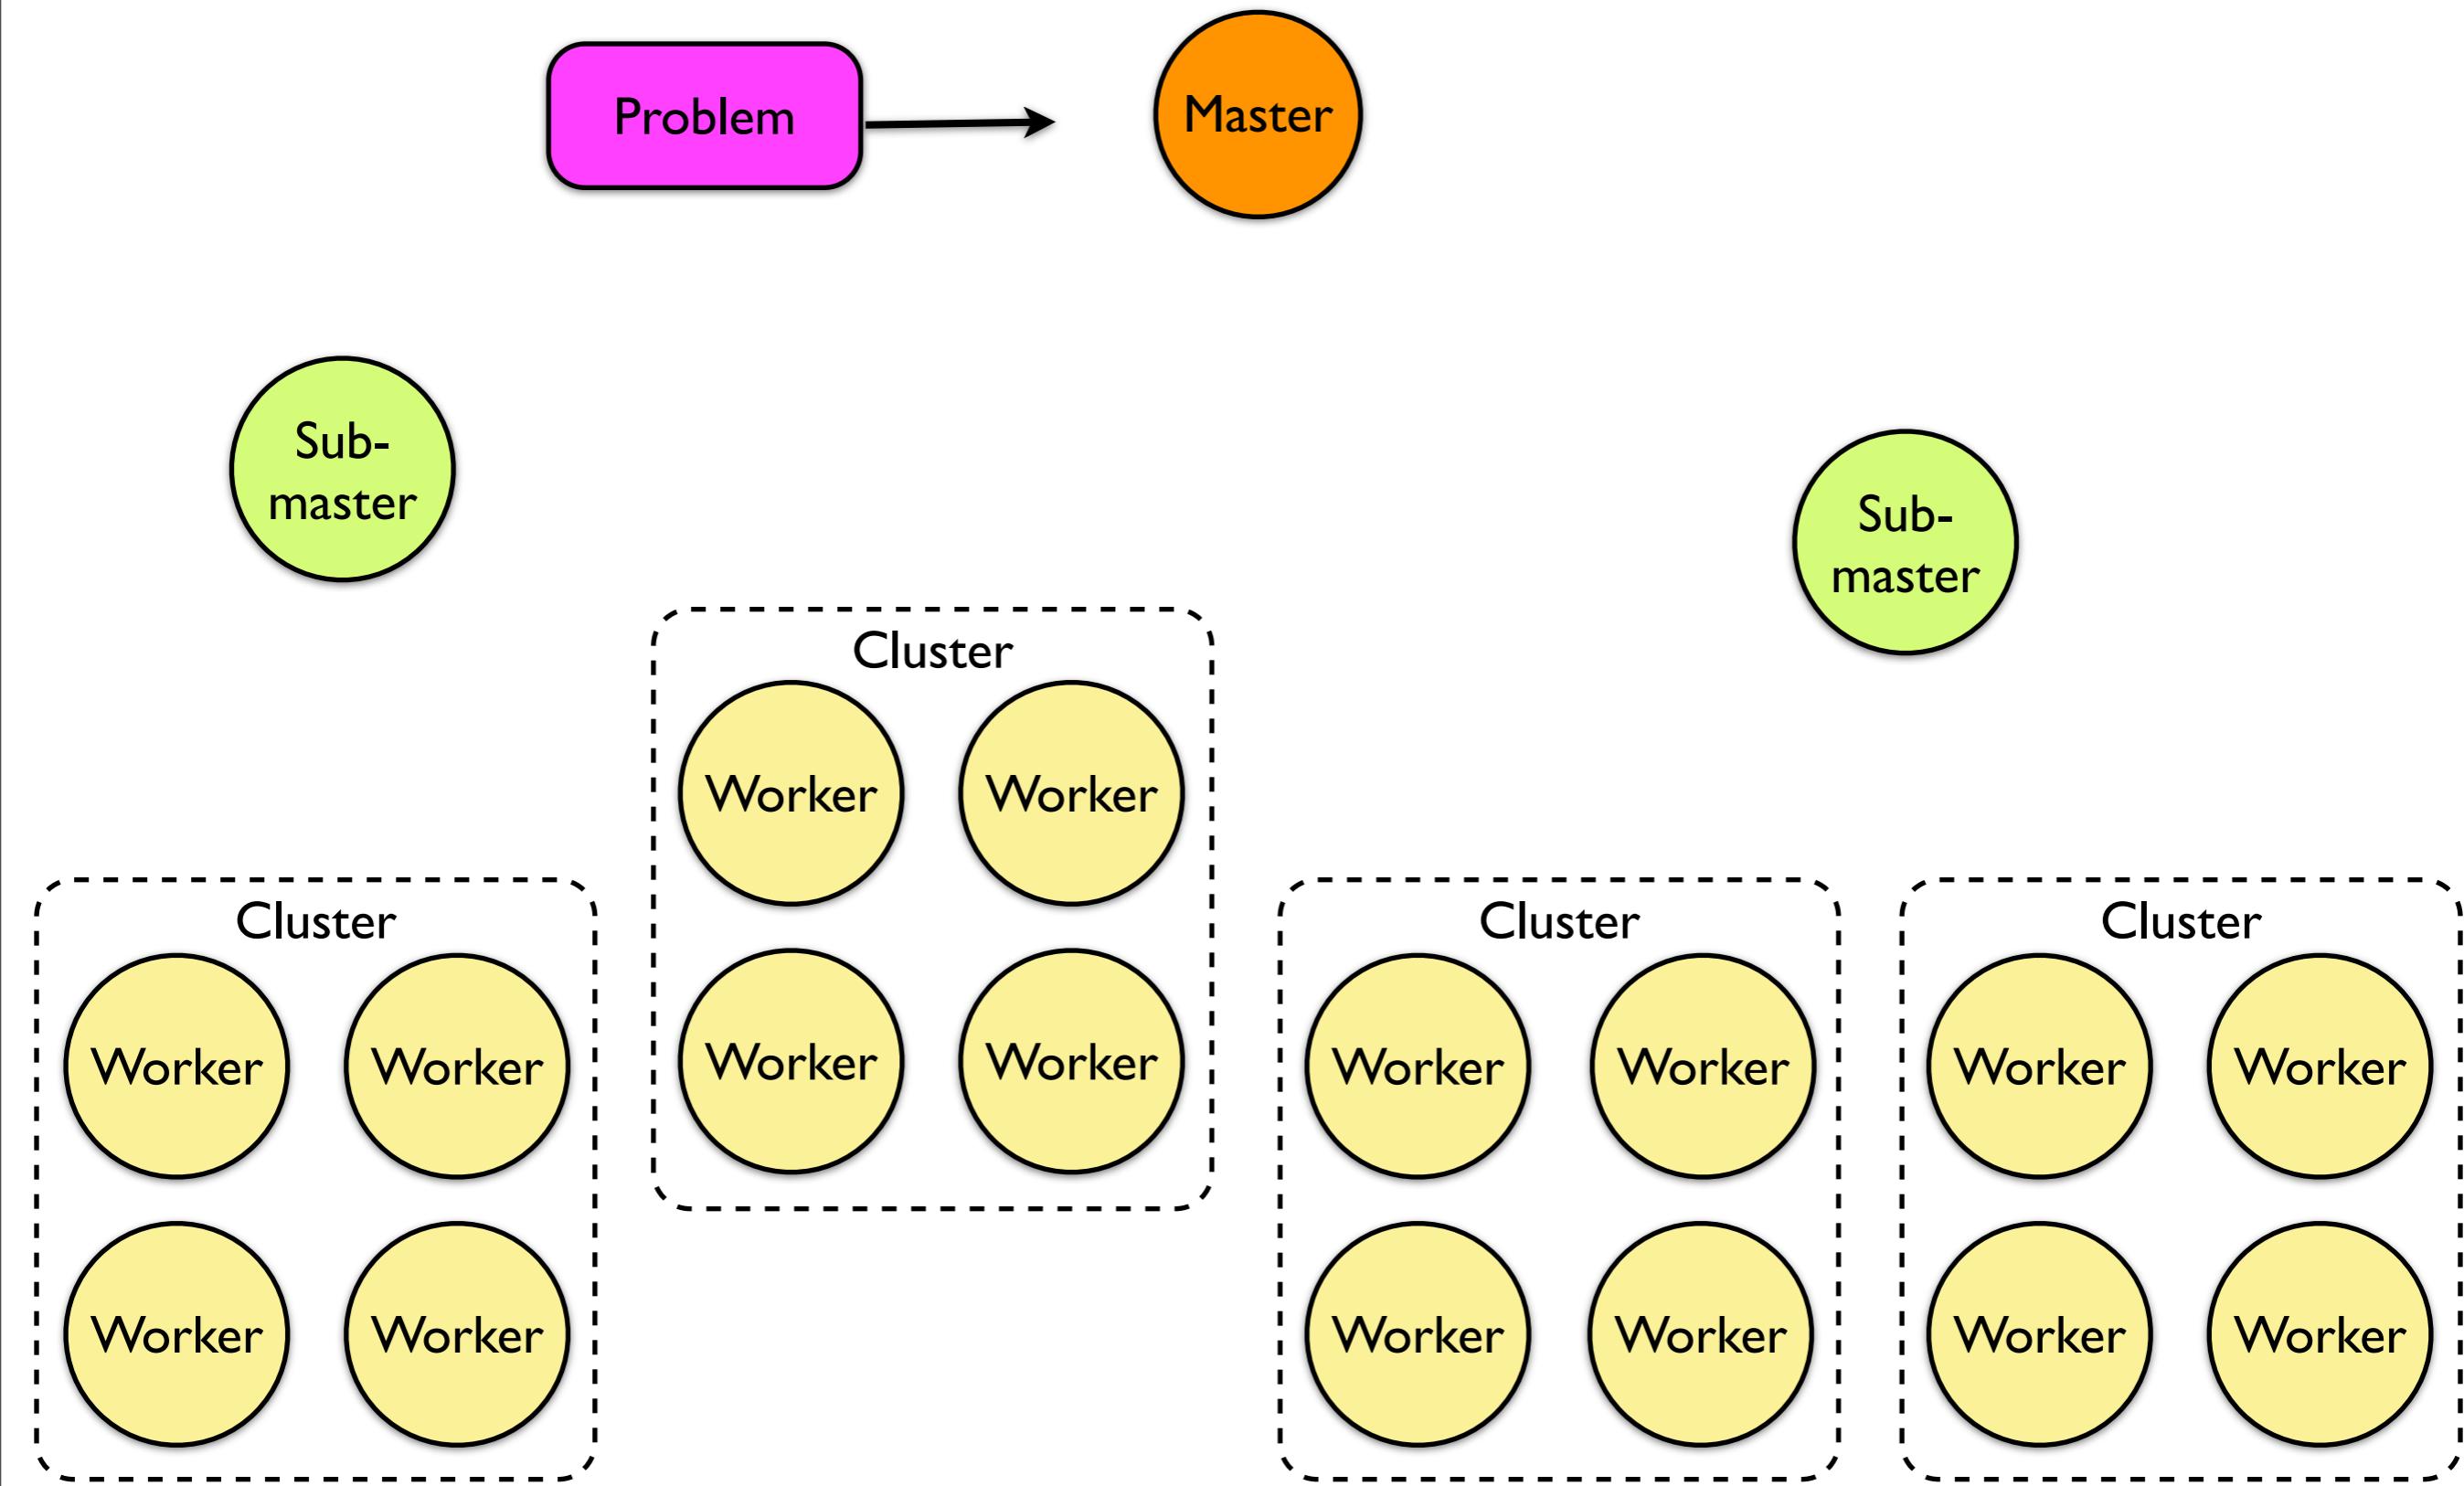
BnB on Grids: Problems and Solutions

Latency	Asynchronous communications
Scalability	Hierarchical master-slave Objective: hide Grid difficulties to users Especially communication problems
Share the best bounds	Efficient parallelism & communication
Faults	Fault-tolerance

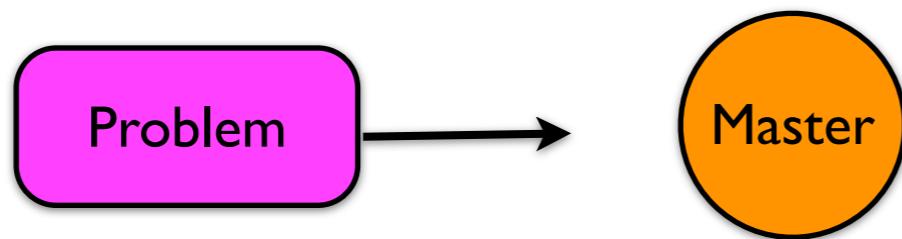
BnB Framework - Architecture



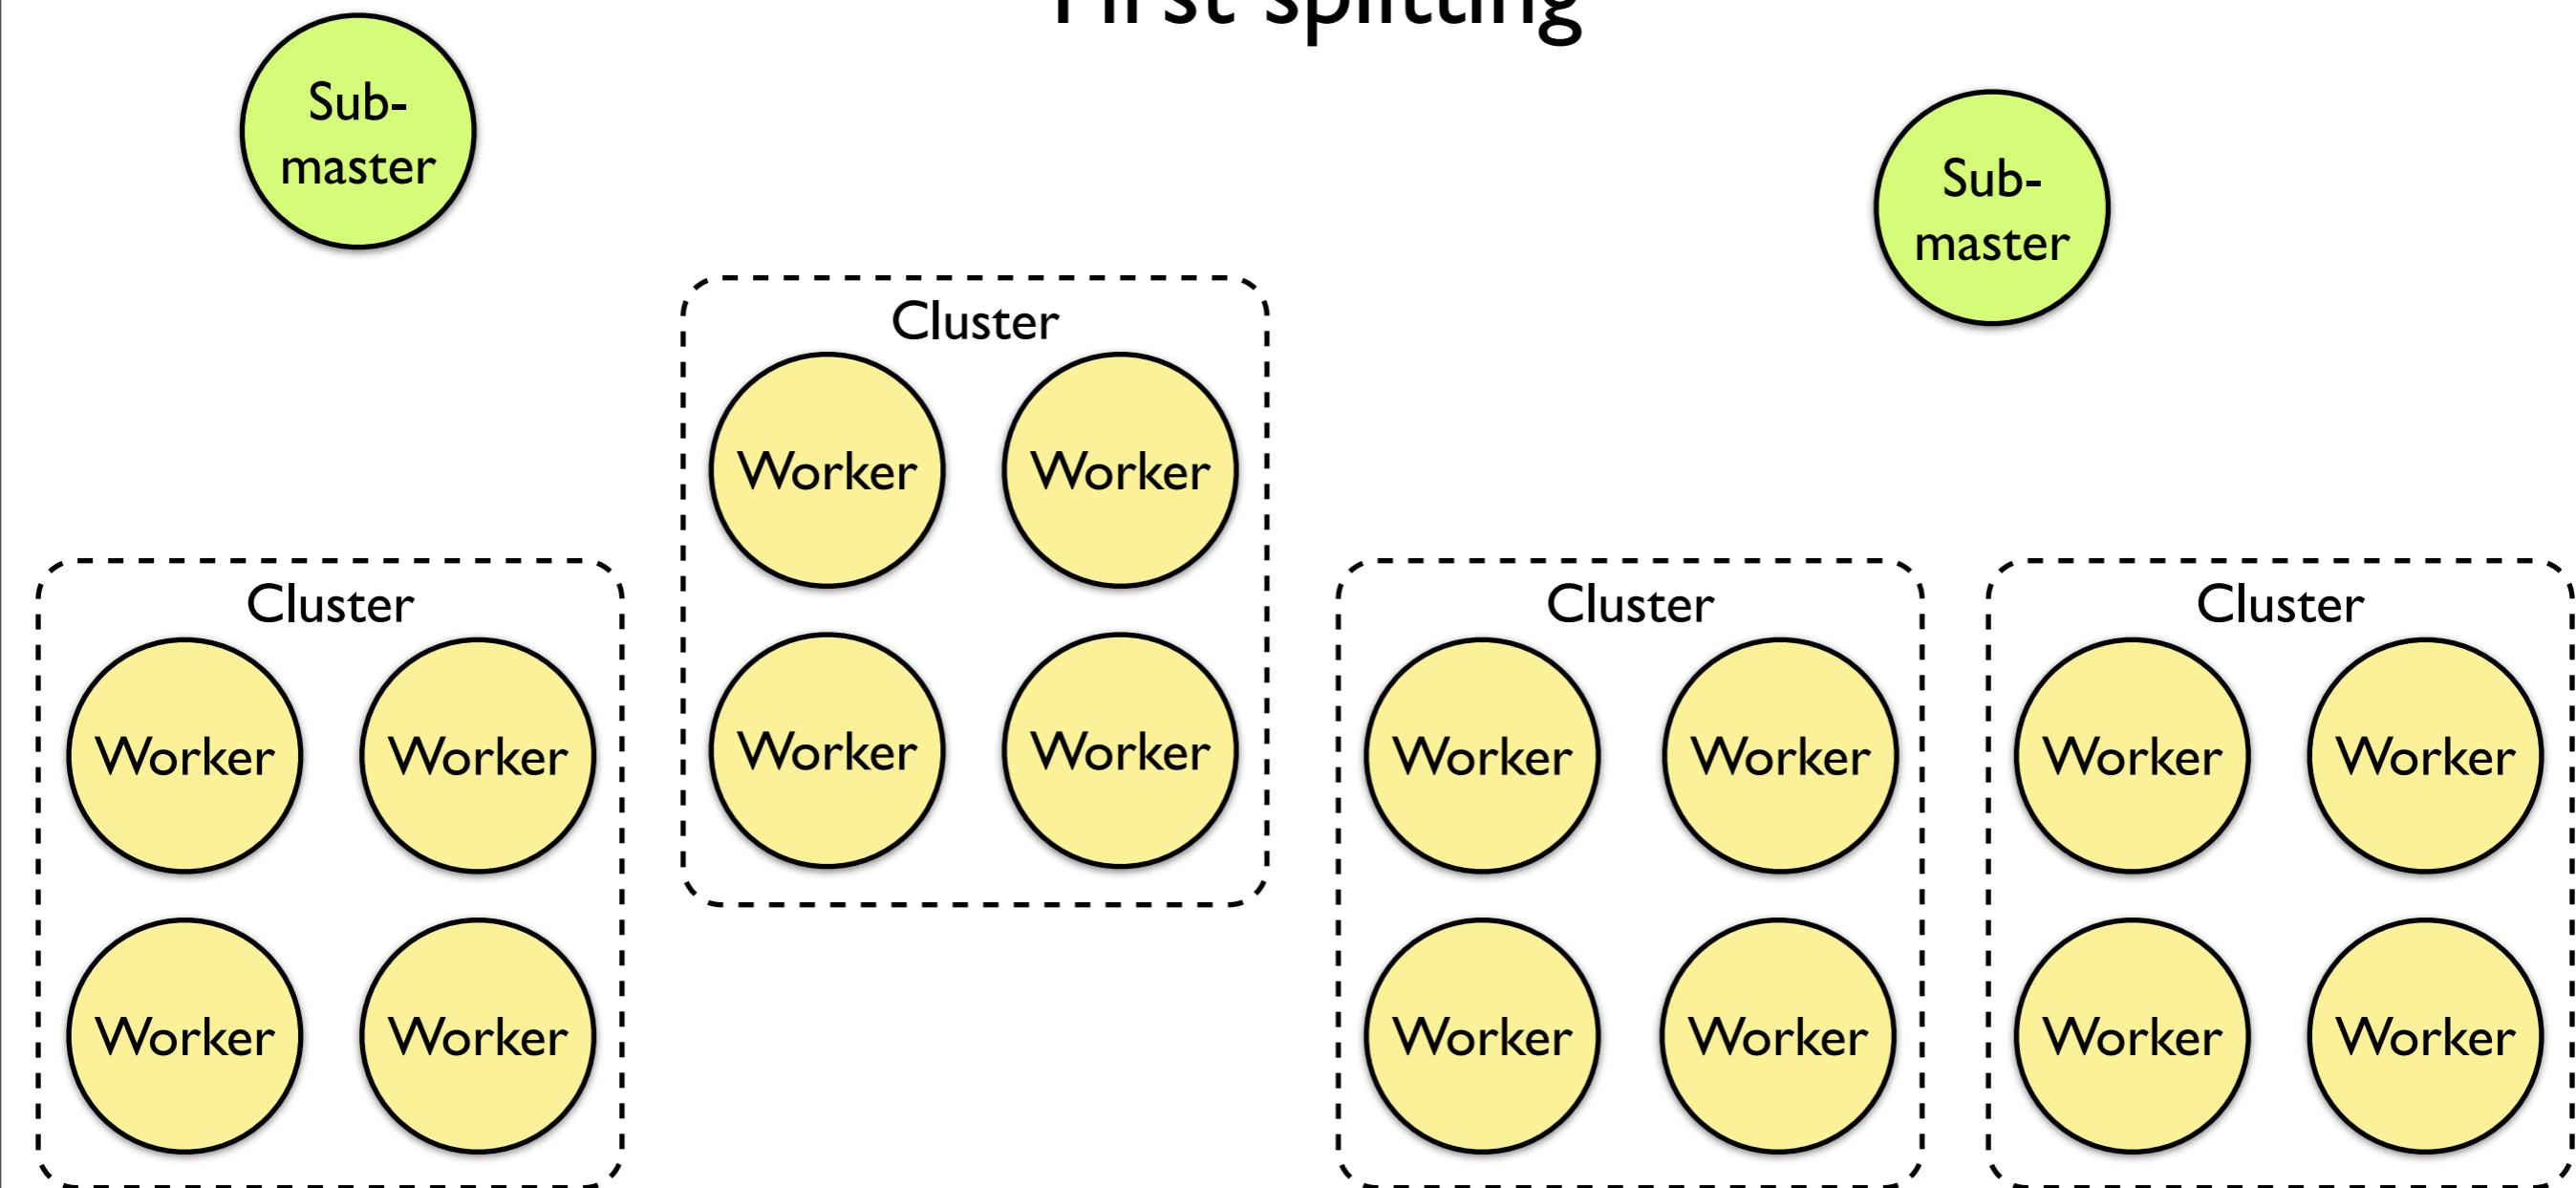
BnB Framework - Architecture



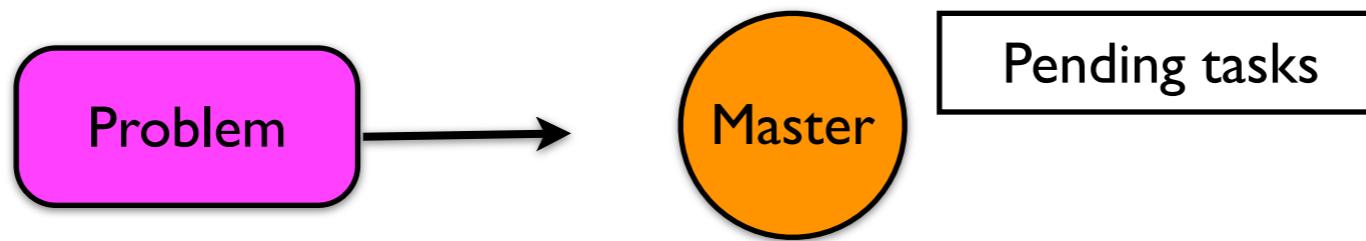
BnB Framework - Architecture



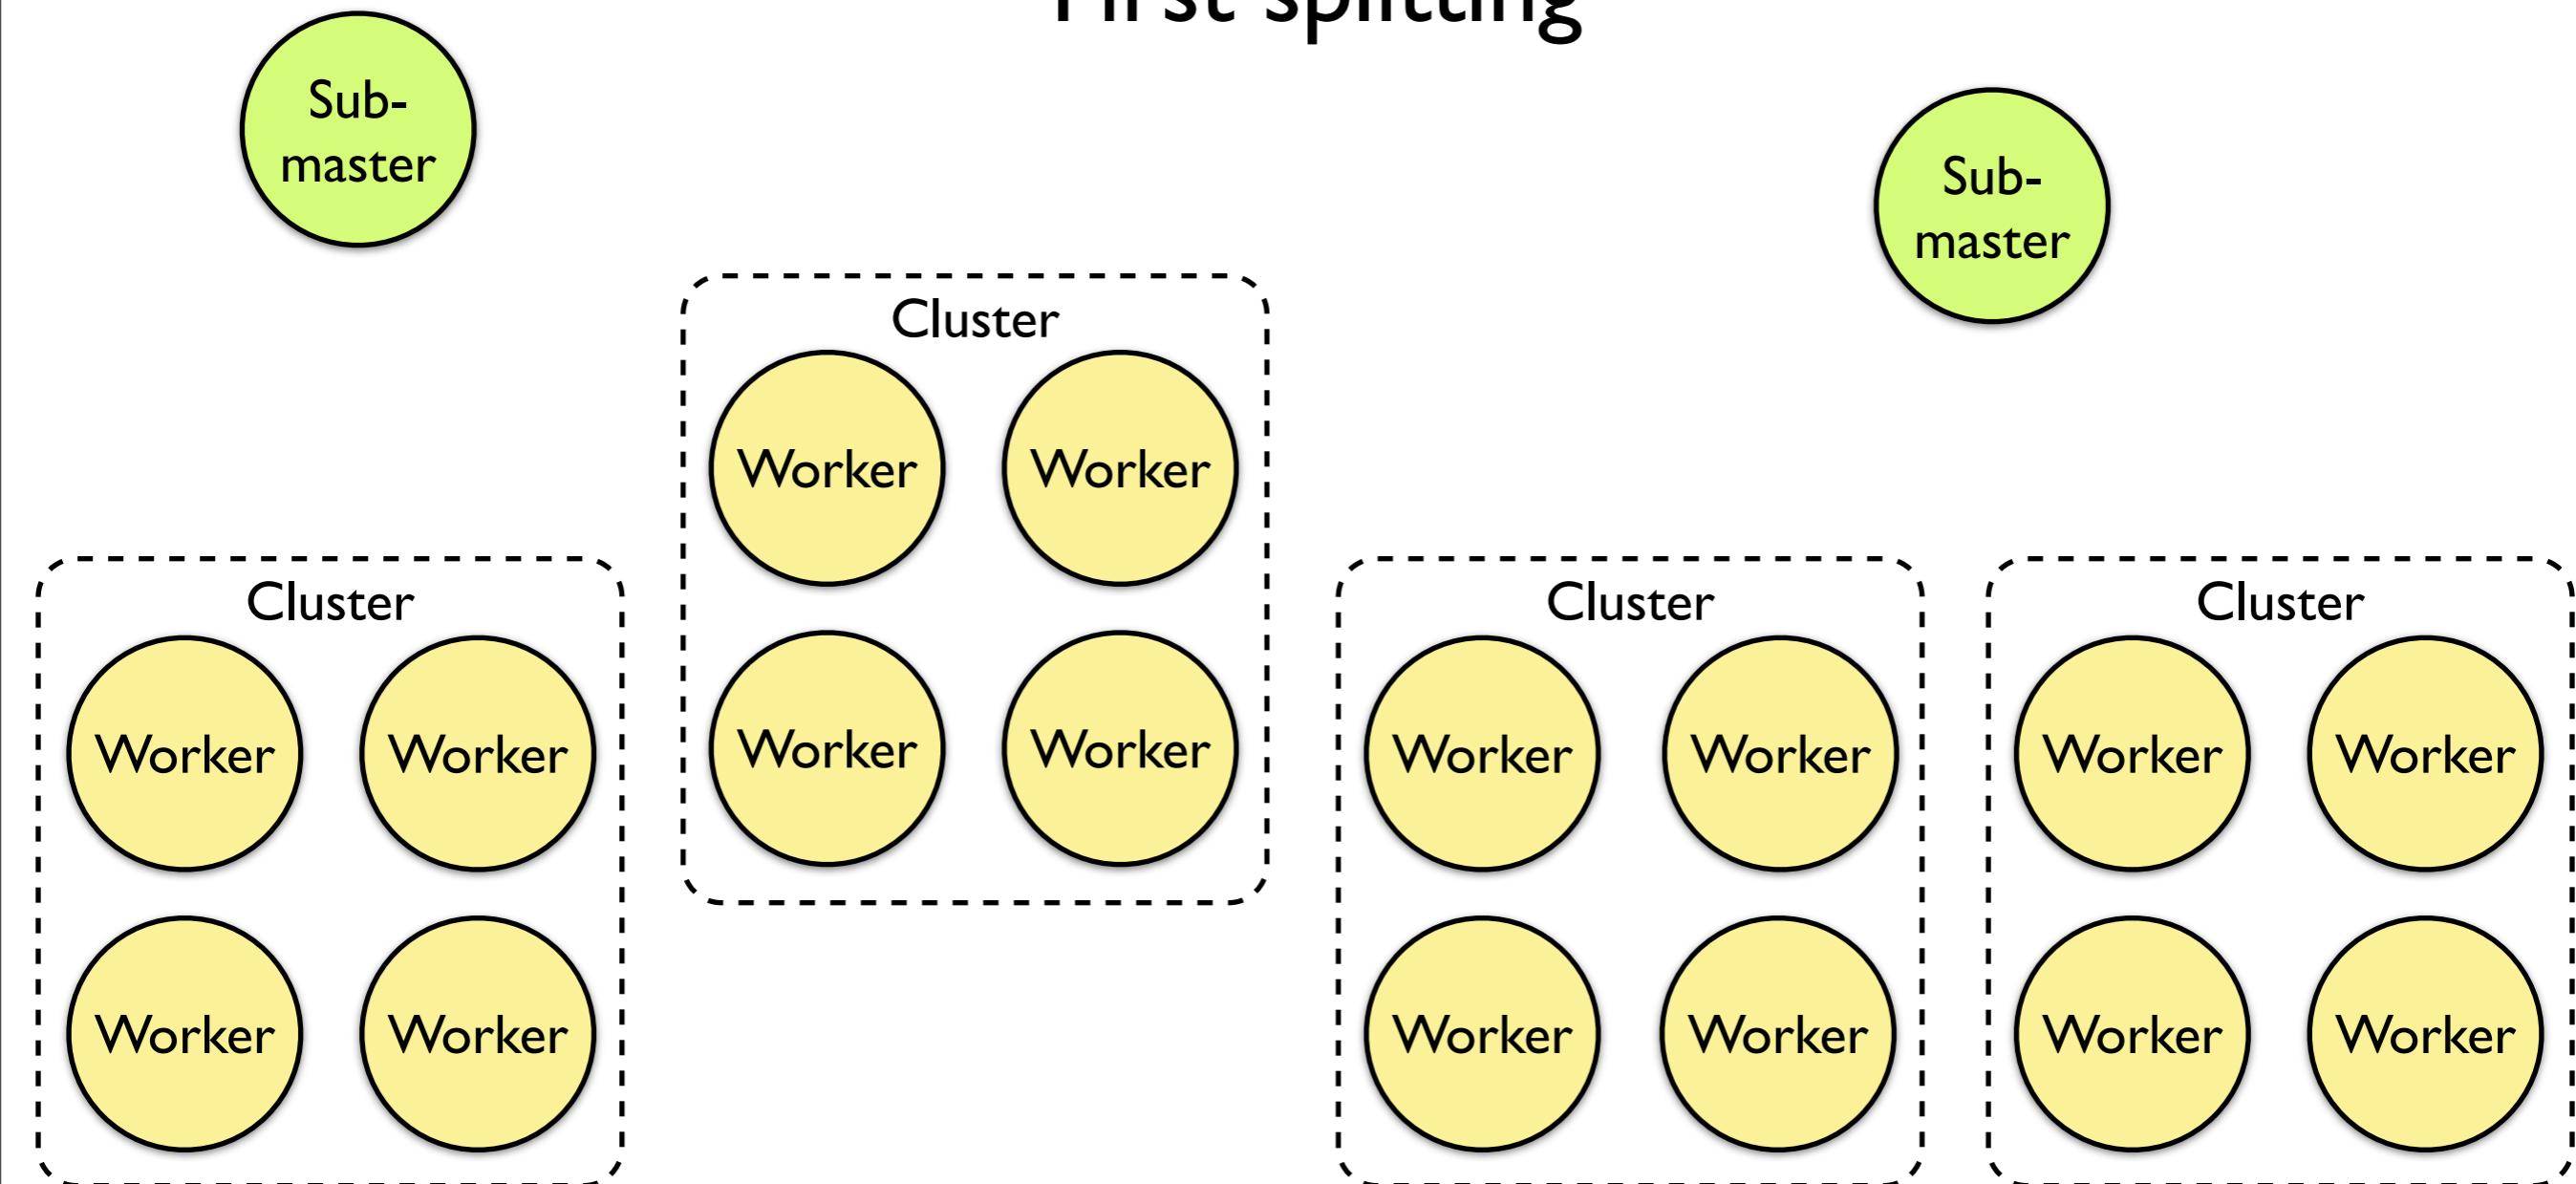
First splitting



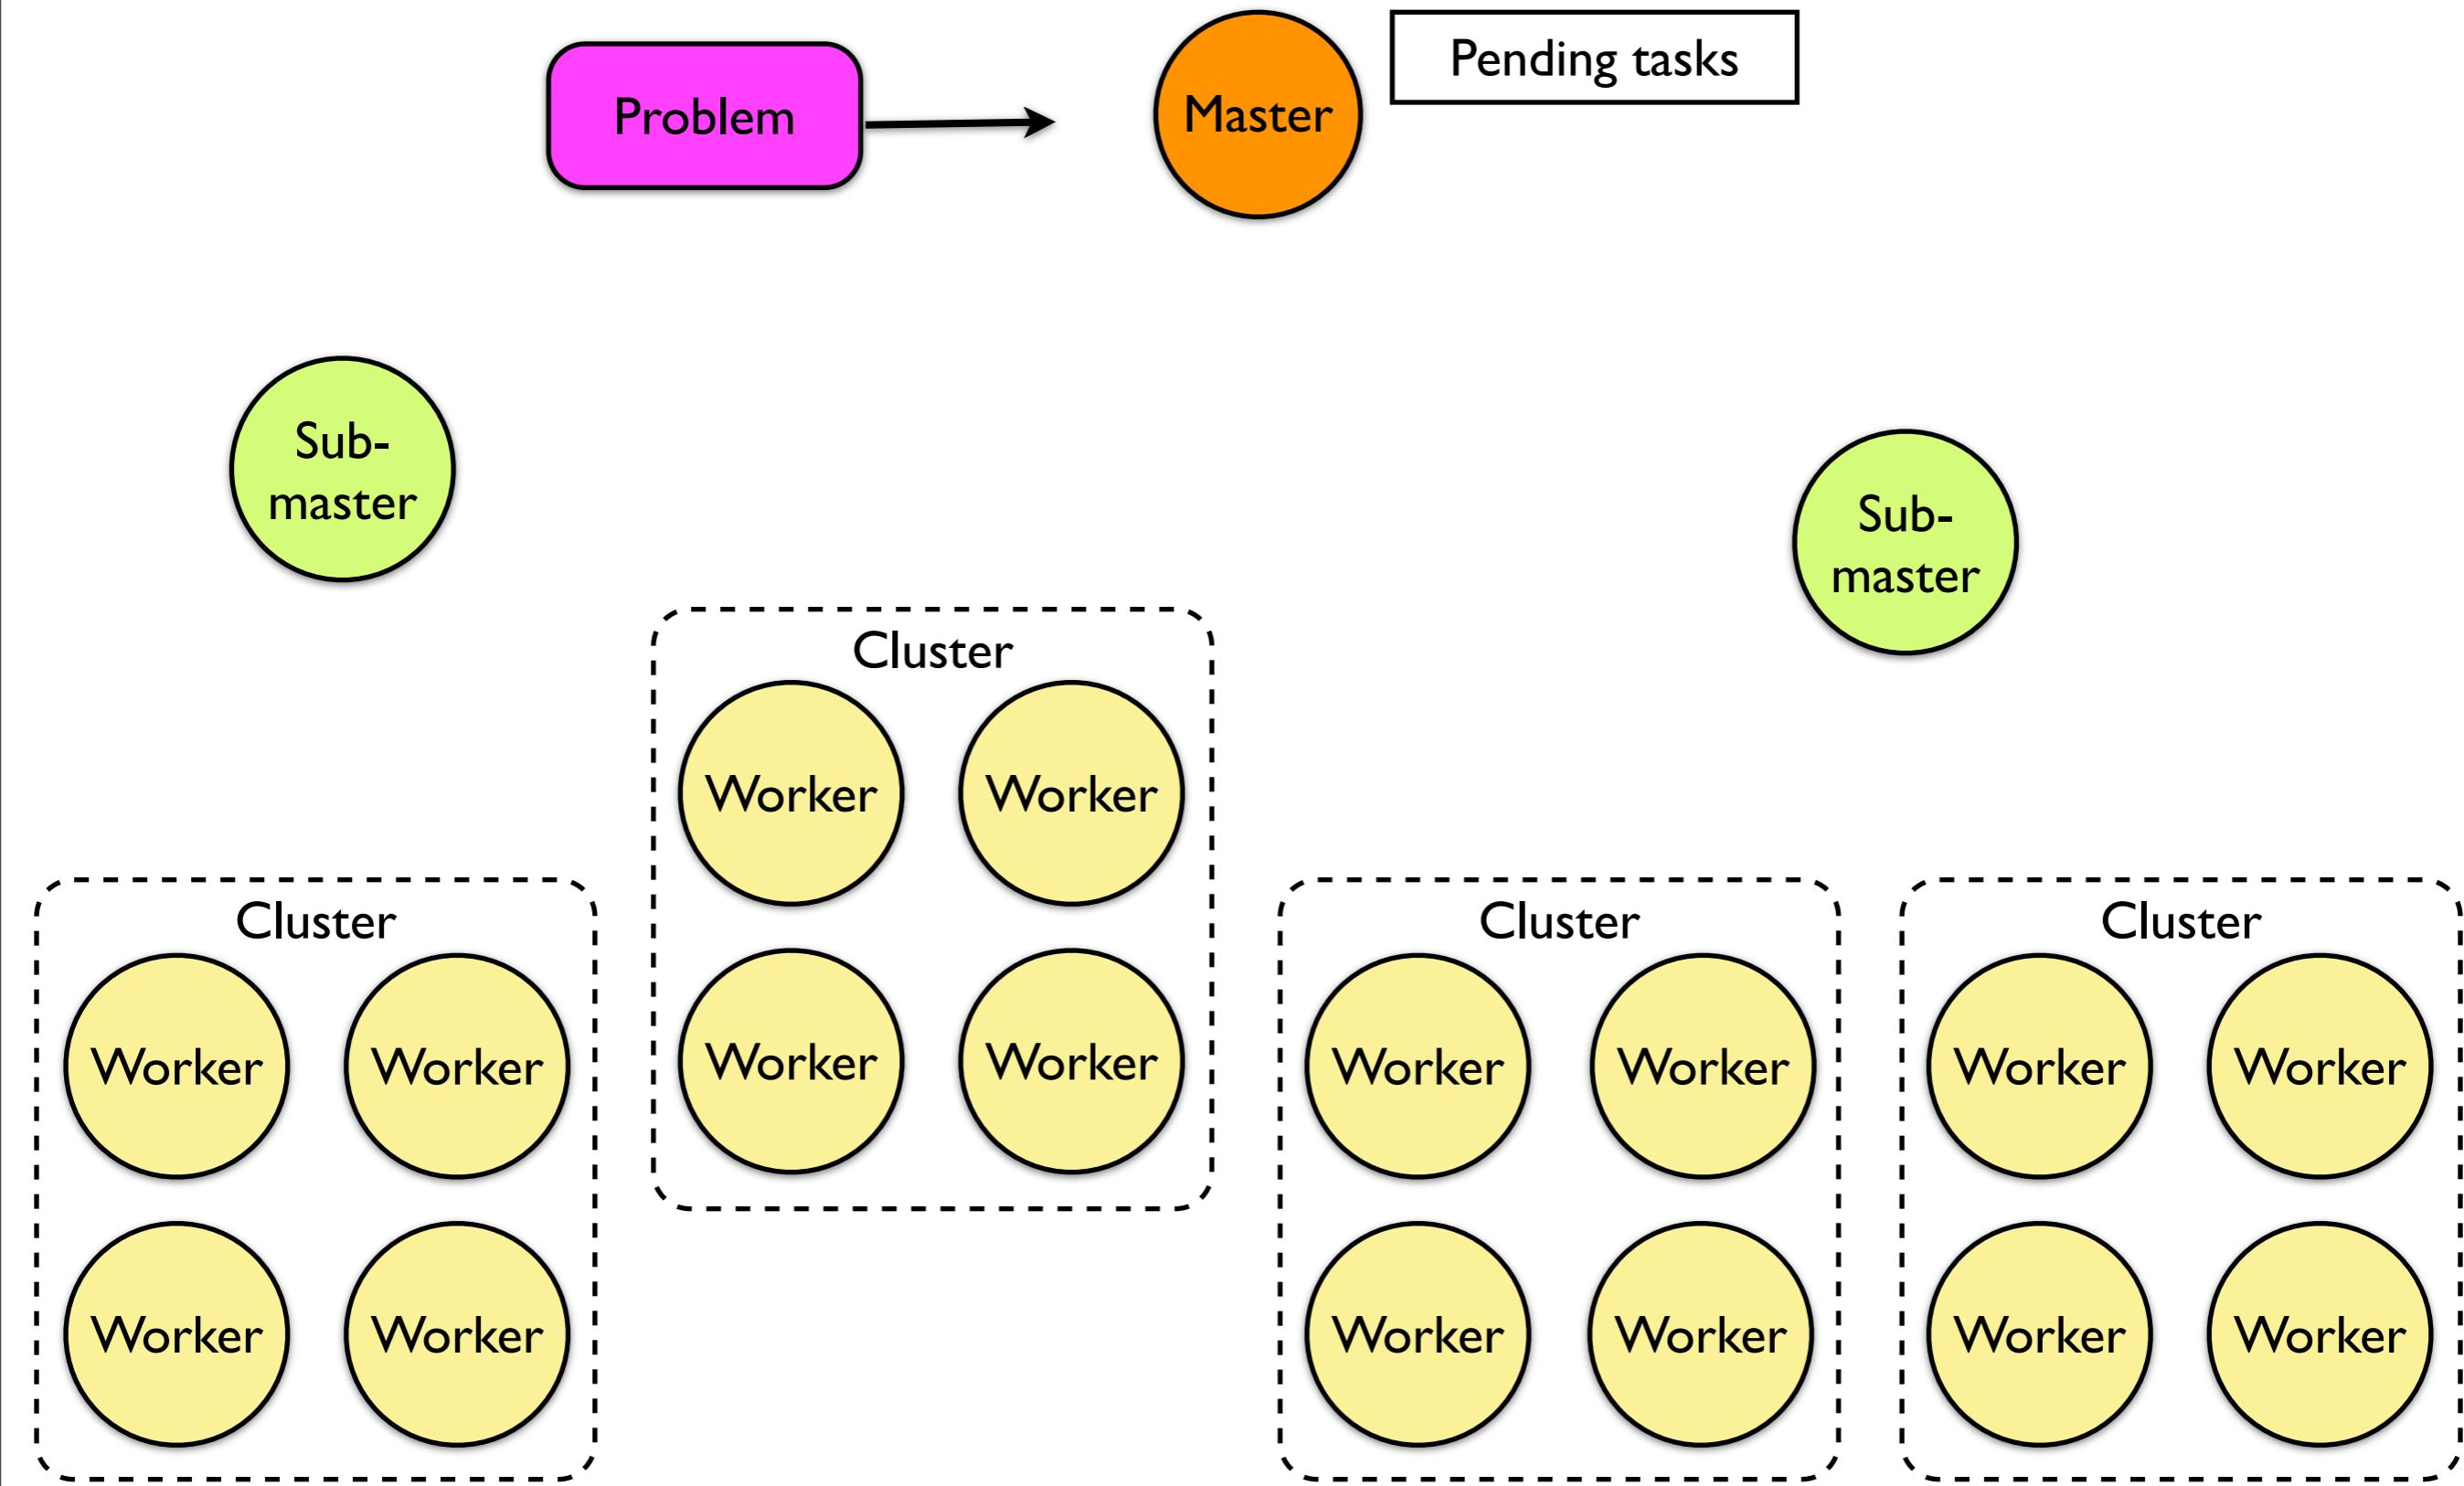
BnB Framework - Architecture



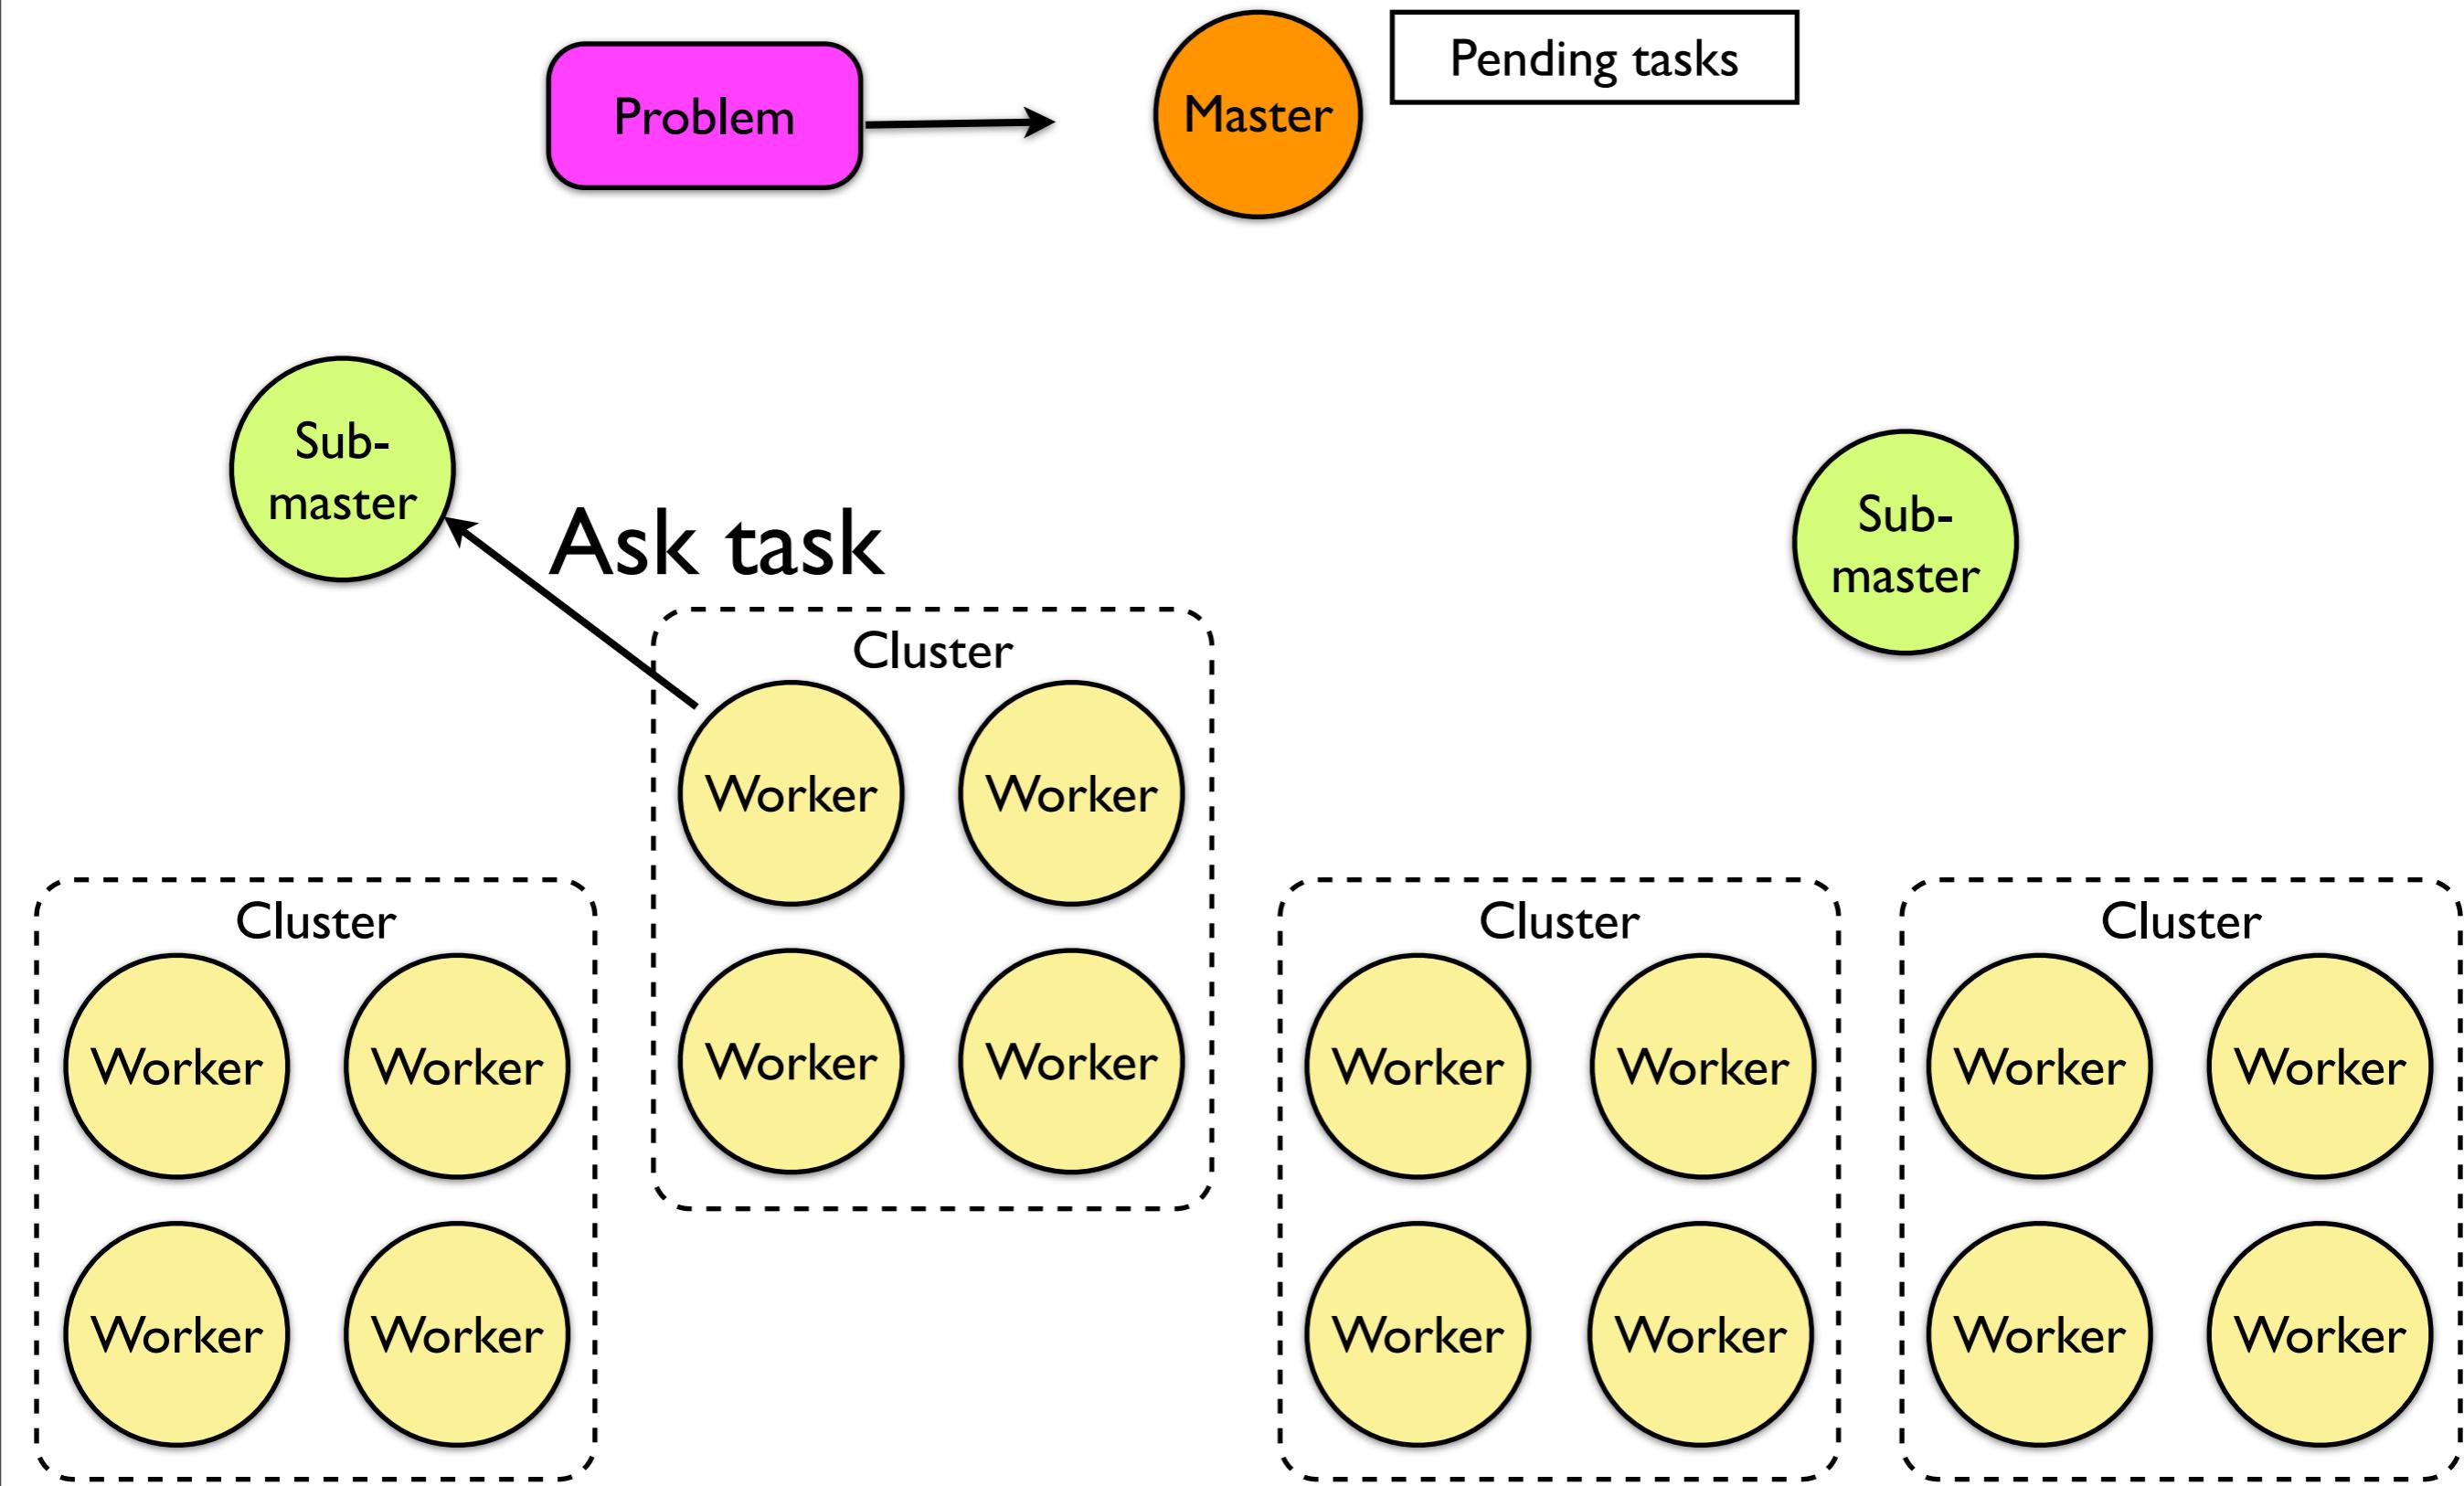
First splitting



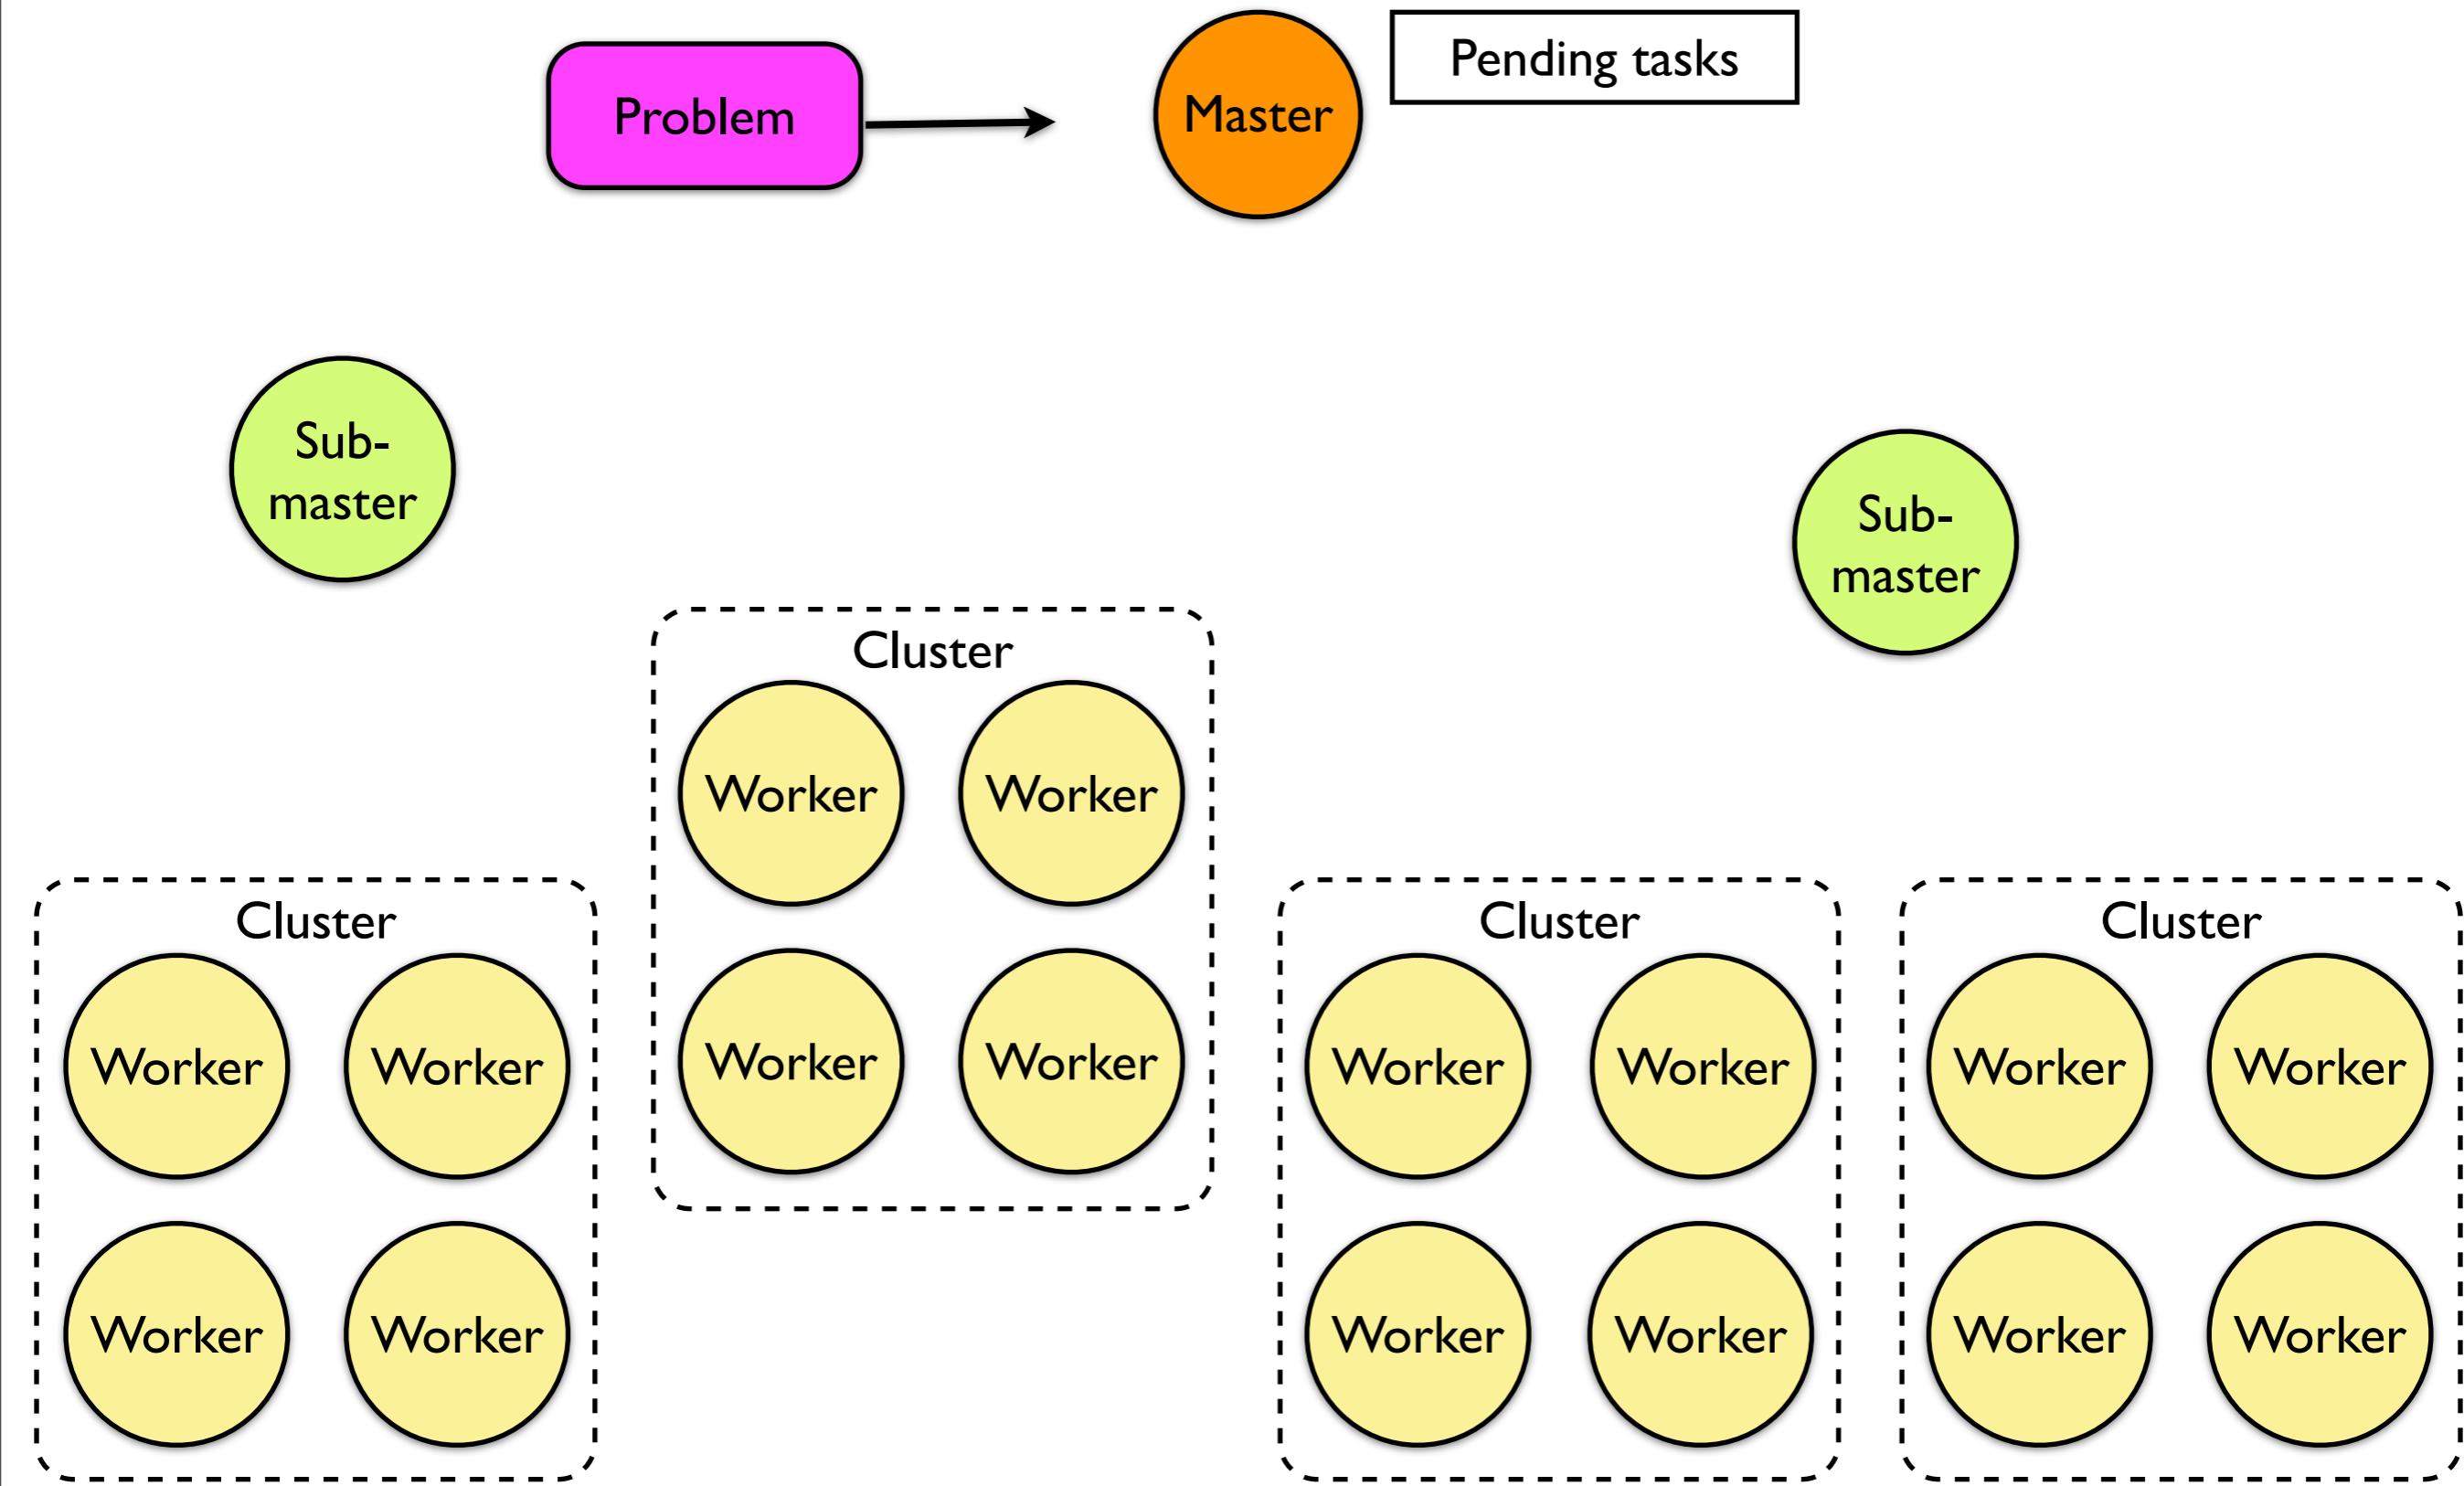
BnB Framework - Architecture



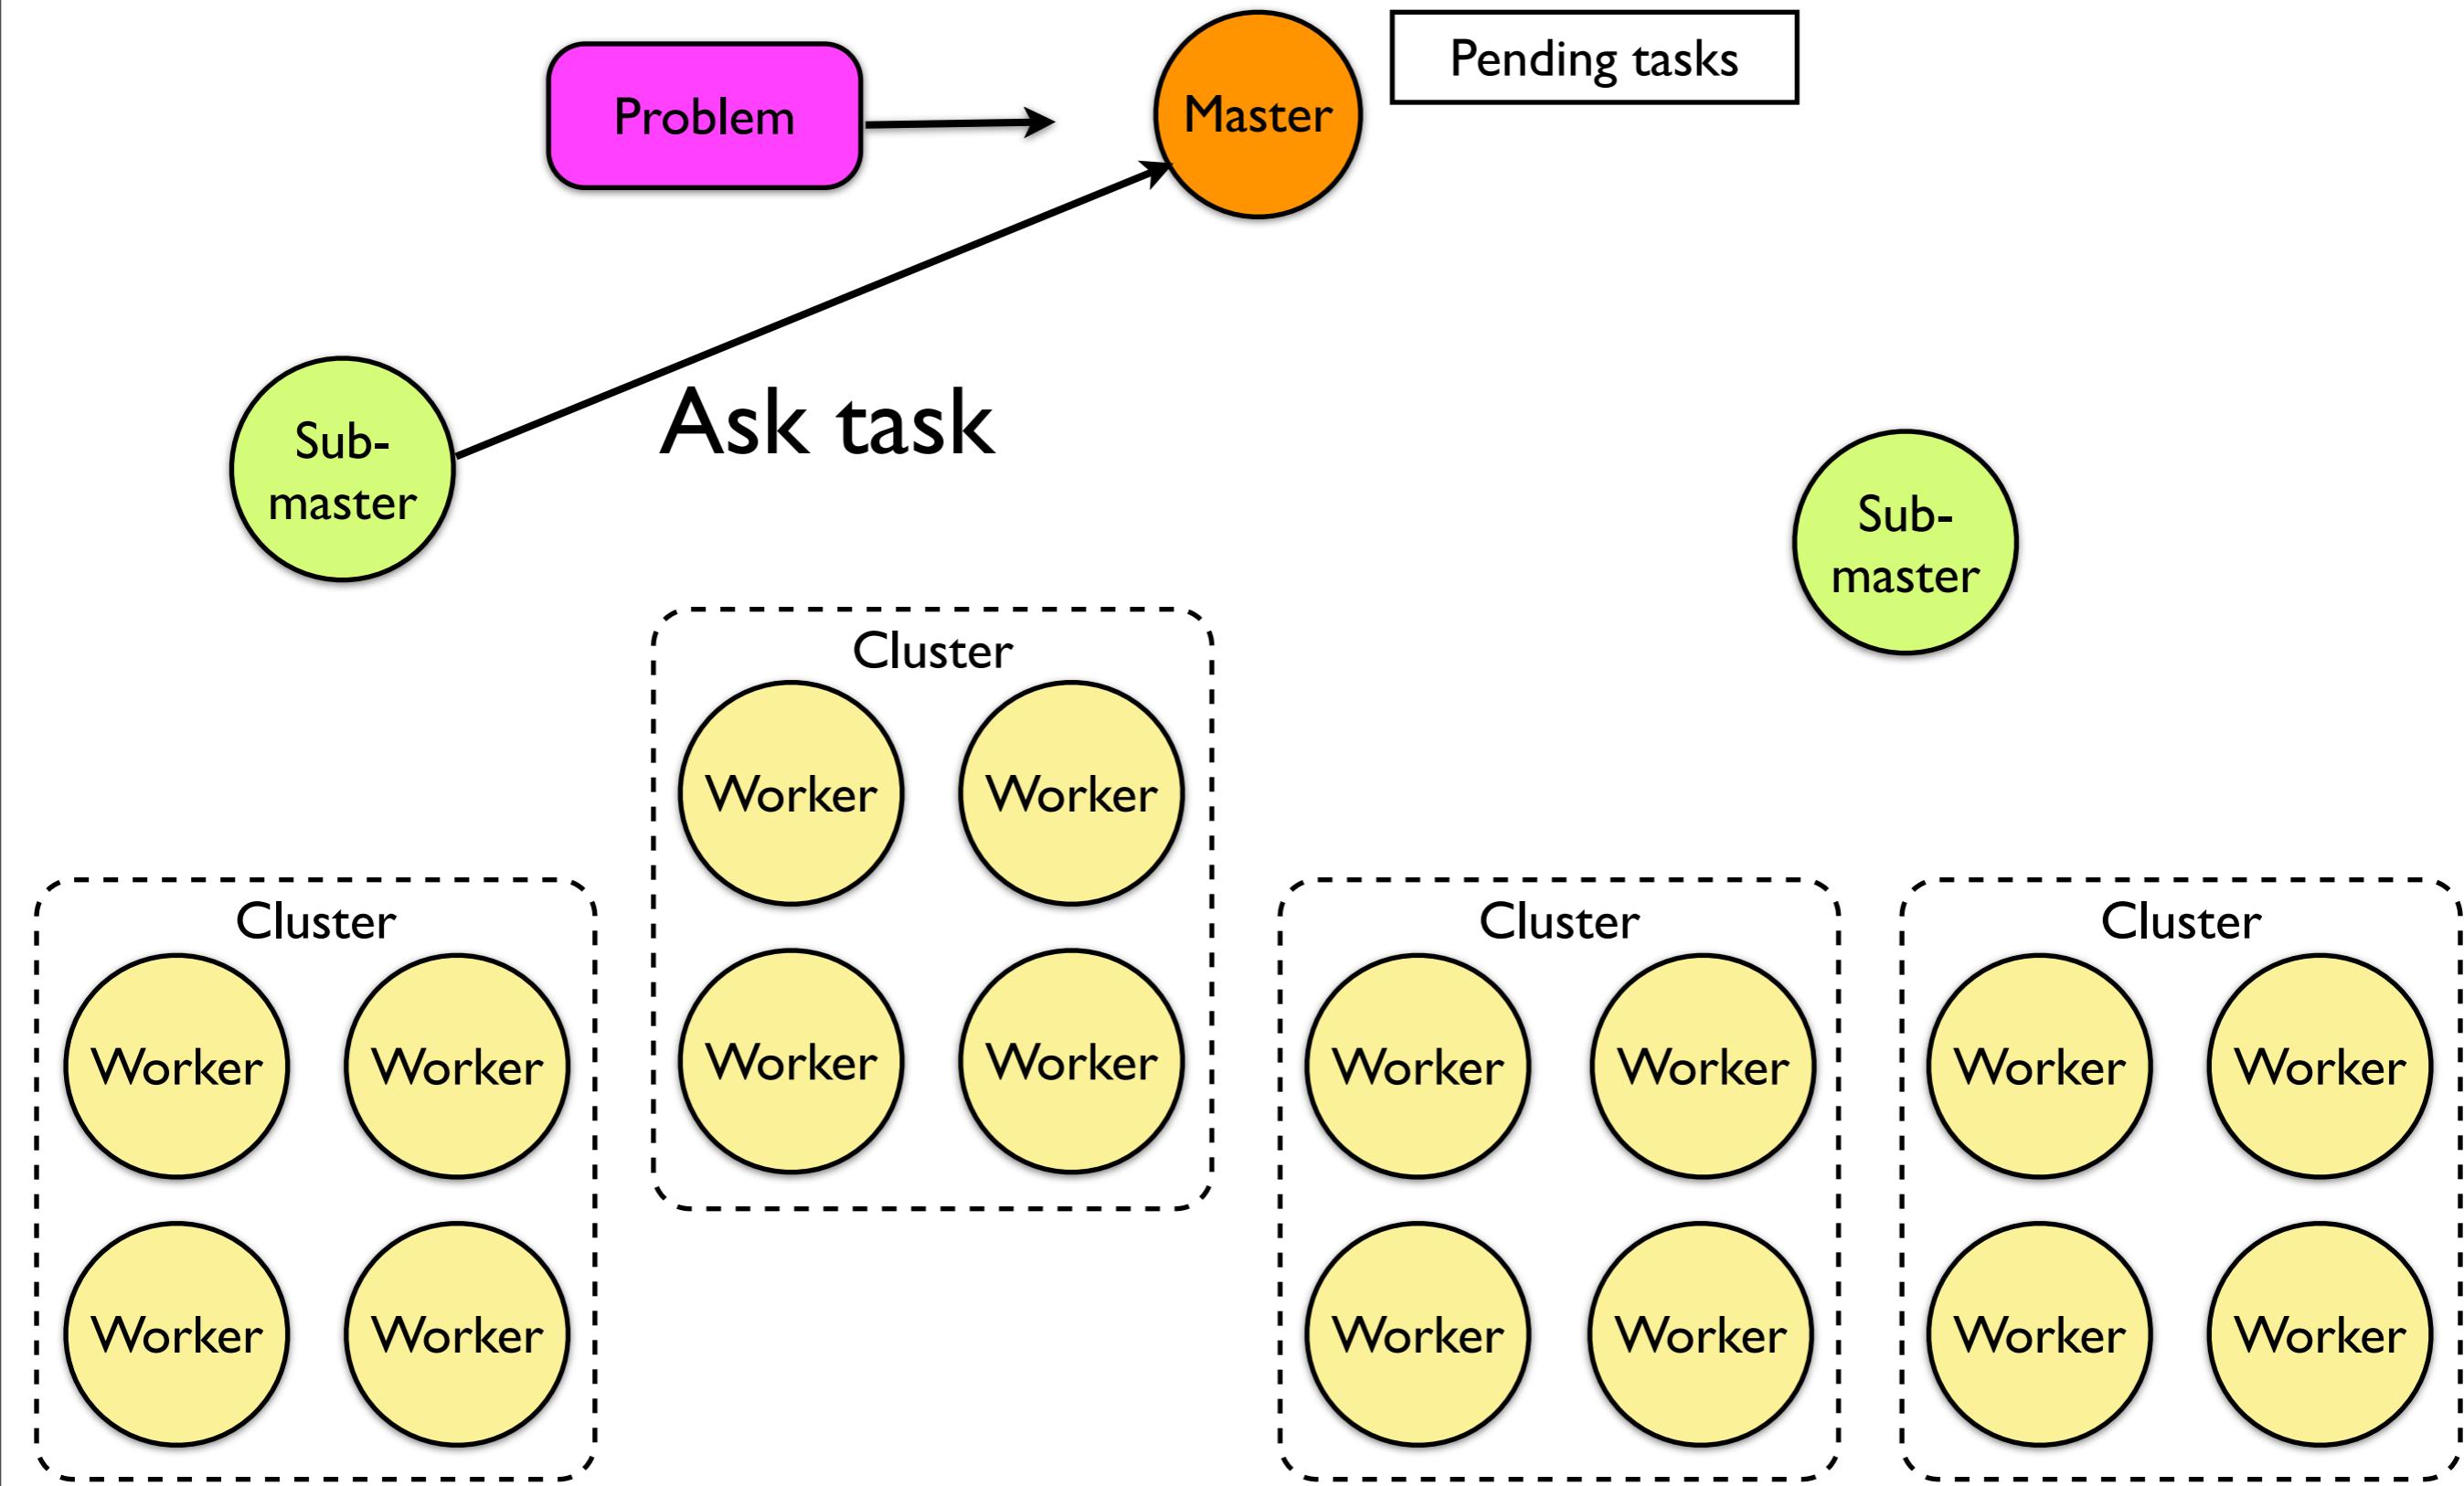
BnB Framework - Architecture



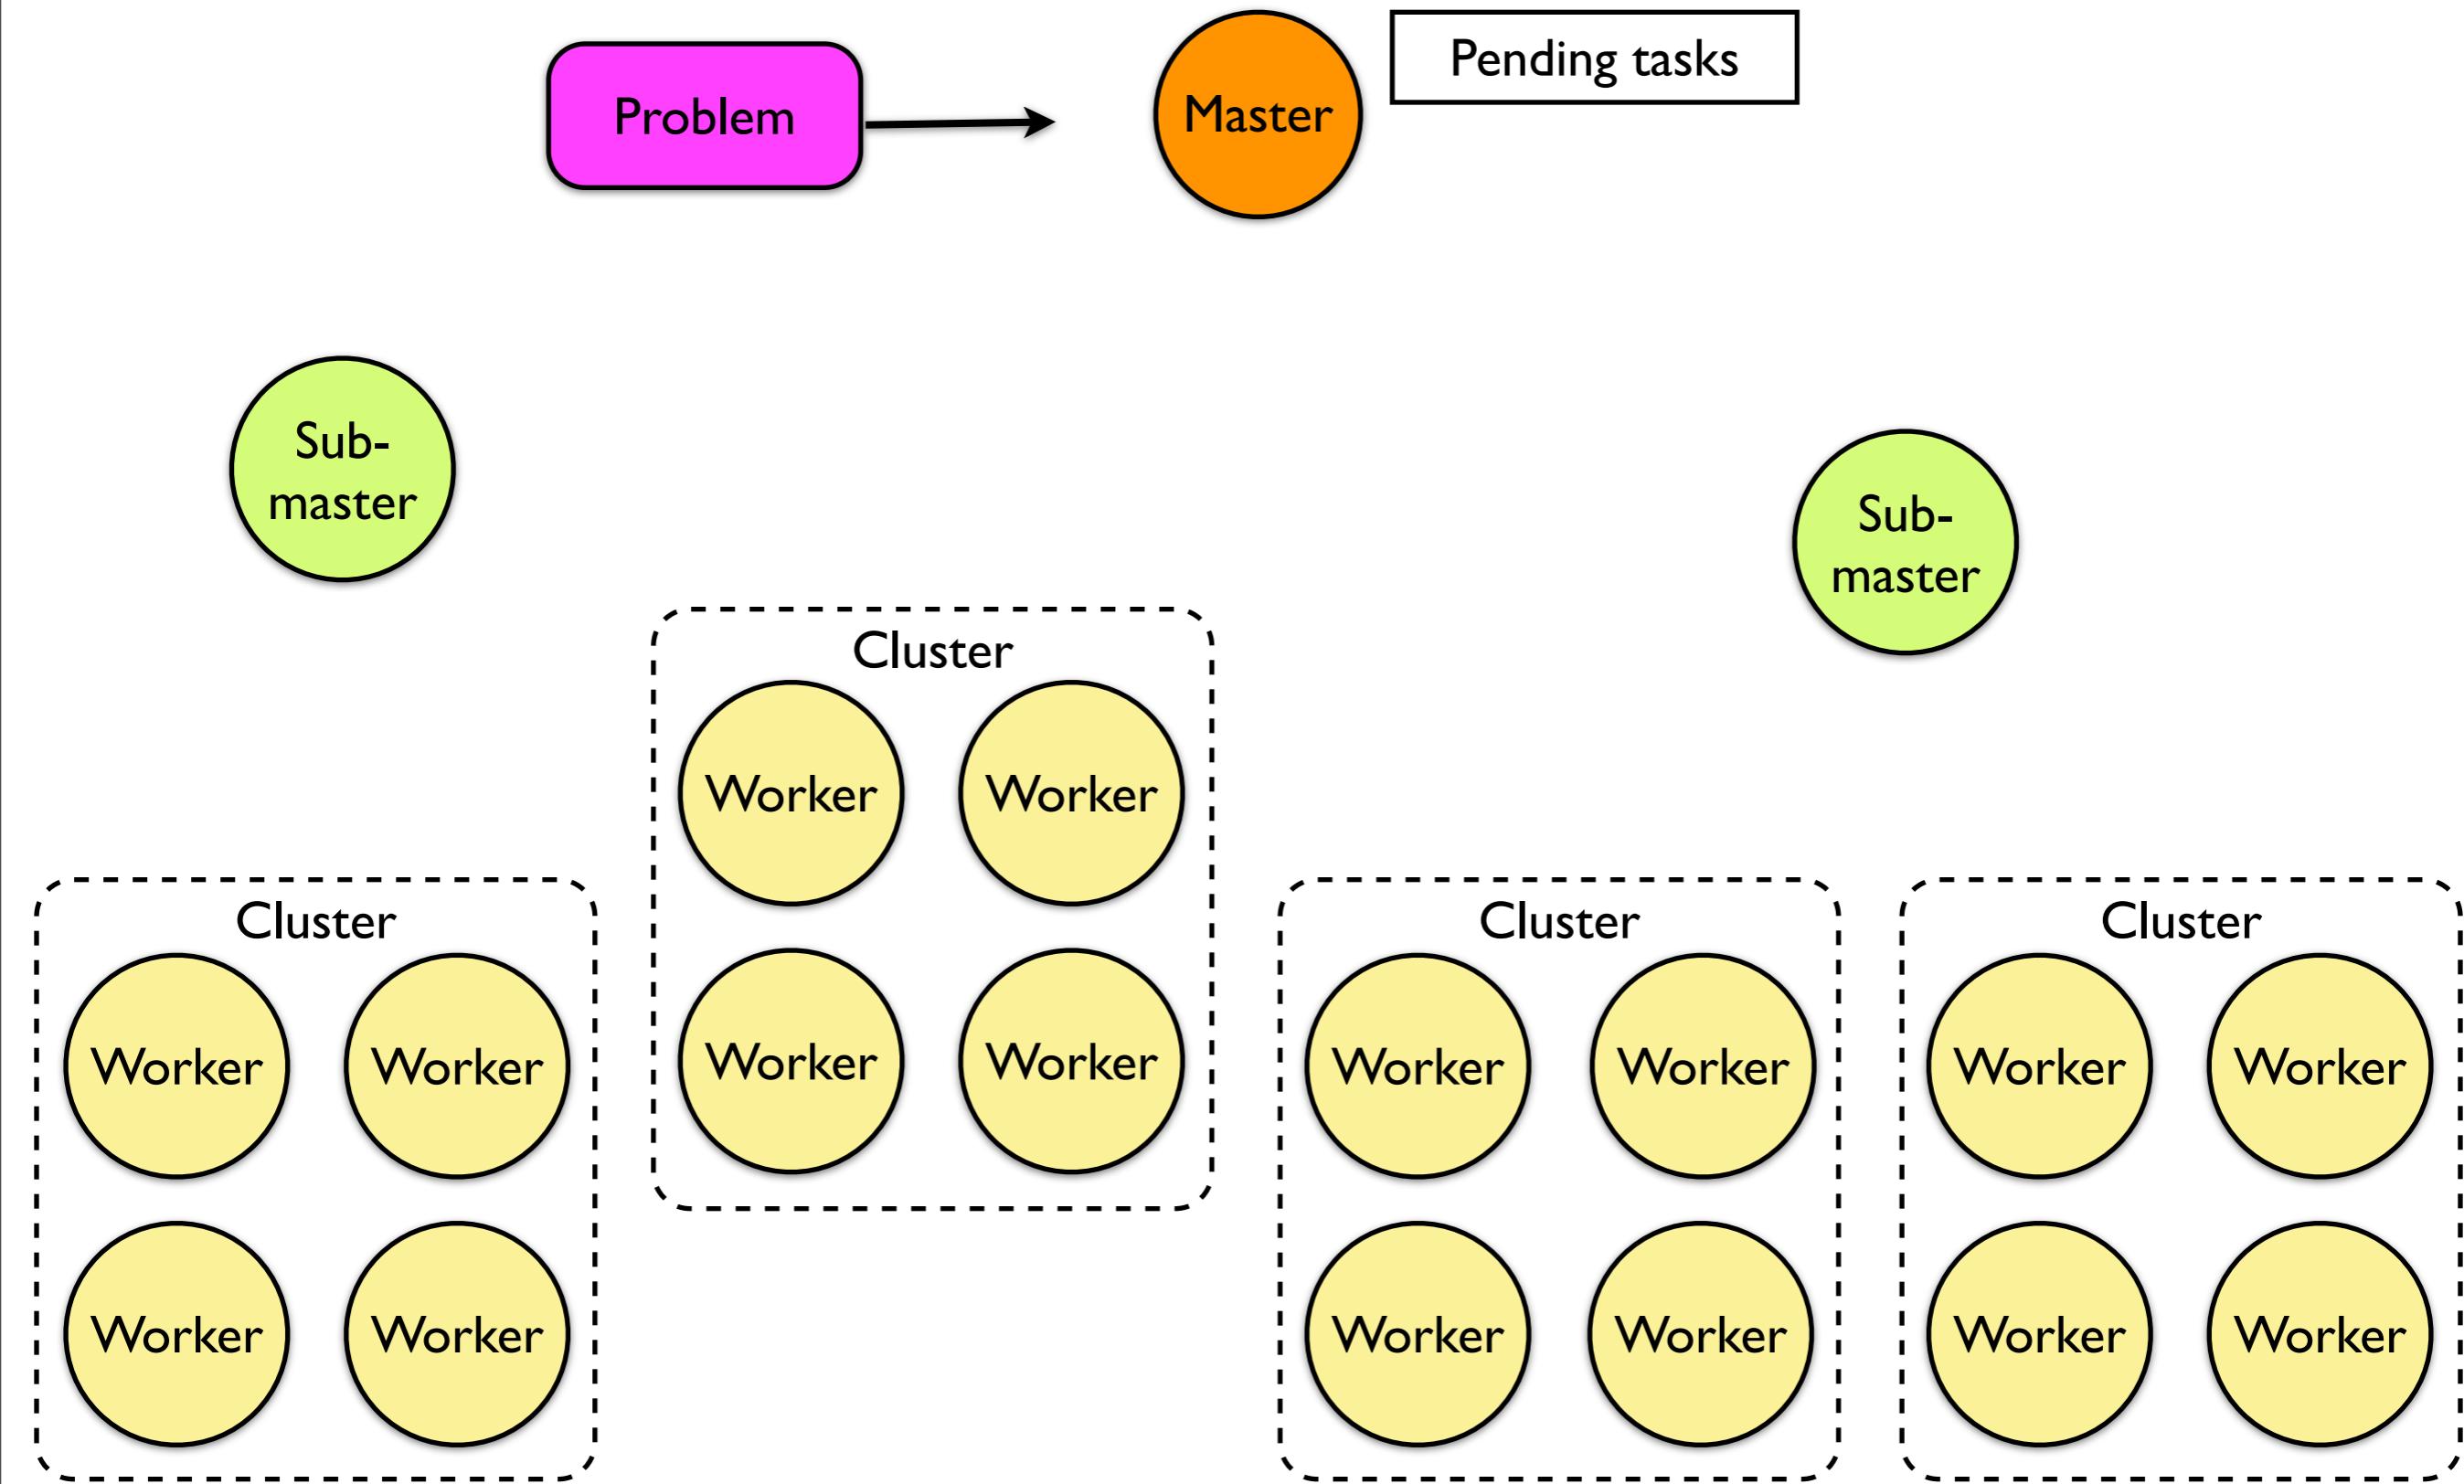
BnB Framework - Architecture



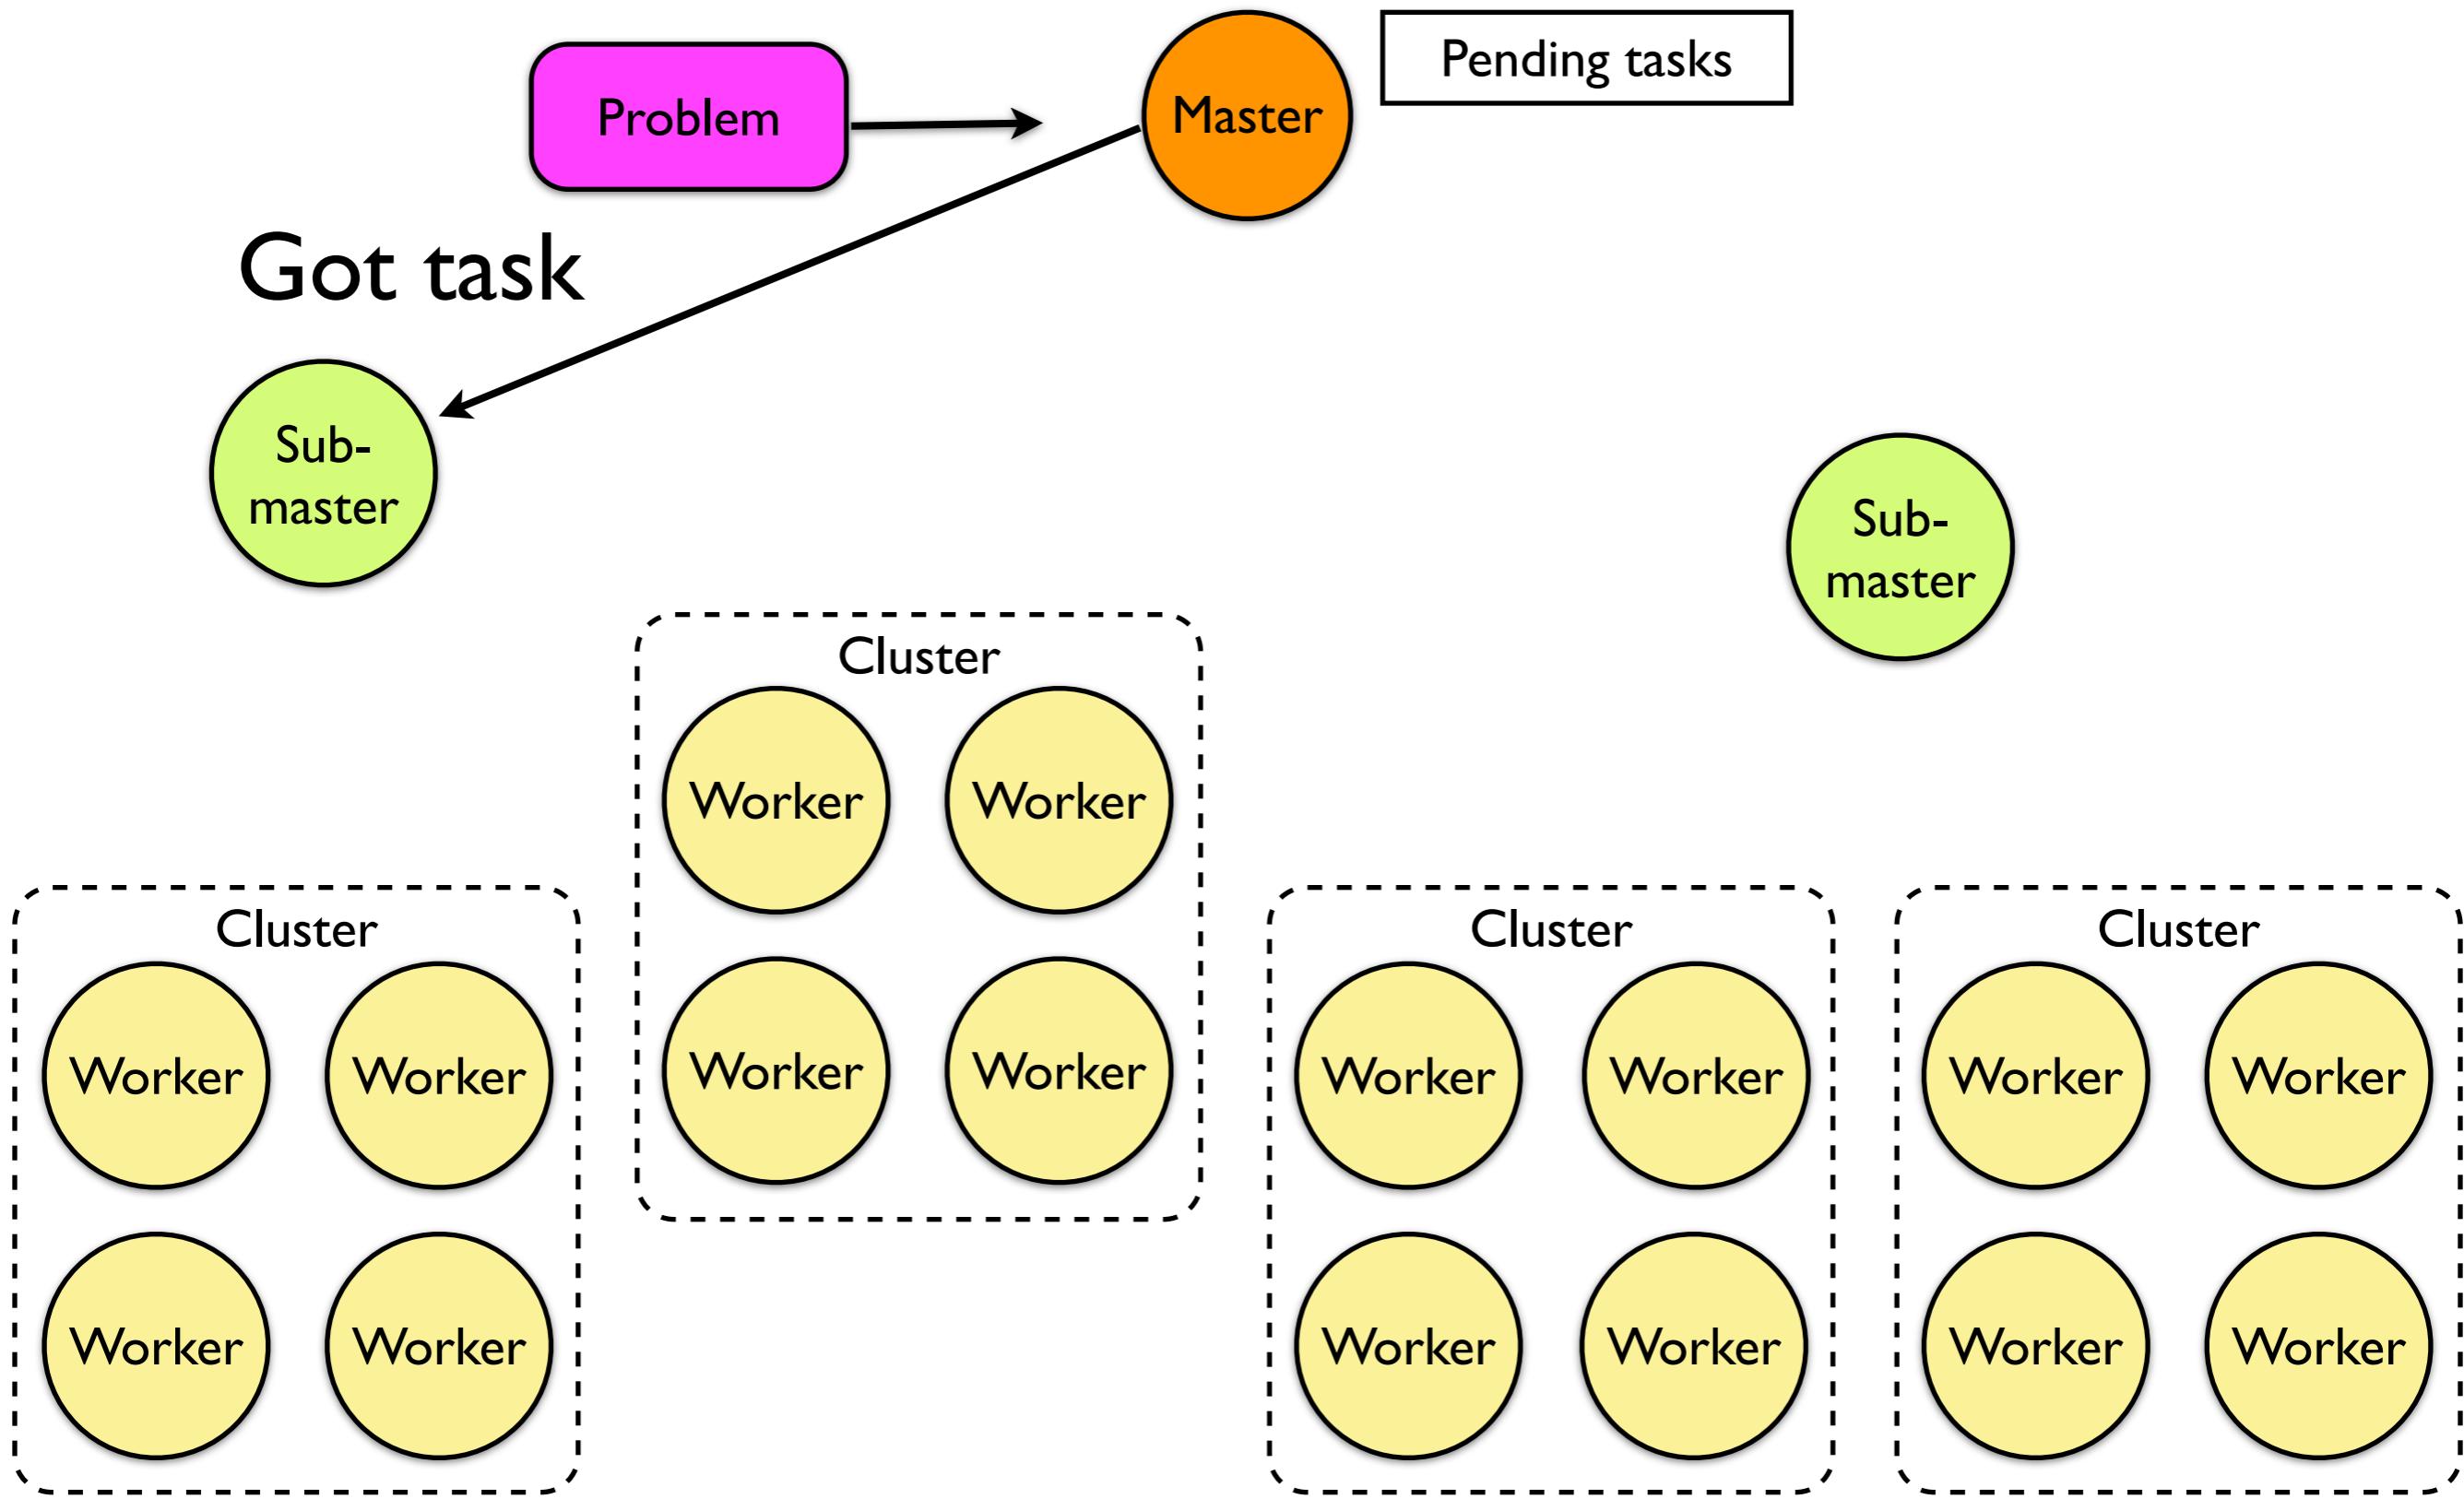
BnB Framework - Architecture



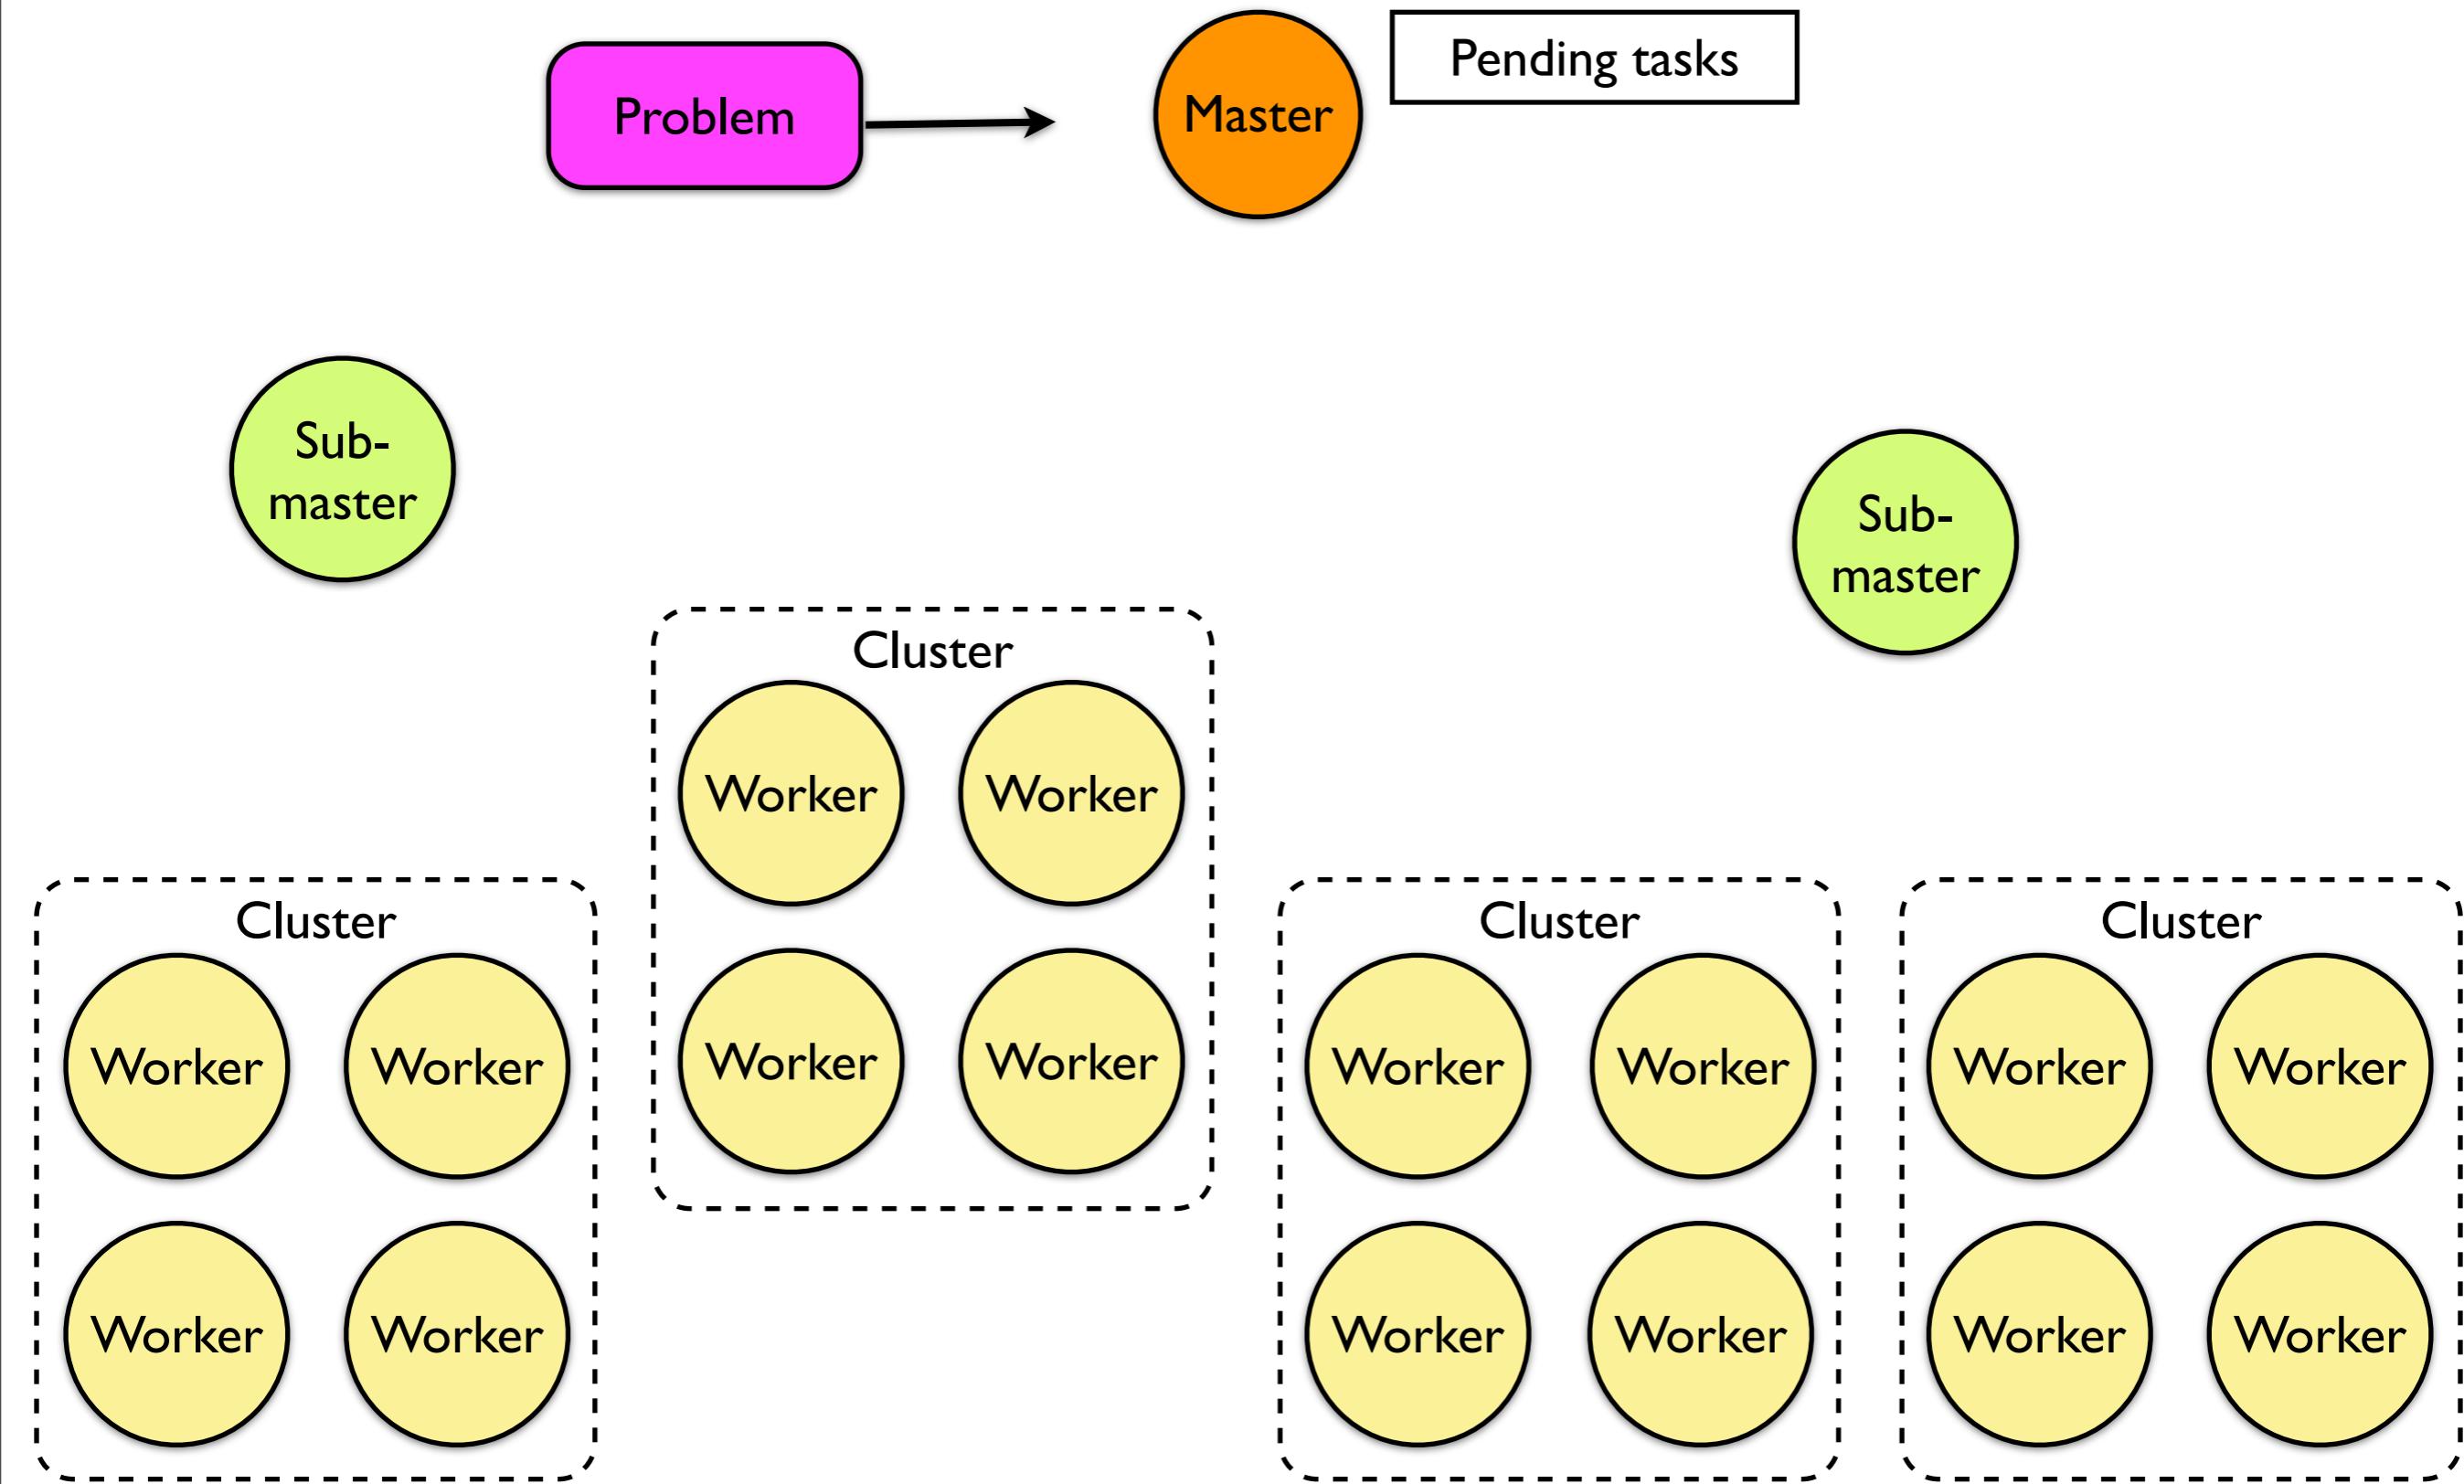
BnB Framework - Architecture



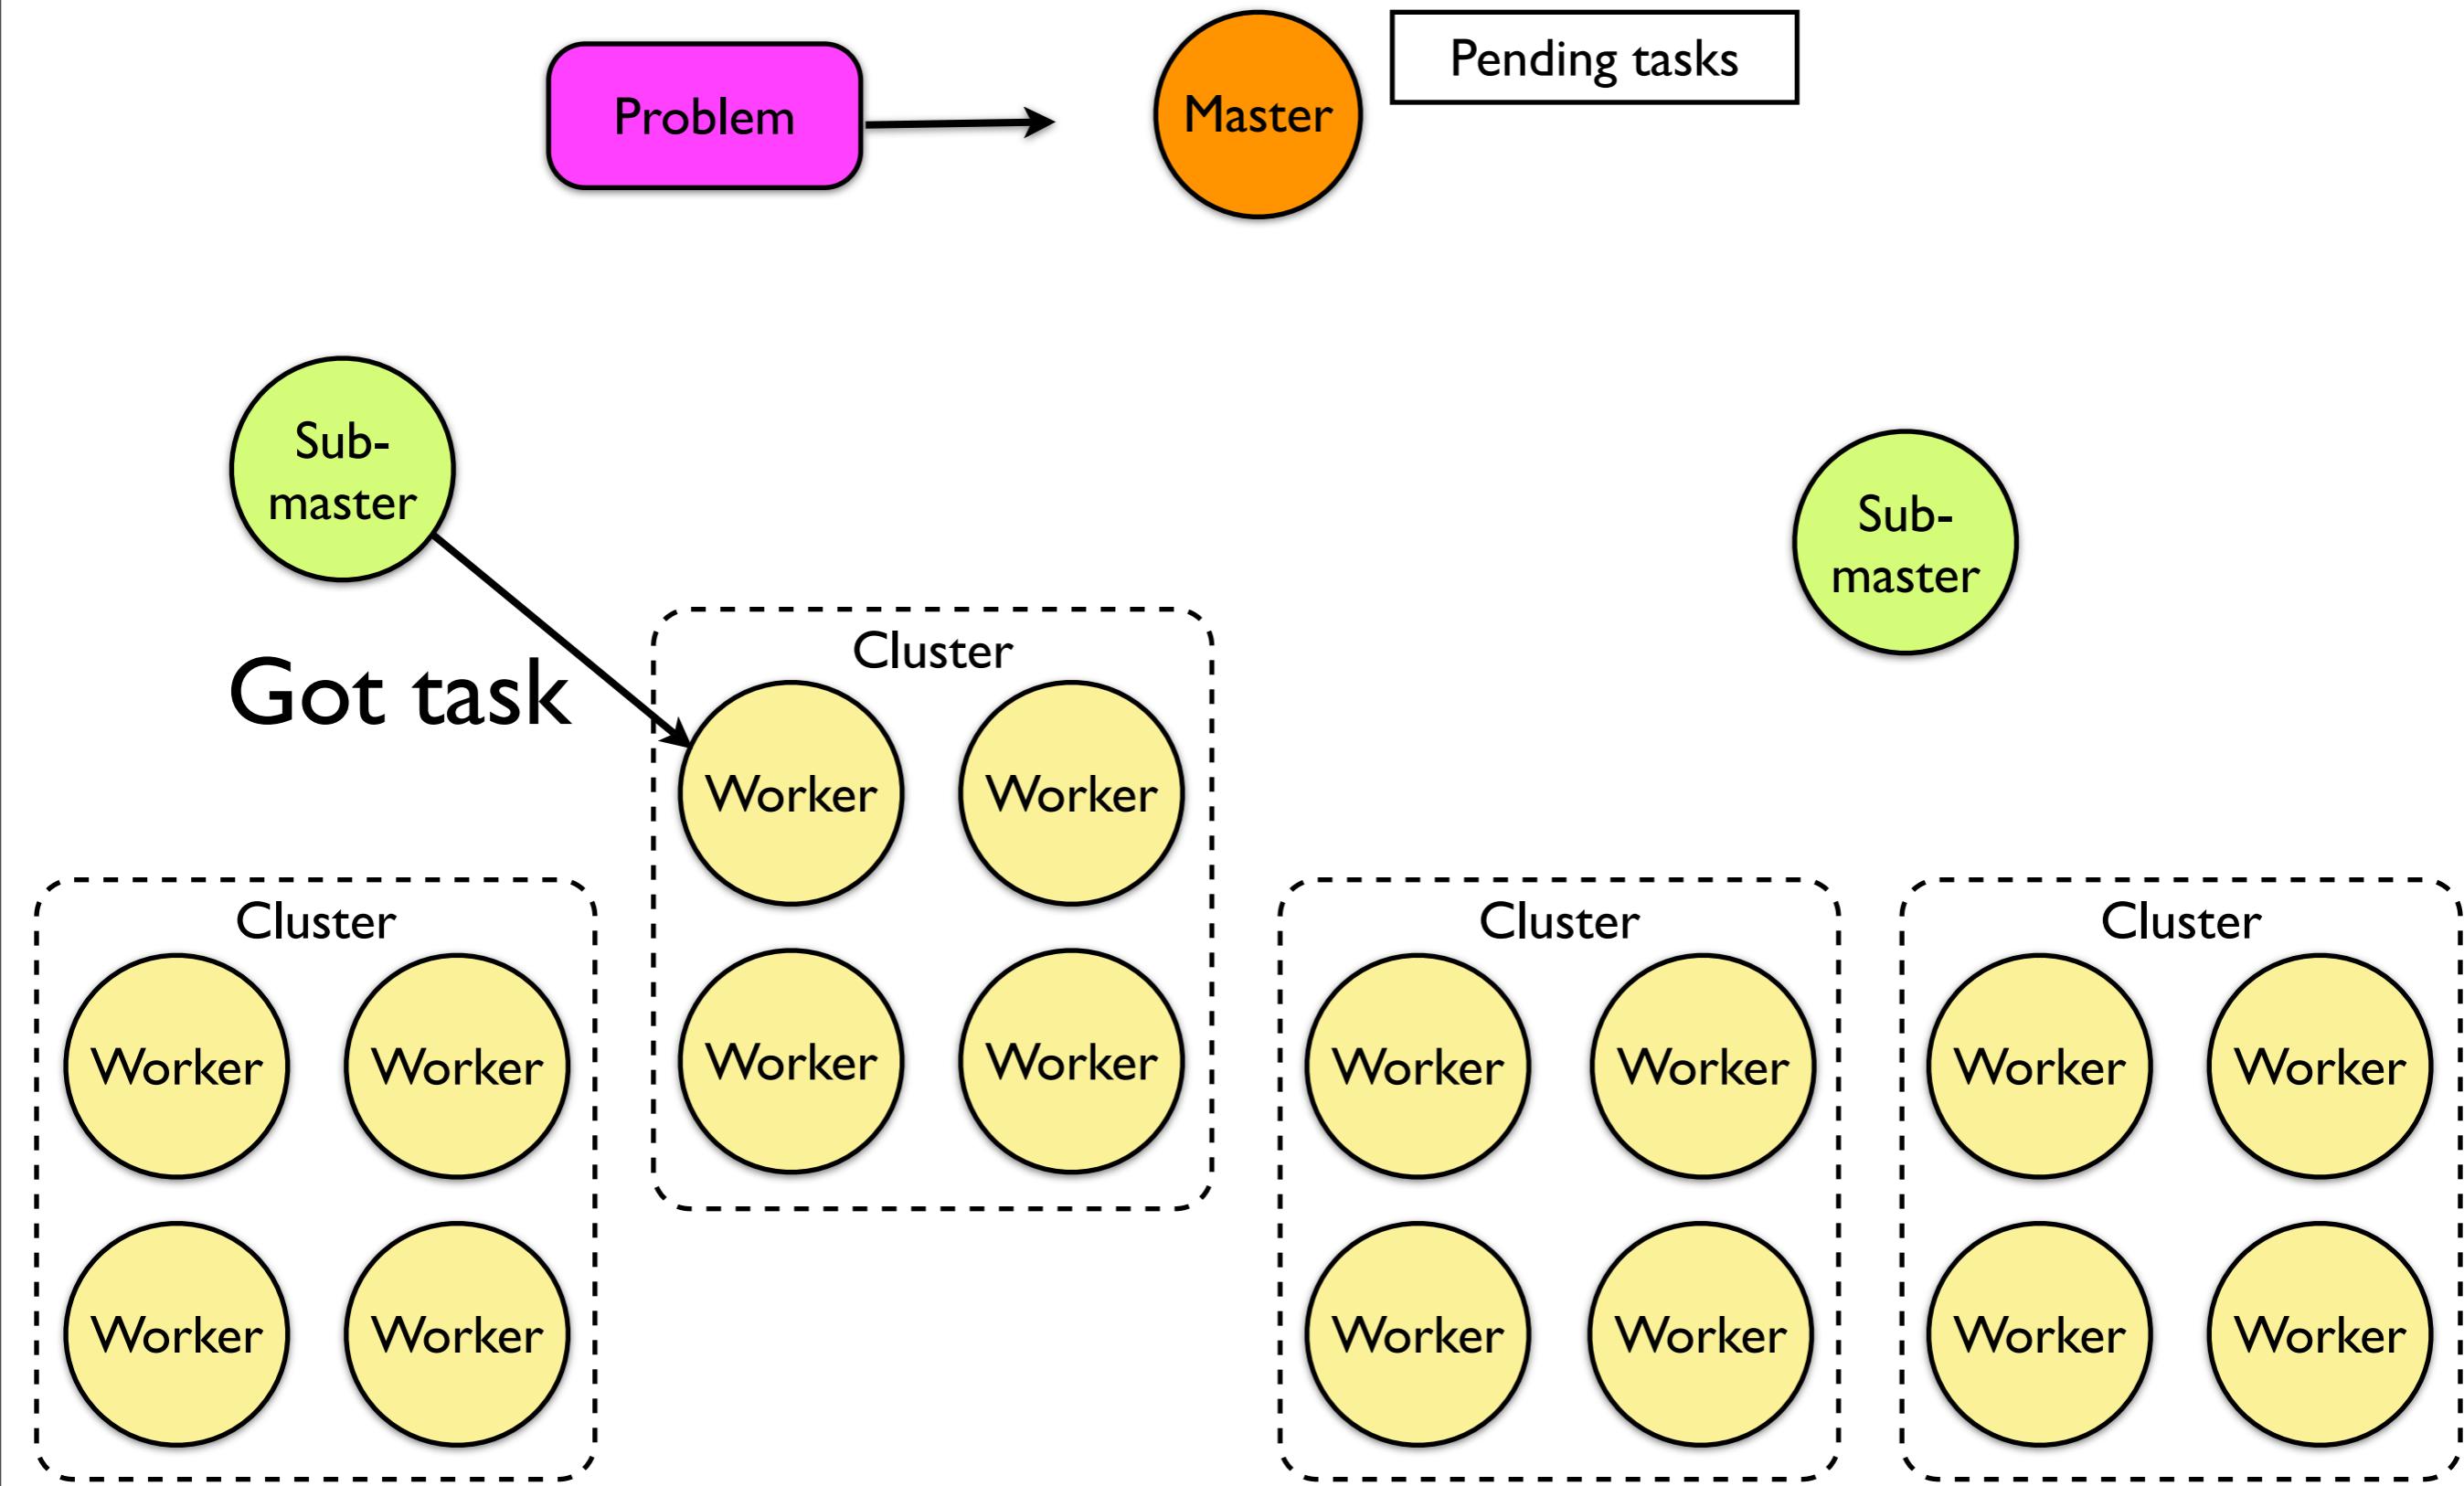
BnB Framework - Architecture



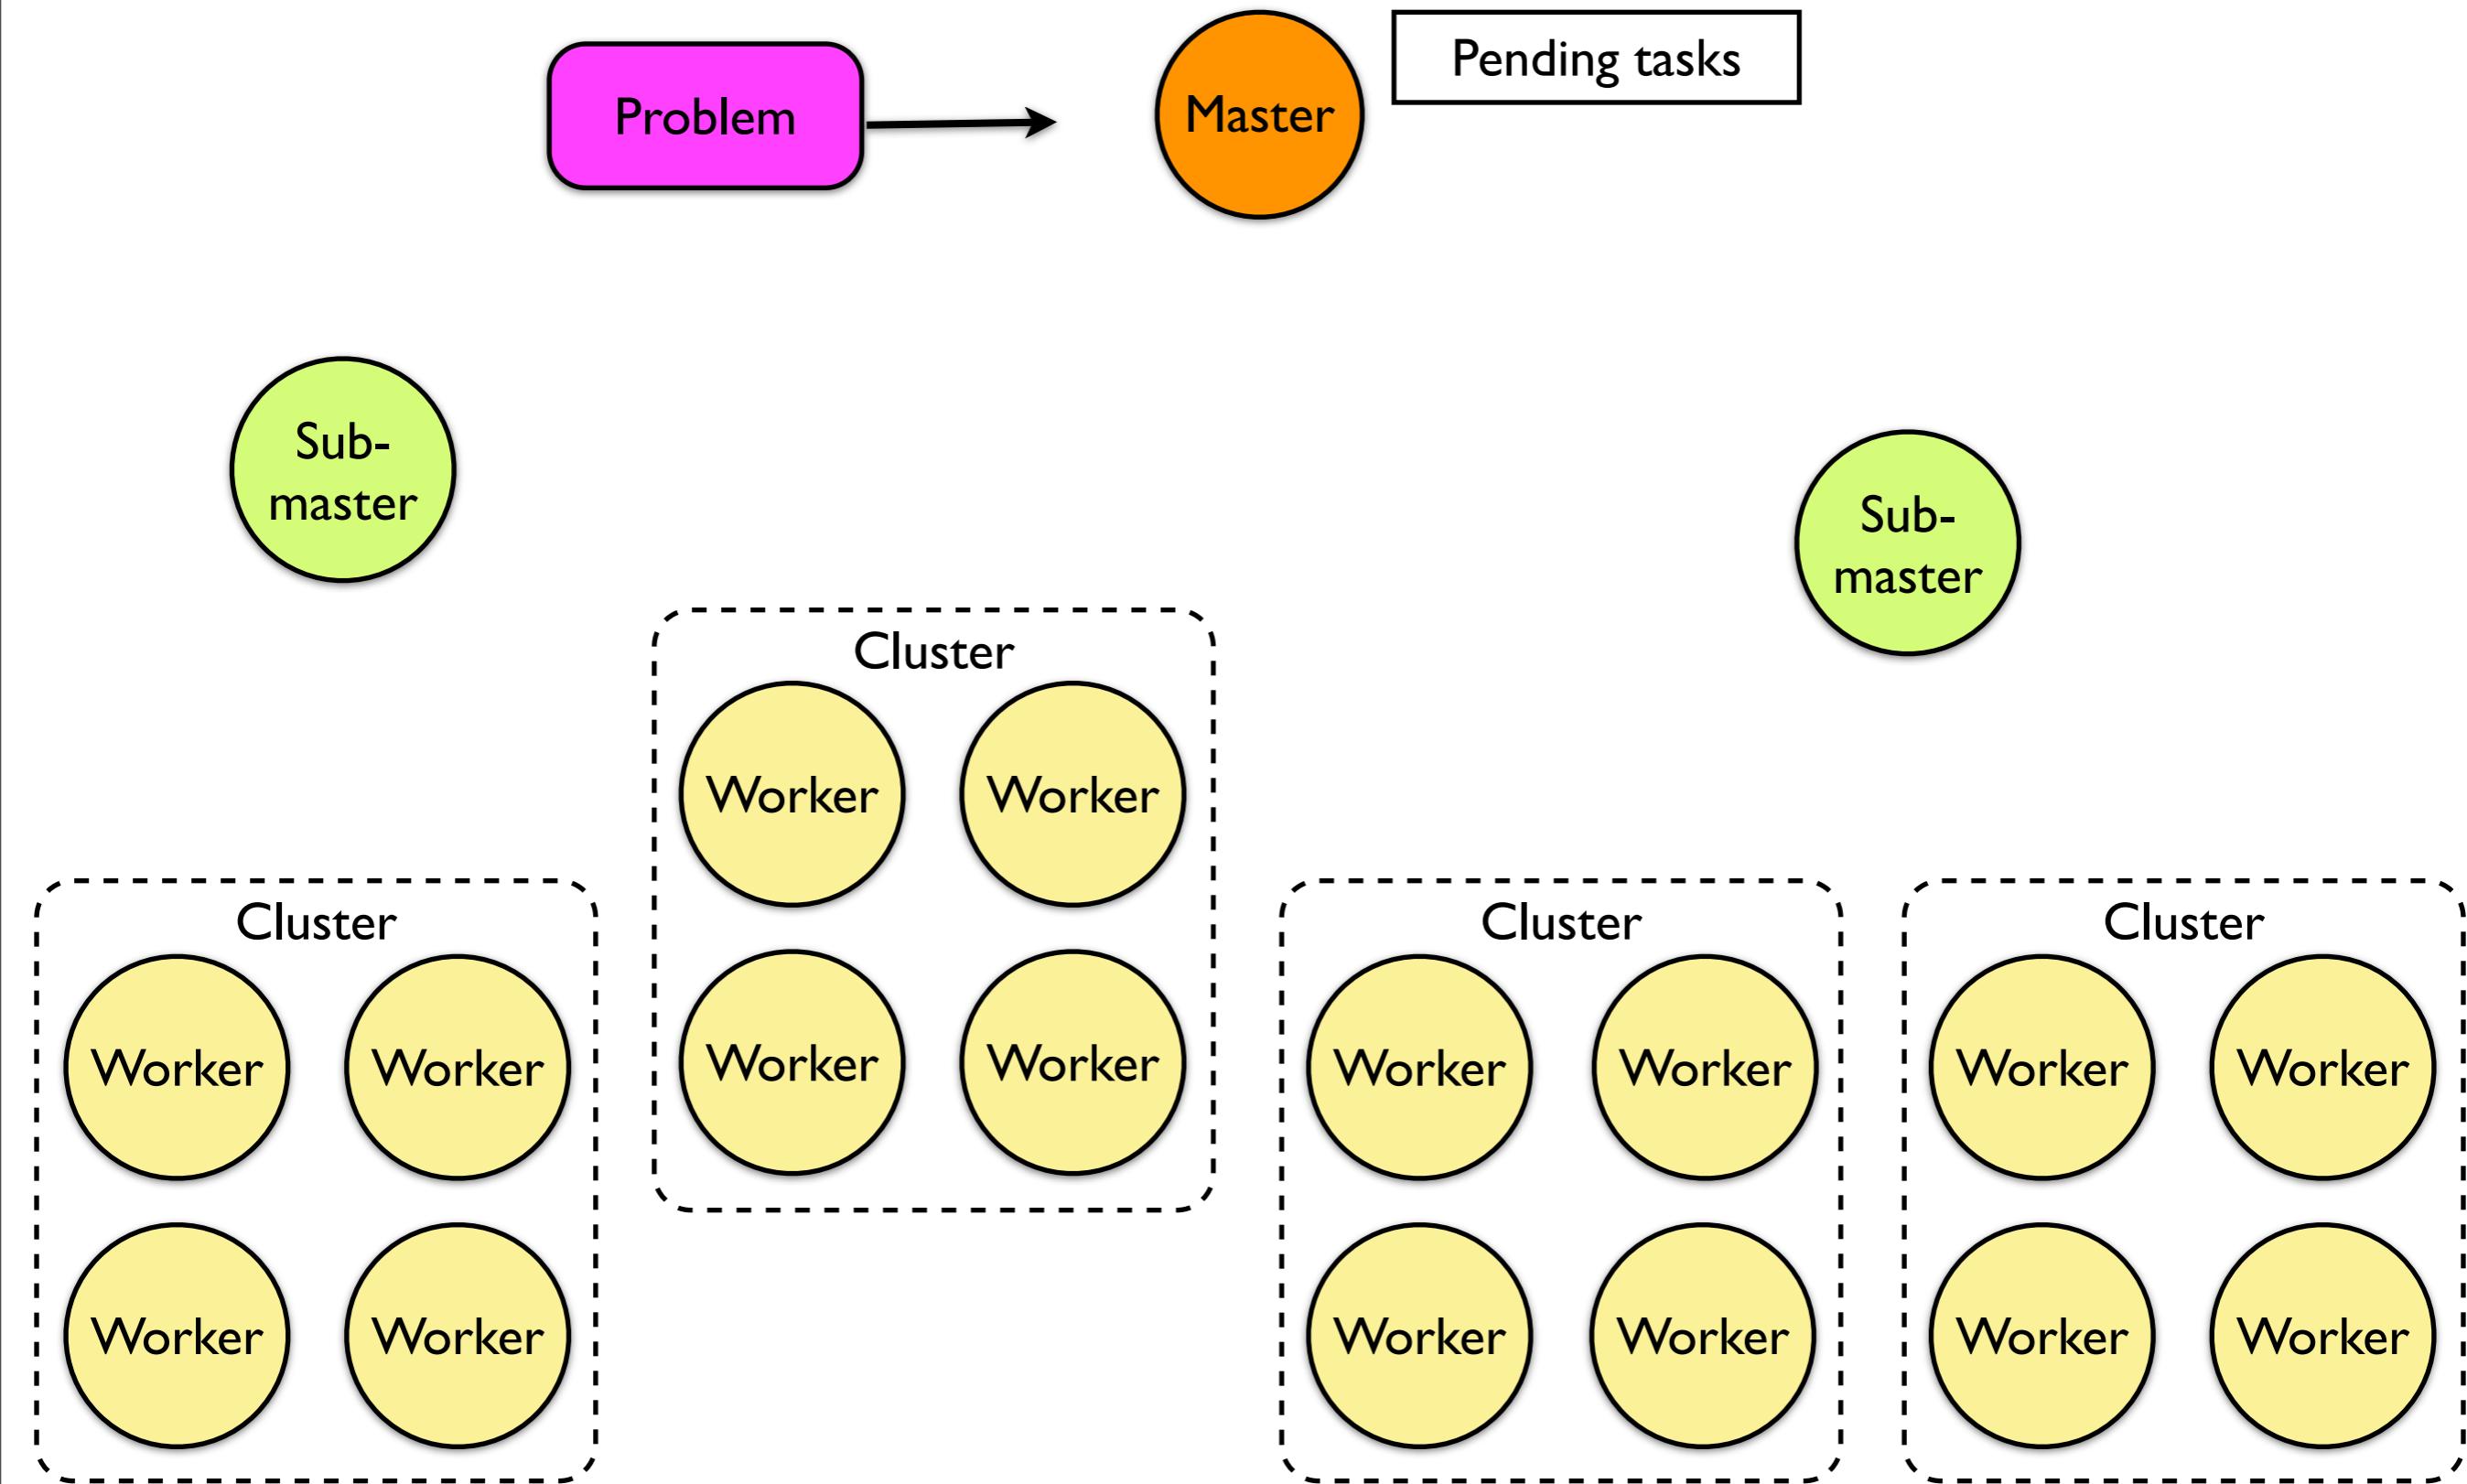
BnB Framework - Architecture



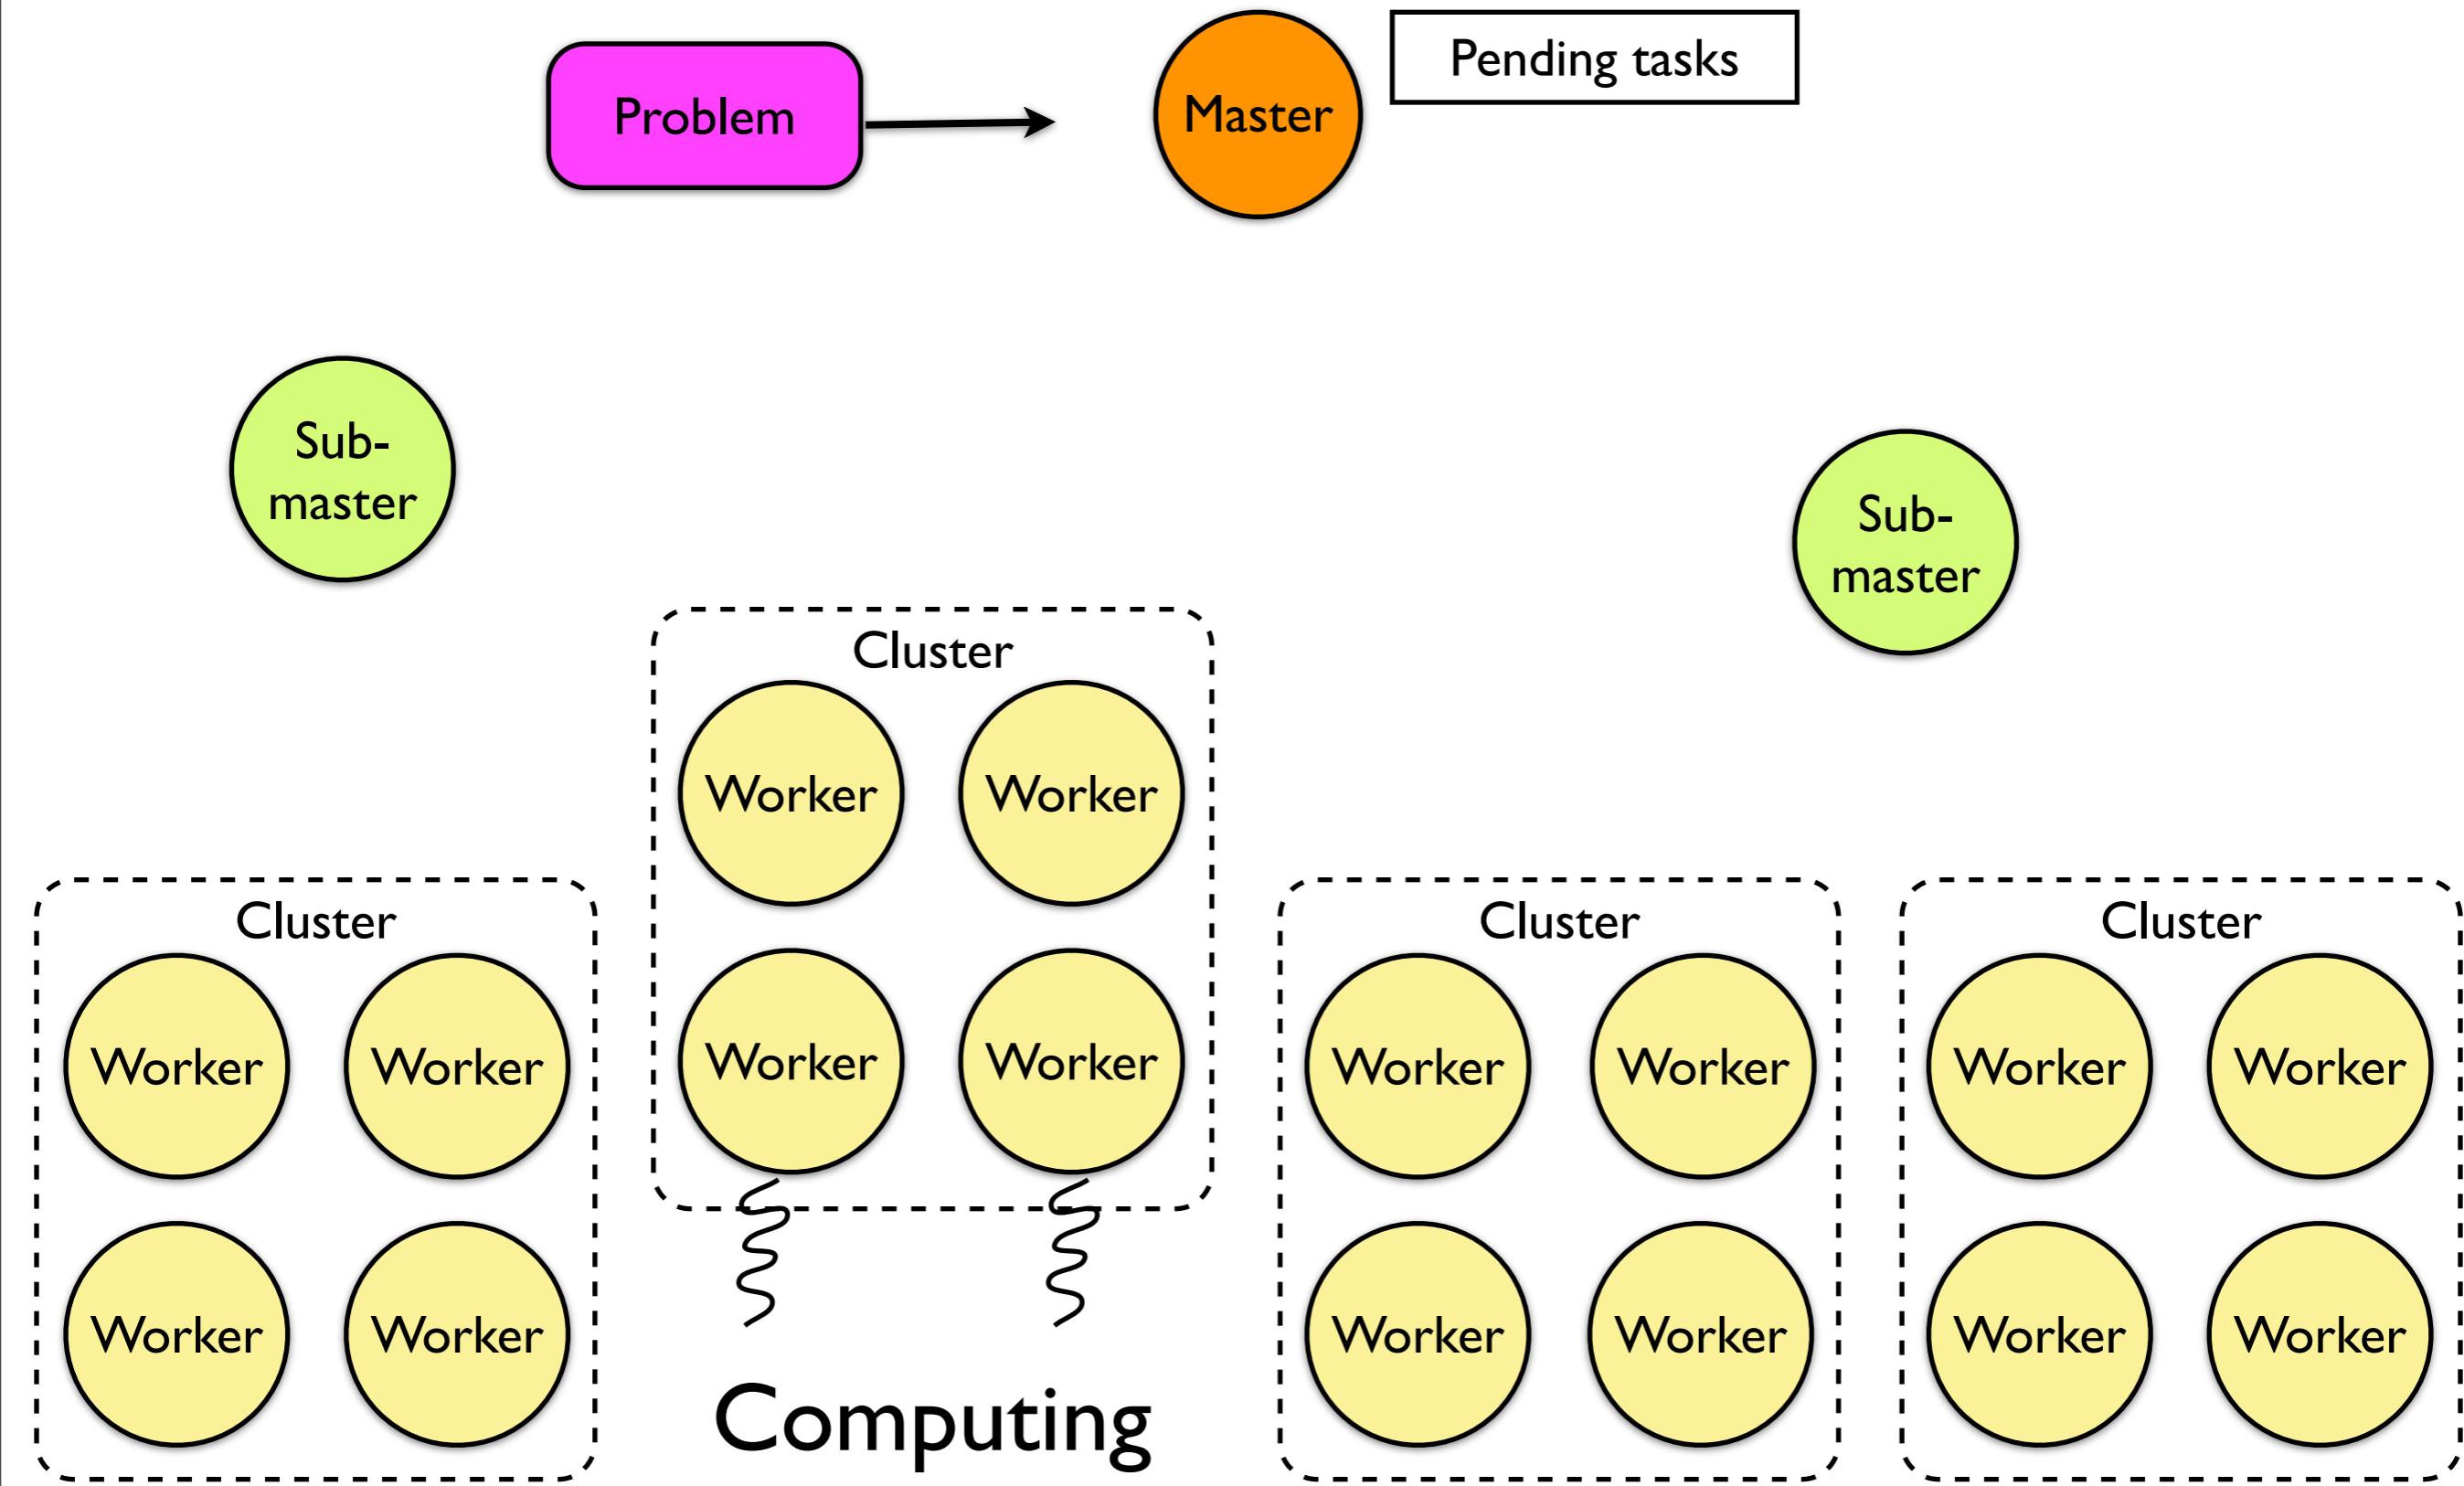
BnB Framework - Architecture



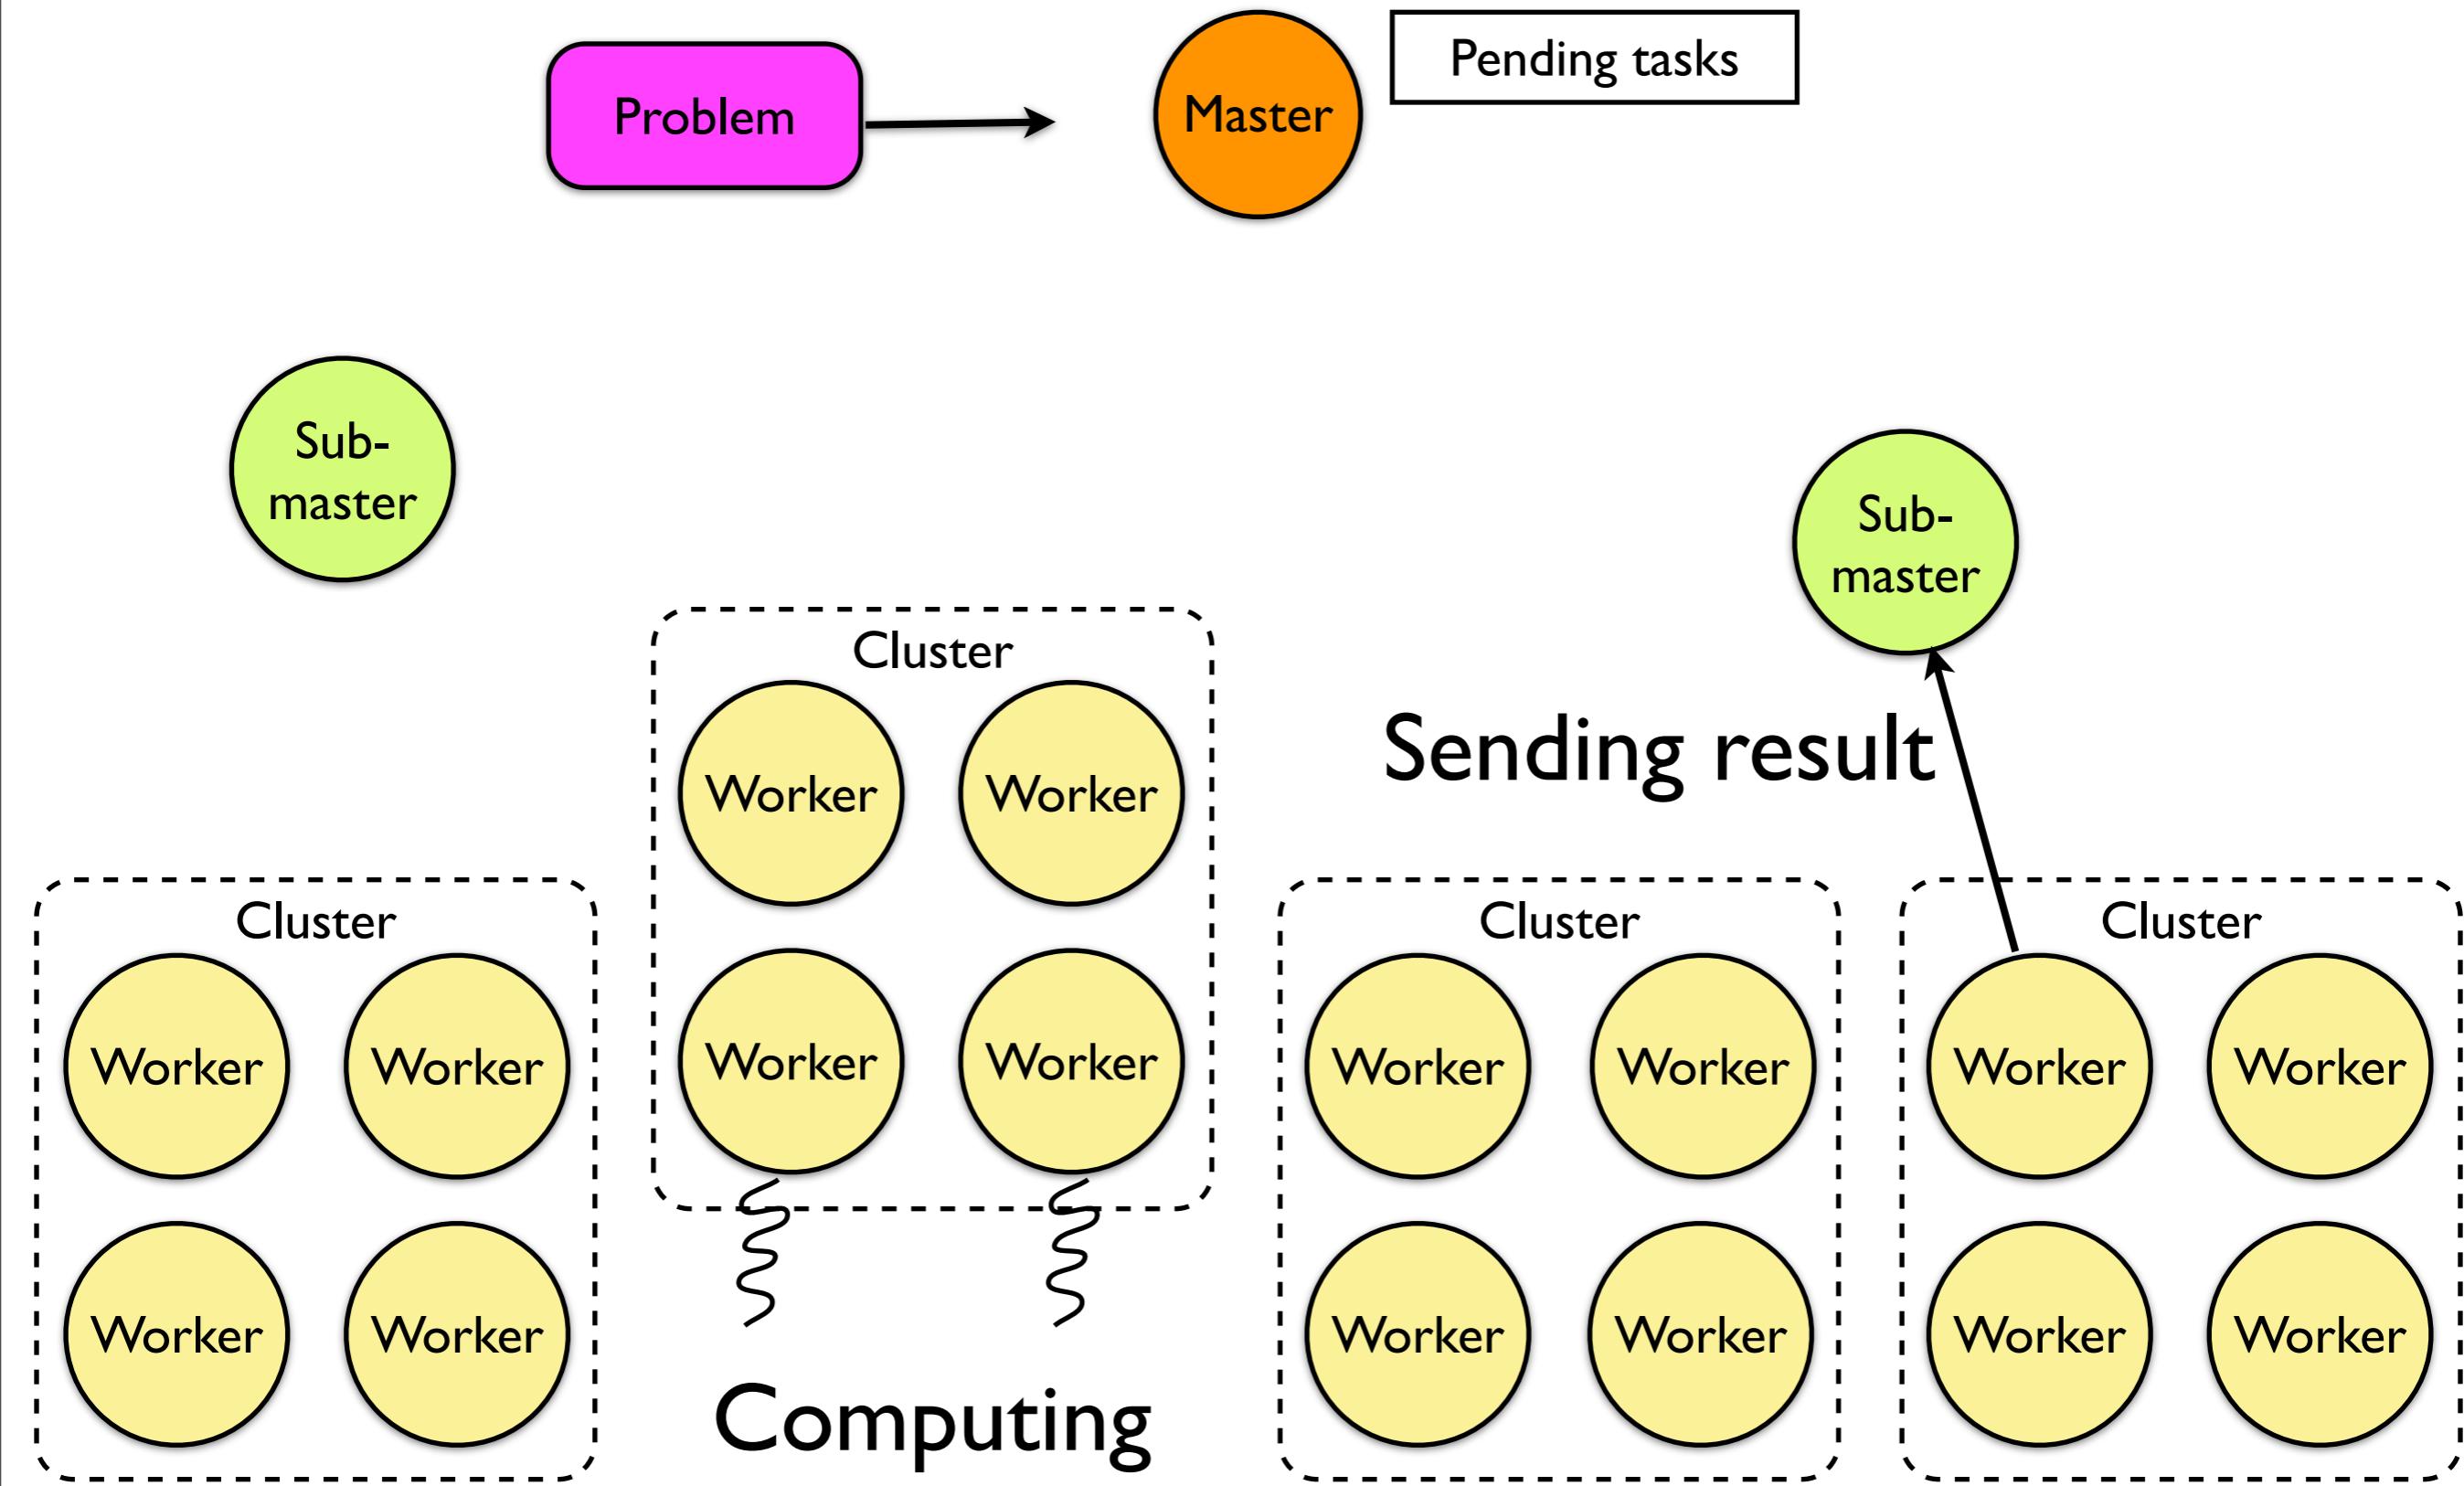
BnB Framework - Architecture



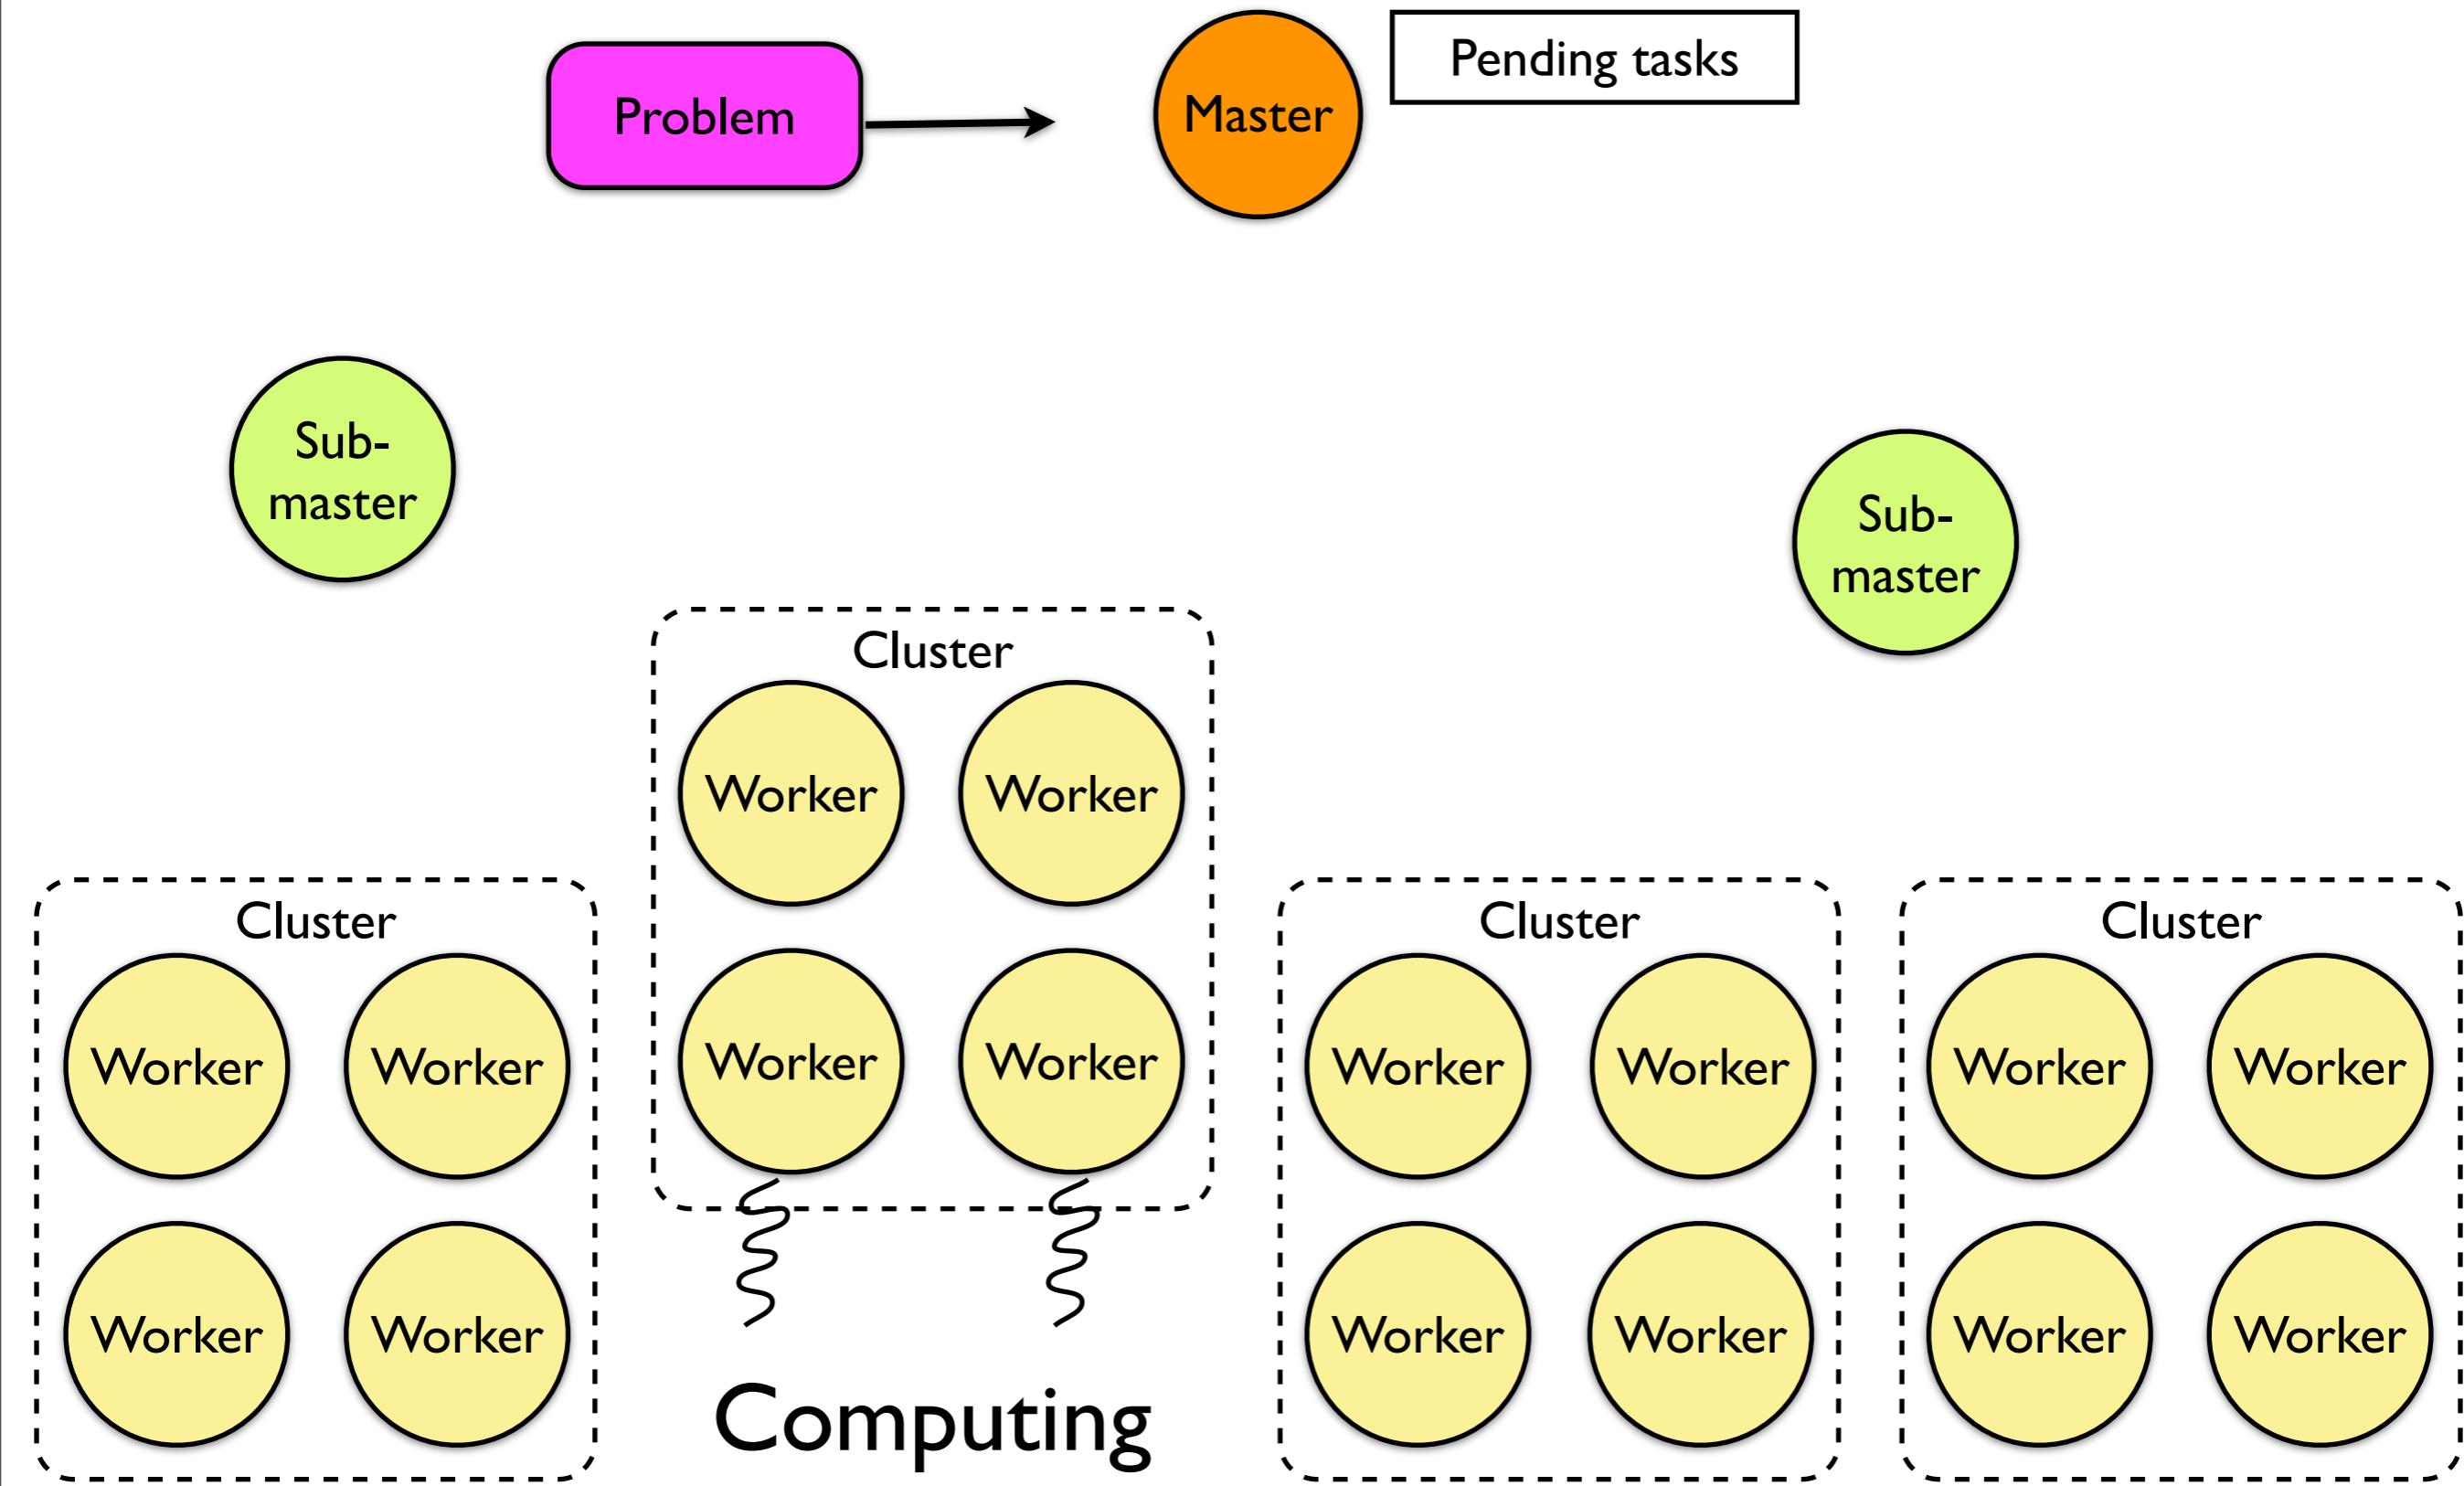
BnB Framework - Architecture



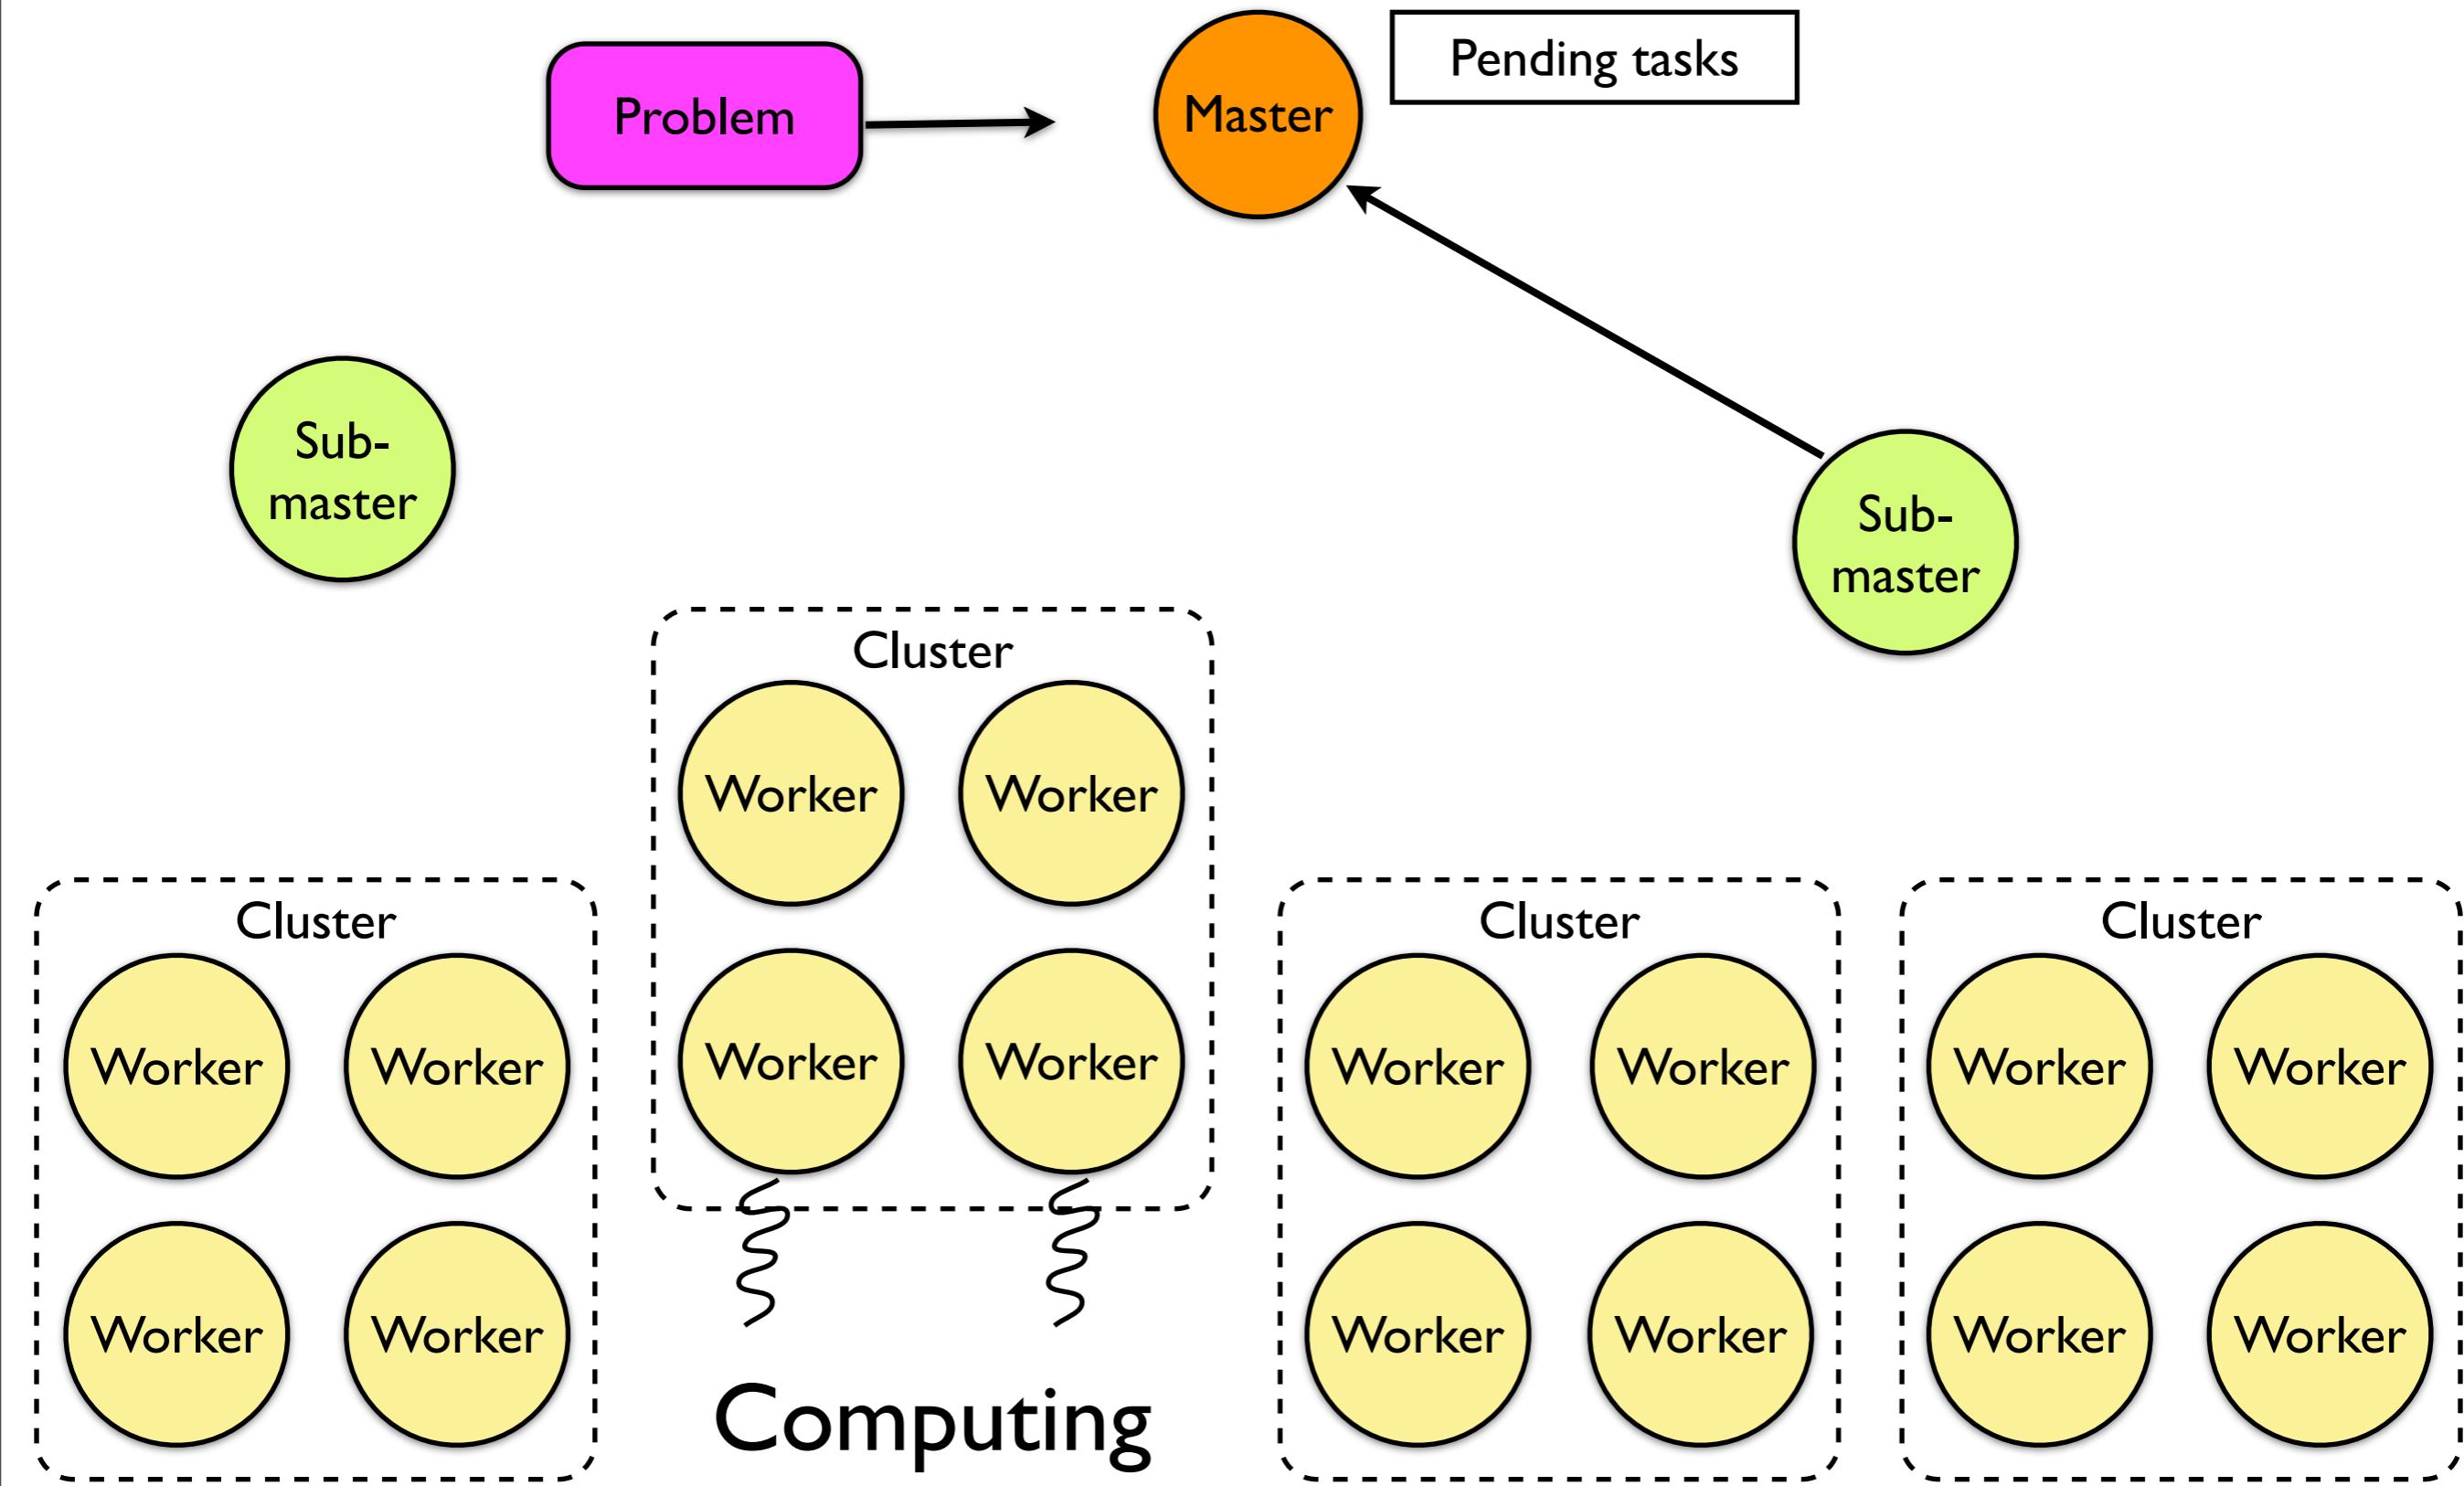
BnB Framework - Architecture



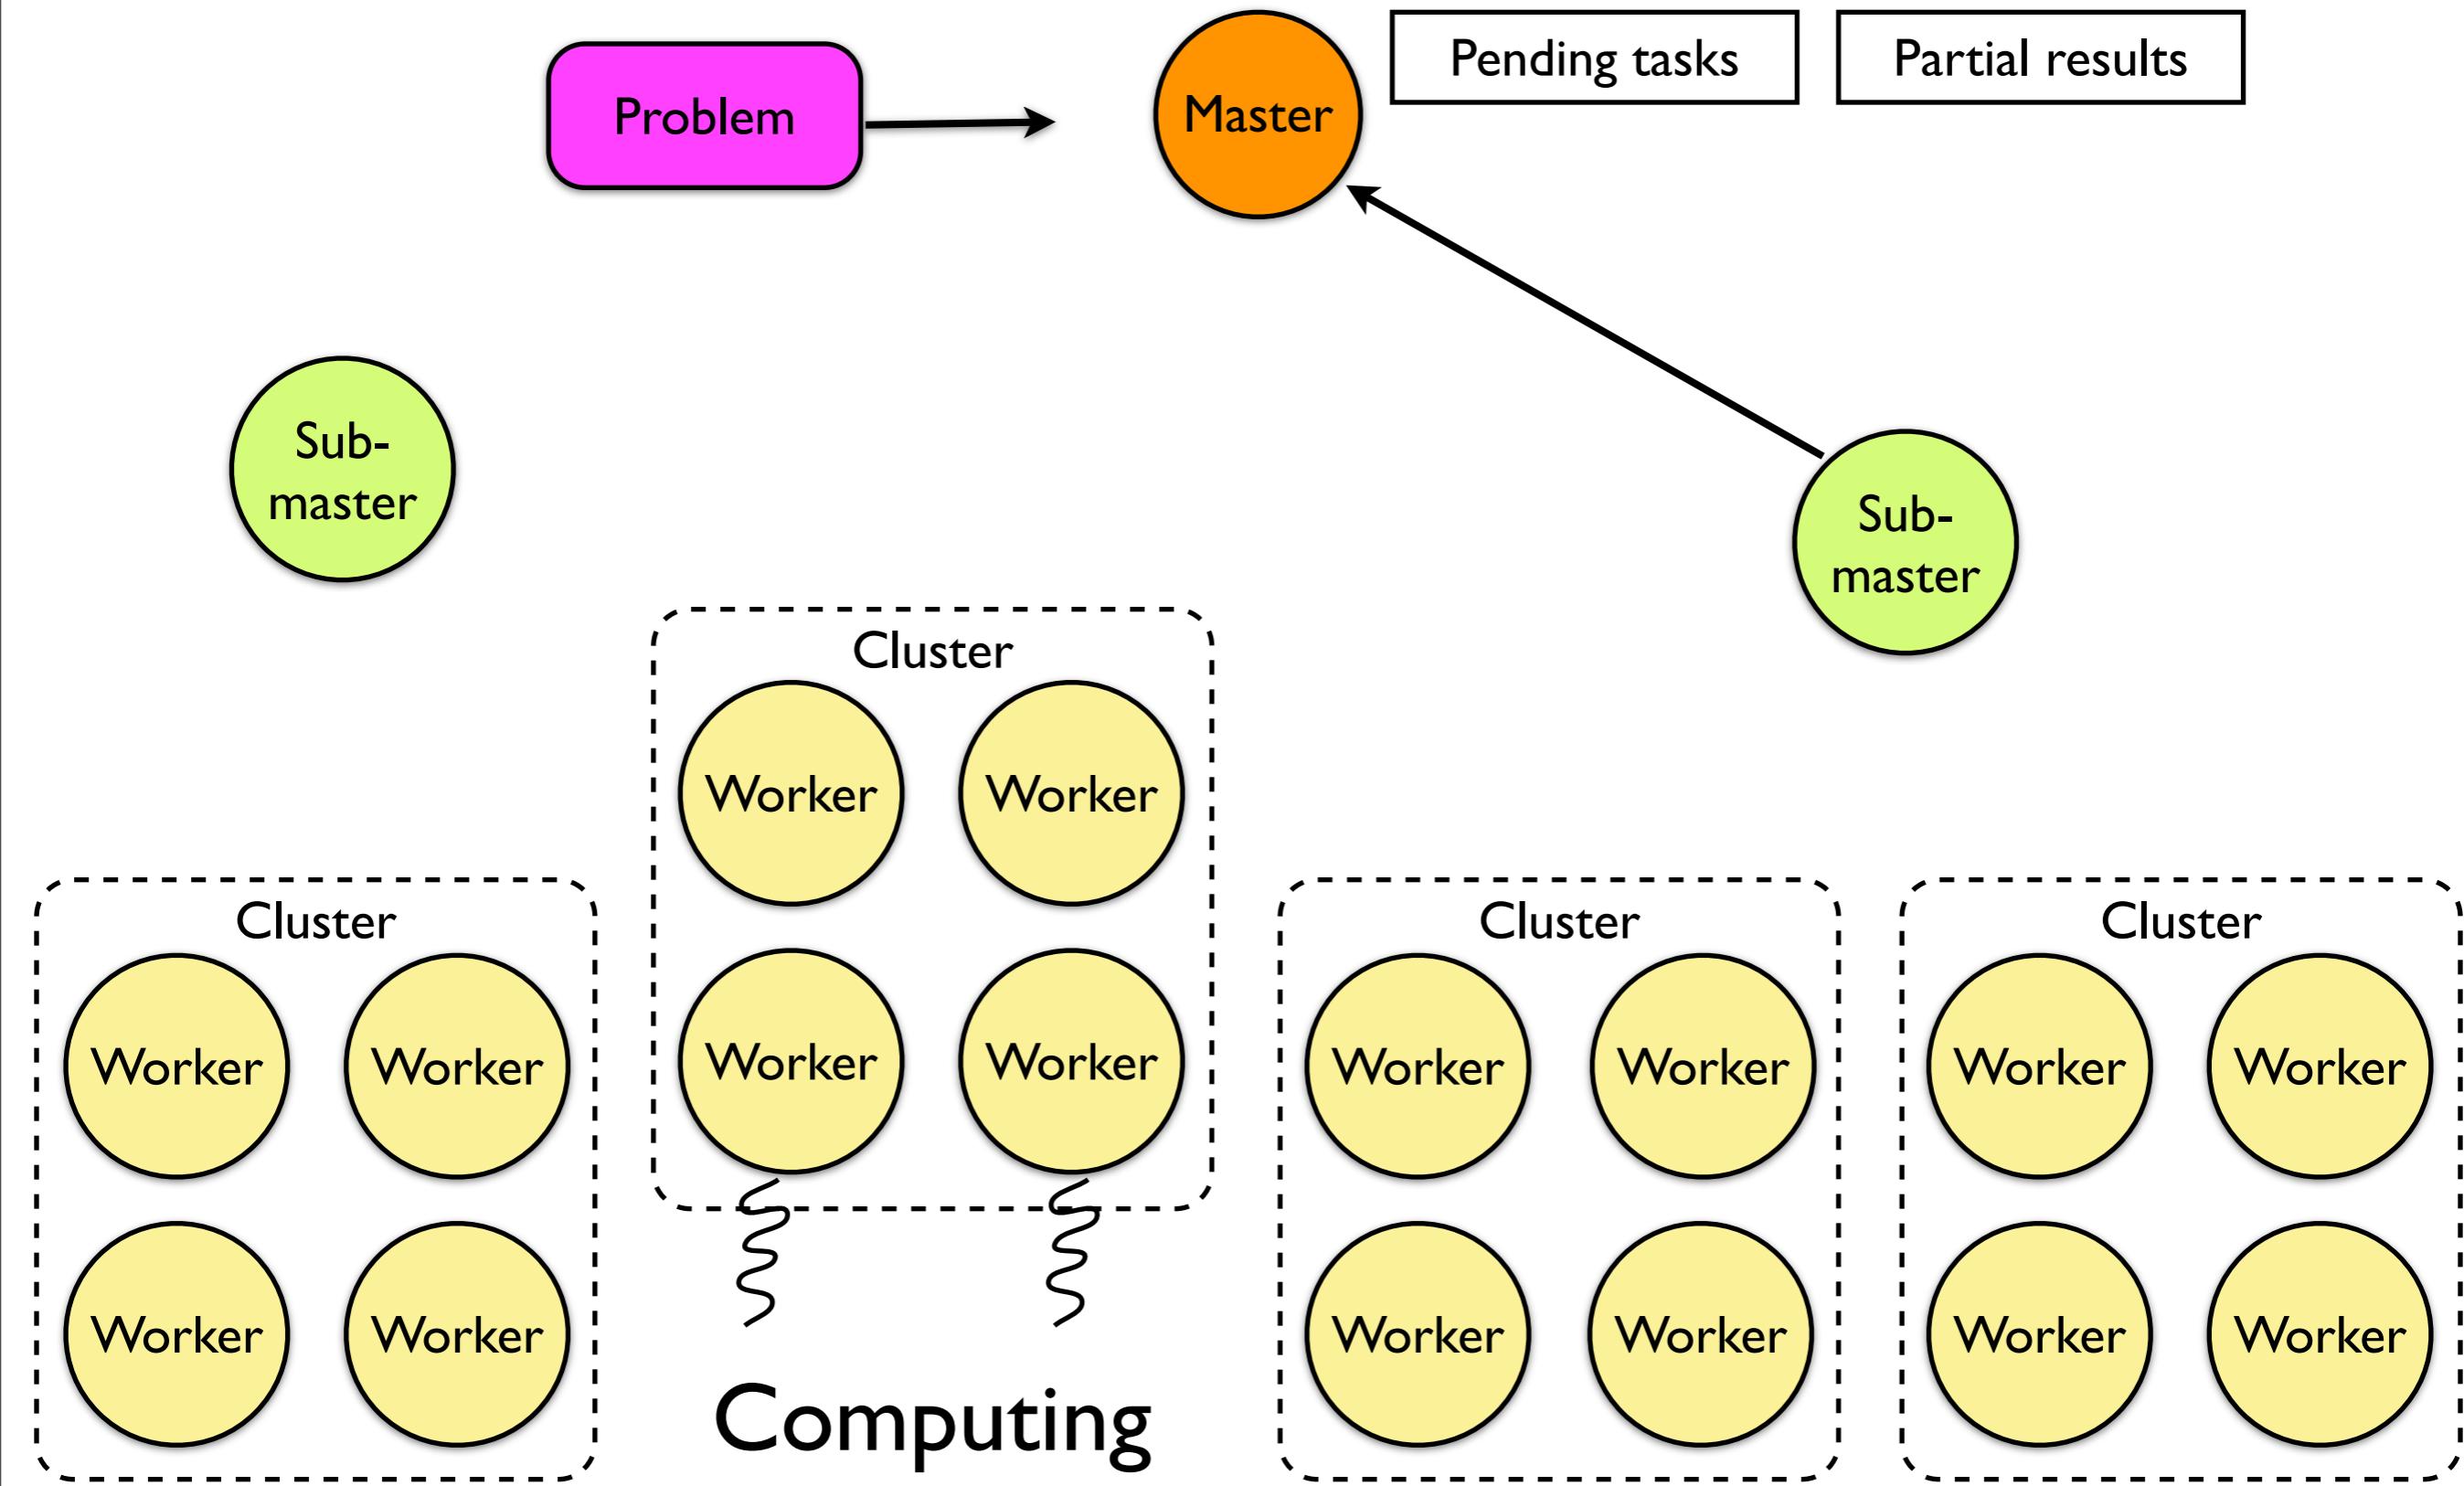
BnB Framework - Architecture



BnB Framework - Architecture



BnB Framework - Architecture



BnB Framework - Solutions

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)
- Solution I: Master keeps the bound

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)
- Solution I: Master keeps the bound
 - previous work shows that not scale [Aida 2003]

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)
- Solution 1: Master keeps the bound
 - previous work shows that not scale [Aida 2003]
- Solution 2: Message framework (Enterprise Service Bus)

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)
- Solution 1: Master keeps the bound
 - previous work shows that not scale [Aida 2003]
- Solution 2: Message framework (Enterprise Service Bus)
 - Grid middleware dependent / Good for SOA

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)
- Solution 1: Master keeps the bound
 - previous work shows that not scale [Aida 2003]
- Solution 2: Message framework (Enterprise Service Bus)
 - Grid middleware dependent / Good for SOA
- Solution 3: Broadcasting

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)
- Solution 1: Master keeps the bound
 - previous work shows that not scale [Aida 2003]
- Solution 2: Message framework (Enterprise Service Bus)
 - Grid middleware dependent / Good for SOA
- Solution 3: Broadcasting
 - 1 to n communication cannot scale

BnB Framework - Solutions

- Context ProActive Java Grid middleware:
 - latency → asynchronous communication
 - underlaying Grid infrastructure → deployment framework (abstraction)
- Implement the tree-based parallel
- Master-worker architecture
- Problem: workers need to share bounds
 - difficult to adapt SPMD for Grids (heterogeneity, distribution, etc.)
- Solution 1: Master keeps the bound
 - previous work shows that not scale [Aida 2003]
- Solution 2: Message framework (Enterprise Service Bus)
 - Grid middleware dependent / Good for SOA
- Solution 3: Broadcasting
 - 1 to n communication cannot scale
 - ✓ hierarchical broadcasting scale [Baduel 05]

Organizing Communications for Broadcasting

Organizing Communications for Broadcasting

Idea: Grids are composed of clusters → organizing Workers in groups

Organizing Communications for Broadcasting

Idea: Grids are composed of clusters → organizing Workers in groups

- clusters are high-performance communication environments

Organizing Communications for Broadcasting

Idea: Grids are composed of clusters → organizing Workers in groups

- clusters are high-performance communication environments

Organizing Communications for Broadcasting

Idea: Grids are composed of clusters \rightarrow organizing Workers in groups

- clusters are high-performance communication environments
- Solution:
 - add a new entity for organizing communications: Leader
 - Leader is a Worker chose by the Master for each group

Organizing Communications for Broadcasting

Idea: Grids are composed of clusters \rightarrow organizing Workers in groups

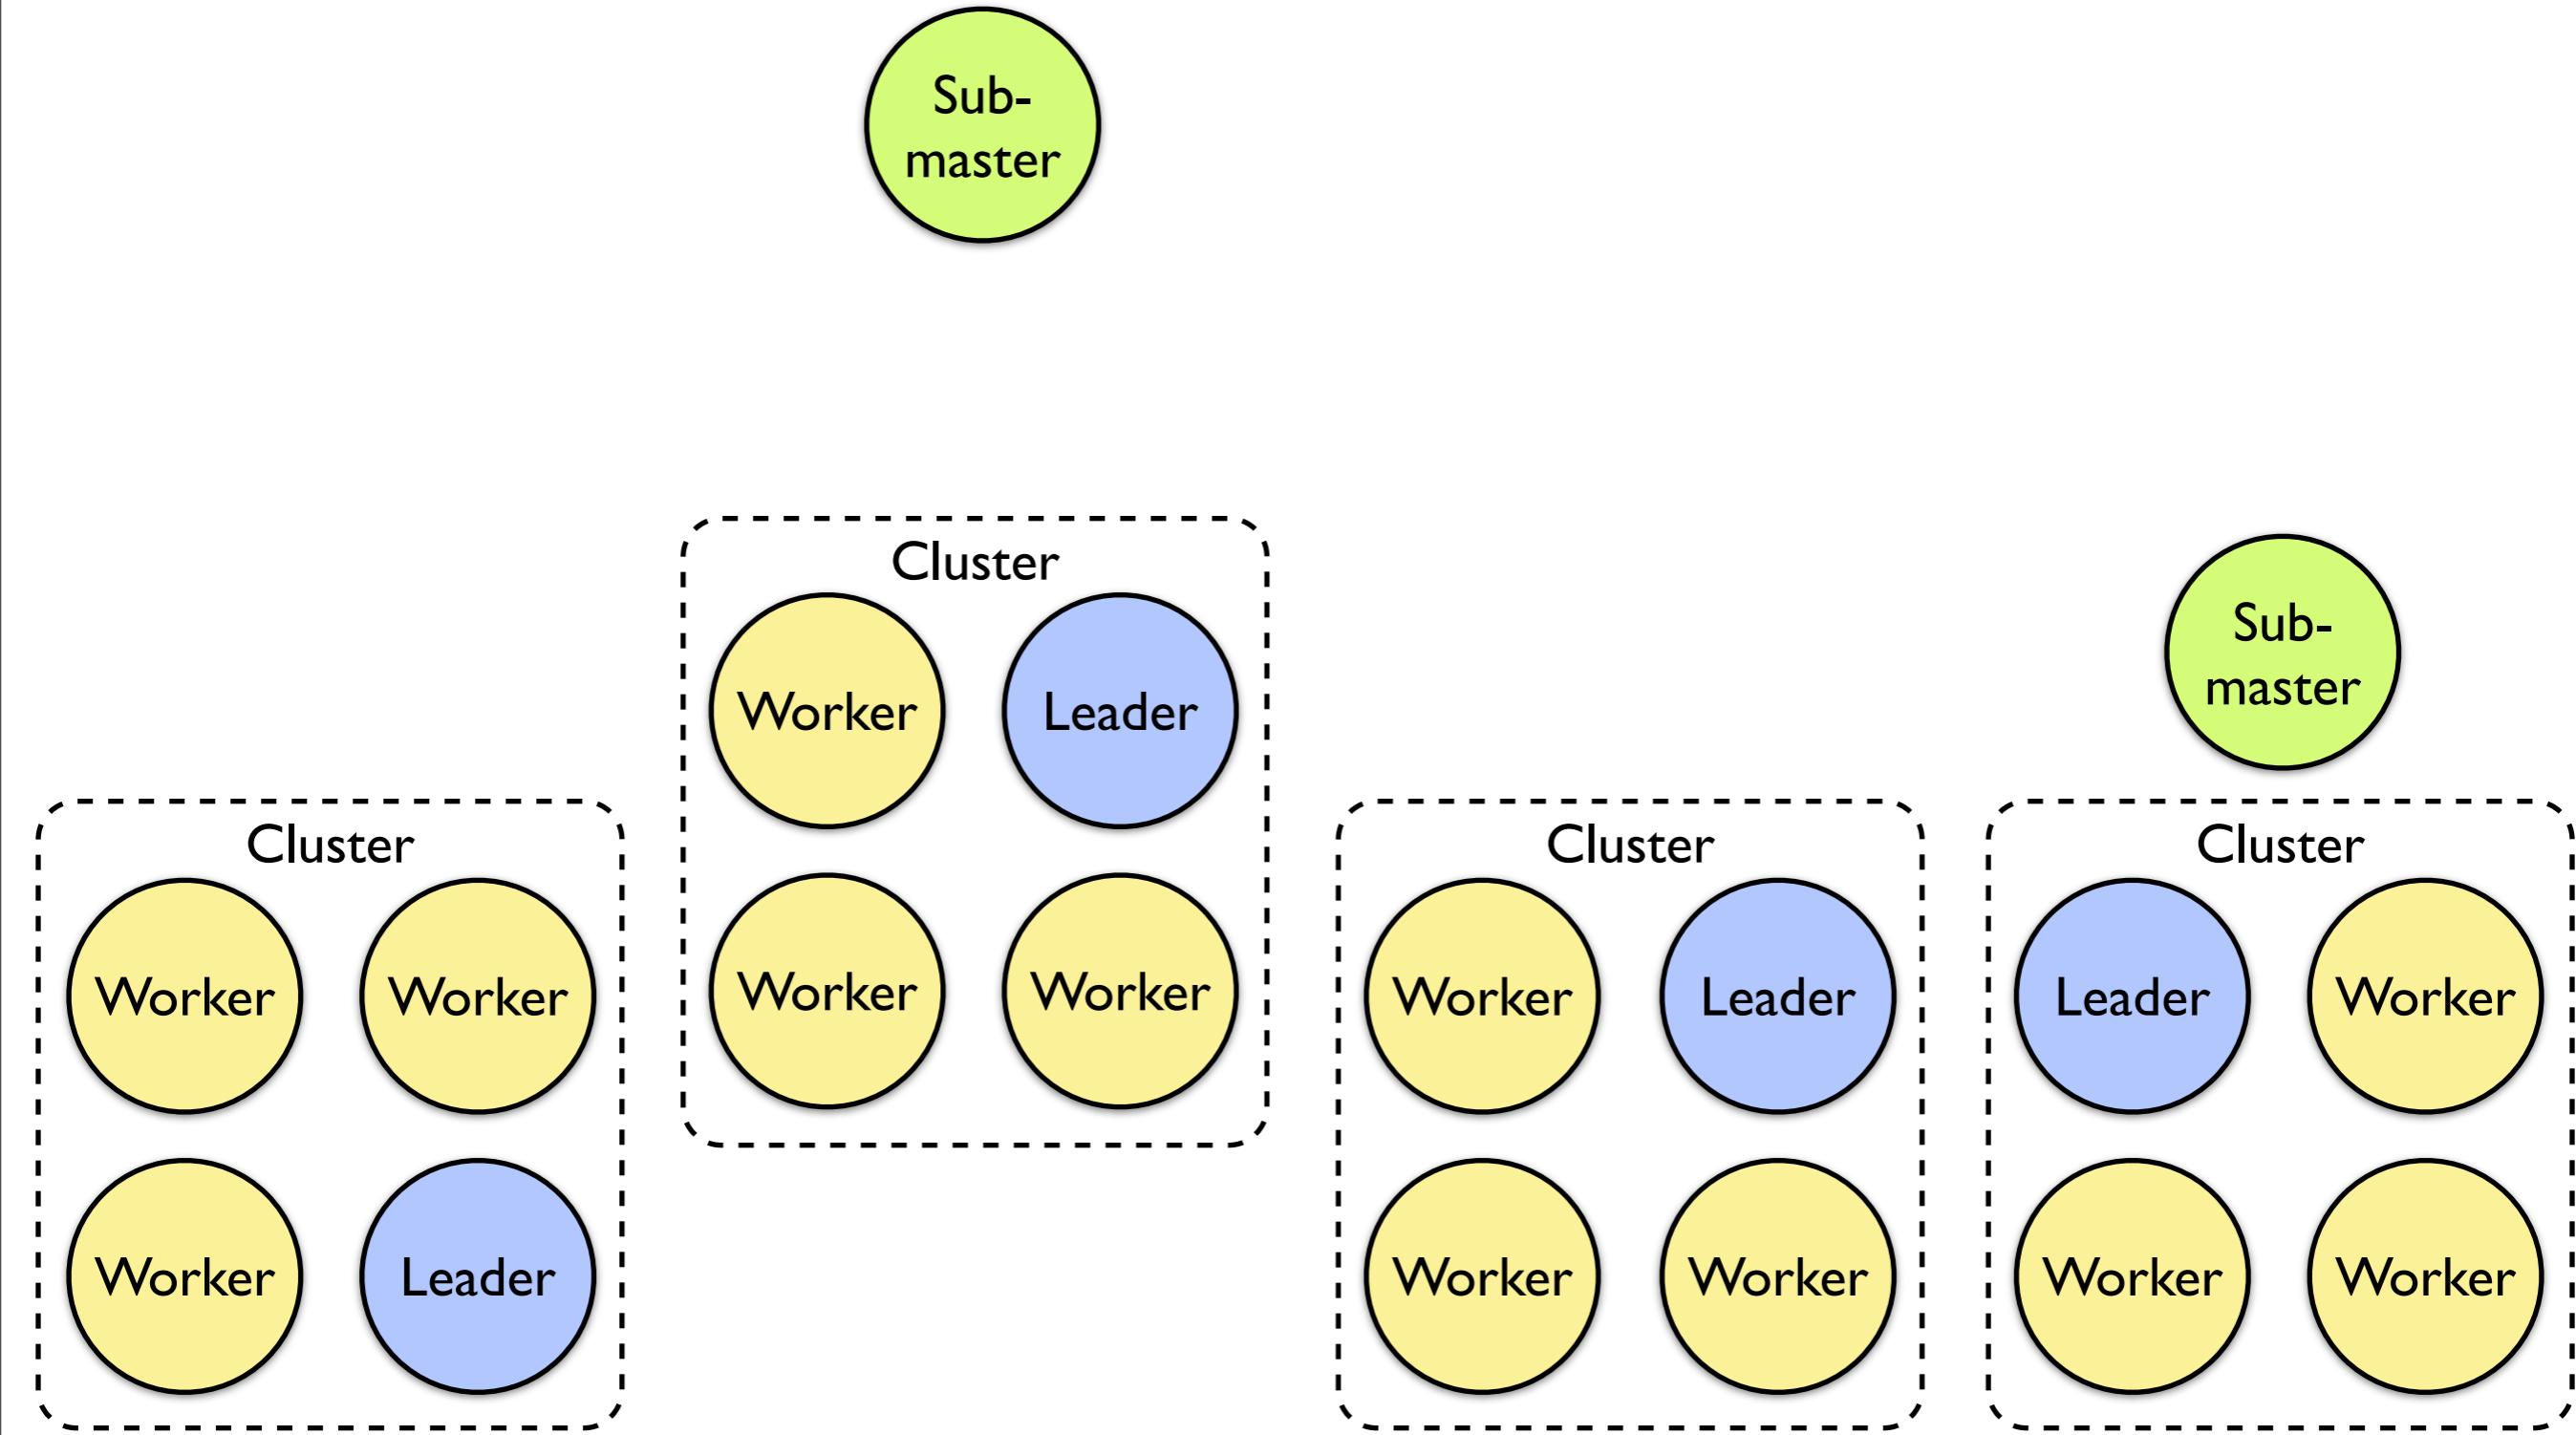
- clusters are high-performance communication environments
- Solution:
 - add a new entity for organizing communications: Leader
 - Leader is a Worker chose by the Master for each group

Organizing Communications for Broadcasting

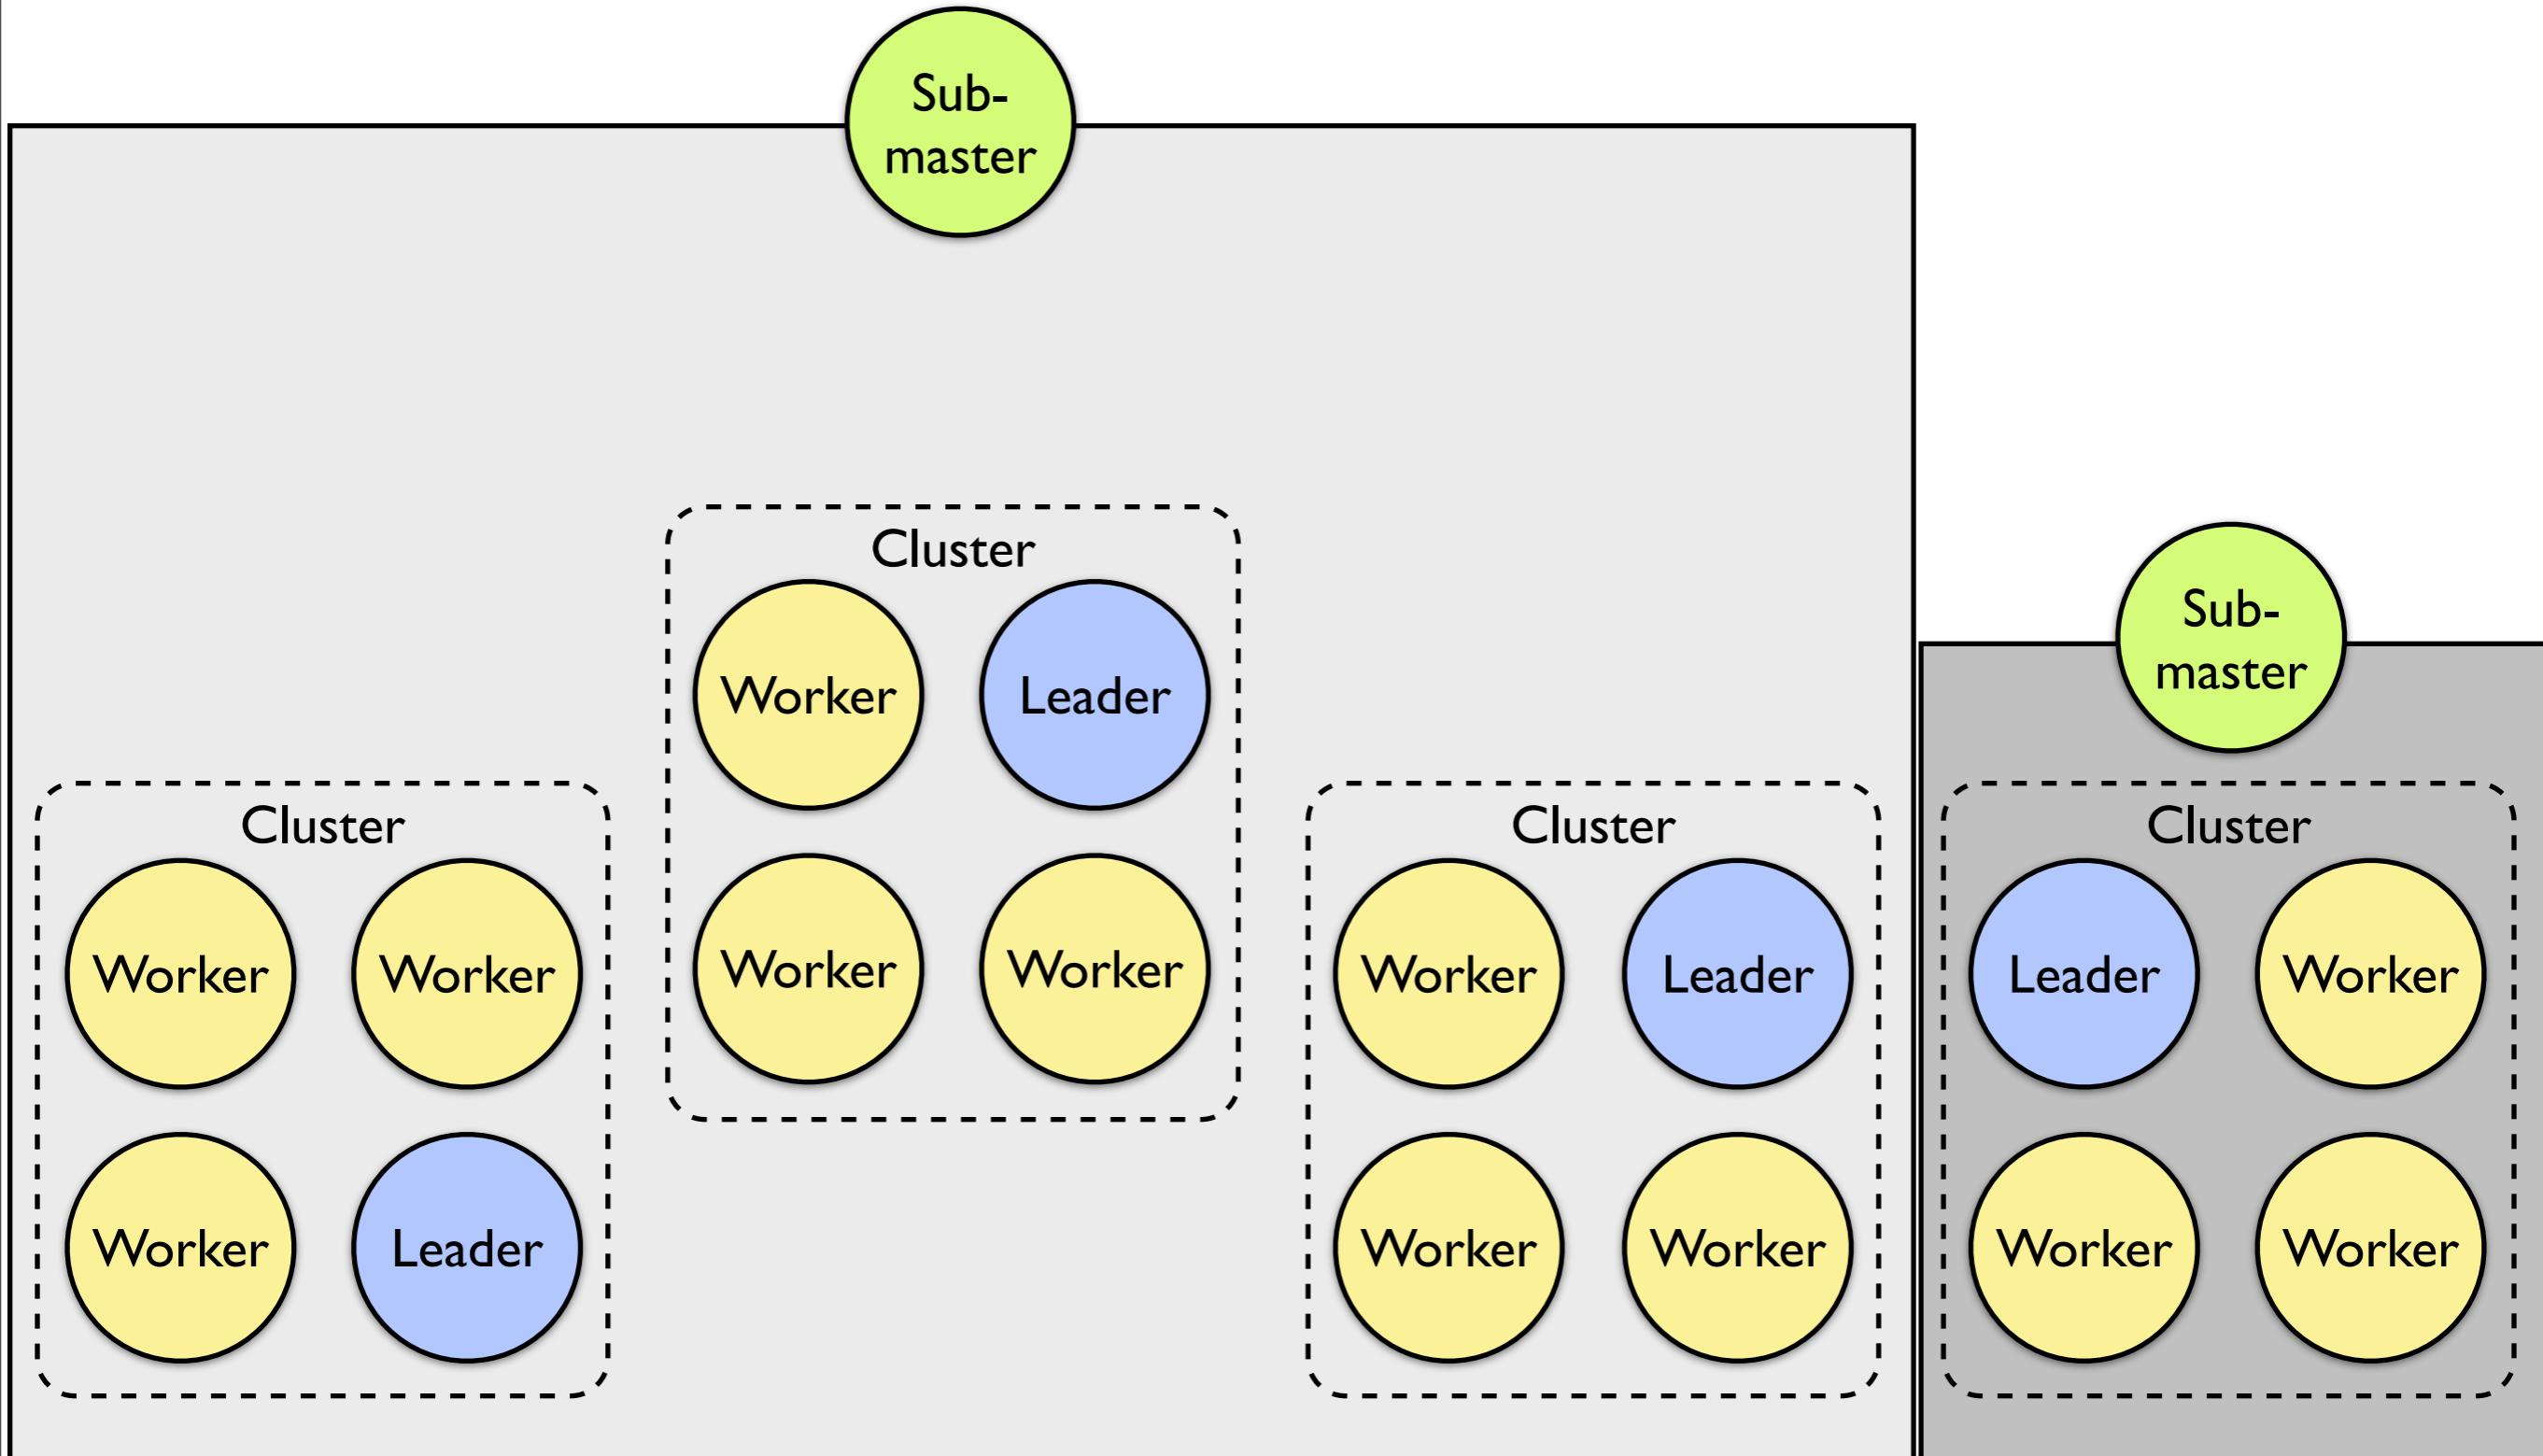
Idea: Grids are composed of clusters \rightarrow organizing Workers in groups

- clusters are high-performance communication environments
- Solution:
 - add a new entity for organizing communications: Leader
 - Leader is a Worker chose by the Master for each group
- Process to update Bounds:
 1. the Worker broadcasts the new Bound inside its group
 2. the group Leader broadcasts the new Bound to all Leaders
 3. each Leader broadcasts the new value inside their groups

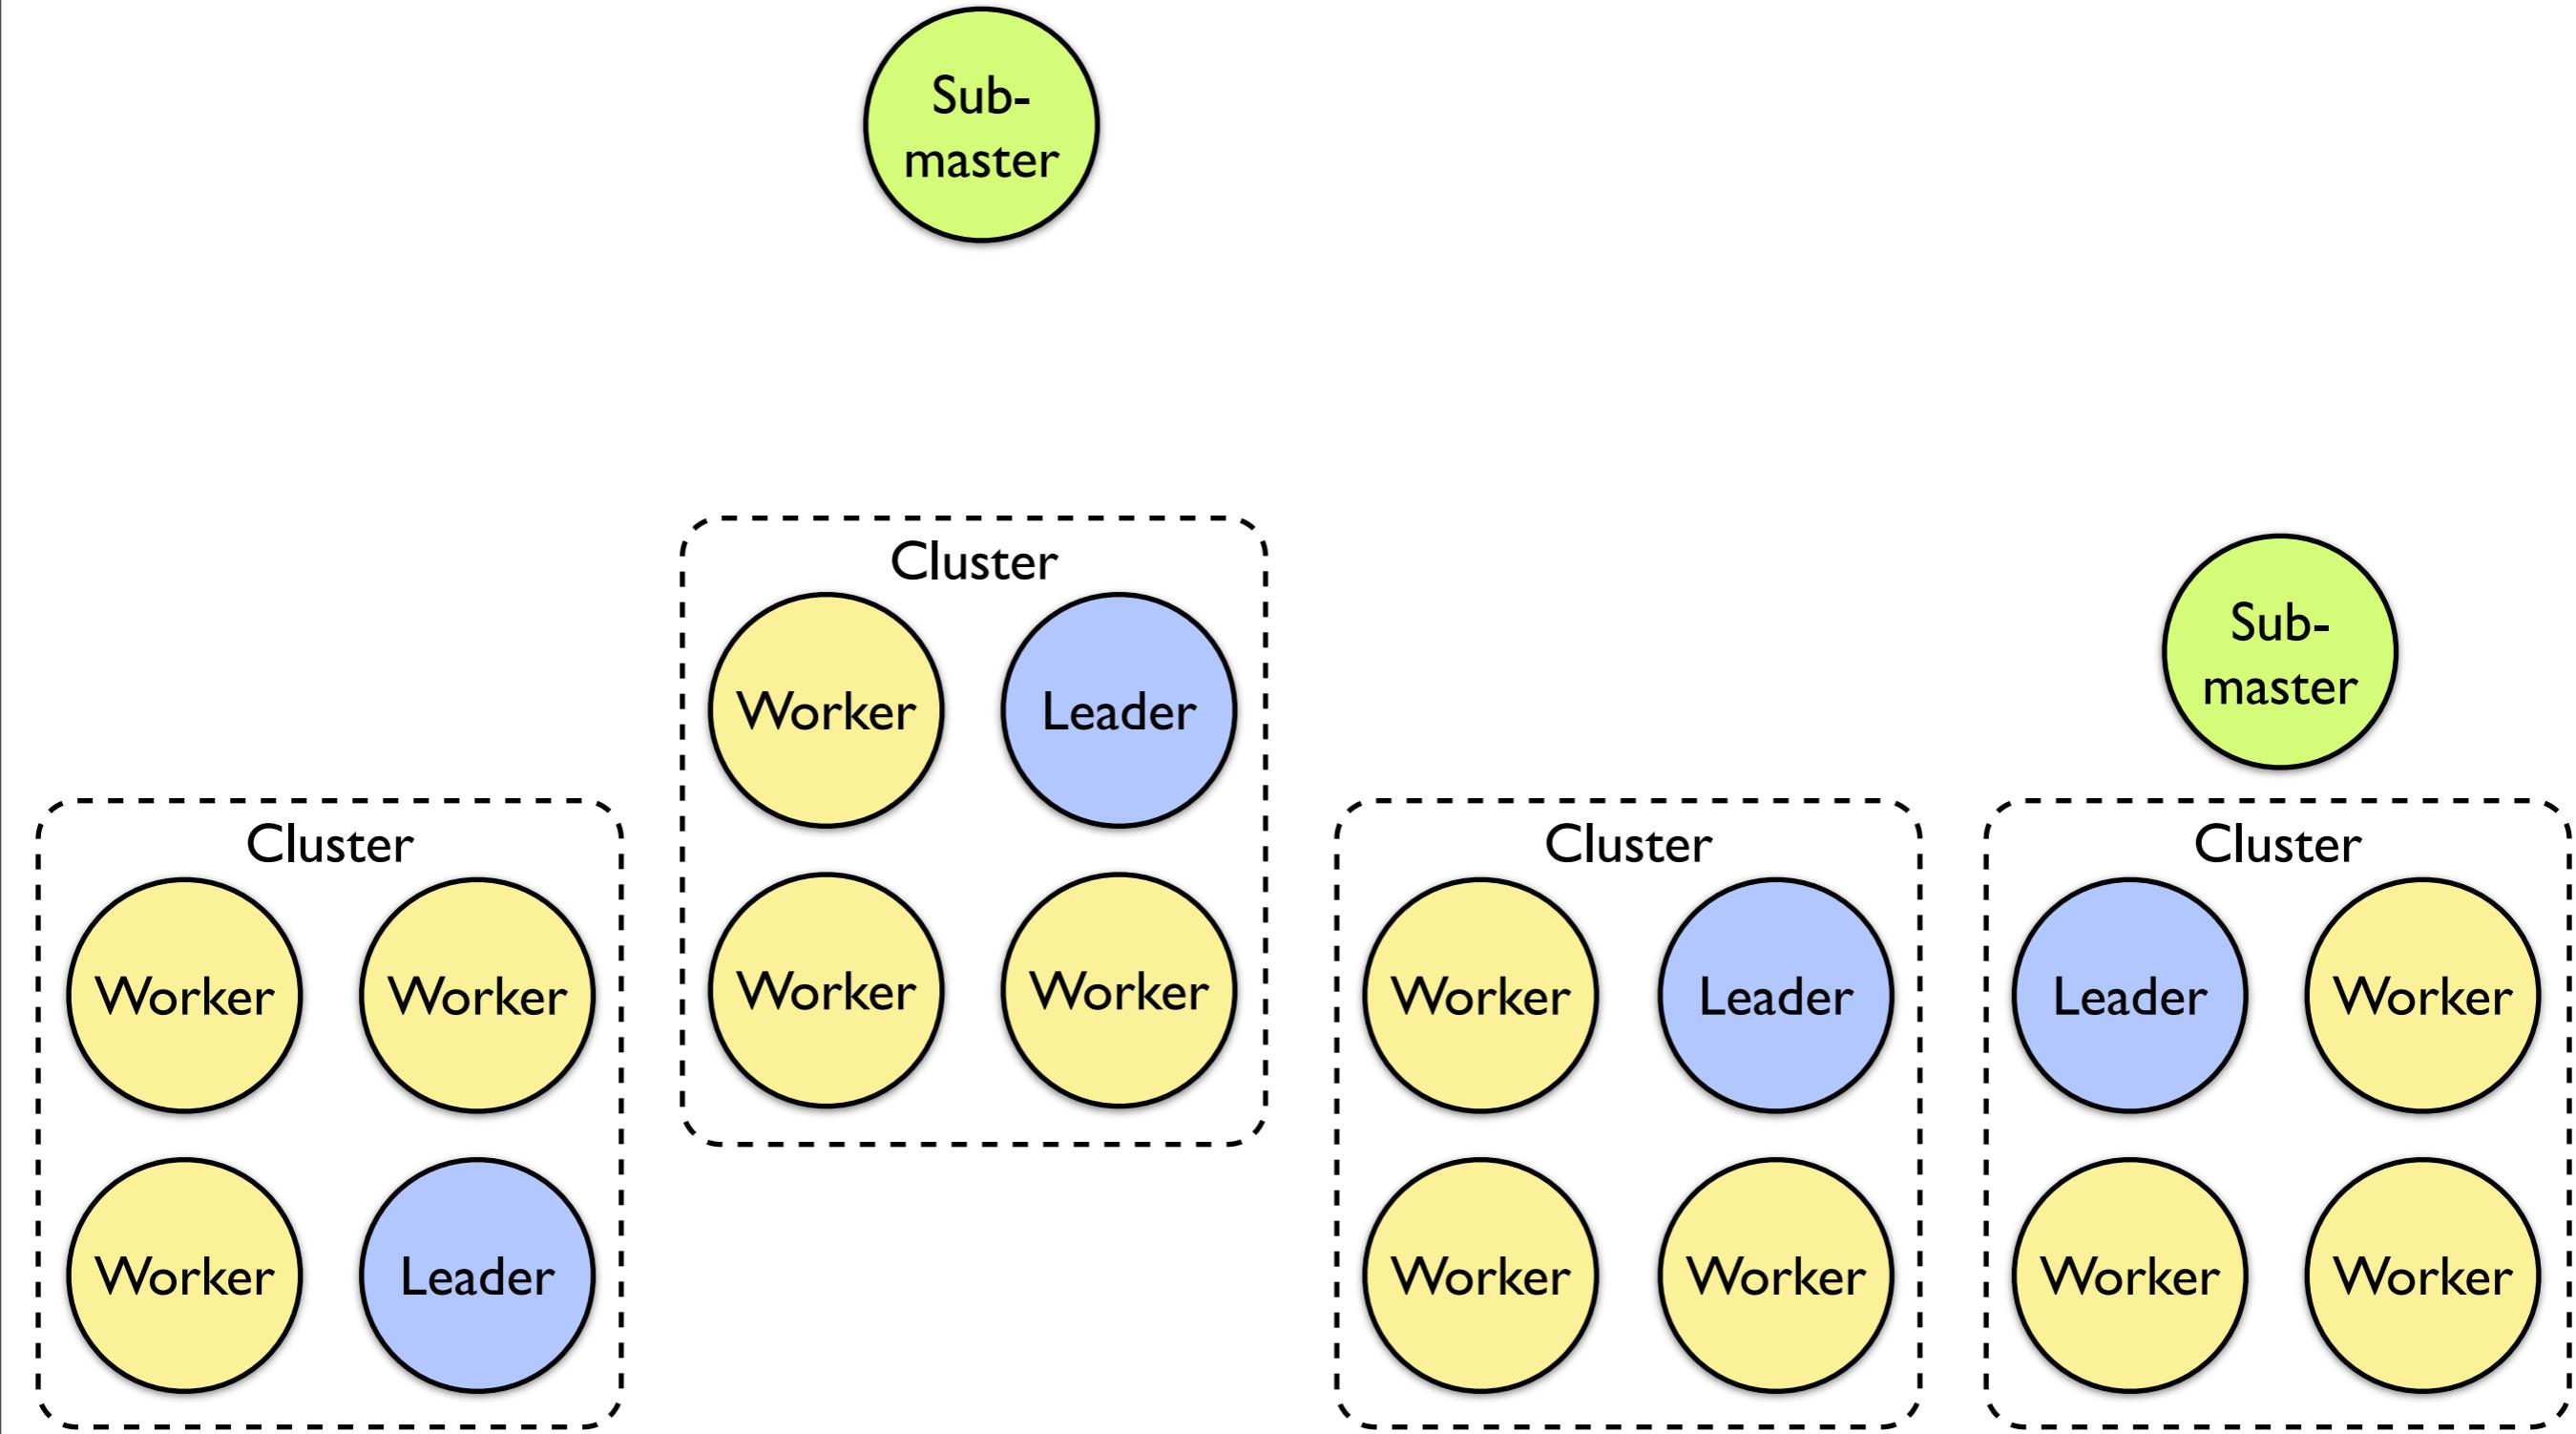
BnB framework - Communications for Sharing Bound



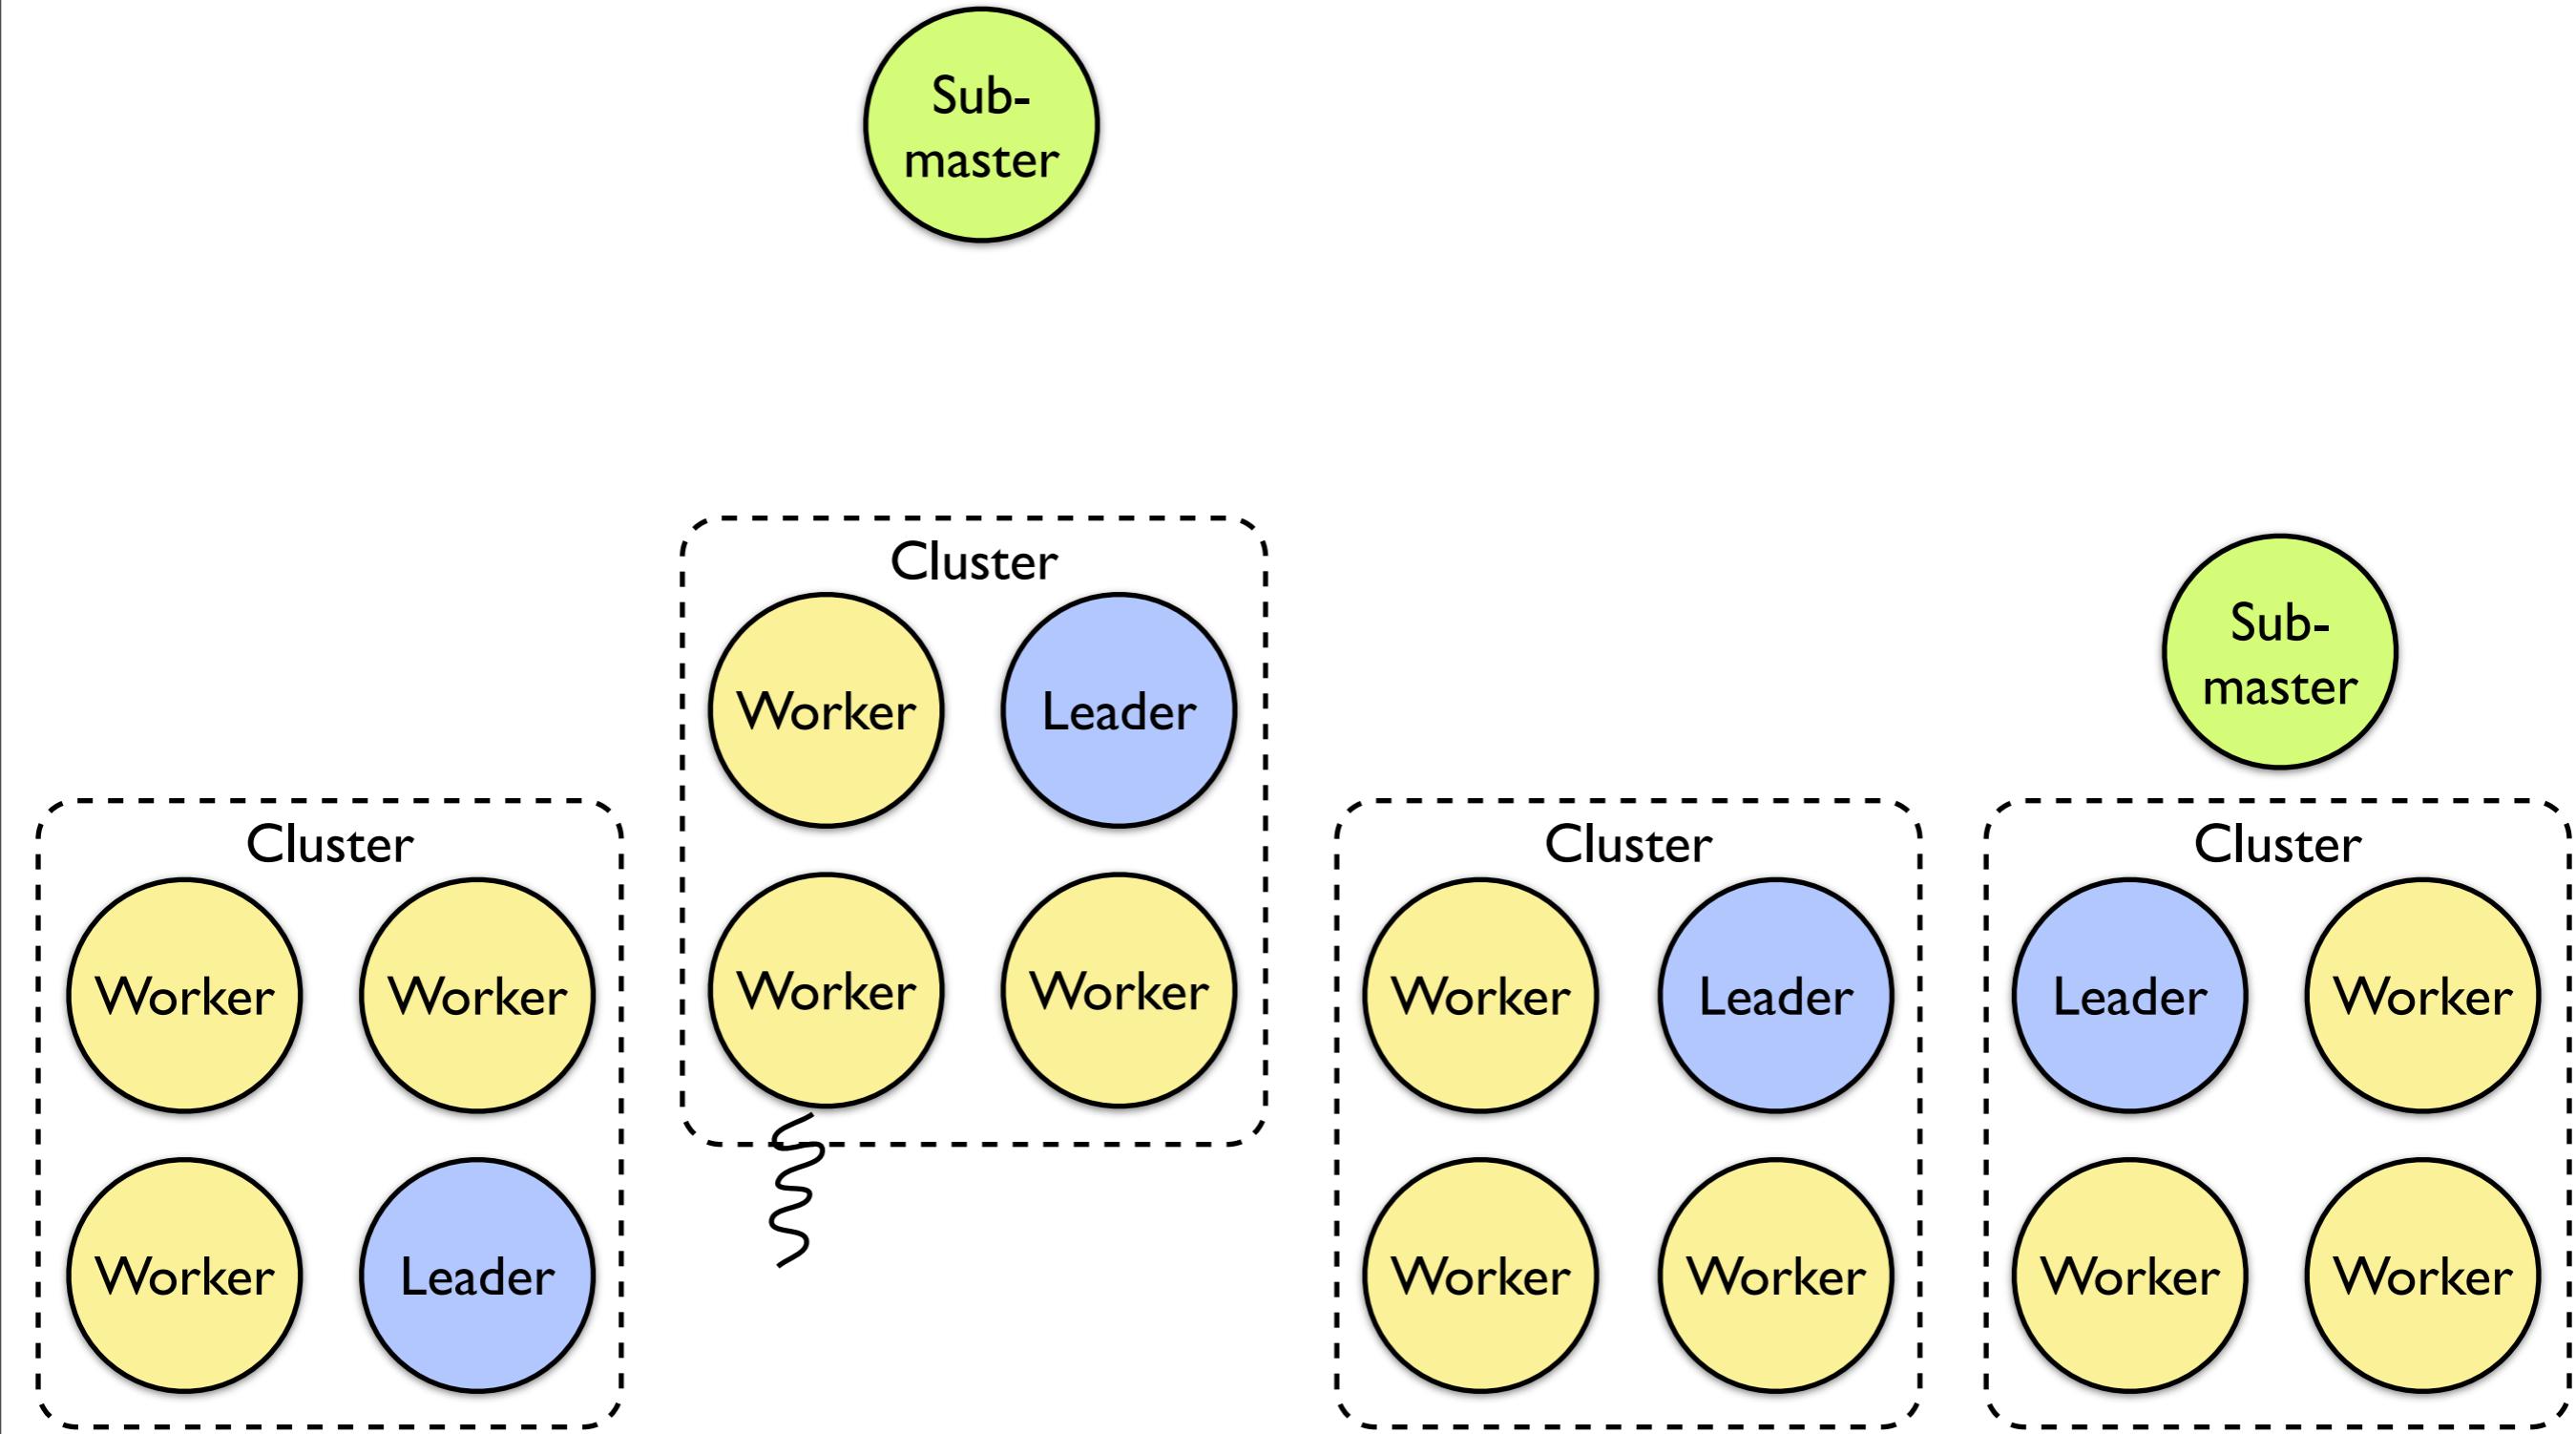
BnB framework - Communications for Sharing Bound



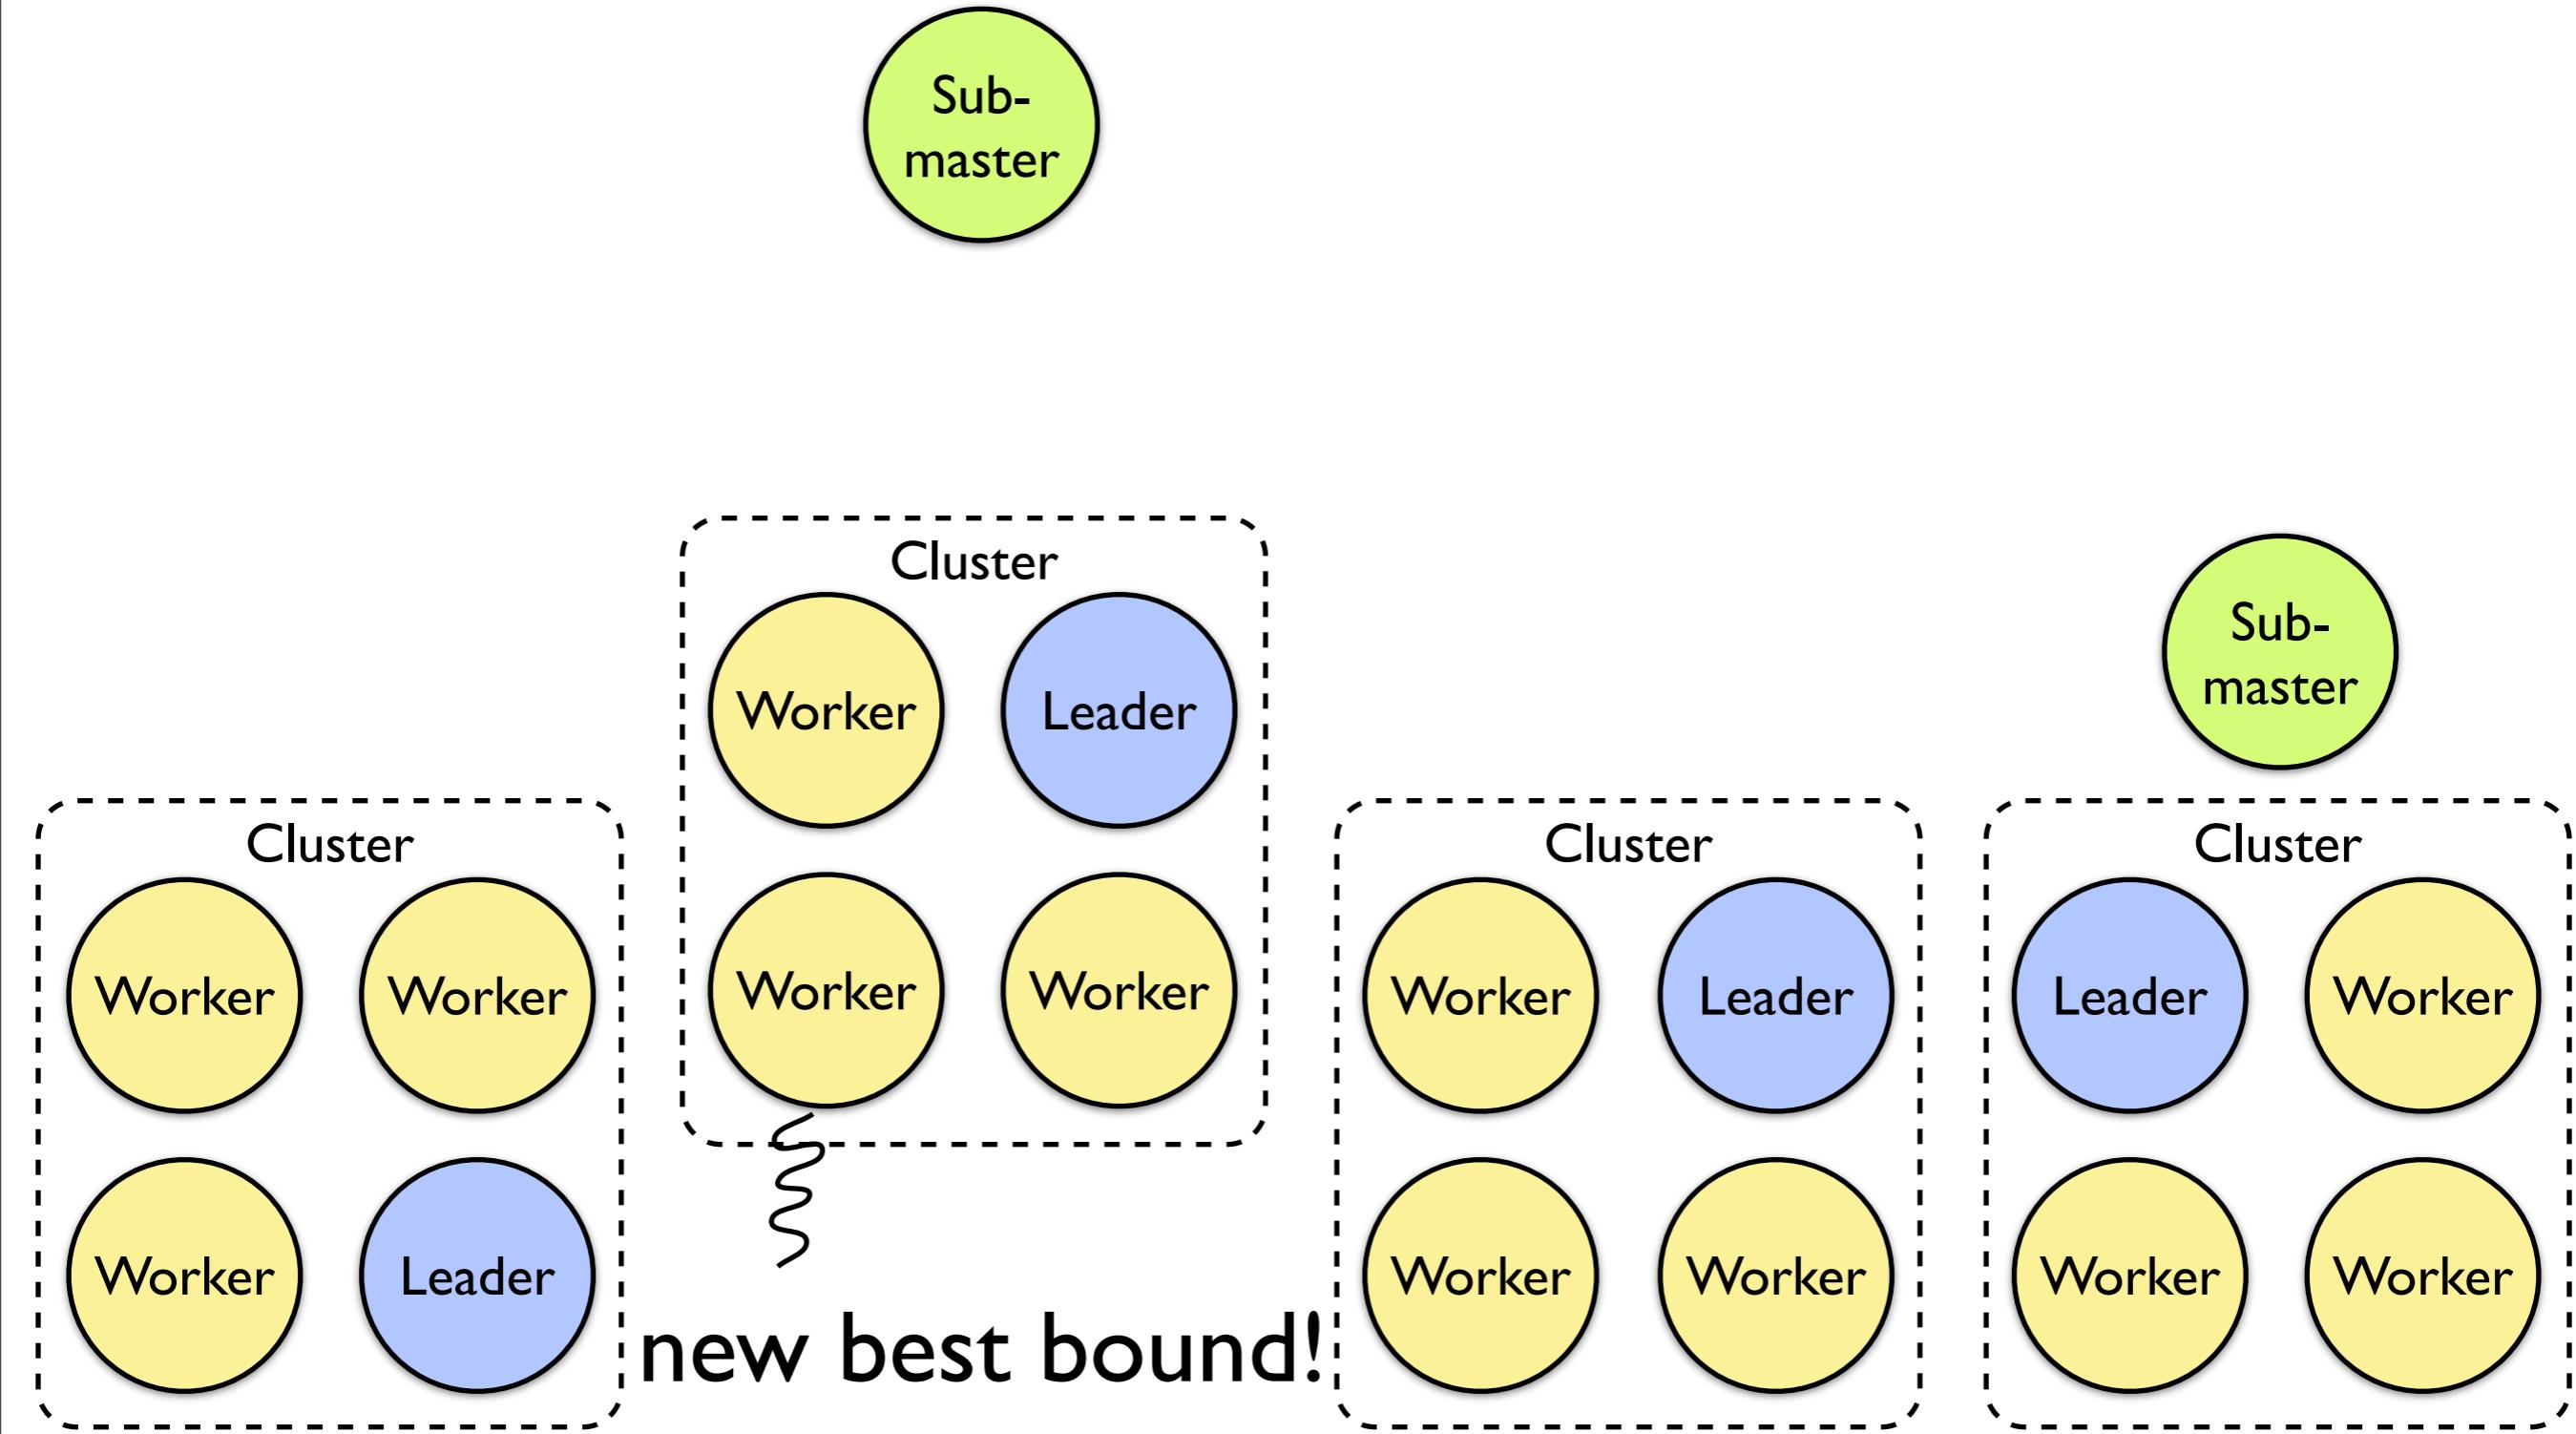
BnB framework - Communications for Sharing Bound



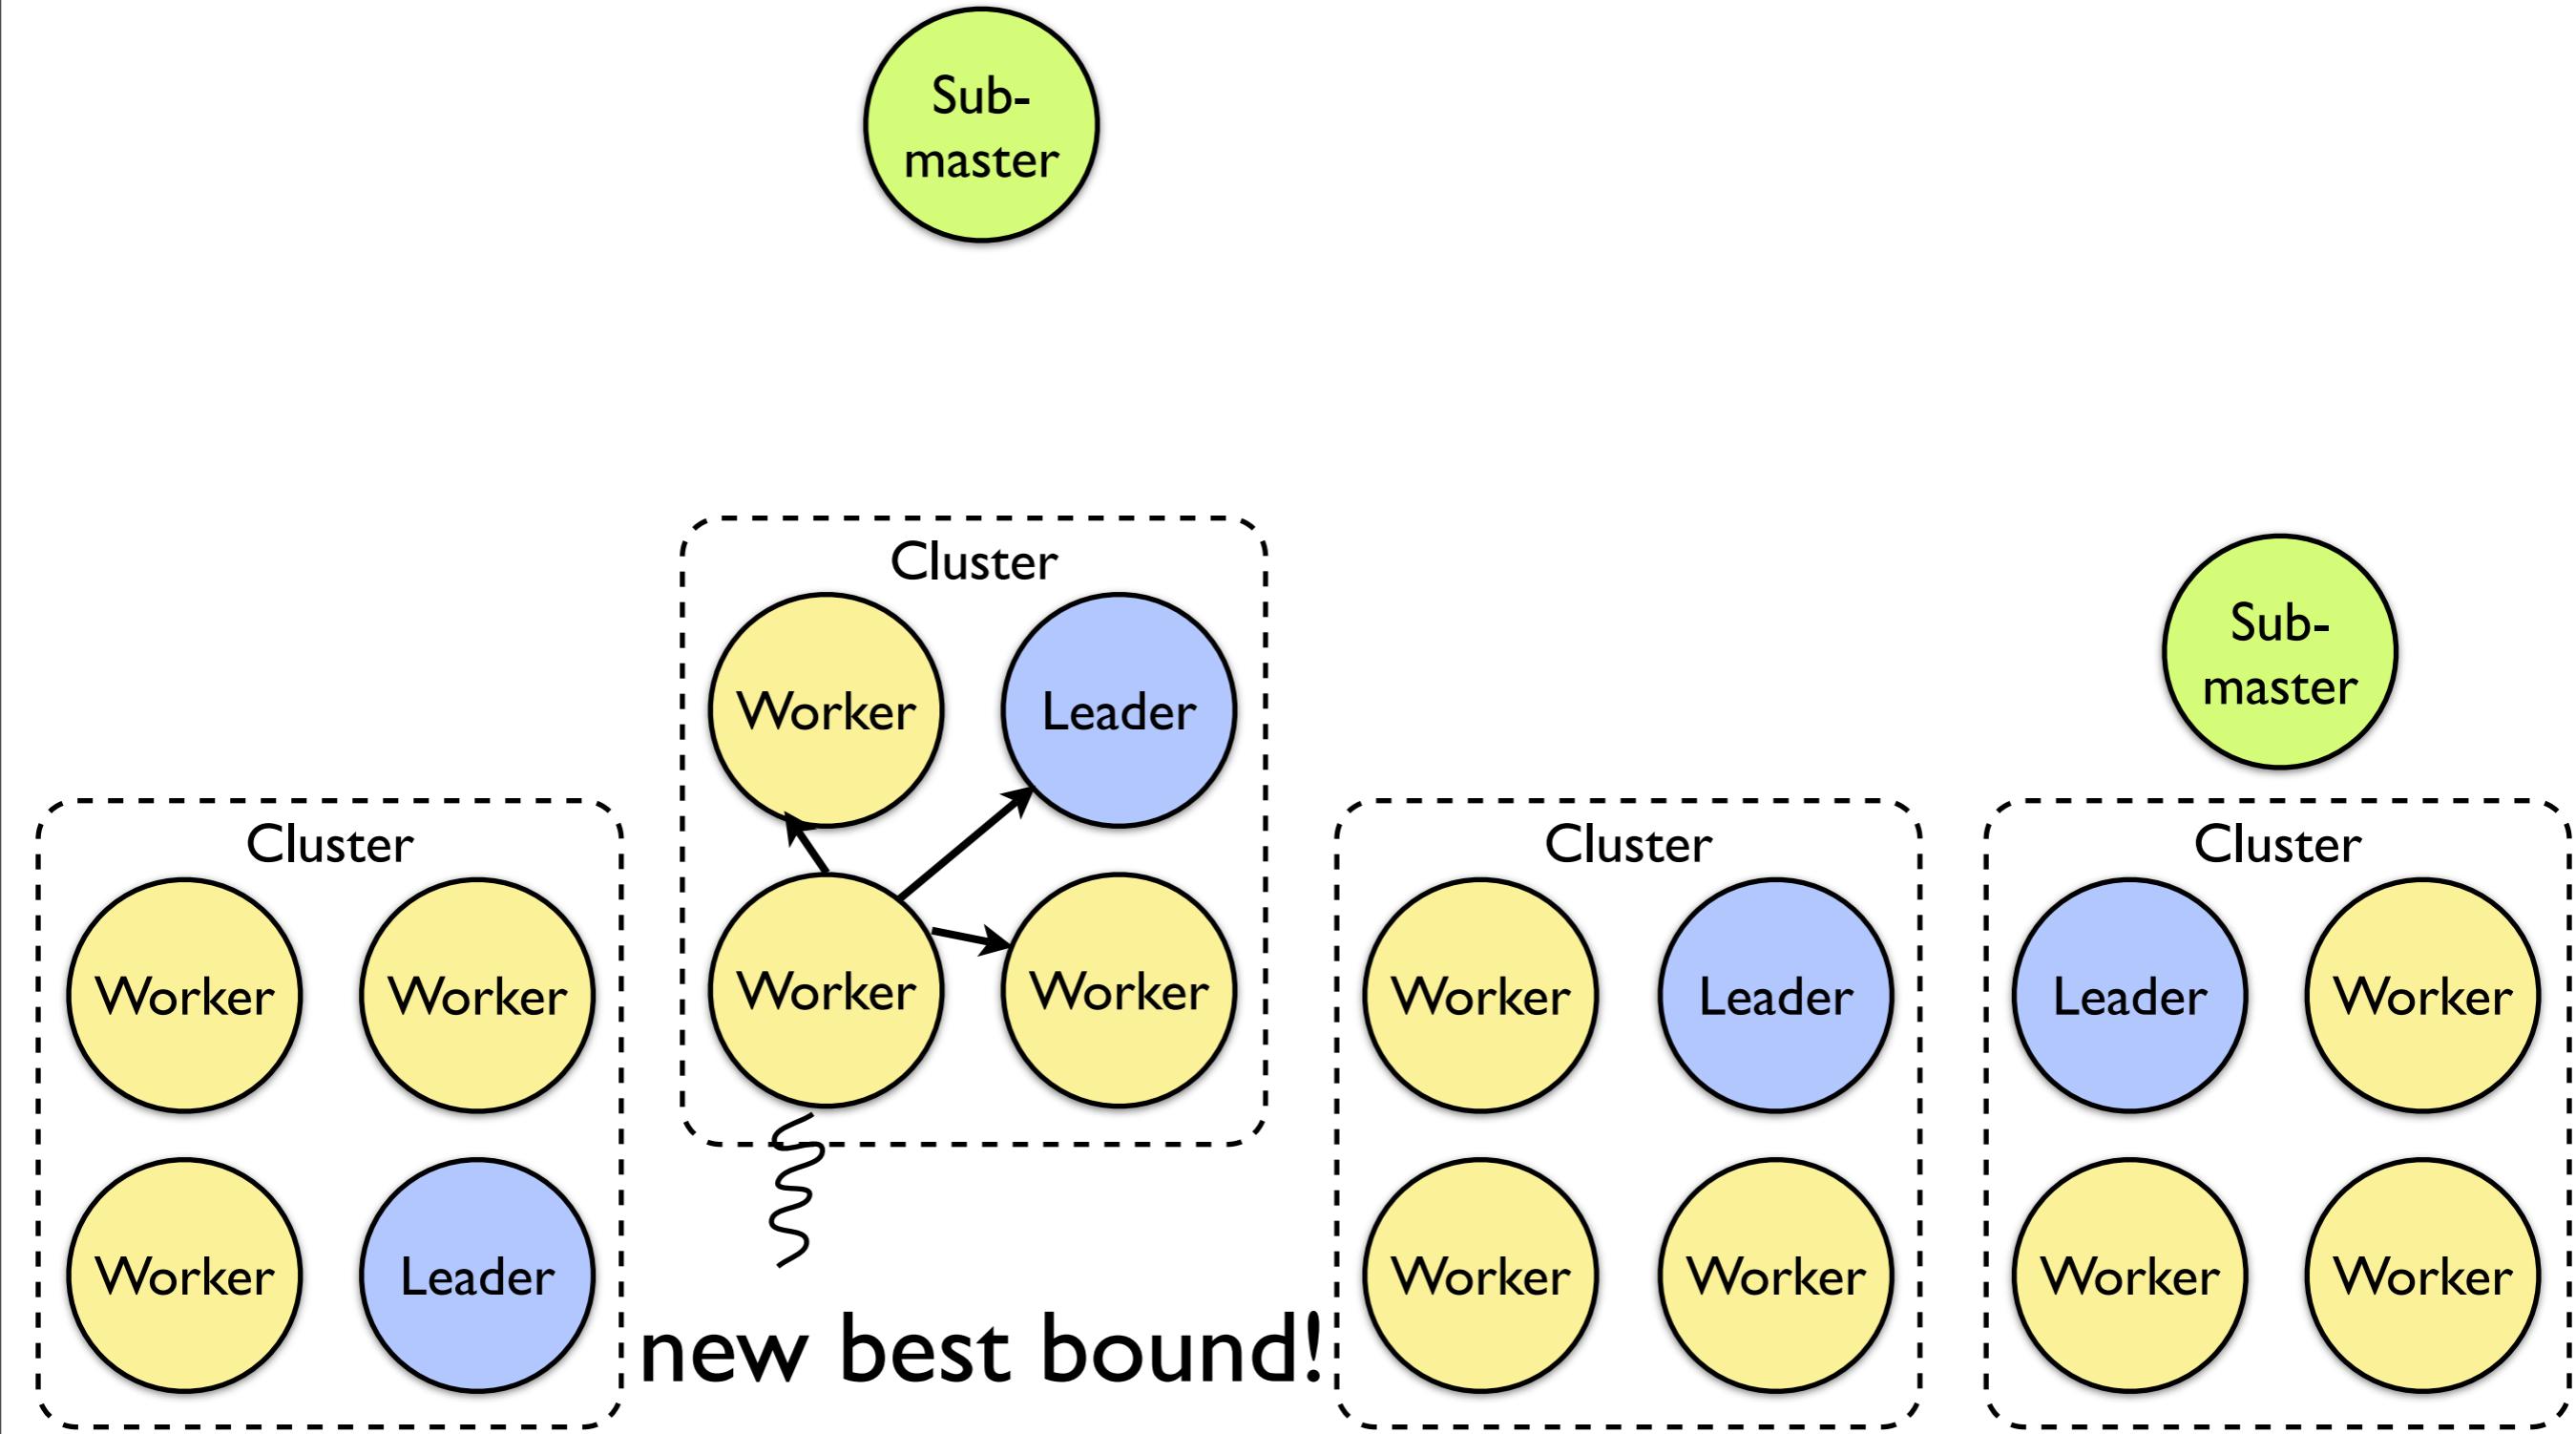
BnB framework - Communications for Sharing Bound



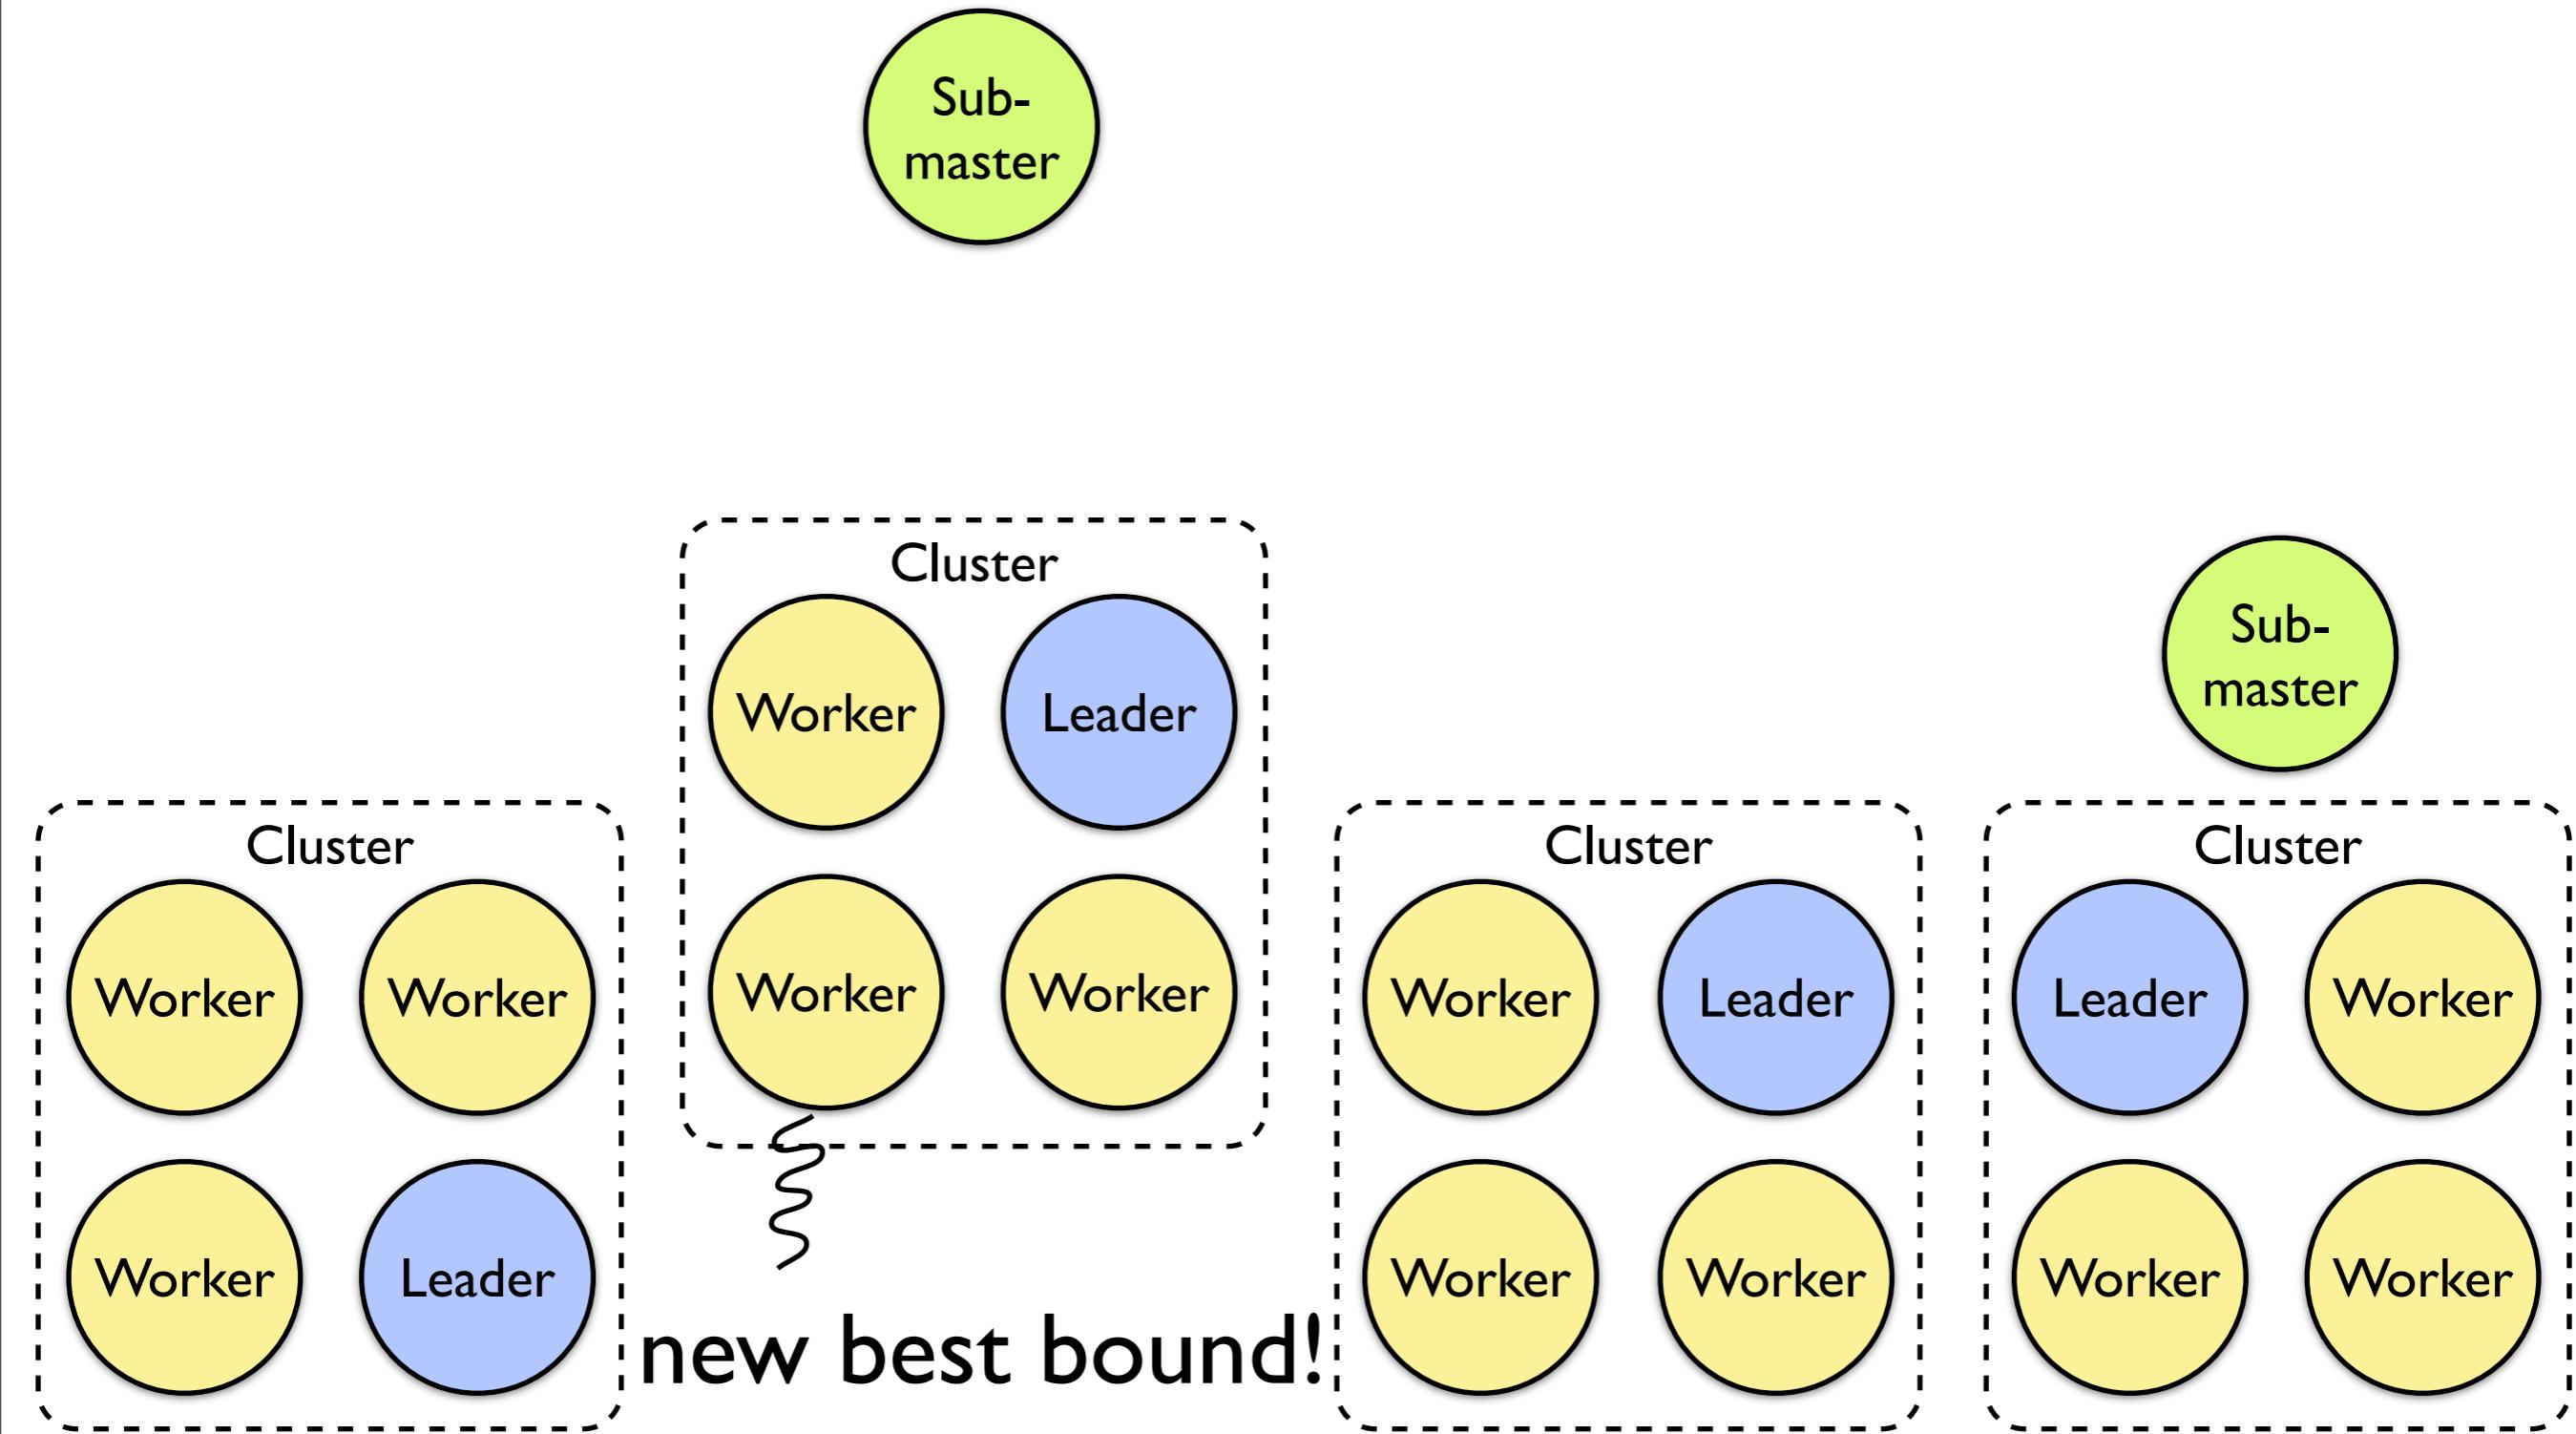
BnB framework - Communications for Sharing Bound



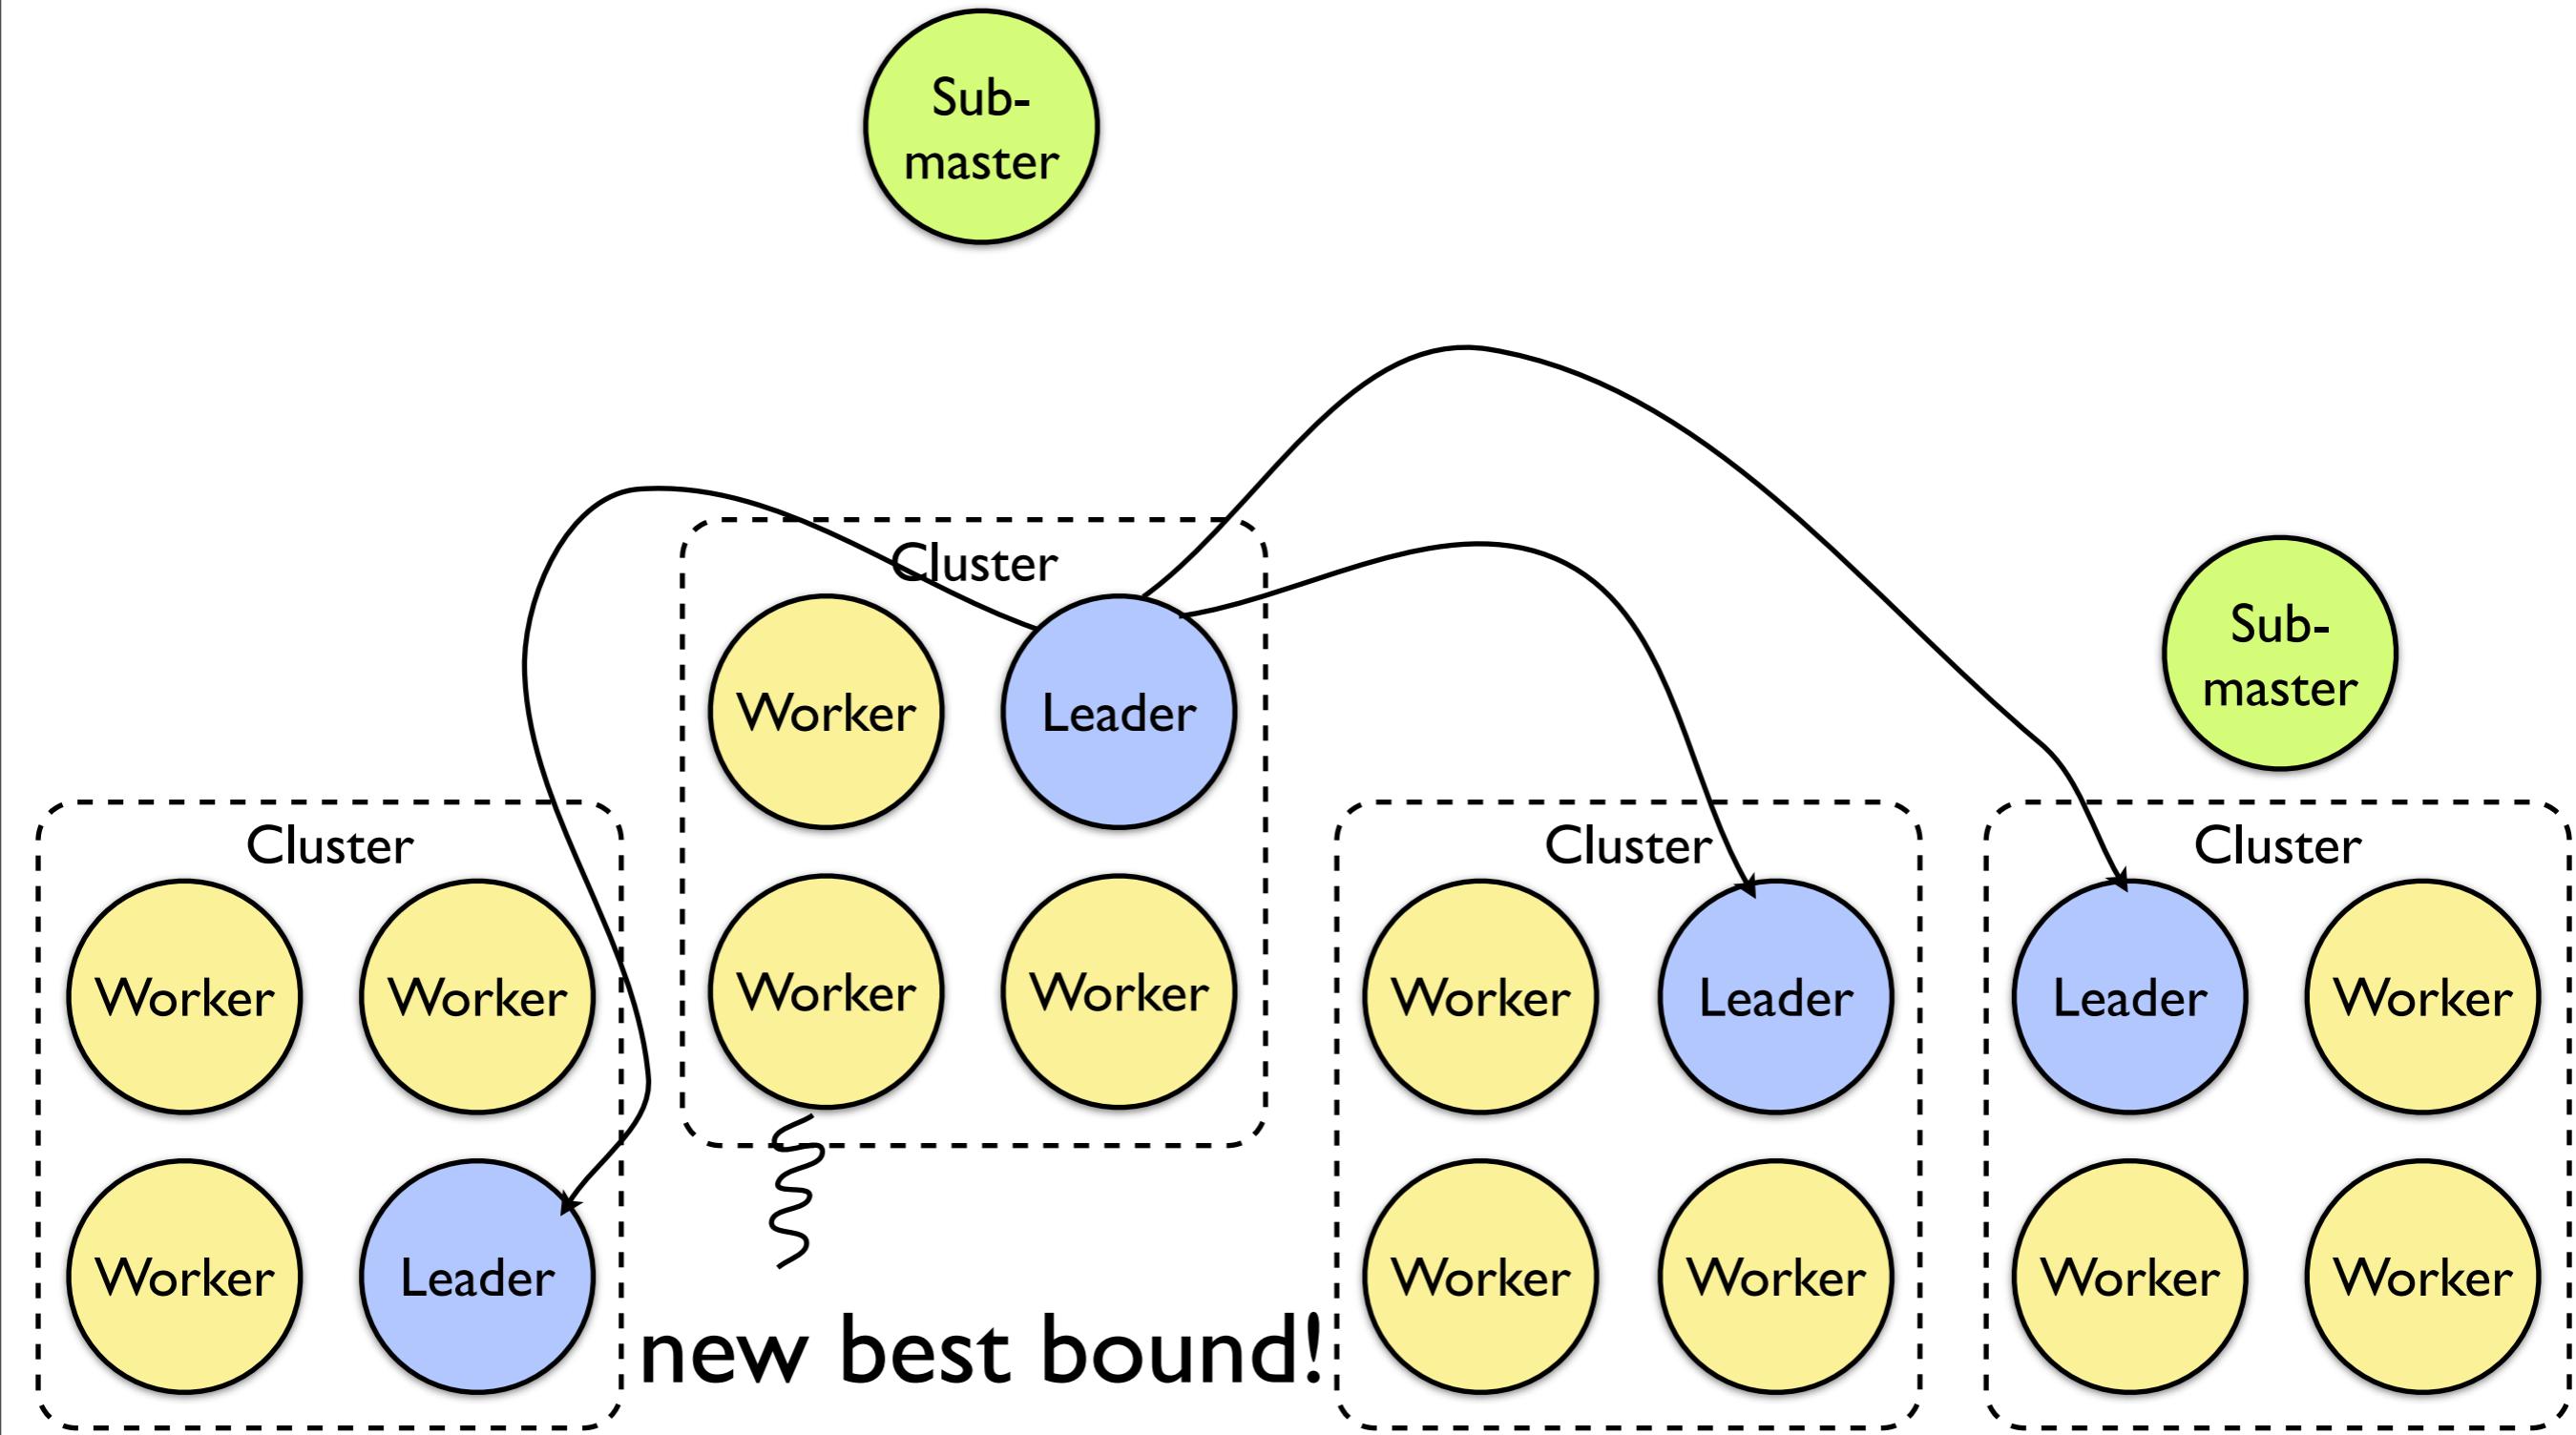
BnB framework - Communications for Sharing Bound



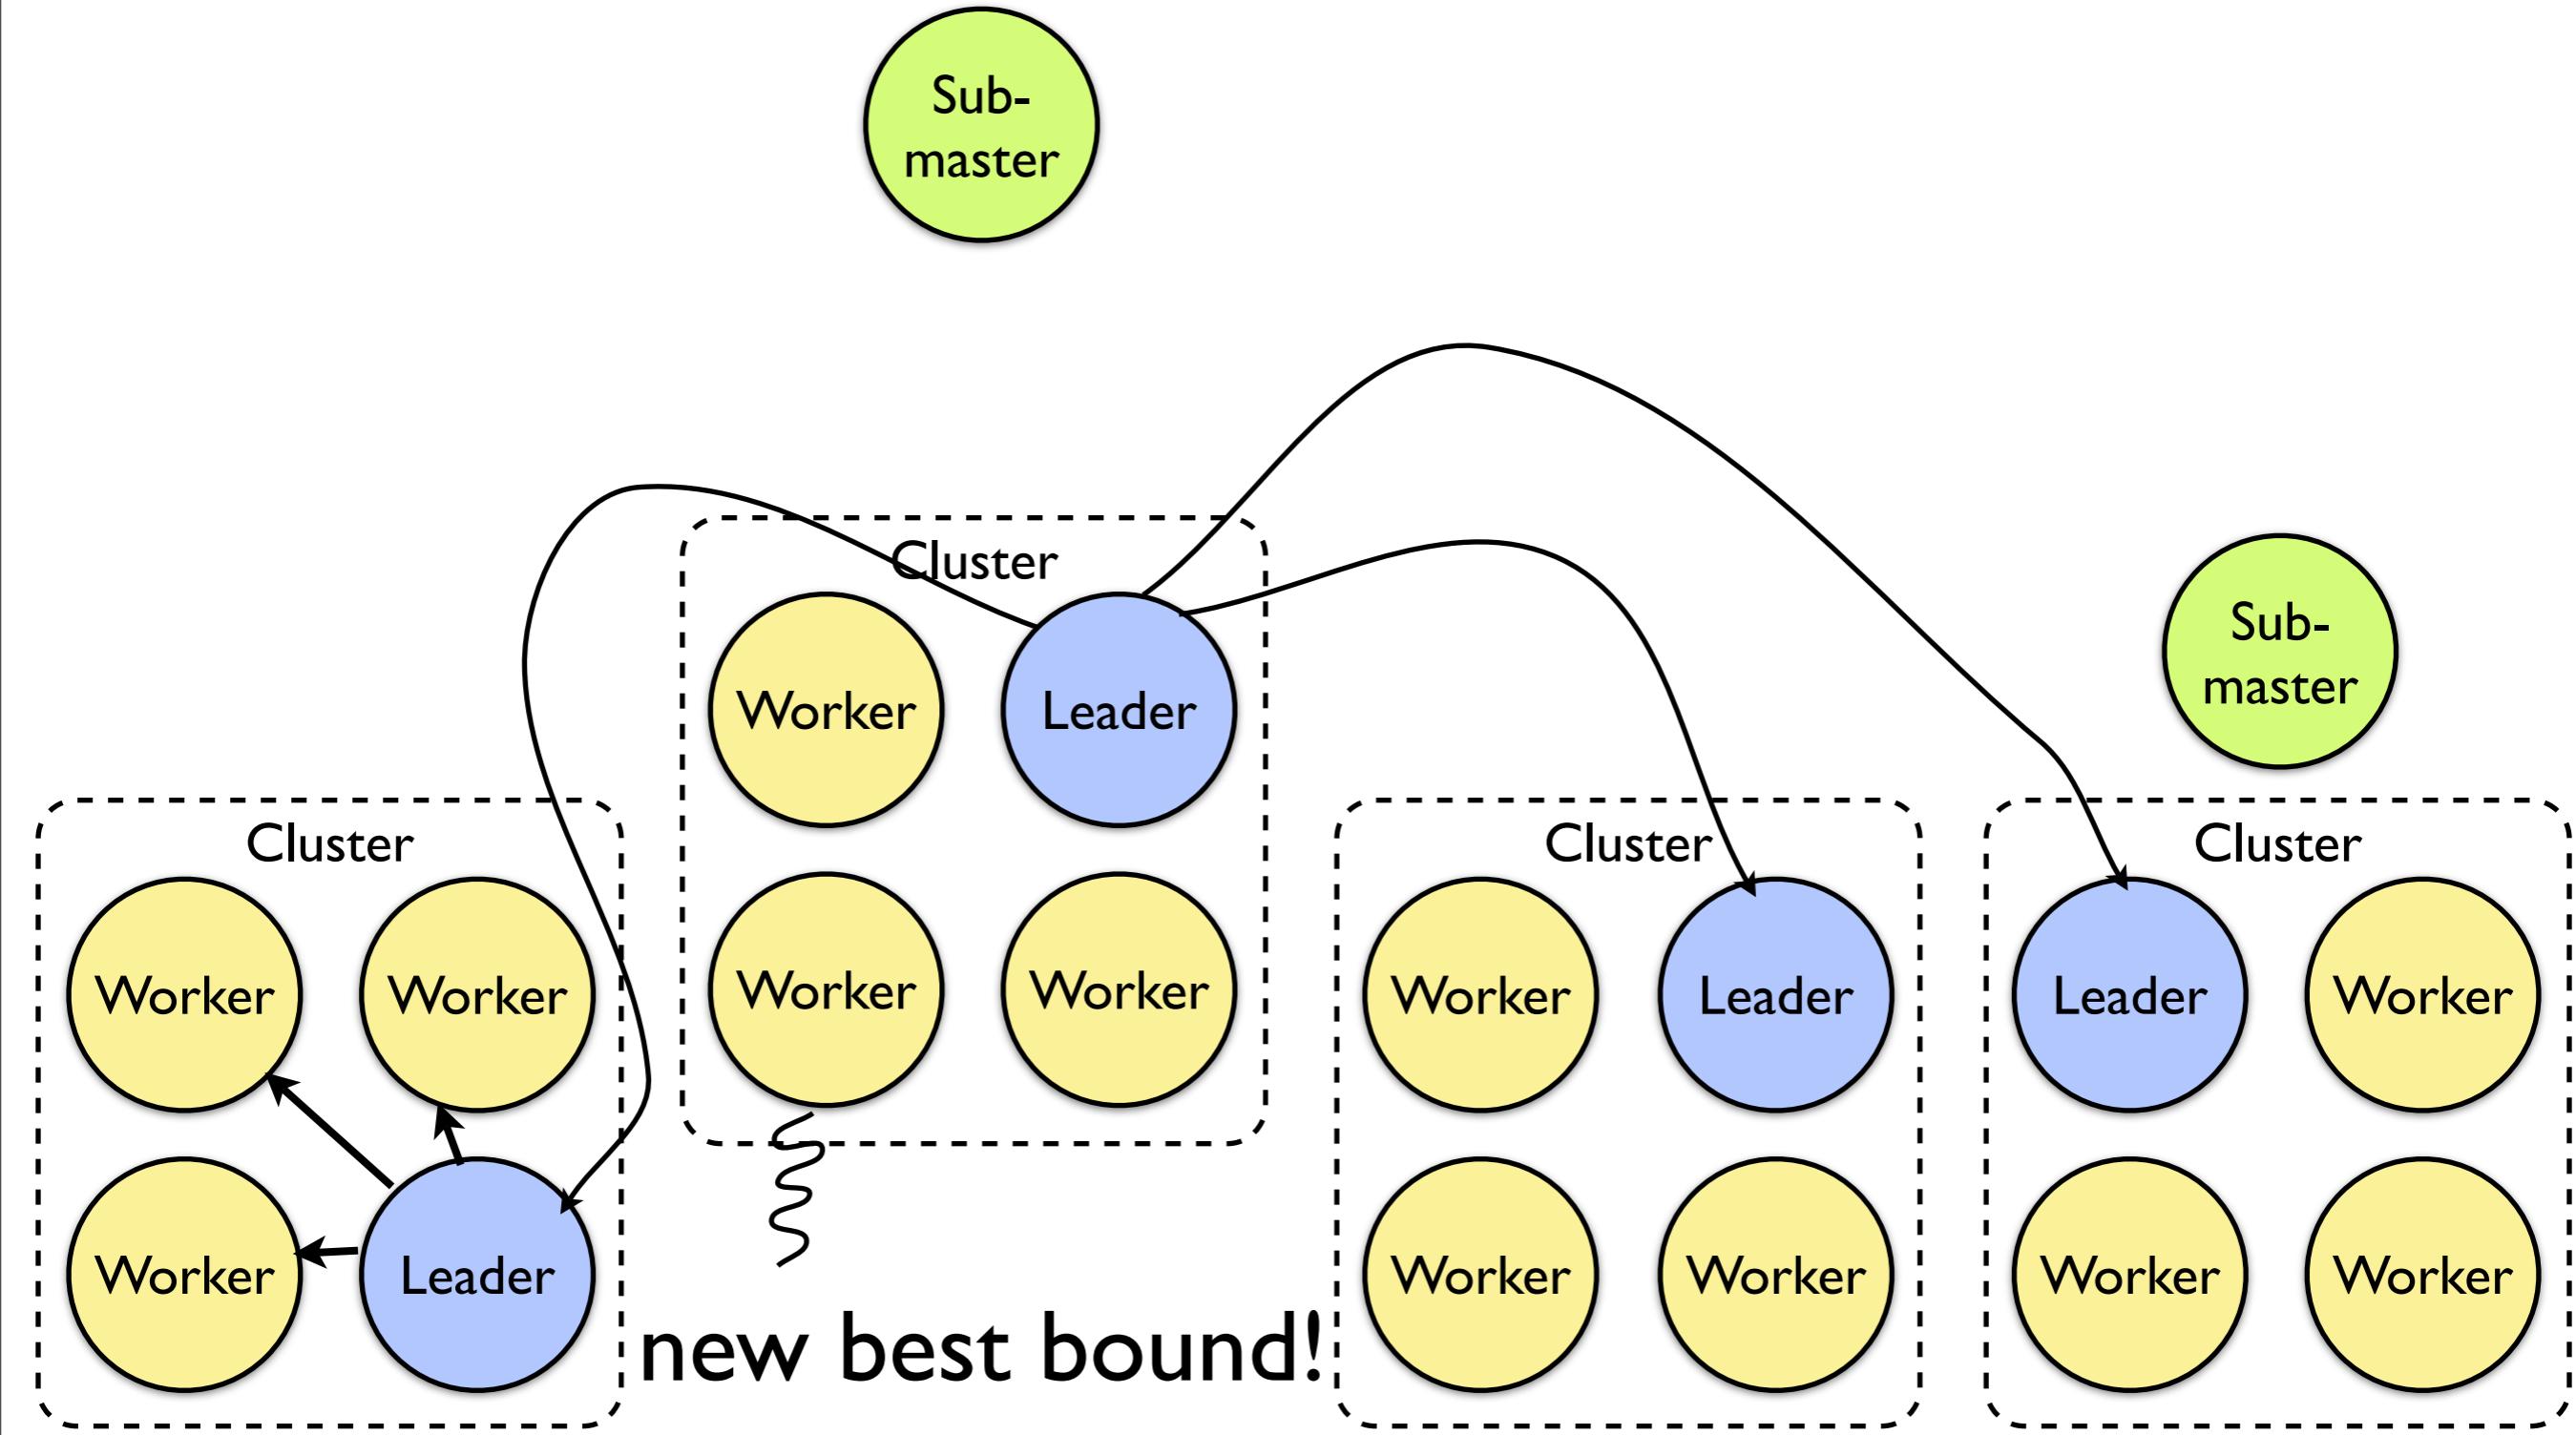
BnB framework - Communications for Sharing Bound



BnB framework - Communications for Sharing Bound

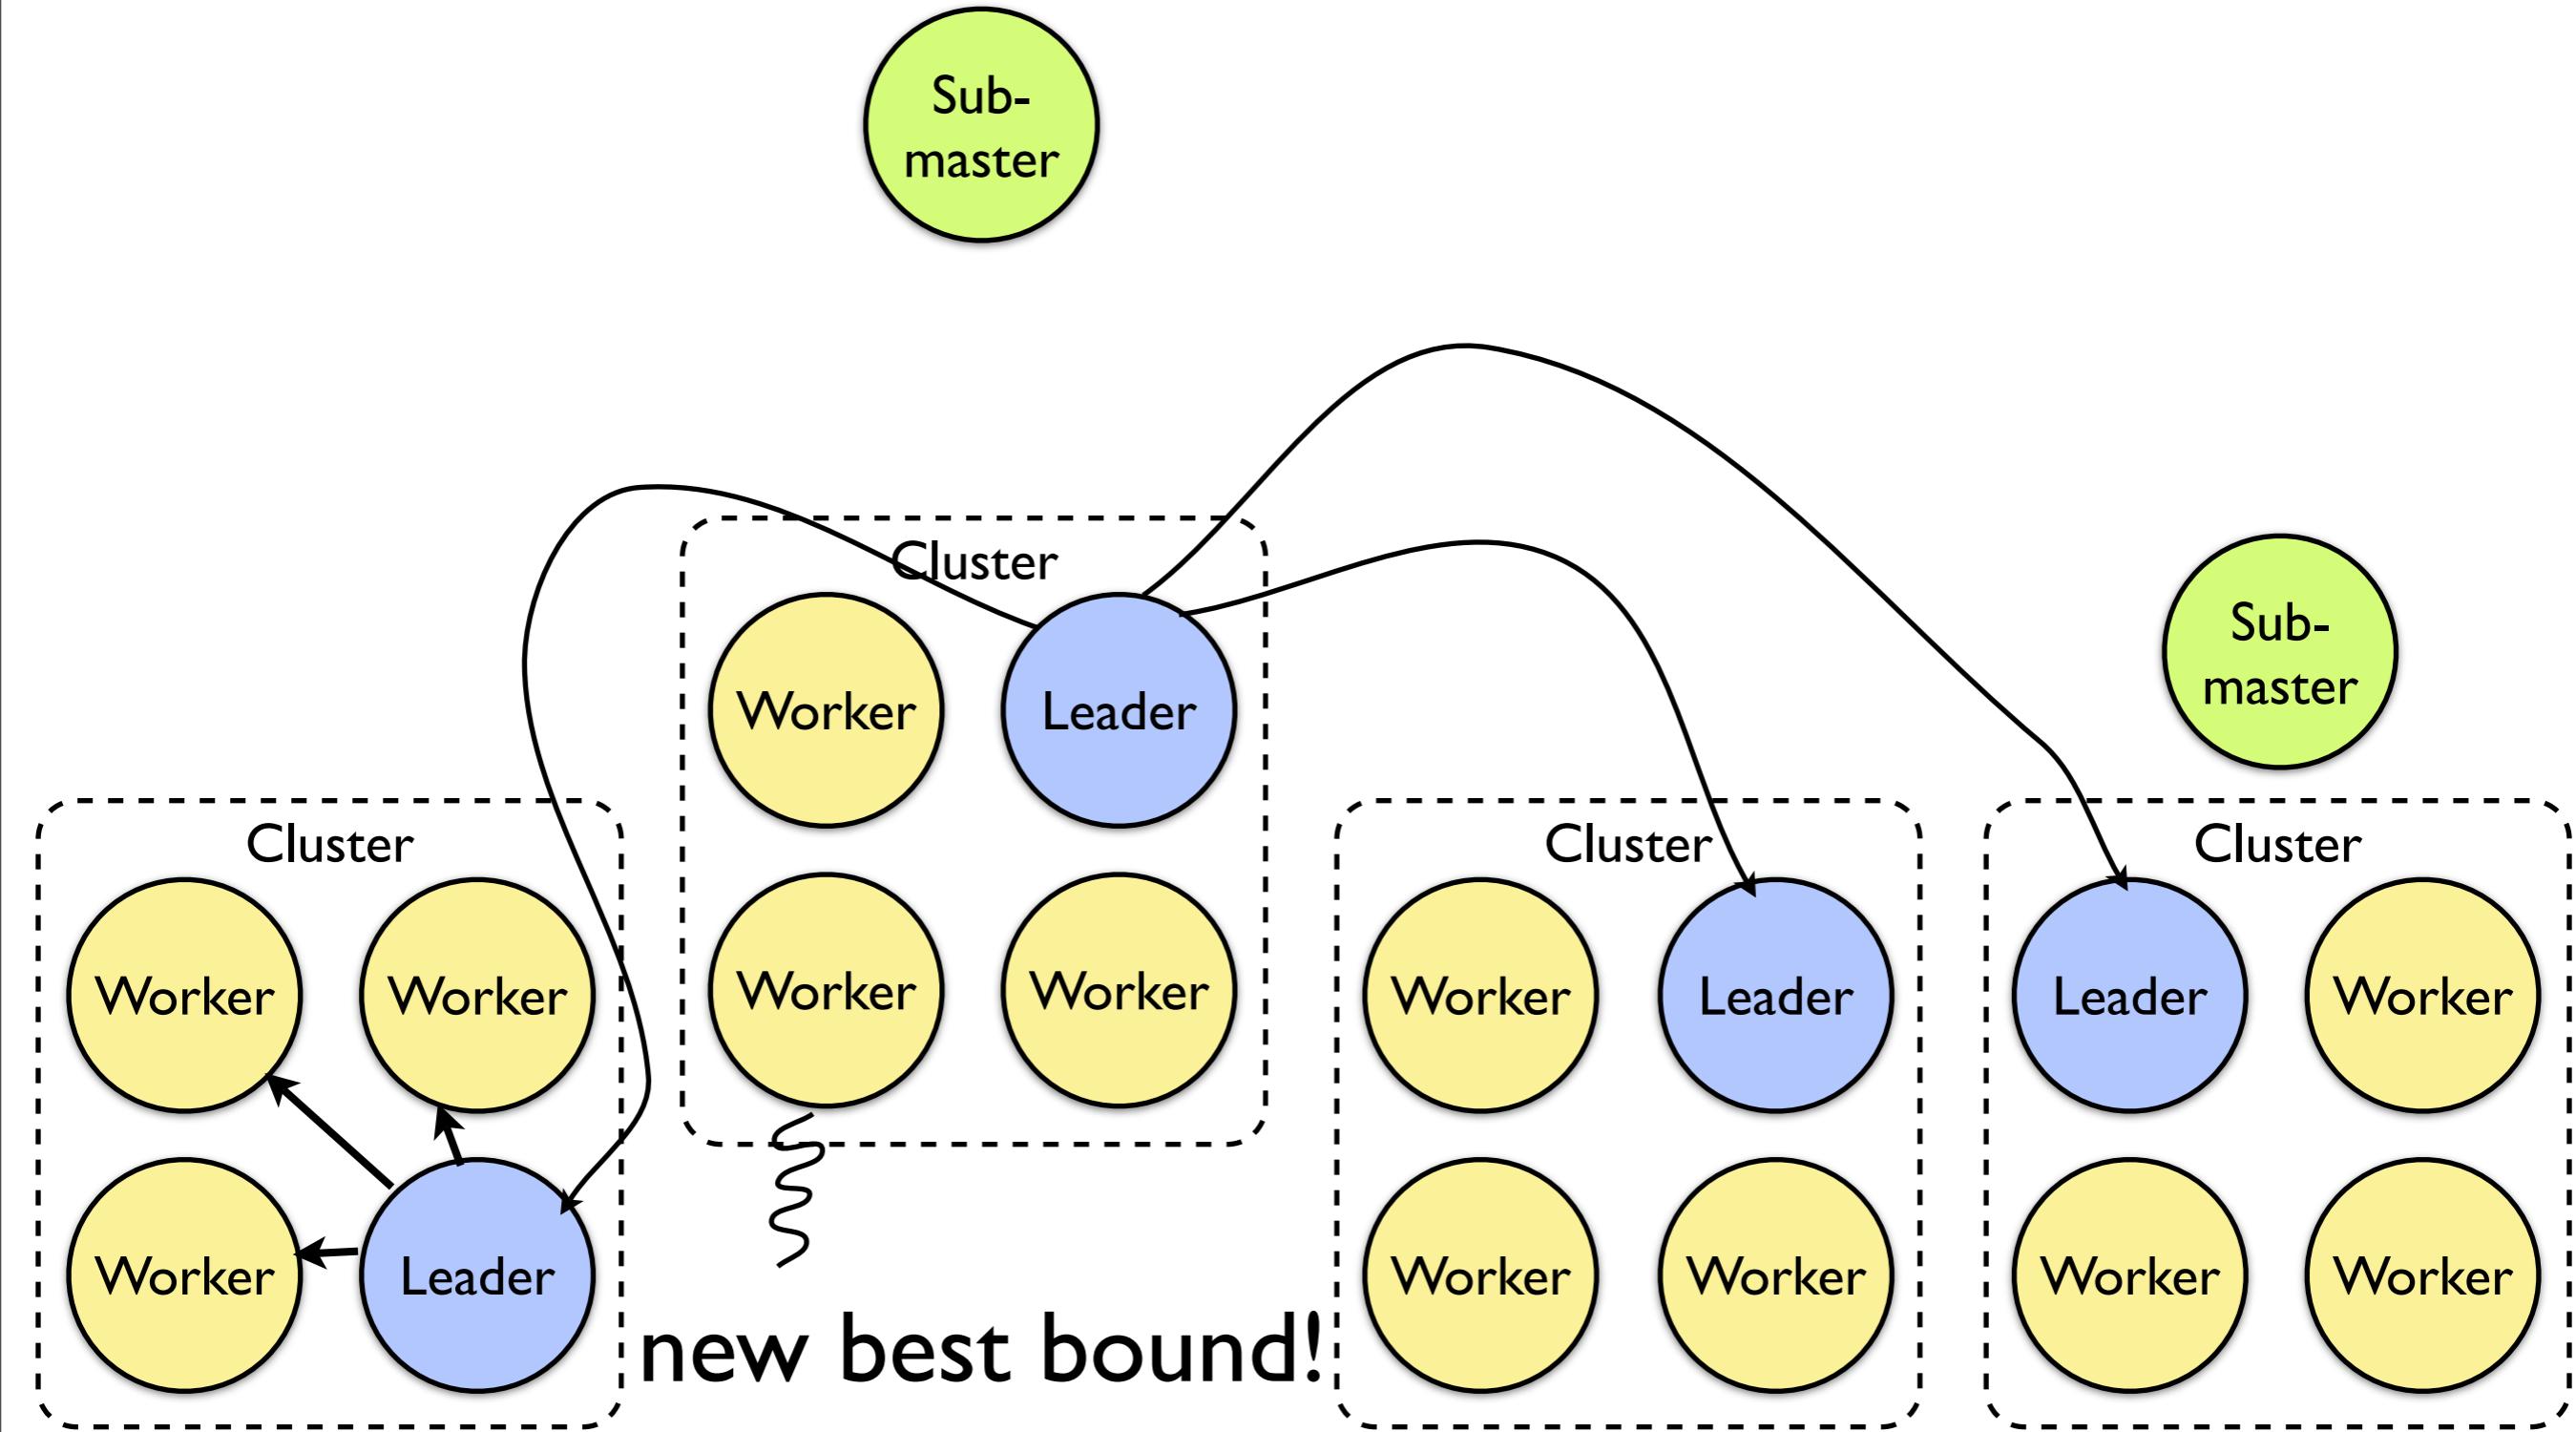


BnB framework - Communications for Sharing Bound



BnB framework - Communications for Sharing Bound

Grid middleware provides communications



Grid'BnB [HiPC07] Features

Grid'BnB [HiPC07] Features

Design

Grid'BnB [HiPC07] Features

- Asynchronous communications
- Hierarchical master-worker with com.
- Dynamic task splitting
- Efficient communications with groups
- Fault-tolerance

Design

Grid'BnB [HiPC07] Features

- Asynchronous communications
- Hierarchical master-worker with com.
- Dynamic task splitting
- Efficient communications with groups
- Fault-tolerance

Design

Users

Grid'BnB [HiPC07] Features

- Asynchronous communications
- Hierarchical master-worker with com.
- Dynamic task splitting
- Efficient communications with groups
- Fault-tolerance

Design

- Hidden parallelism and Grid difficulties
- API for COPs
- Ease of deployment
- Principally tree-based
- Implementing and testing search strategies
- Focus on objective function

Users

Grid'BnB [HiPC07] Features

- Asynchronous communications
- Hierarchical master-worker with com.
- Dynamic task splitting
- Efficient communications with groups
-

Design

Validate and Test Grid'BnB by experiments

- Hidden parallelism and Grid difficulties
- API for COPs
- Ease of deployment
- Principally tree-based
- Implementing and testing search strategies
- Focus on objective function

Users

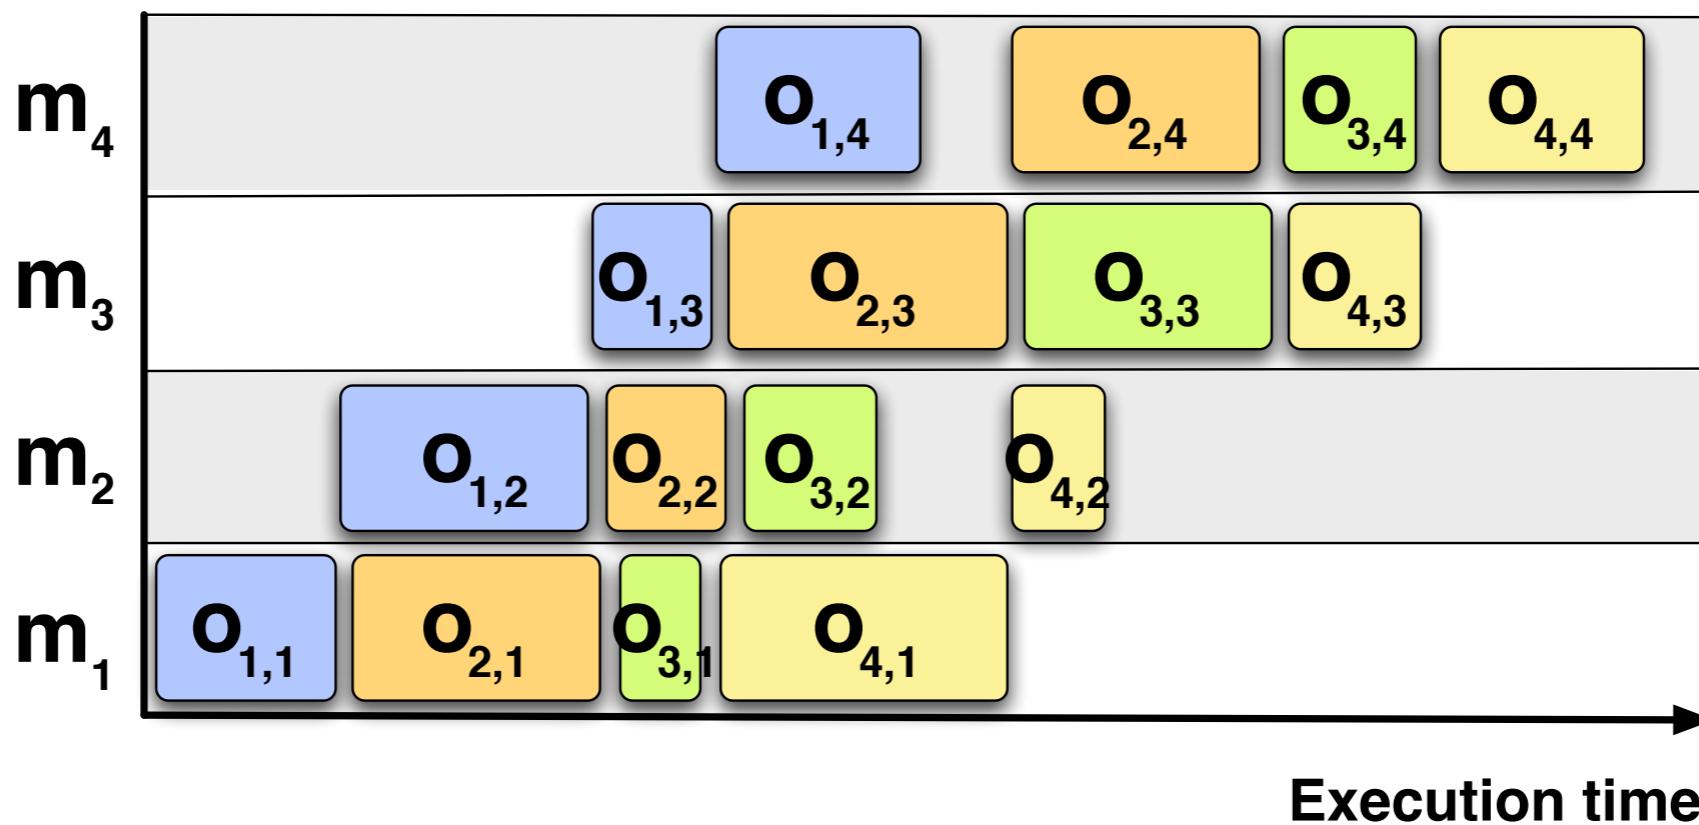
Flow-Shop Experiments

- NP-hard permutation optimization problem

$$J = \{j_1, j_2, \dots, j_n\}$$

$$j_i = \{o_{i1}, o_{i2}, \dots, o_{im}\}$$

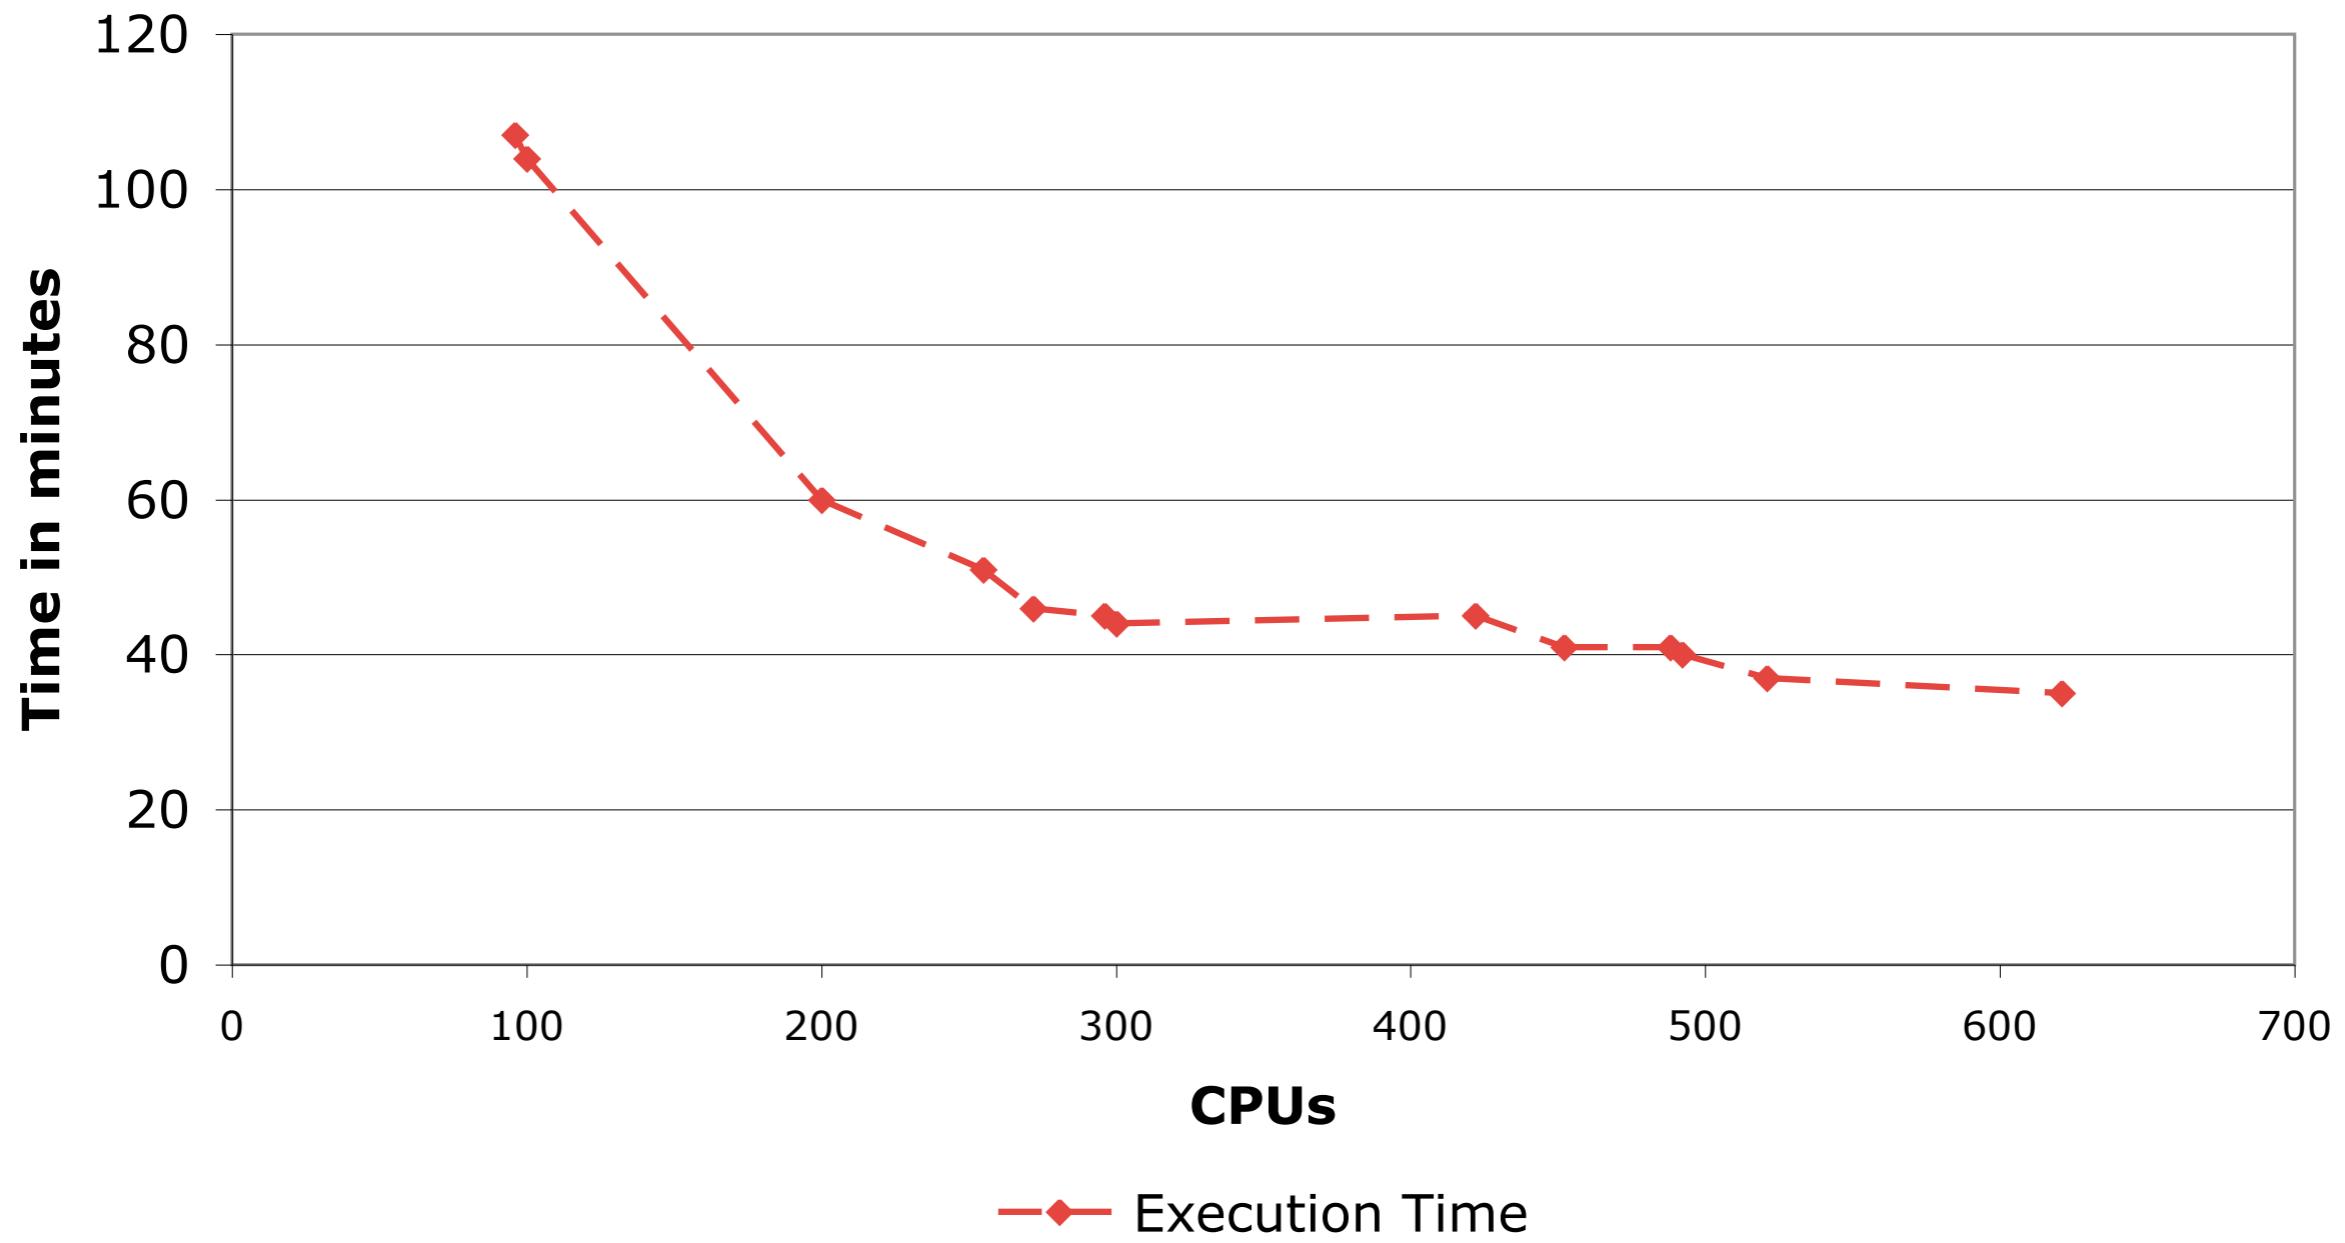
$$M = \{m_1, m_2, \dots, m_m\}$$



Grid Experimentations

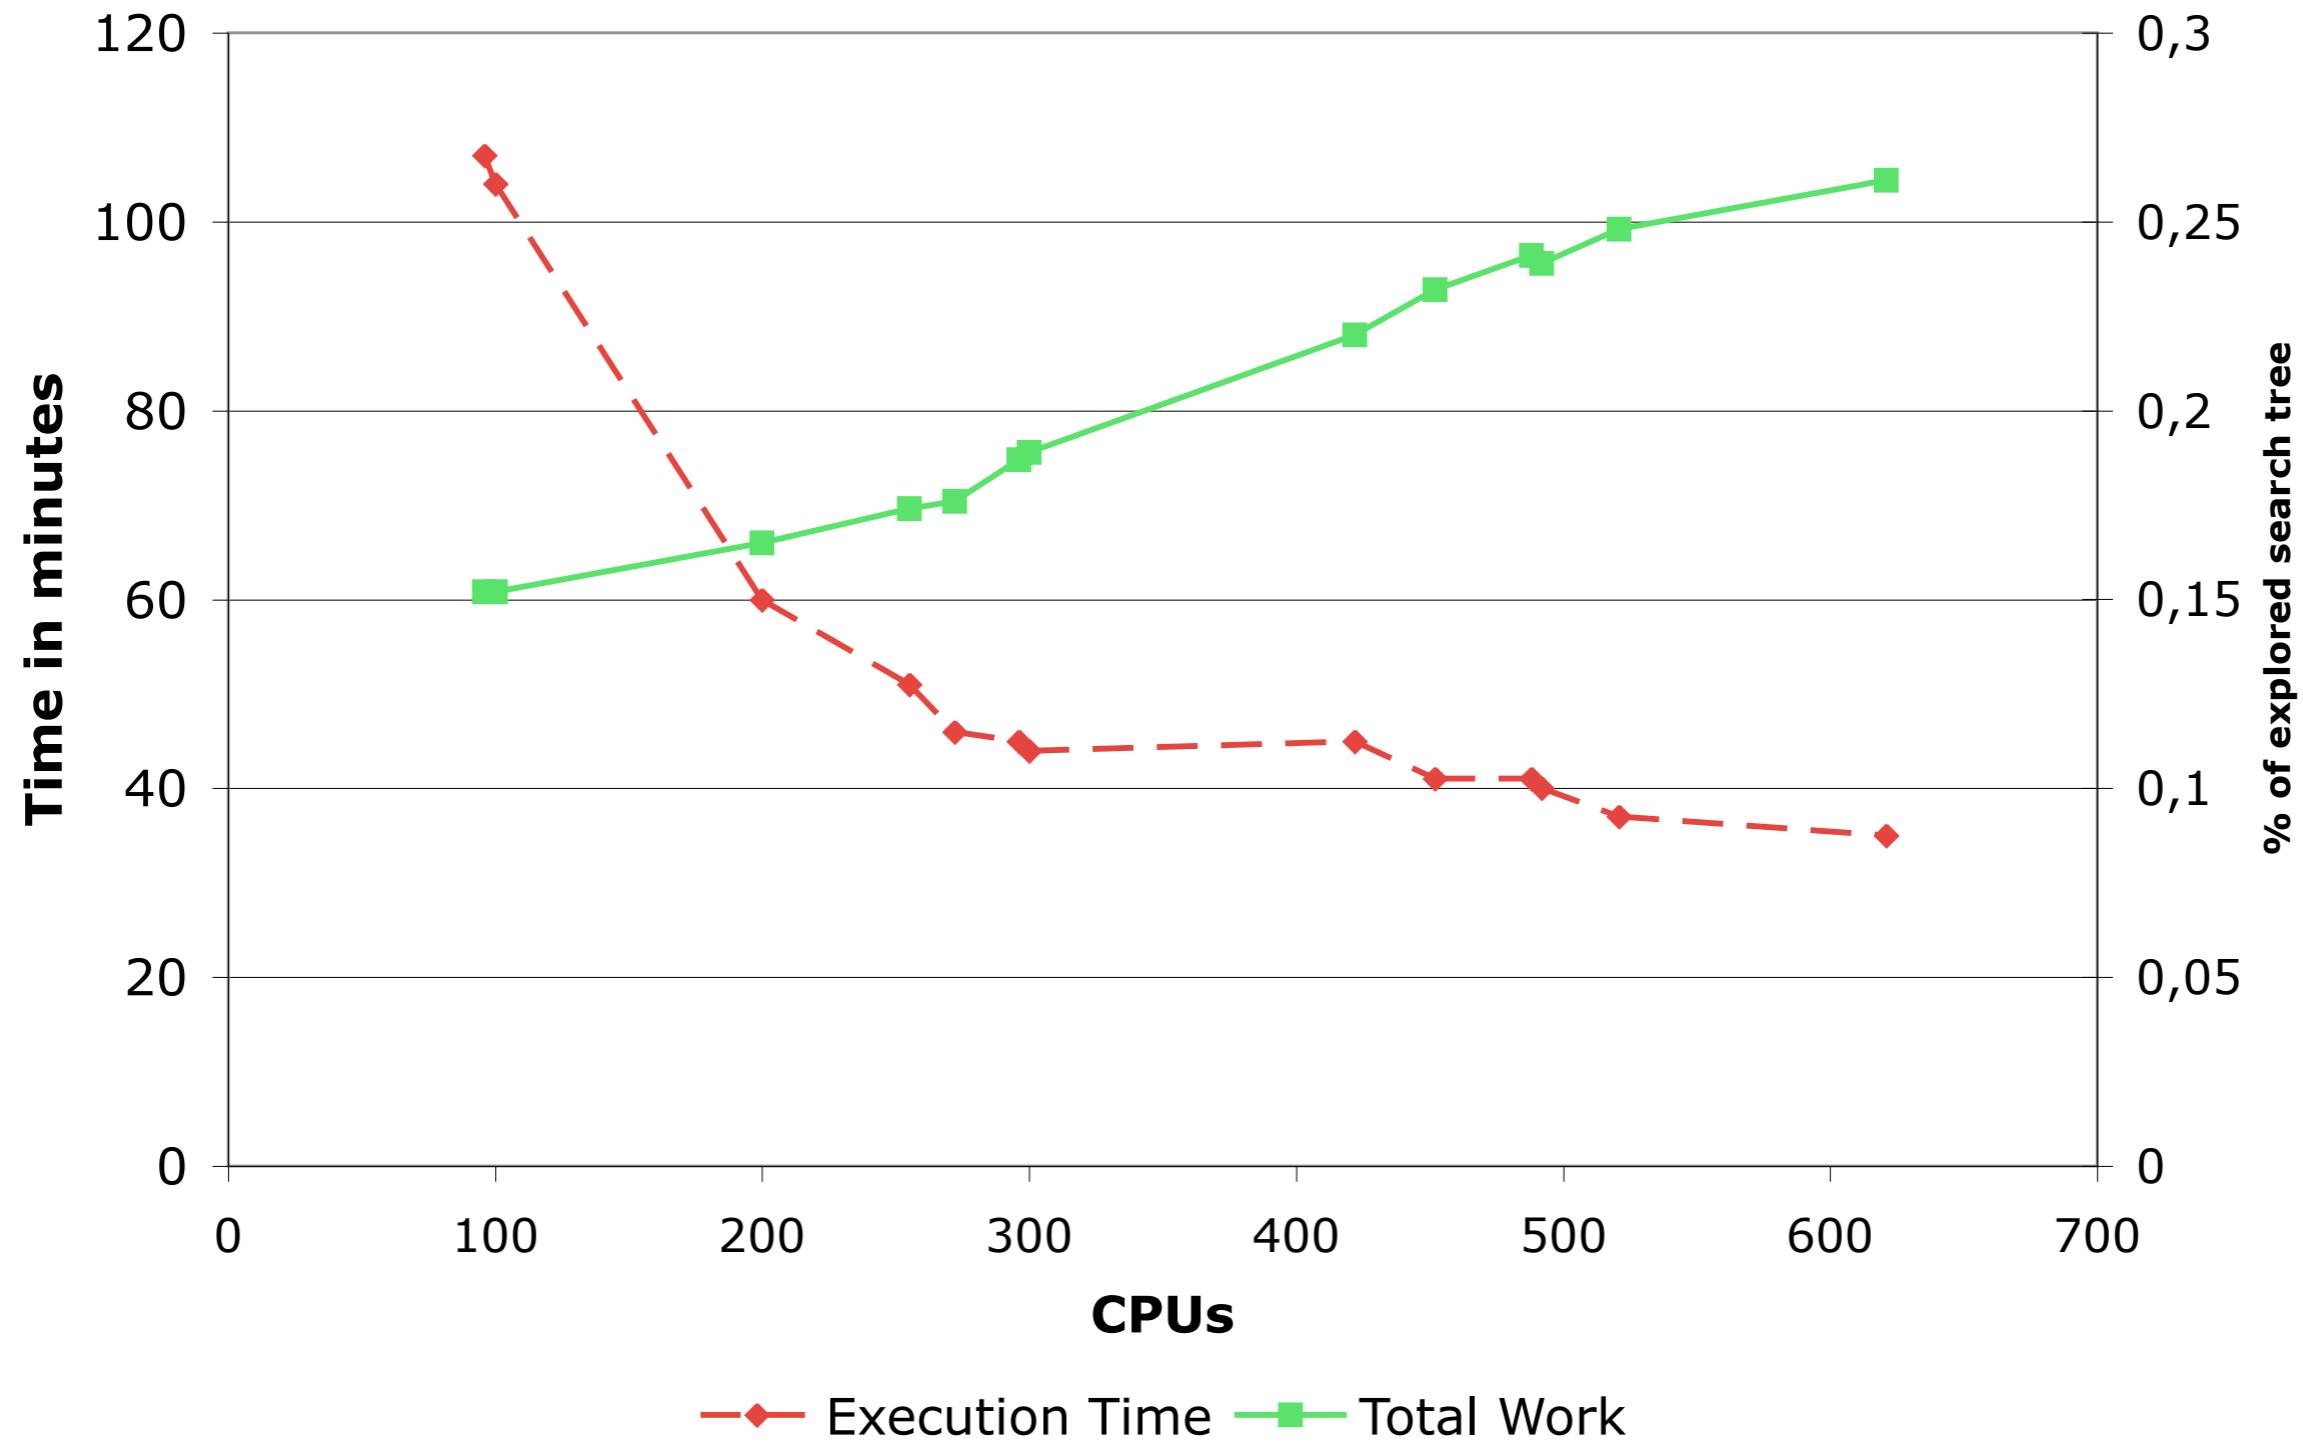
up to 621 CPUs on 5 sites

Flow-Shop: 17 Jobs / 17 Machines



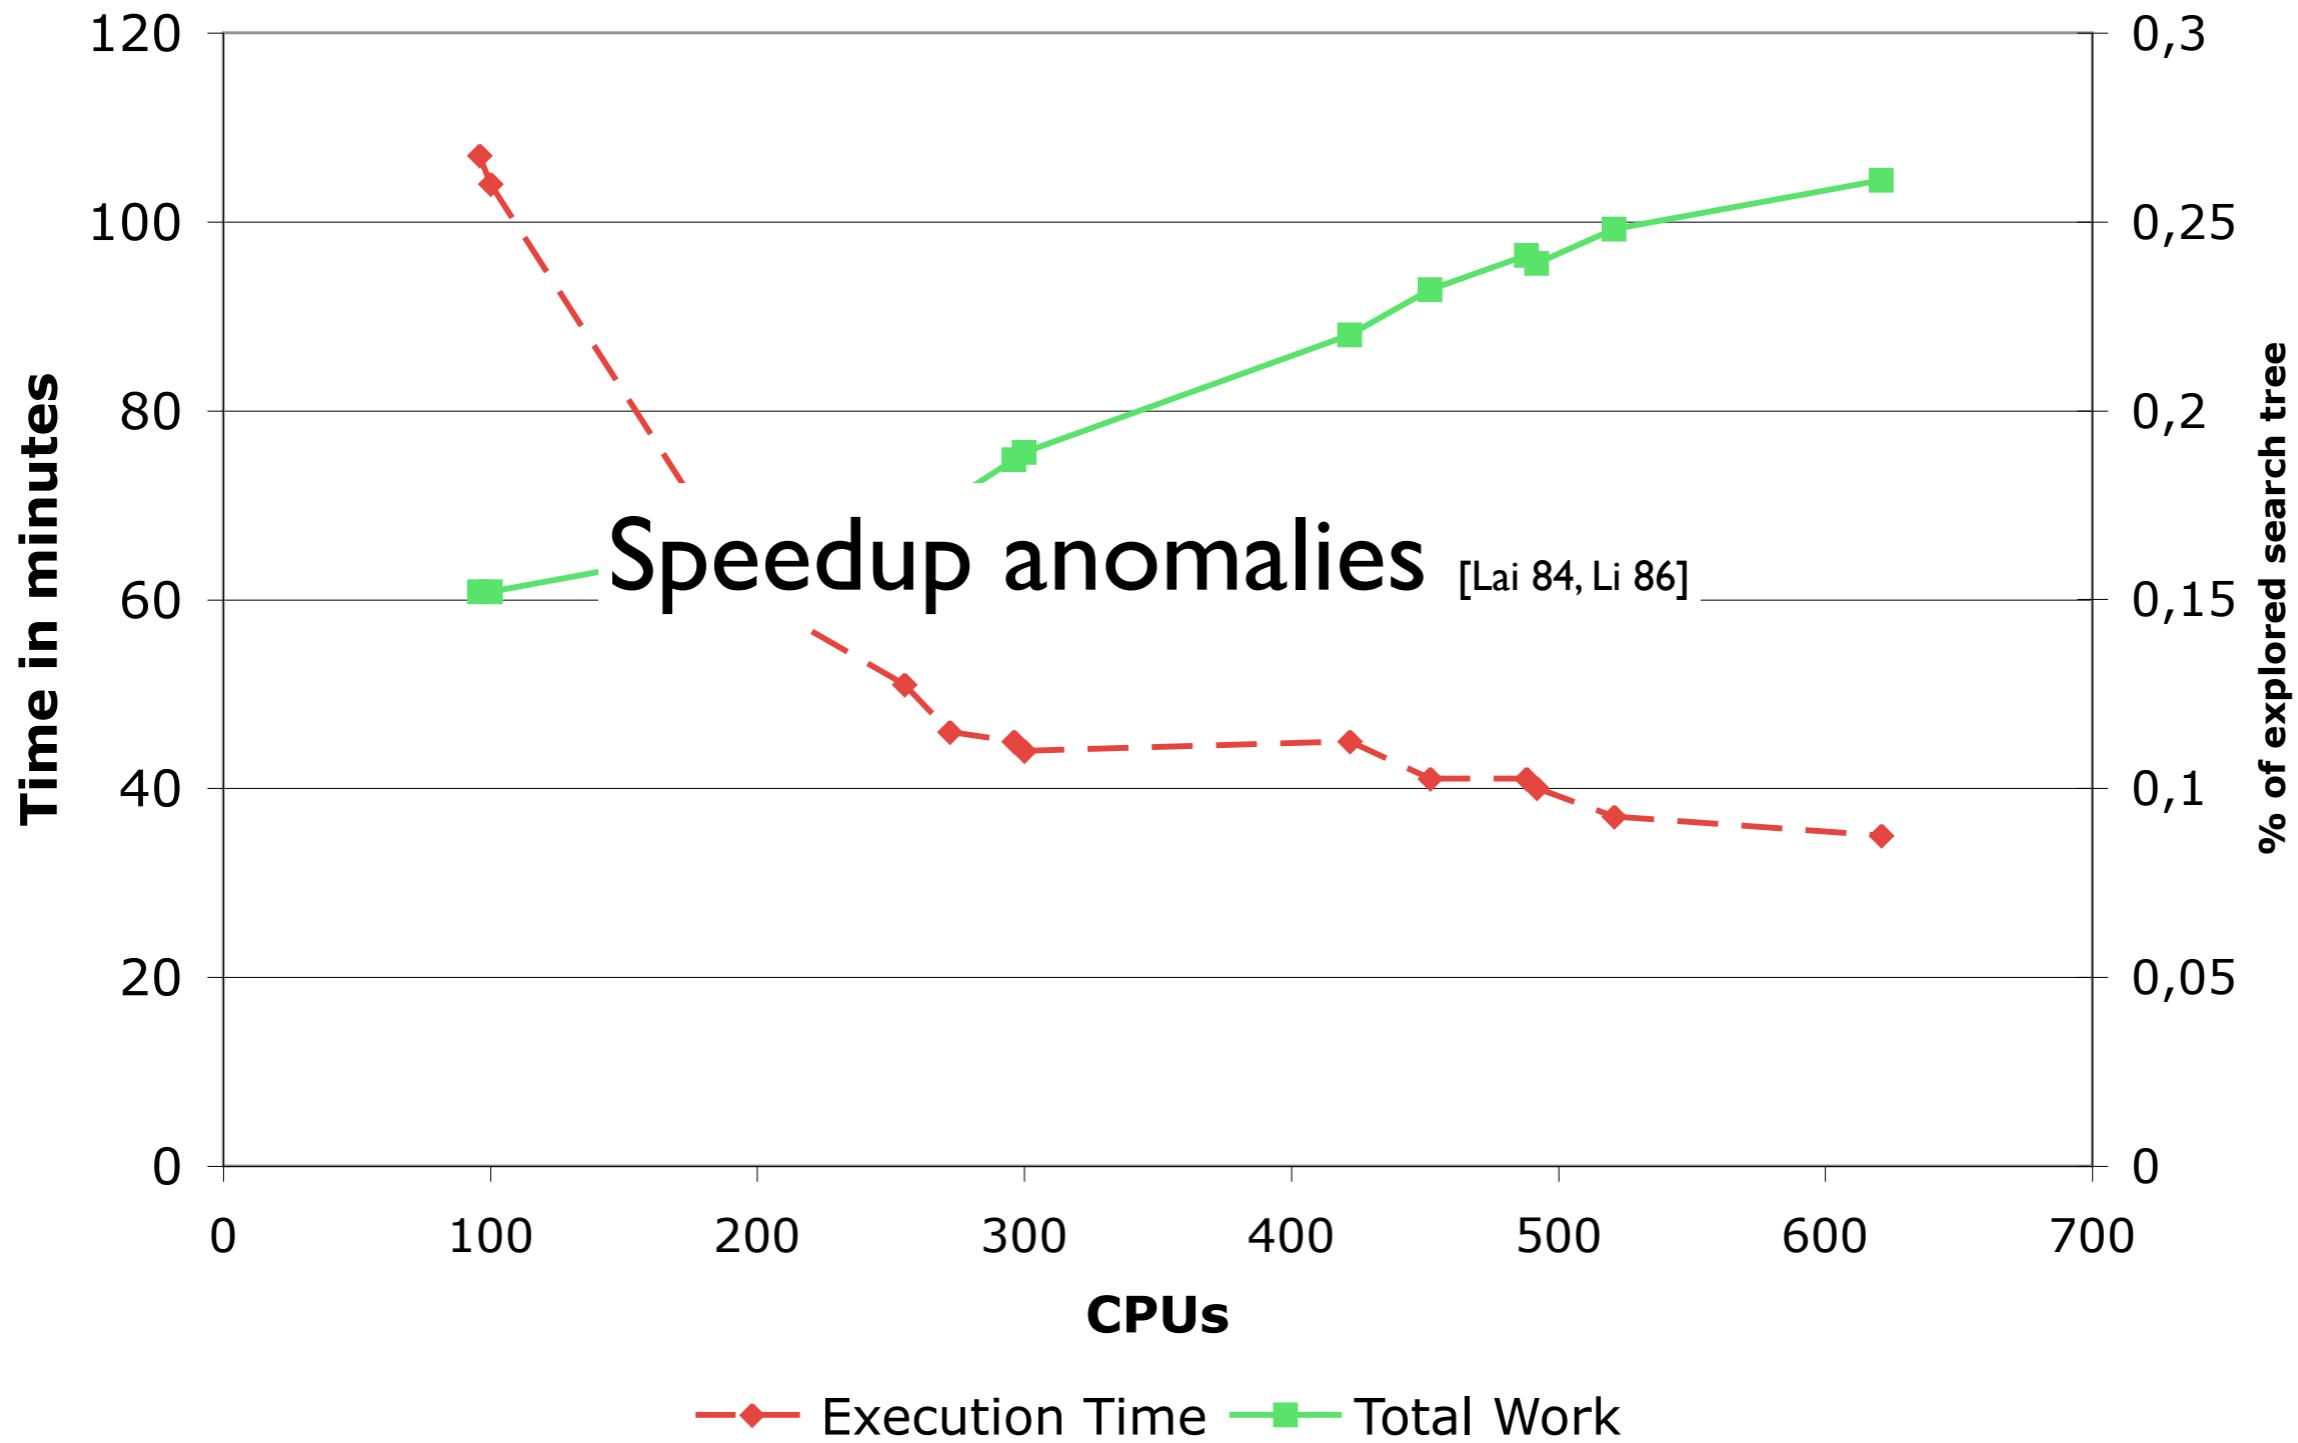
Grid Experimentations

up to 621 CPUs on 5 sites
Flow-Shop: 17 Jobs / 17 Machines



Grid Experimentations

up to 621 CPUs on 5 sites
Flow-Shop: 17 Jobs / 17 Machines

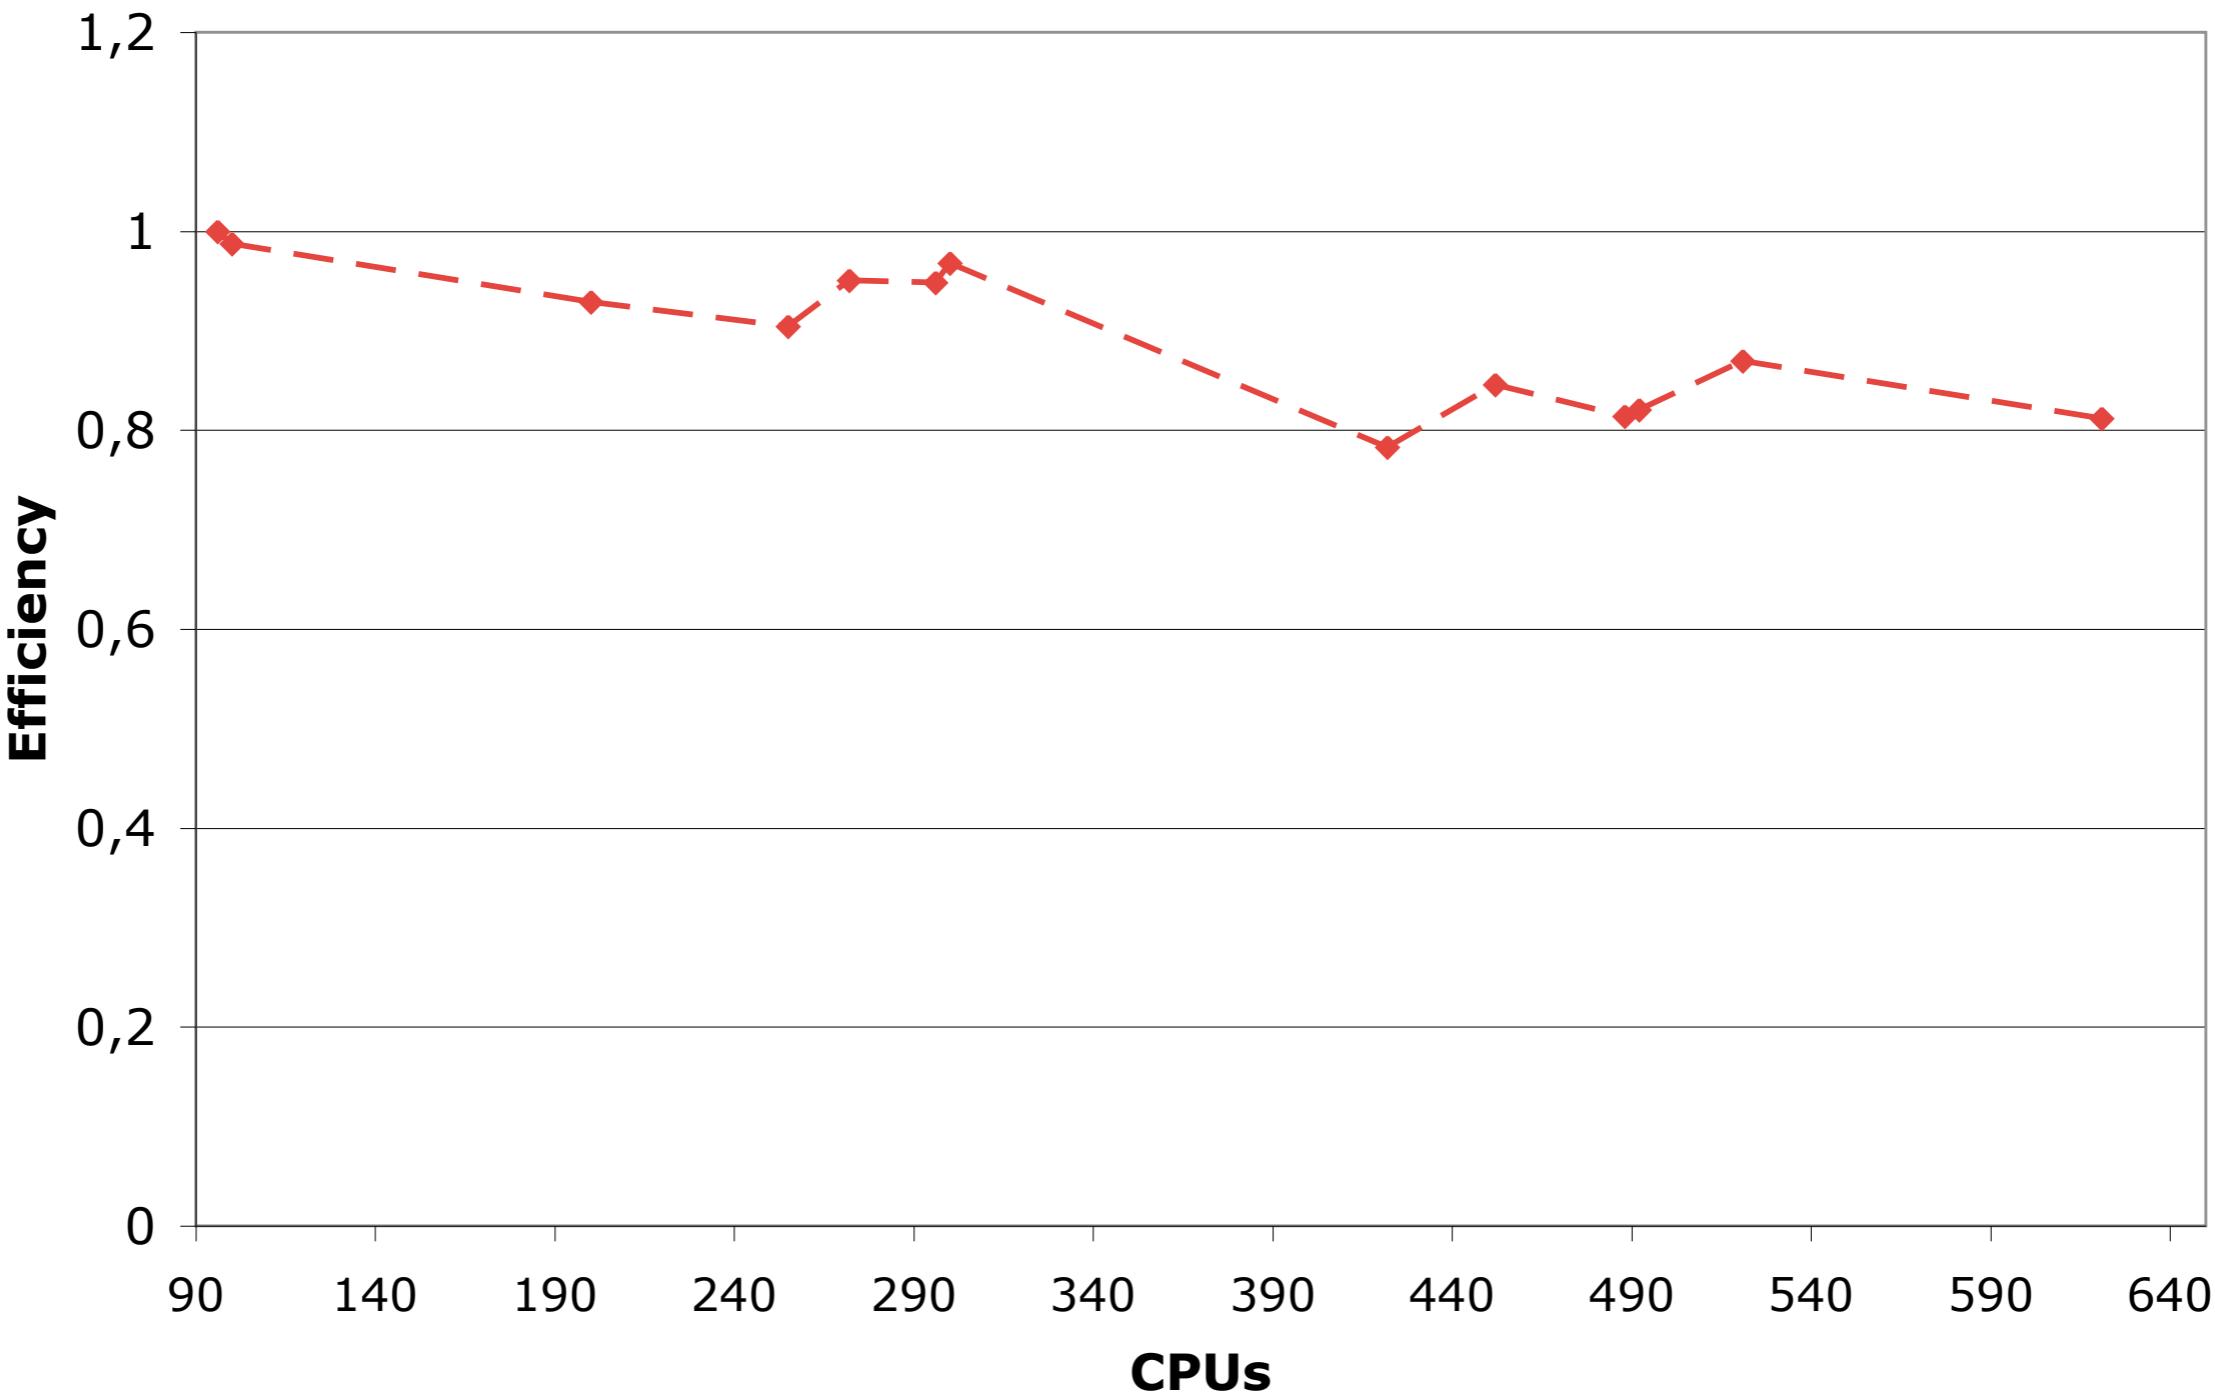


Speedup Anomalies & Efficiency

- Parallel tree-based speedup may sometimes quite spectacular (> or < linear) [Mans 95]
- Speedup Anomalies in BnB [Roucairol 87, Lai 84, Li 86]
 - speedup depends of how the tree is dynamically built
- Efficiency: is a related measure computed as the speedup divided by the number of processors.

Speedup Anomalies & Efficiency

Flow-Shop: 17 Jobs / 17 Machines



Grid'BnB: Results

Grid'BnB: Results

- Experimentally validate our BnB framework for Grids
 - **validity** of organizing communications
 - **scalability** on Grid (up to 621 CPUs on 5 sites)

Grid'BnB: Results

- Experimentally validate our BnB framework for Grids
 - **validity** of organizing communications
 - **scalability** on Grid (up to 621 CPUs on 5 sites)

Grid'BnB: Results

- Experimentally validate our BnB framework for Grids
 - **validity** of organizing communications
 - **scalability** on Grid (up to 621 CPUs on 5 sites)
- **Problems:**
 - **deployment** on Grids is difficult
 - **dynamically** acquiring new resources is difficult
 - **popularity** of Grid'5000
 - **cannot mix** Grid'5000 and under-utilized lab desktops

Grid'BnB: Results

- Experimentally validate our BnB framework for Grids
 - **validity** of organizing communications
 - **scalability** on Grid (up to 621 CPI Is on 5 sites)
 - Grid middleware needs a better supporting of dynamic infrastructure
- **deployment** on Grids is difficult
- **dynamically** acquiring new resources is difficult
- **popularity** of Grid'5000
- **cannot mix** Grid'5000 and under-utilized lab desktops

Agenda

- About Me / Previous Experience
- PhD Main Contributions
- Branch-and-Bound Framework for Grids
- **Desktop Grid with Peer-to-Peer**
- Mixing Desktops & Clusters
- Perspectives
- Possible Support & Conclusion

Peer-to-Peer as Grid Middleware

- Grid Computing and Peer-to-Peer share a common goal:
 - sharing resources [Foster 03, Goux 00]
- Grid related work: [Globus]
 - ✓ providing computational resources
 - installing/deploying Grid middleware is difficult
- P2P related work: [Gnutella, Freenet, DHT]
 - focusing on sharing data & mono-application
 - ✓ dynamic & easy to deploy

Peer-to-Peer as Grid Middleware

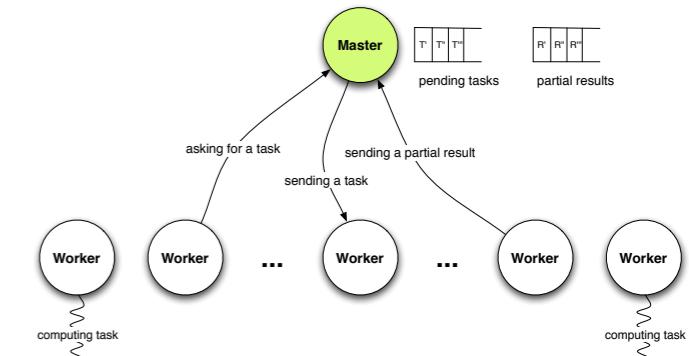
- Grid Computing and Peer-to-Peer share a common goal:
 - sharing resources [Foster 03, Goux 00]
- Grid related work: [Globus]
 - ✓ providing computational resources
 - installing/deploying Grid middleware is difficult
- P2P related work: [Gnutella, Freenet, DHT]
 - focusing on sharing data & mono-application
 - ✓ dynamic & easy to deploy

Objective: provide a P2P infrastructure for Grids and sharing computational resources

Which P2P Architecture?

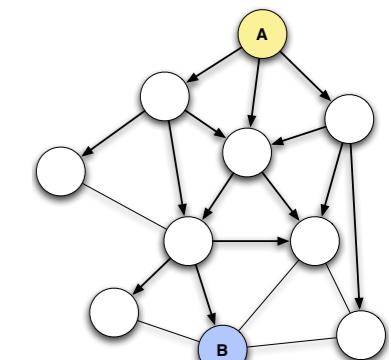
Master-Worker (SETI@home)

- centralized
- targets only desktops
- good for embarrassingly parallel



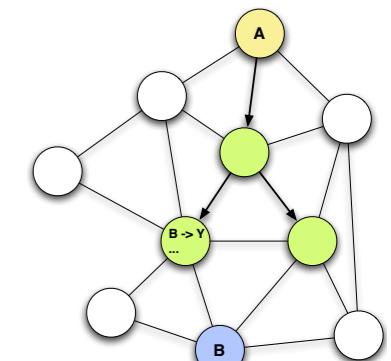
Pure/Unstructured Peer-to-Peer (Gnutella)

- flooding problems
- ✓ supports high-churn (good for Desktop Grids)
- ✓ supports many kind of application (data, computational)



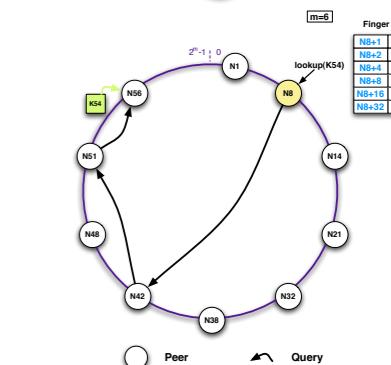
Hybrid Peer-to-Peer (JXTA)

- uses central servers
- ✓ limits the flooding
- has to manage churn



Structured Peer-to-Peer (Chord)

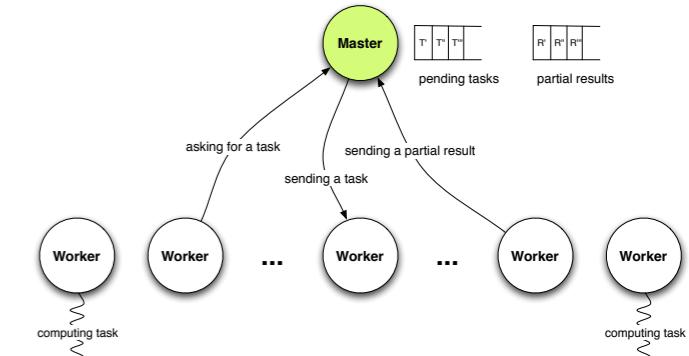
- high cost for managing churn
- efficient for data sharing (Distributed Hash Table)



Which P2P Architecture?

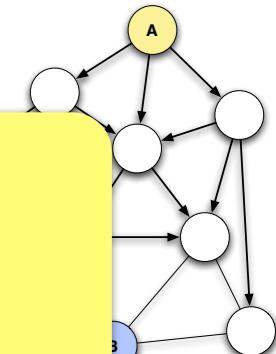
Master-Worker (SETI@home)

- centralized
- targets only desktops
- good for embarrassingly parallel



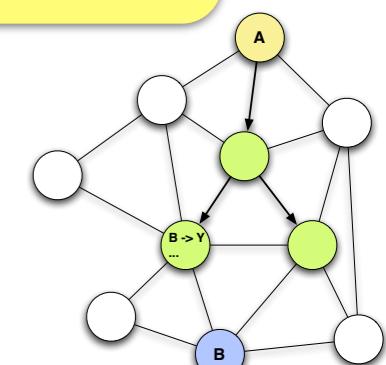
Pure/Unstructured Peer-to-Peer (Gnutella)

- ✓ Pure Peer-to-Peer is the most adapted
- ✓ Needs to avoid the flooding problem



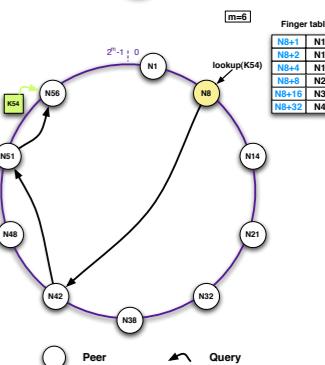
Hybrid Peer-to-Peer (JXTA)

- uses central servers
- ✓ limits the flooding
- has to manage churn

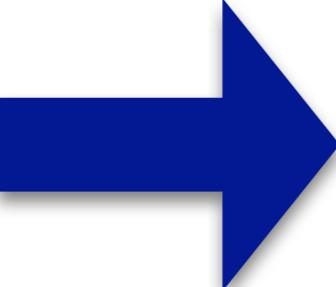


Structured Peer-to-Peer (Chord)

- high cost for managing churn
- efficient for data sharing (Distributed Hash Table)



Contribution & Positioning

Branch-and-Bound API 

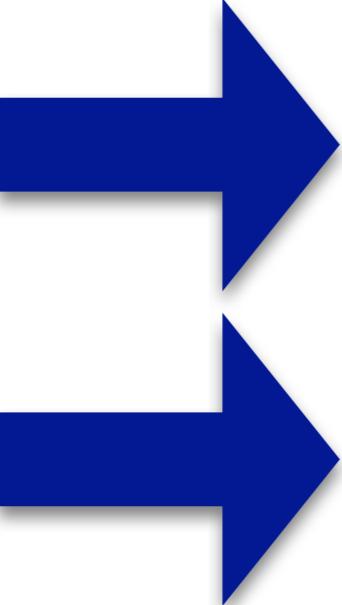
Grid Applications & Portals

Grid Programming

**Grid Middleware
Infrastructure**

Grid Fabric

Contribution & Positioning

Branch-and-Bound API 
P2P Infrastructure

Grid Applications & Portals

Grid Programming

**Grid Middleware
Infrastructure**

Grid Fabric

The Peer-to-Peer Infrastructure [CMST06]

The Peer-to-Peer Infrastructure [CMST06]

- Pure Peer-to-Peer overlay network
- Using it as Grid middleware infrastructure
- The proposed solution to avoid the flooding problem:
 - 3 node-request protocols:
 - 1 node: Random walk algorithm
 - n nodes: Breadth-First-Search (BFS) algorithm with acknowledgement
 - max nodes: BFS without acknowledgement
 - Best-effort

INRIA Sophia P2P Desktop Grid

- Need to validate the infrastructure with desktops
- 260 desktops at INRIA Sophia lab
- No disturbing normal users:
 - running in low priority
 - working schedules:
 - 24/24 ≈ 50 machines (INRIA-2424)
 - night/weekend ≈ 260 machines (INRIA-ALL)

INRIA Sophia P2P Desktop Grid

- Need to validate the infrastructure with desktops
- 260 desktops at INRIA Sophia lab
- No disturbing normal users:
 - running in low priority
 - working schedules:
 - 24/24 ≈ 50 machines (INRIA-2424)
 - night/weekend ≈ 260 machines (INRIA-ALL)

Deployed our P2P infrastructure as permanent
Desktop Grid at INRIA Sophia

INRIA Sophia DOD Desktop Grid

Ne

260

Nc

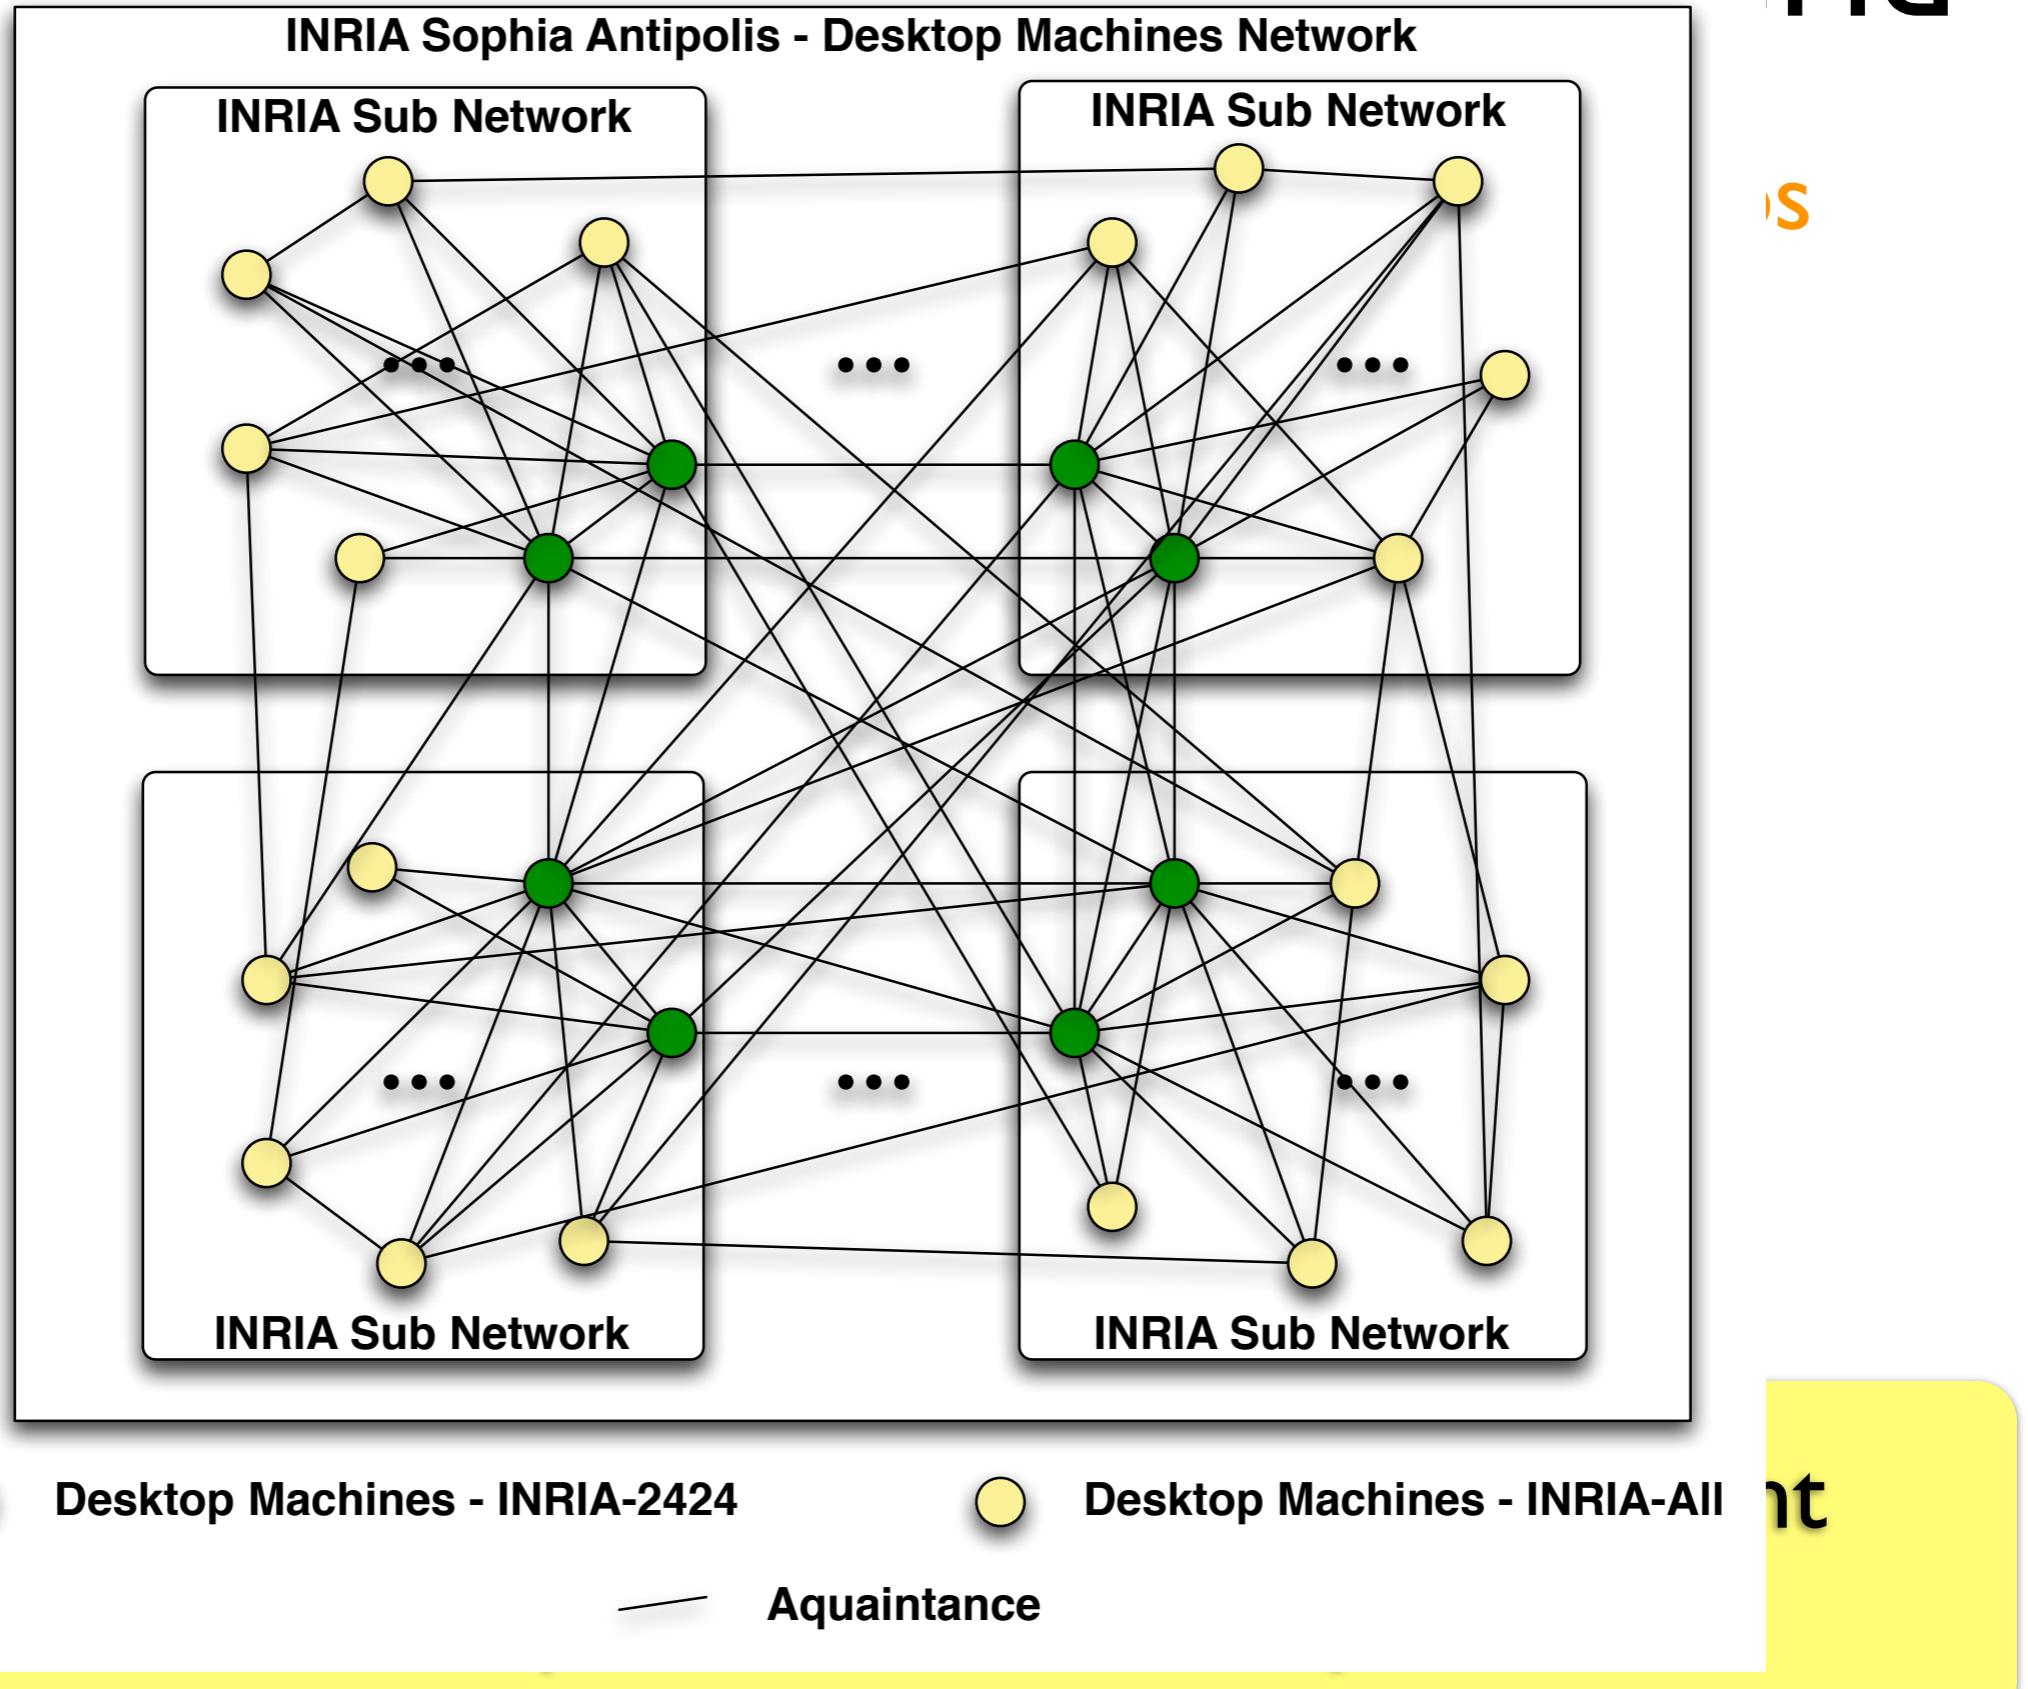
r

v

l

b

S



Long Running Experiments with the P2P Desktop Grid

- Context: ETSI Grid Plugtests contest \Rightarrow n-queens
- n-queens:
- embarrassingly parallel / CPU intensive / master-worker

Long Running Experiments with the P2P Desktop Grid

- Context: ETSI Grid Plugtests contest \Rightarrow **n-queens**
- **n-queens:**
- embarrassingly parallel / CPU intensive / master-worker

How many solution for 25 queens ?

n-queens Experiment Results

Total # of Tasks	12,125,199
Task Computation	$\approx 138''$
Computation Time	$\approx 185 \text{ days}$
Cumulated Time	$\approx 53 \text{ years}$
# of Desktop Machines	260
Total of Solution Found	2,207,893,435,808,352 $\approx 2 \text{ quadrillions}$

● World Record [Sloane Integers Sequence A000170]

● What we learn from this experiments:

● validate the **workability** of the infrastructure

● validate the **robustness** of the infrastructure

● hard to forecast machine's performances

Agenda

- About Me / Previous Experience
- PhD Main Contributions
- Branch-and-Bound Framework for Grids
- Desktop Grid with Peer-to-Peer
- **Mixing Desktops & Clusters**
- Perspectives
- Possible Support & Conclusion

Mixing Desktops & Clusters [PARCO07]

Grid'BnB

Parallel BnB Framework for Grid

P2P Infrastructure

as Grid Middleware

Mixing Desktops & Clusters [PARCO07]

Grid'BnB
Parallel BnB Framework for Grid

+

P2P Infrastructure
as Grid Middleware

**API for solving COPs &
Dynamic Grid Infrastructure**

Mixing Desktops & Clusters [PARCO07]

Grid'BnB

Parallel BnB Framework for Grid

+

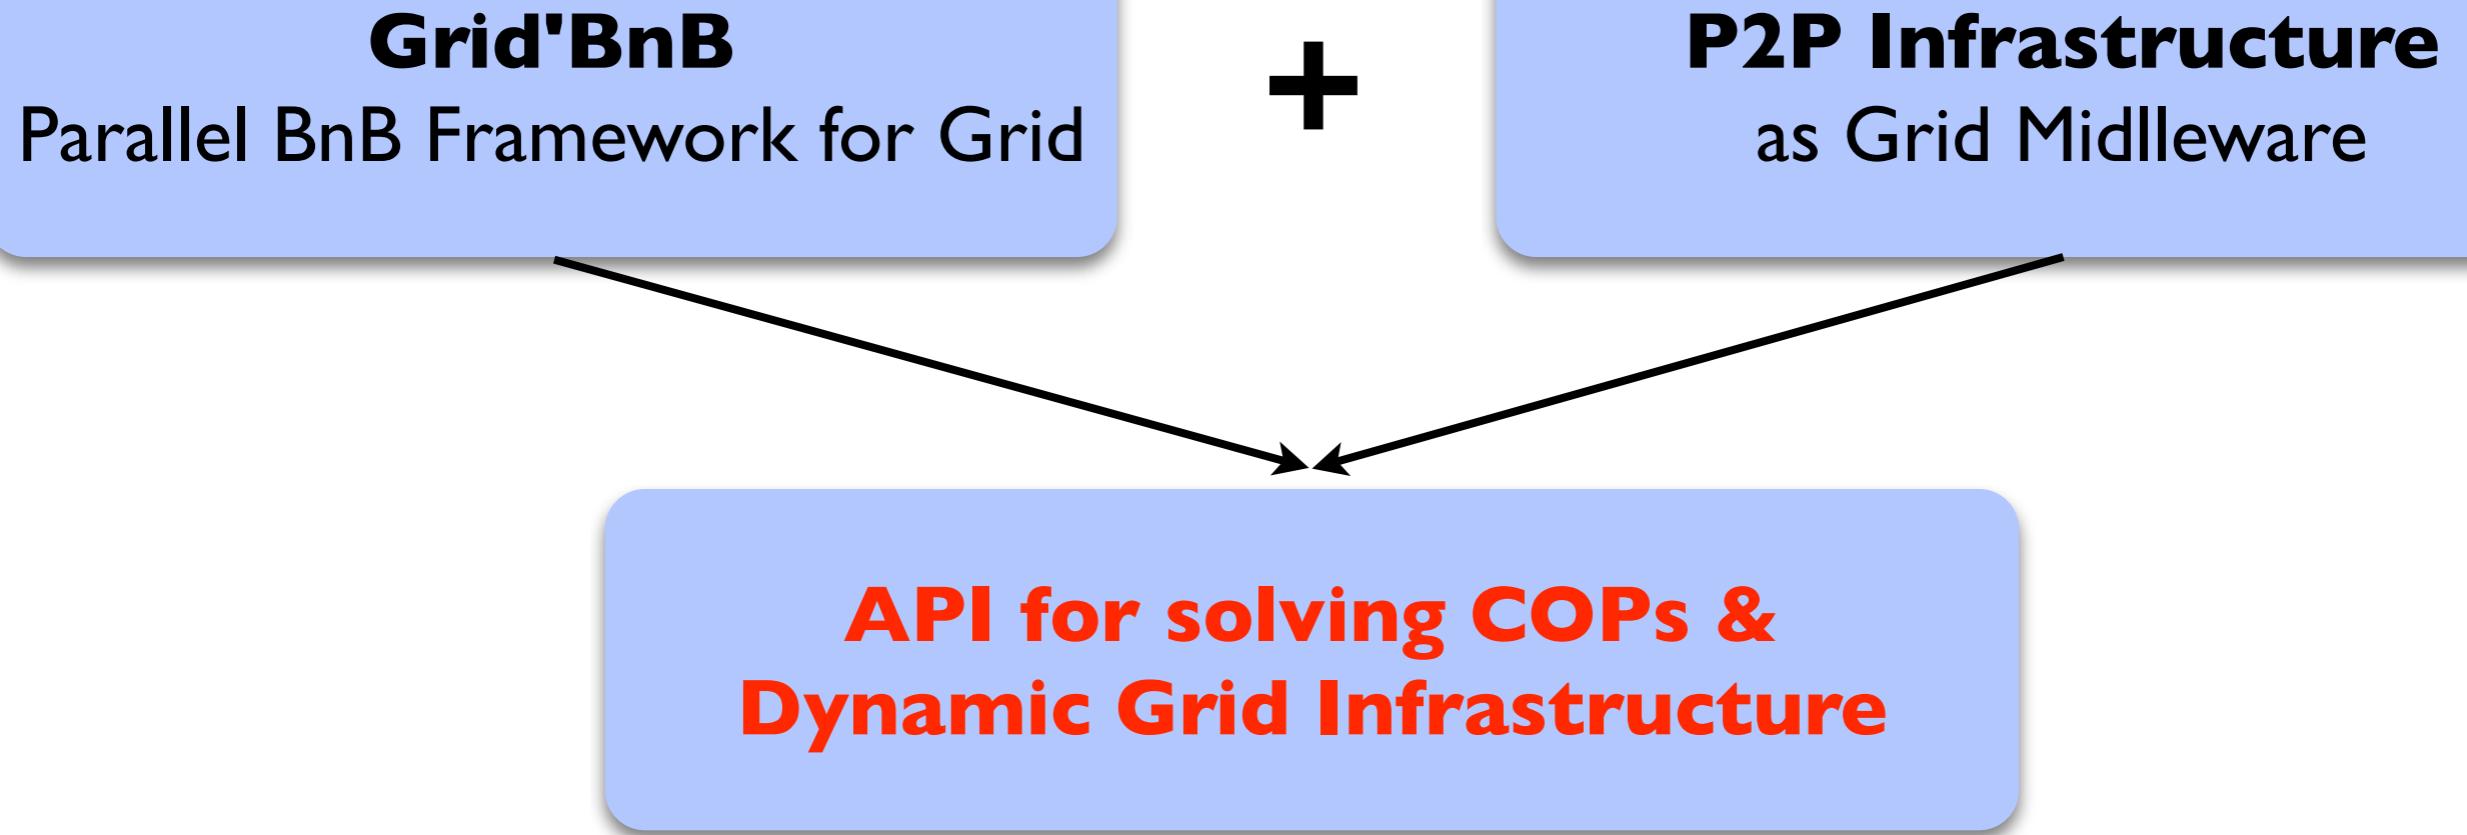
P2P Infrastructure

as Grid Middleware

**API for solving COPs &
Dynamic Grid Infrastructure**

Validate by experiments on a Grid of Desktops & Clusters

Mixing Desktops & Clusters [PARCO07]



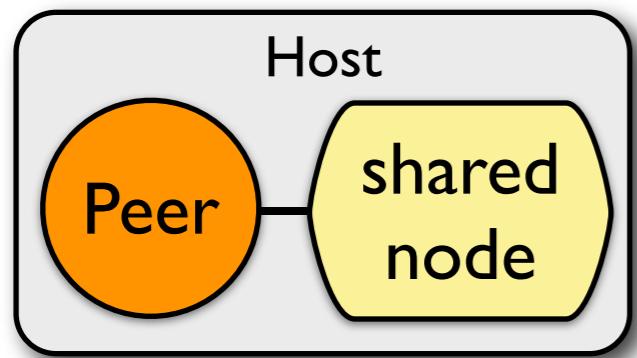
Validate by experiments on a Grid of Desktops & Clusters

New Problems:

firewalls → forwarder

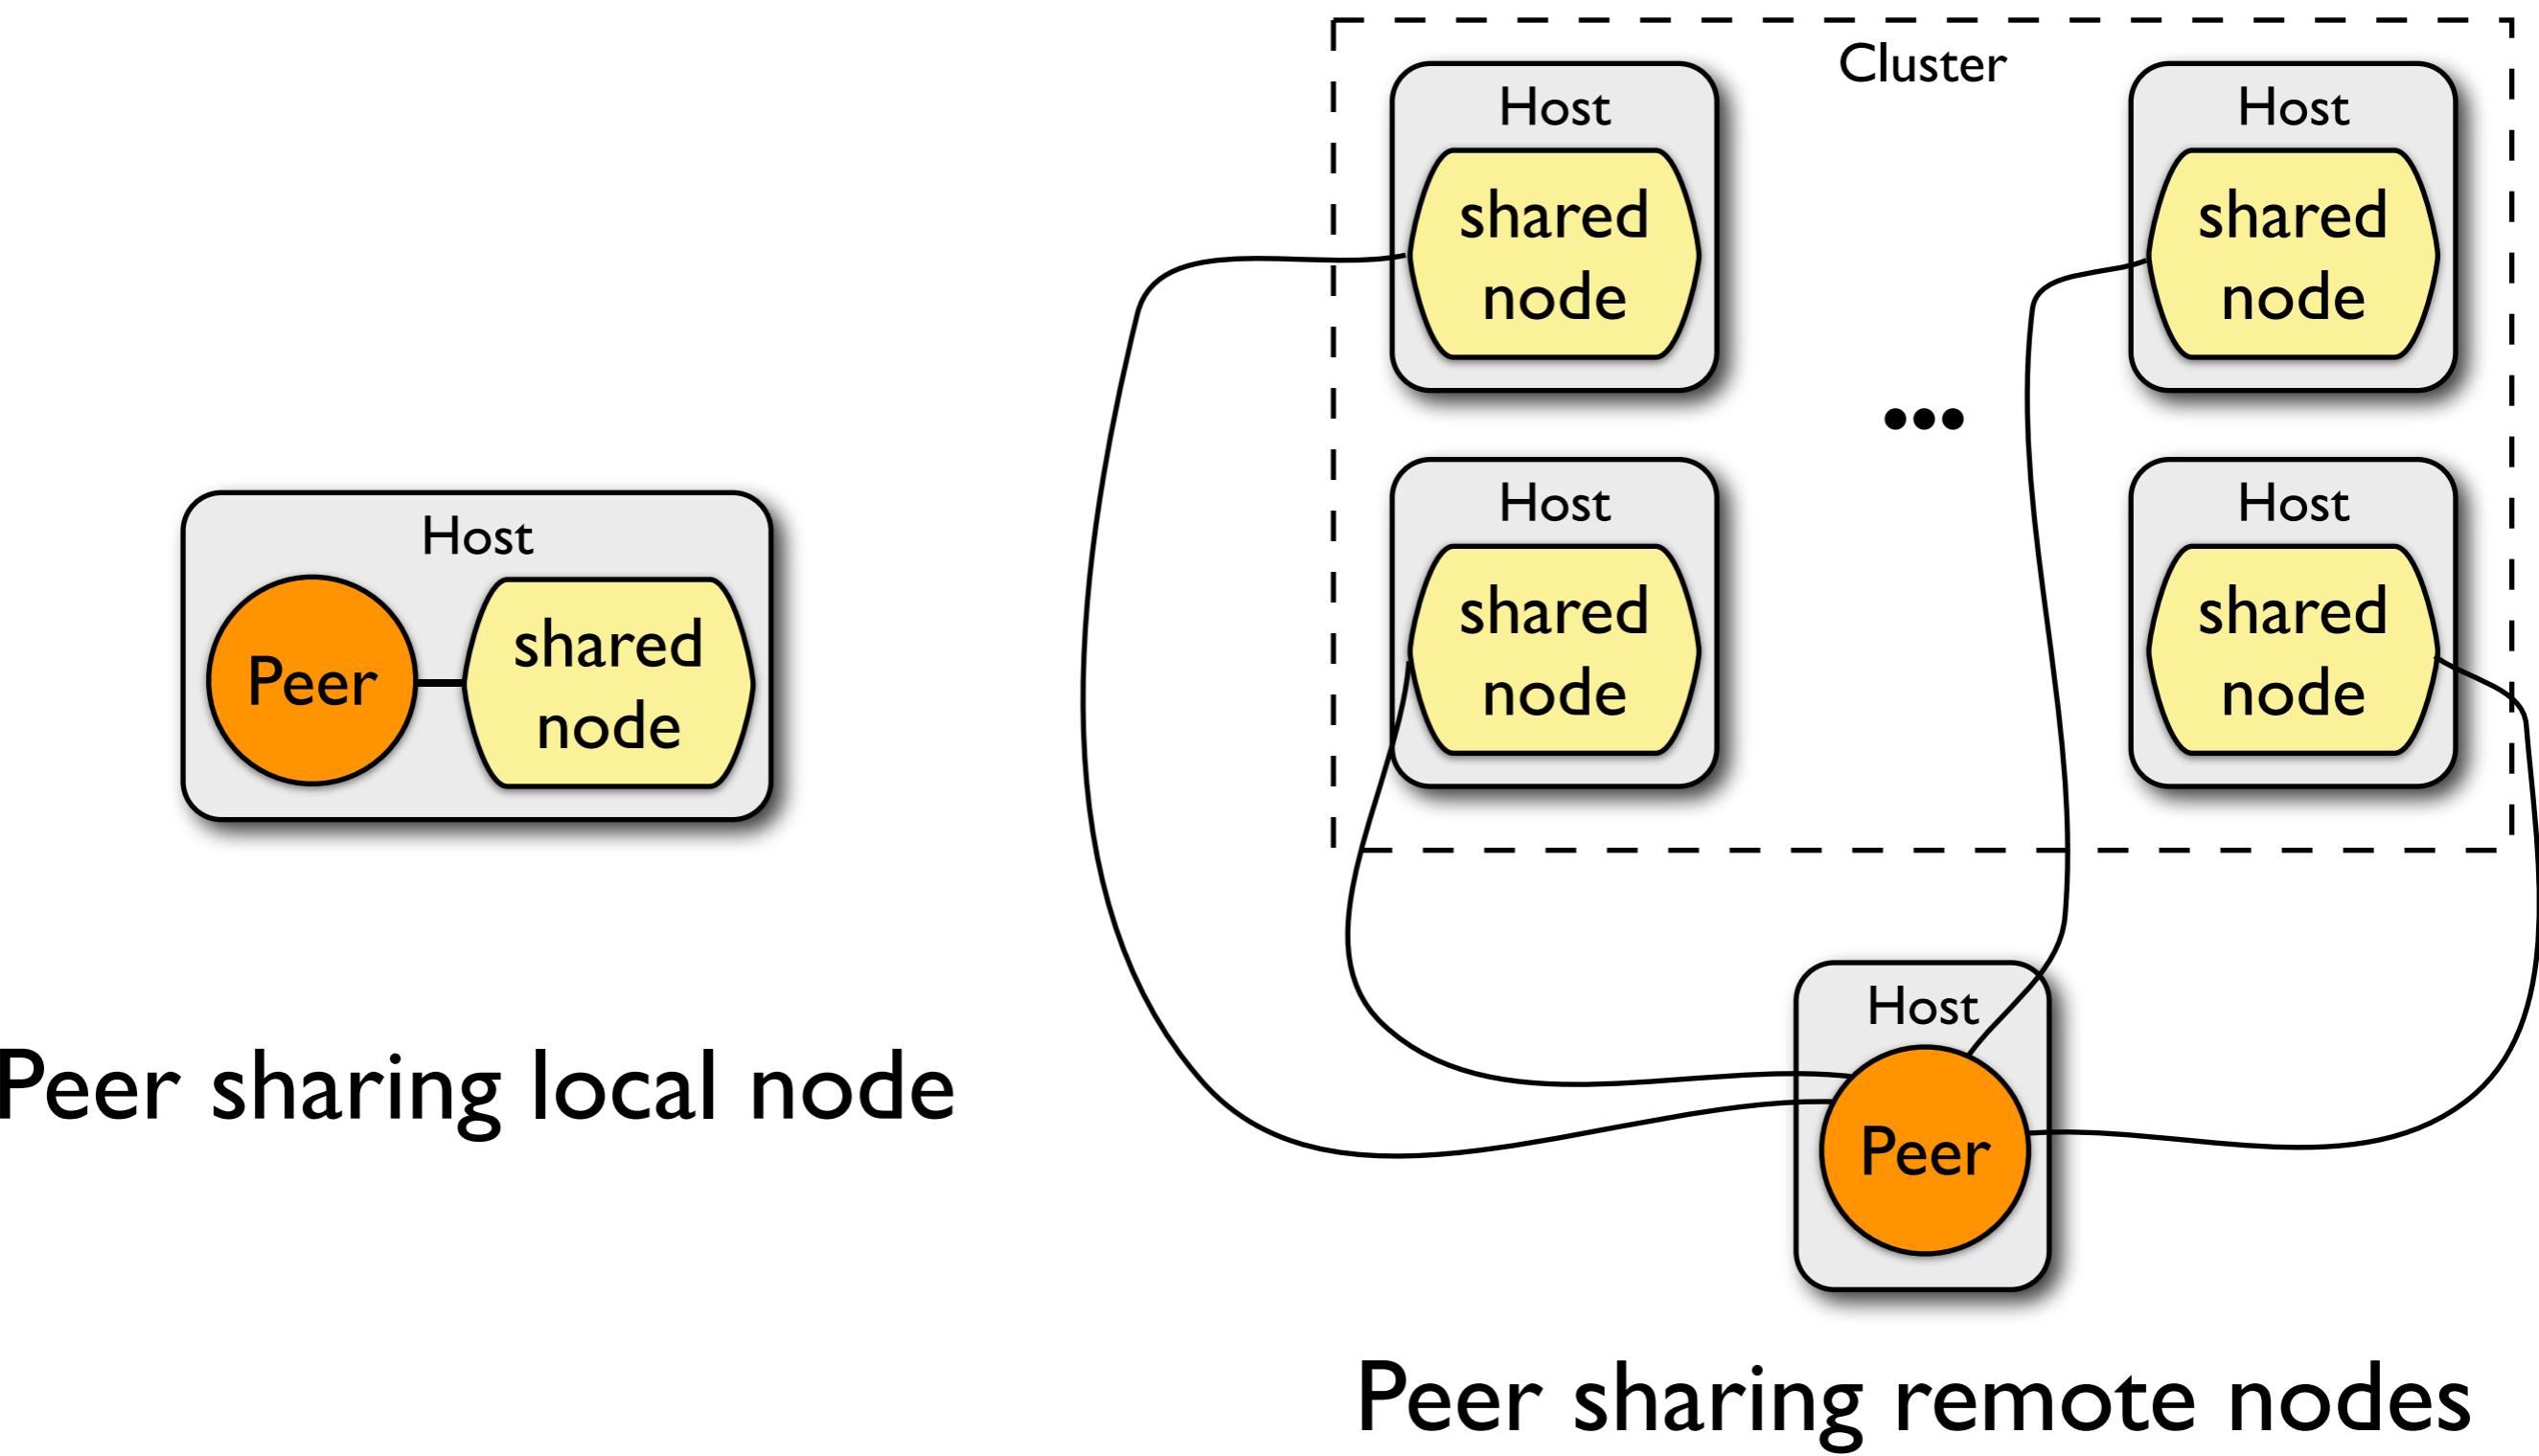
sharing clusters with the P2P infrastructure

Sharing Cluster's Nodes with P2P



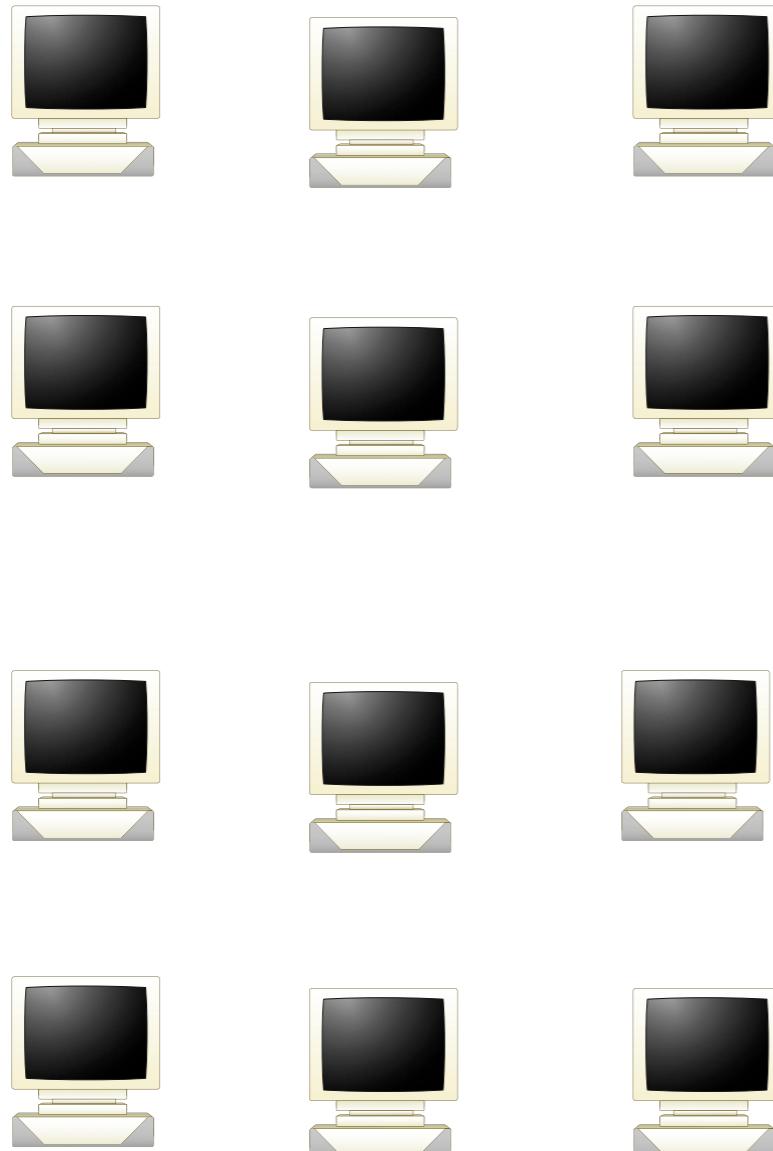
Peer sharing local node

Sharing Cluster's Nodes with P2P

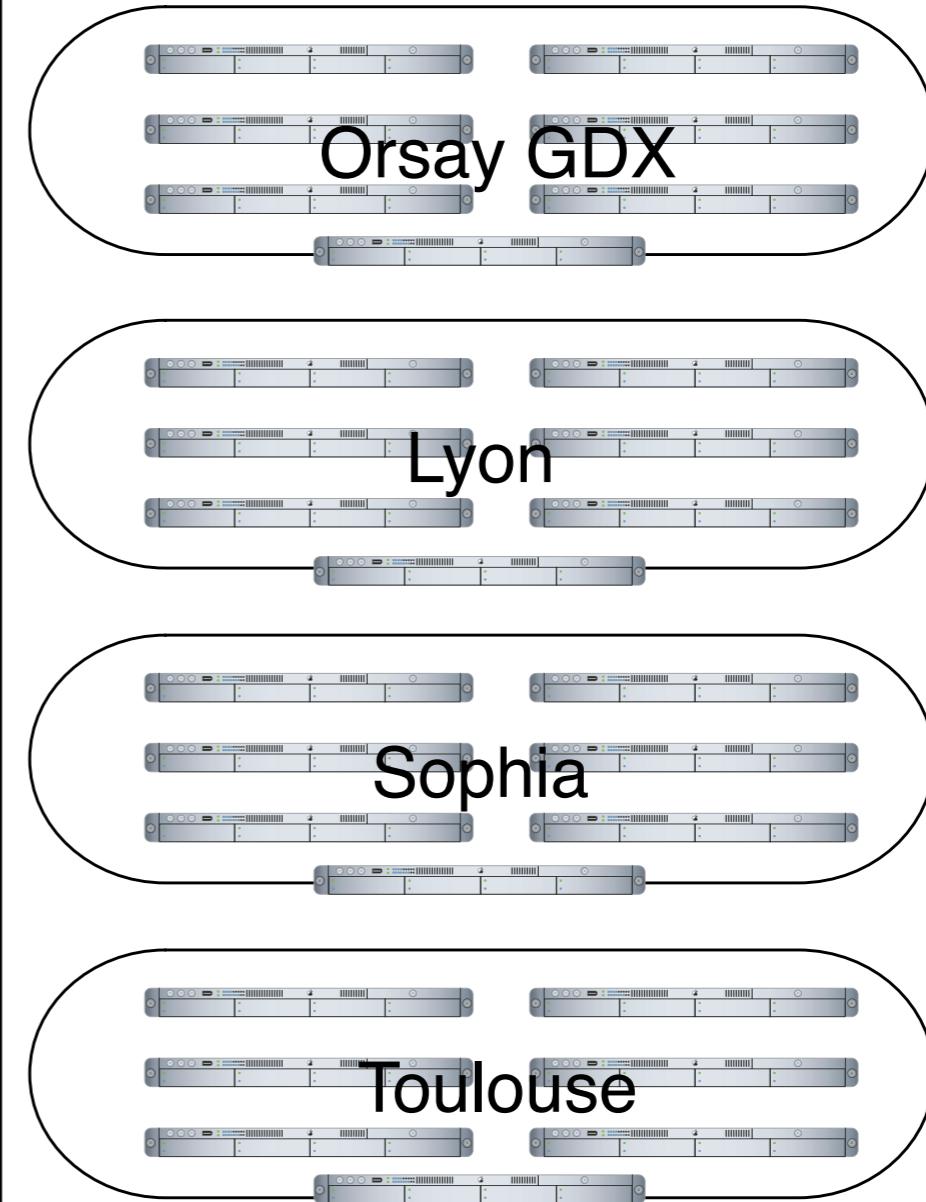


Testbed

INRIA Sophia - Desktops

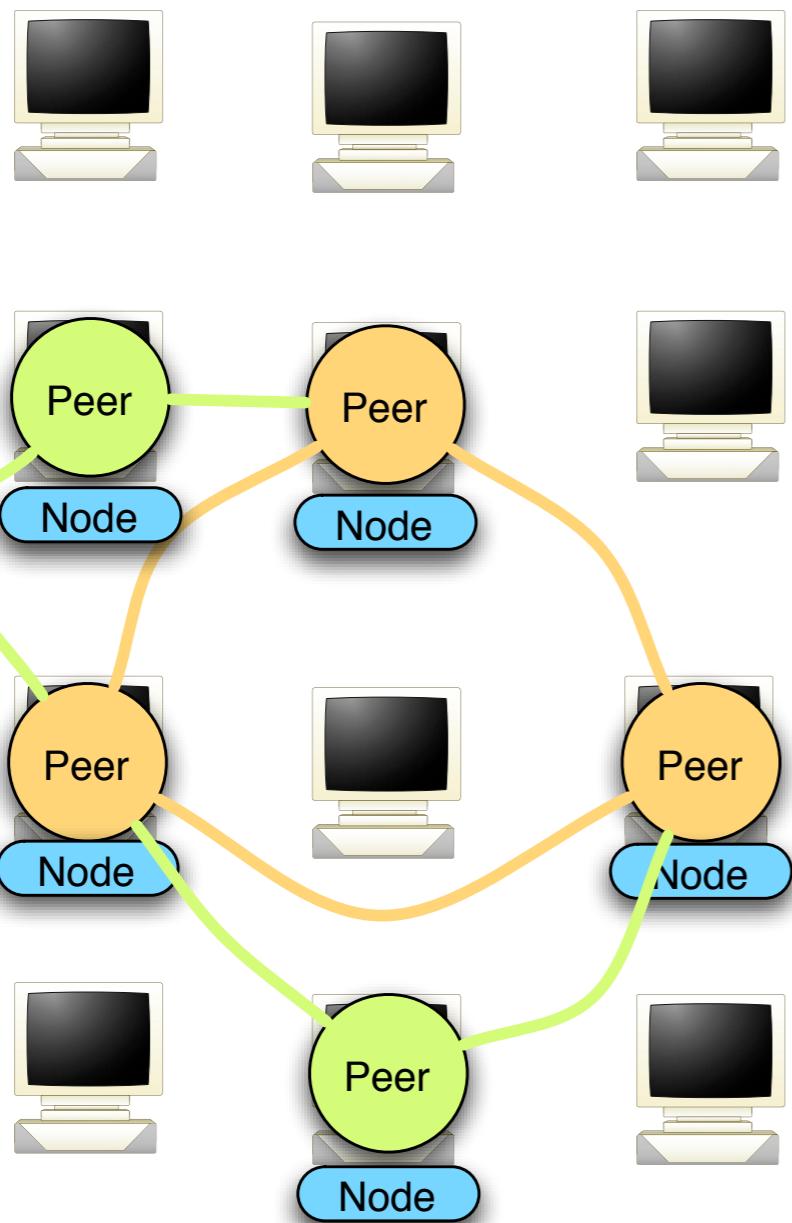


G5K Platform

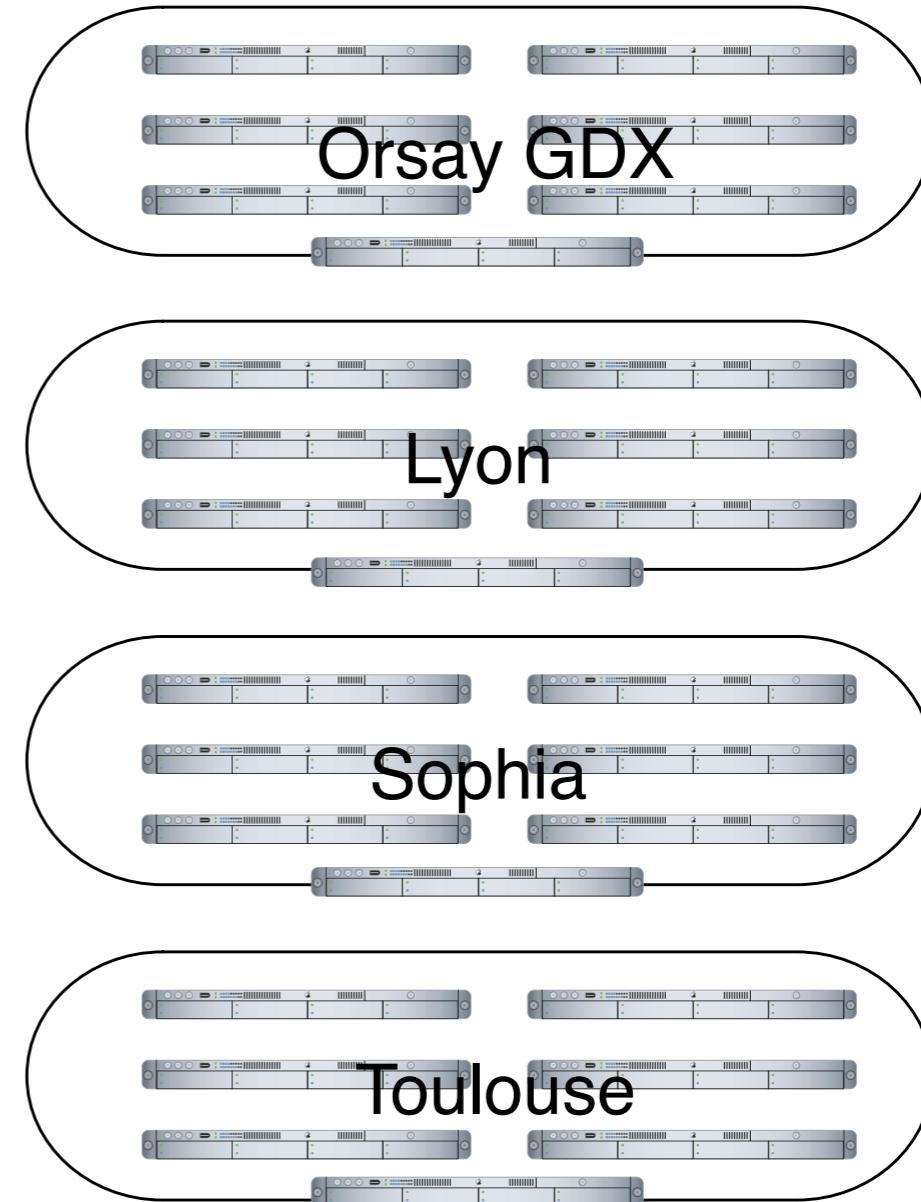


Testbed

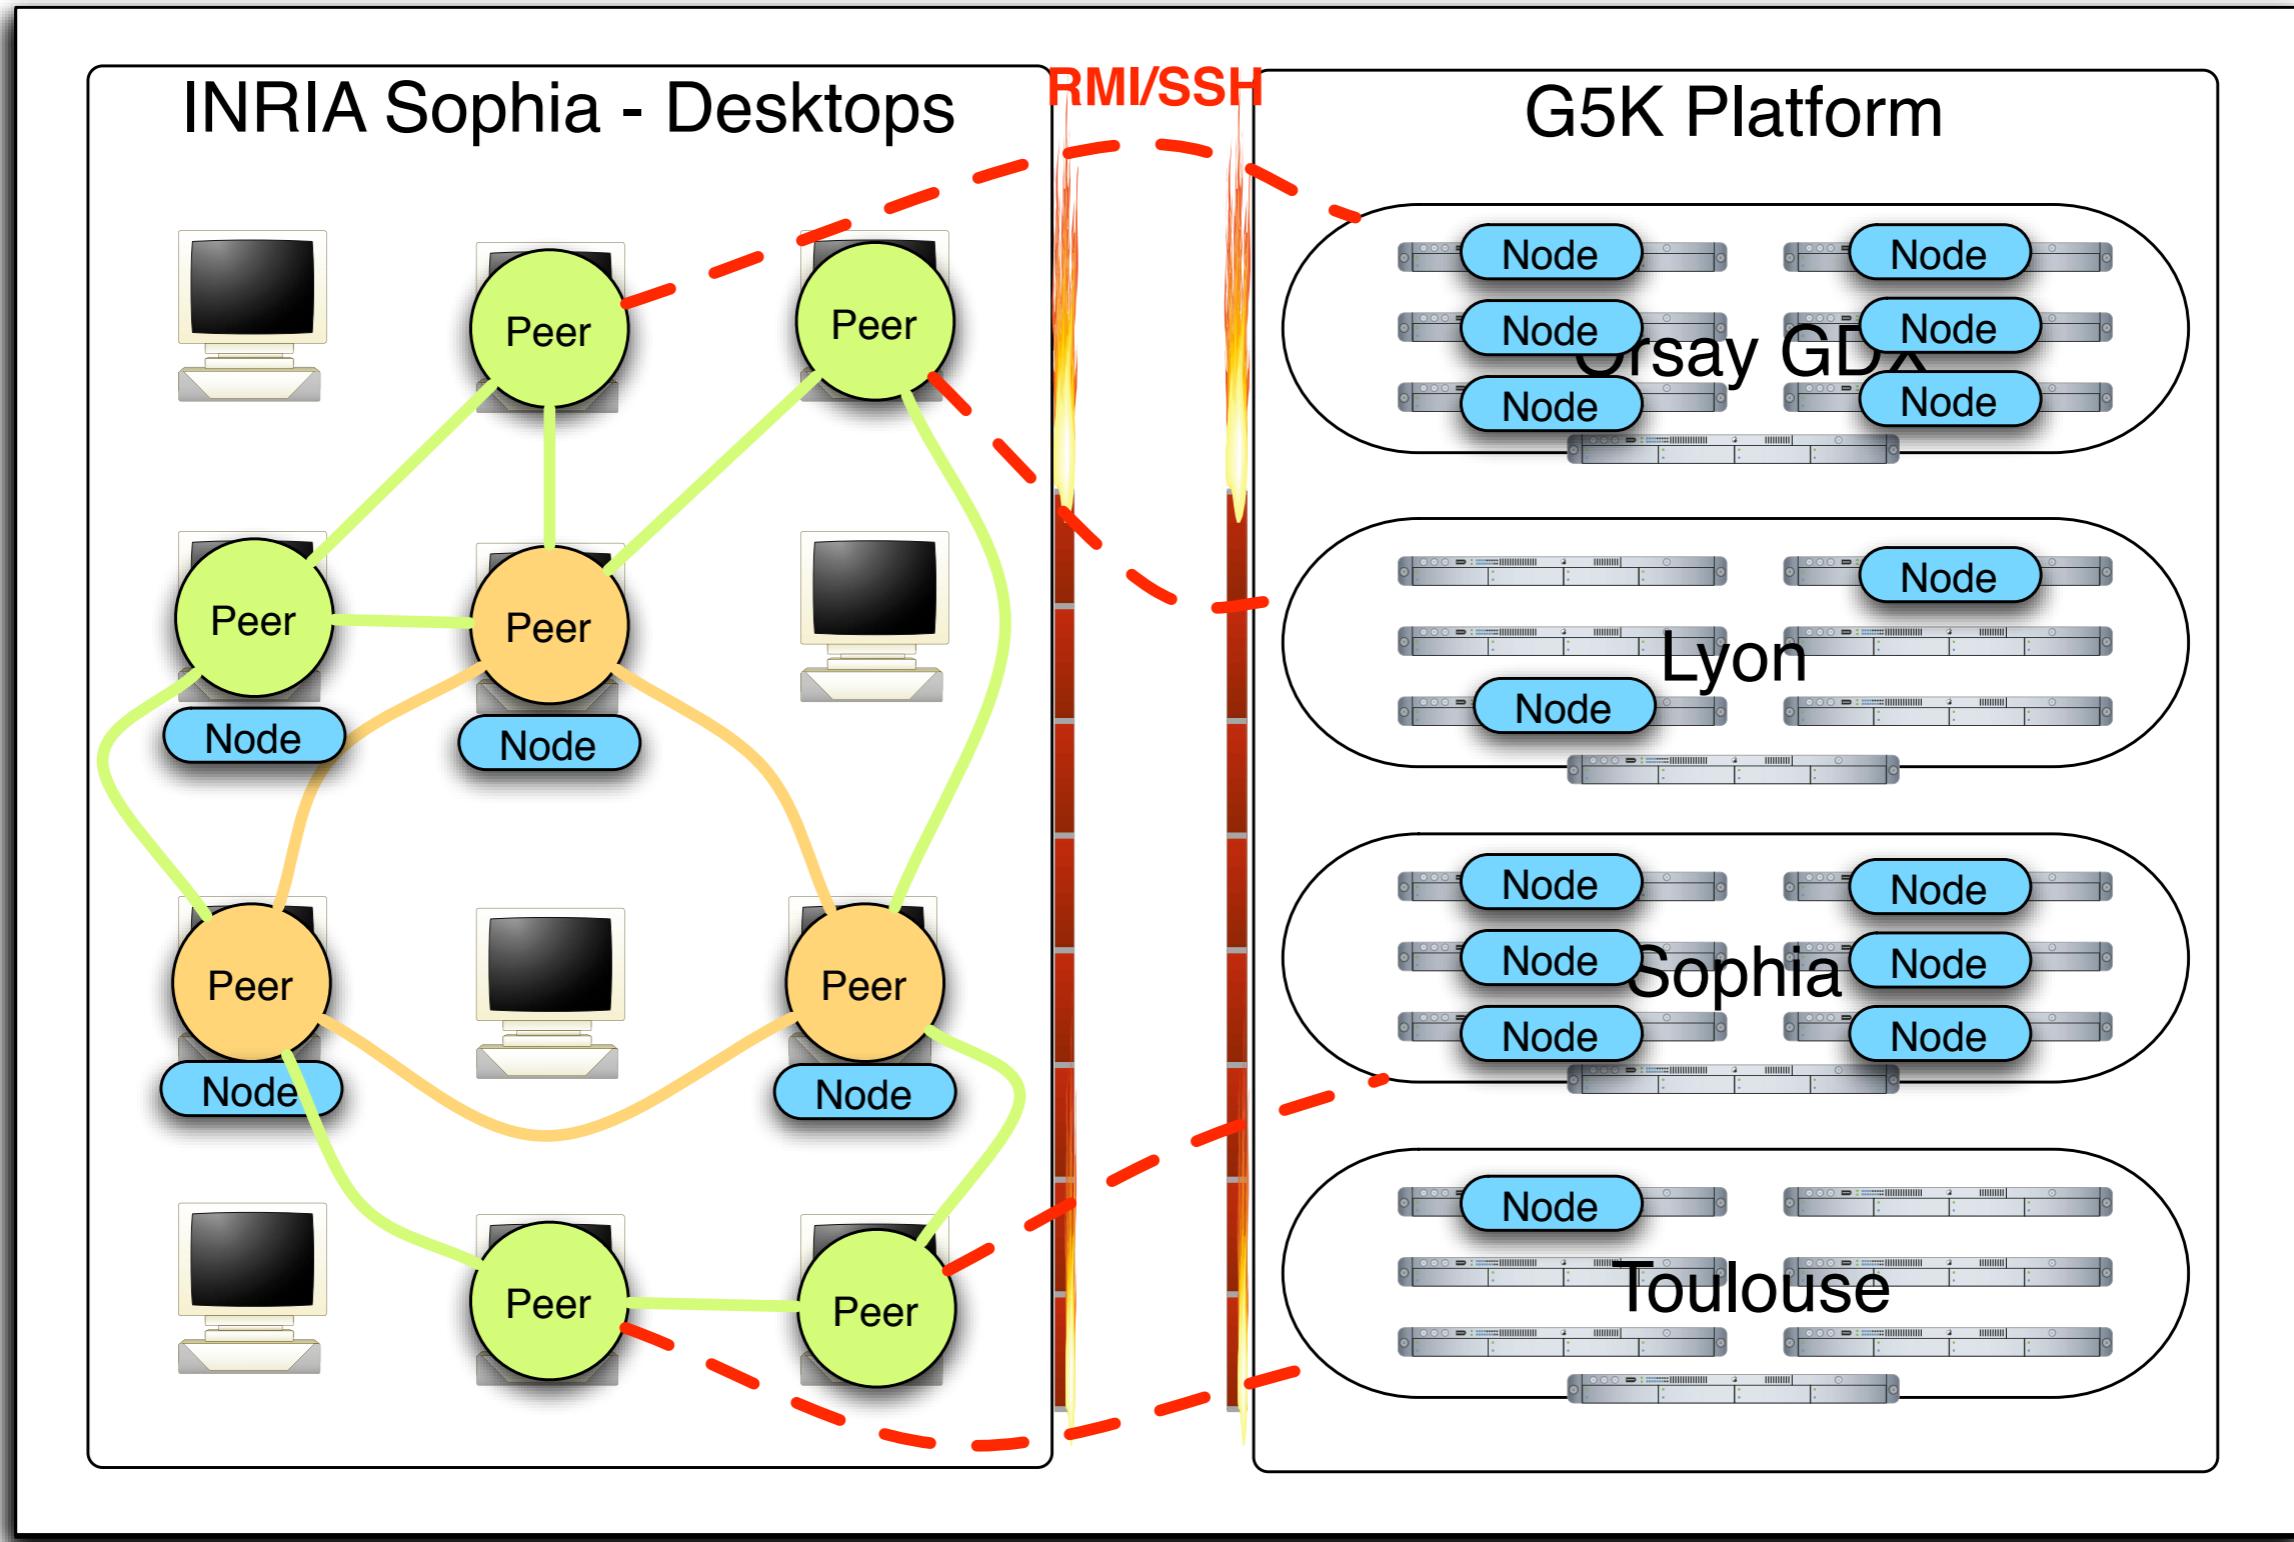
INRIA Sophia - Desktops



G5K Platform



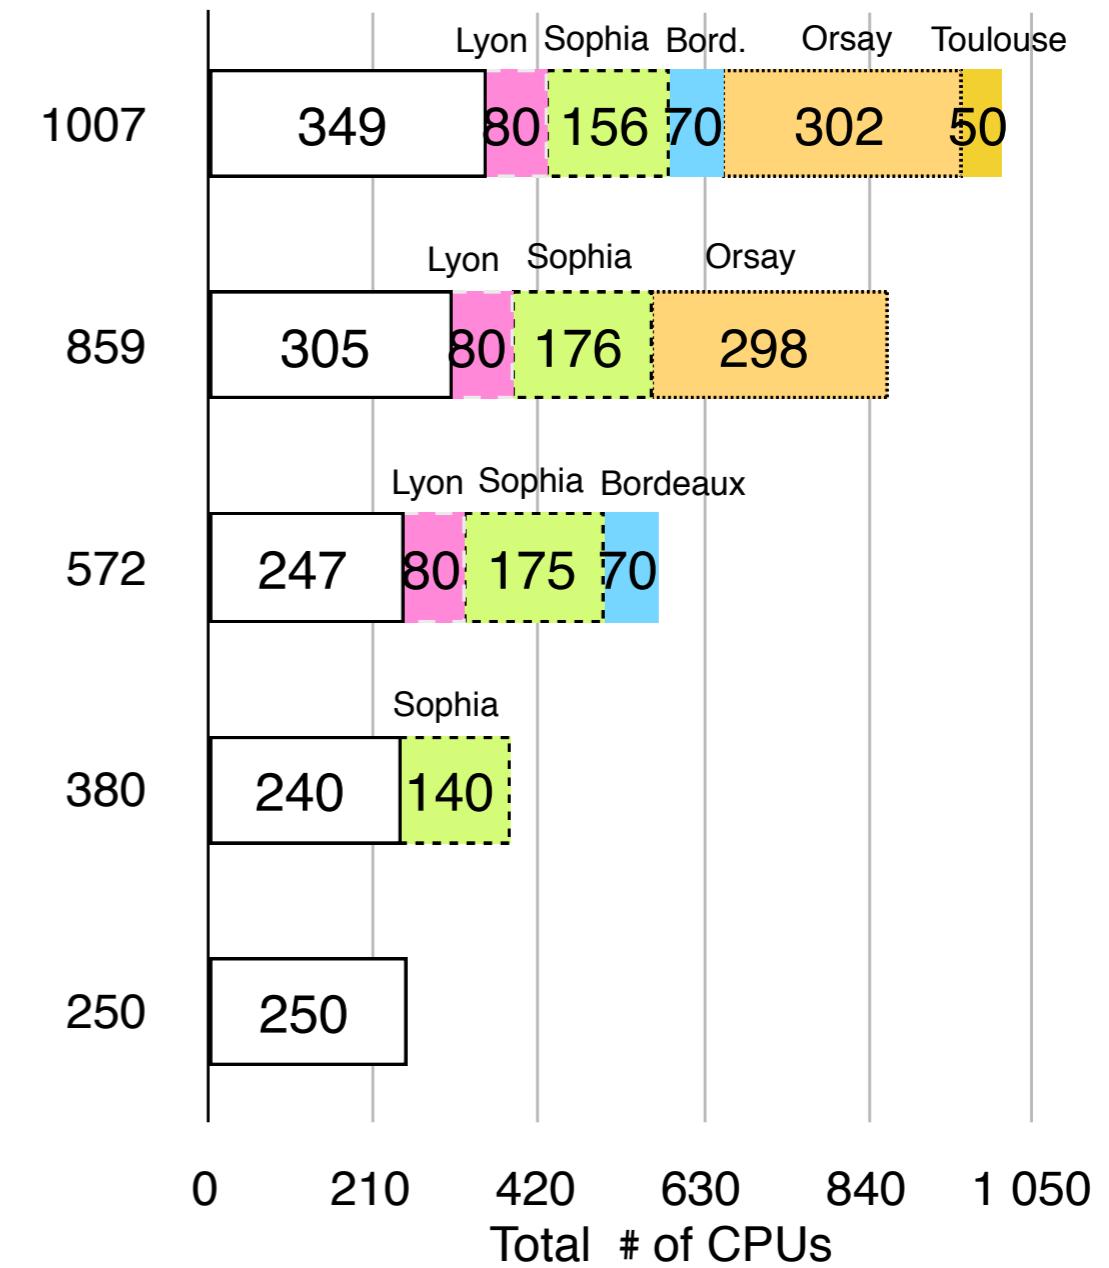
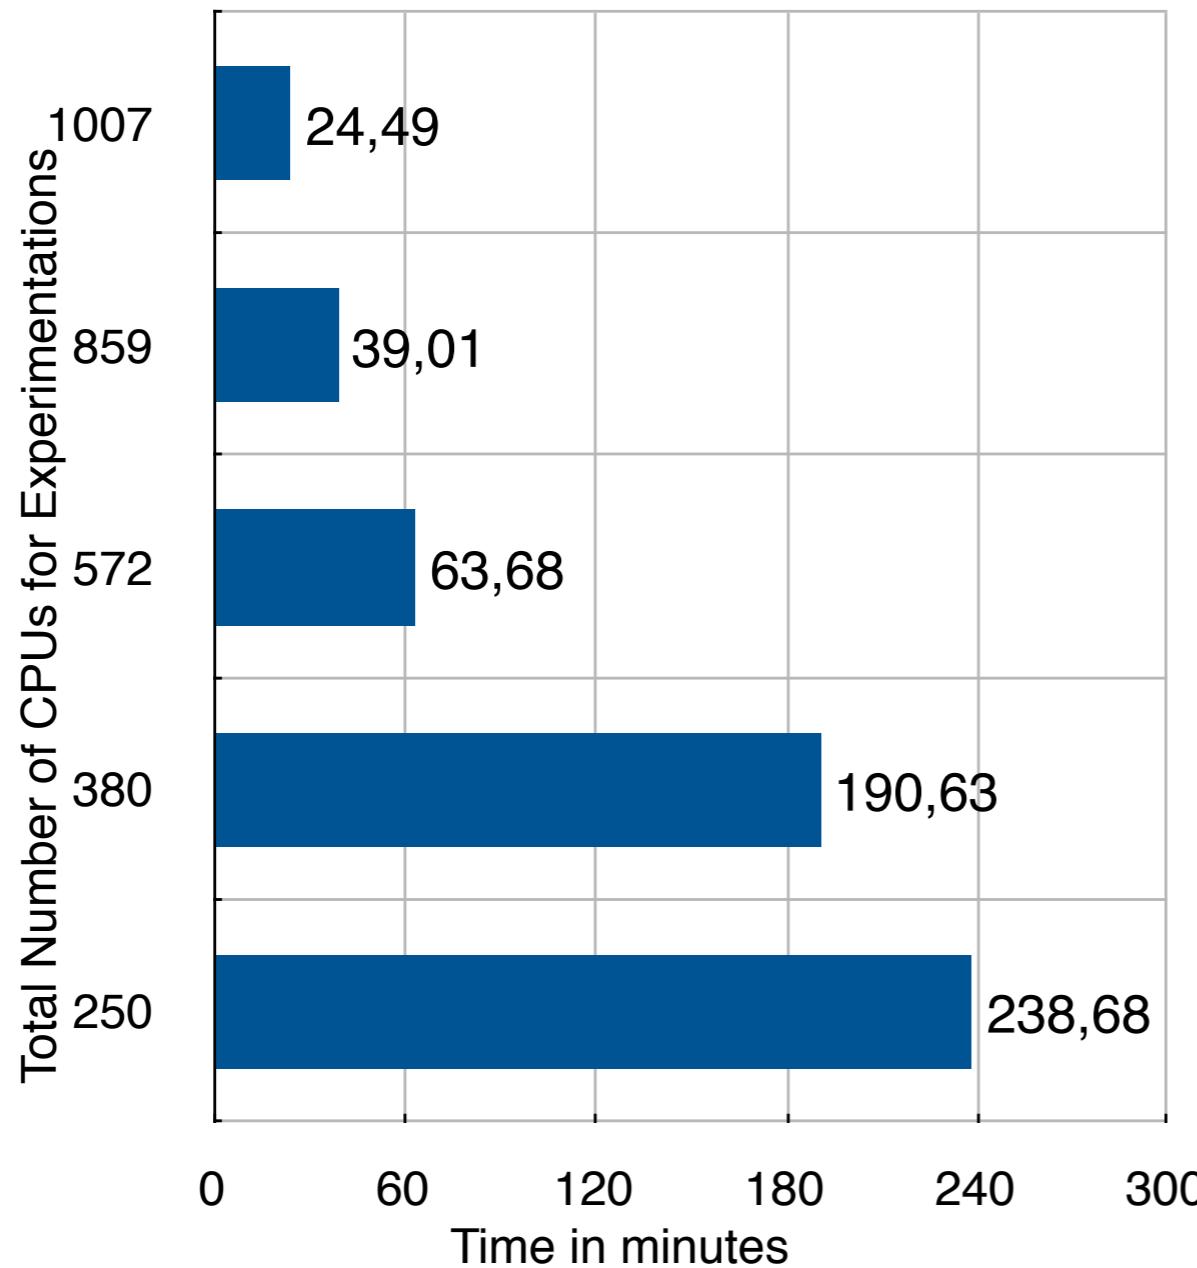
Testbed



Large-Scale Experiments

- **Goal:** validate the infrastructure by experiments
- With **n-queens:**
 - no communication between workers
 - test the **workability** of the infrastructure
- With **flow-shop:**
 - communications between workers
 - test **solving** COPs

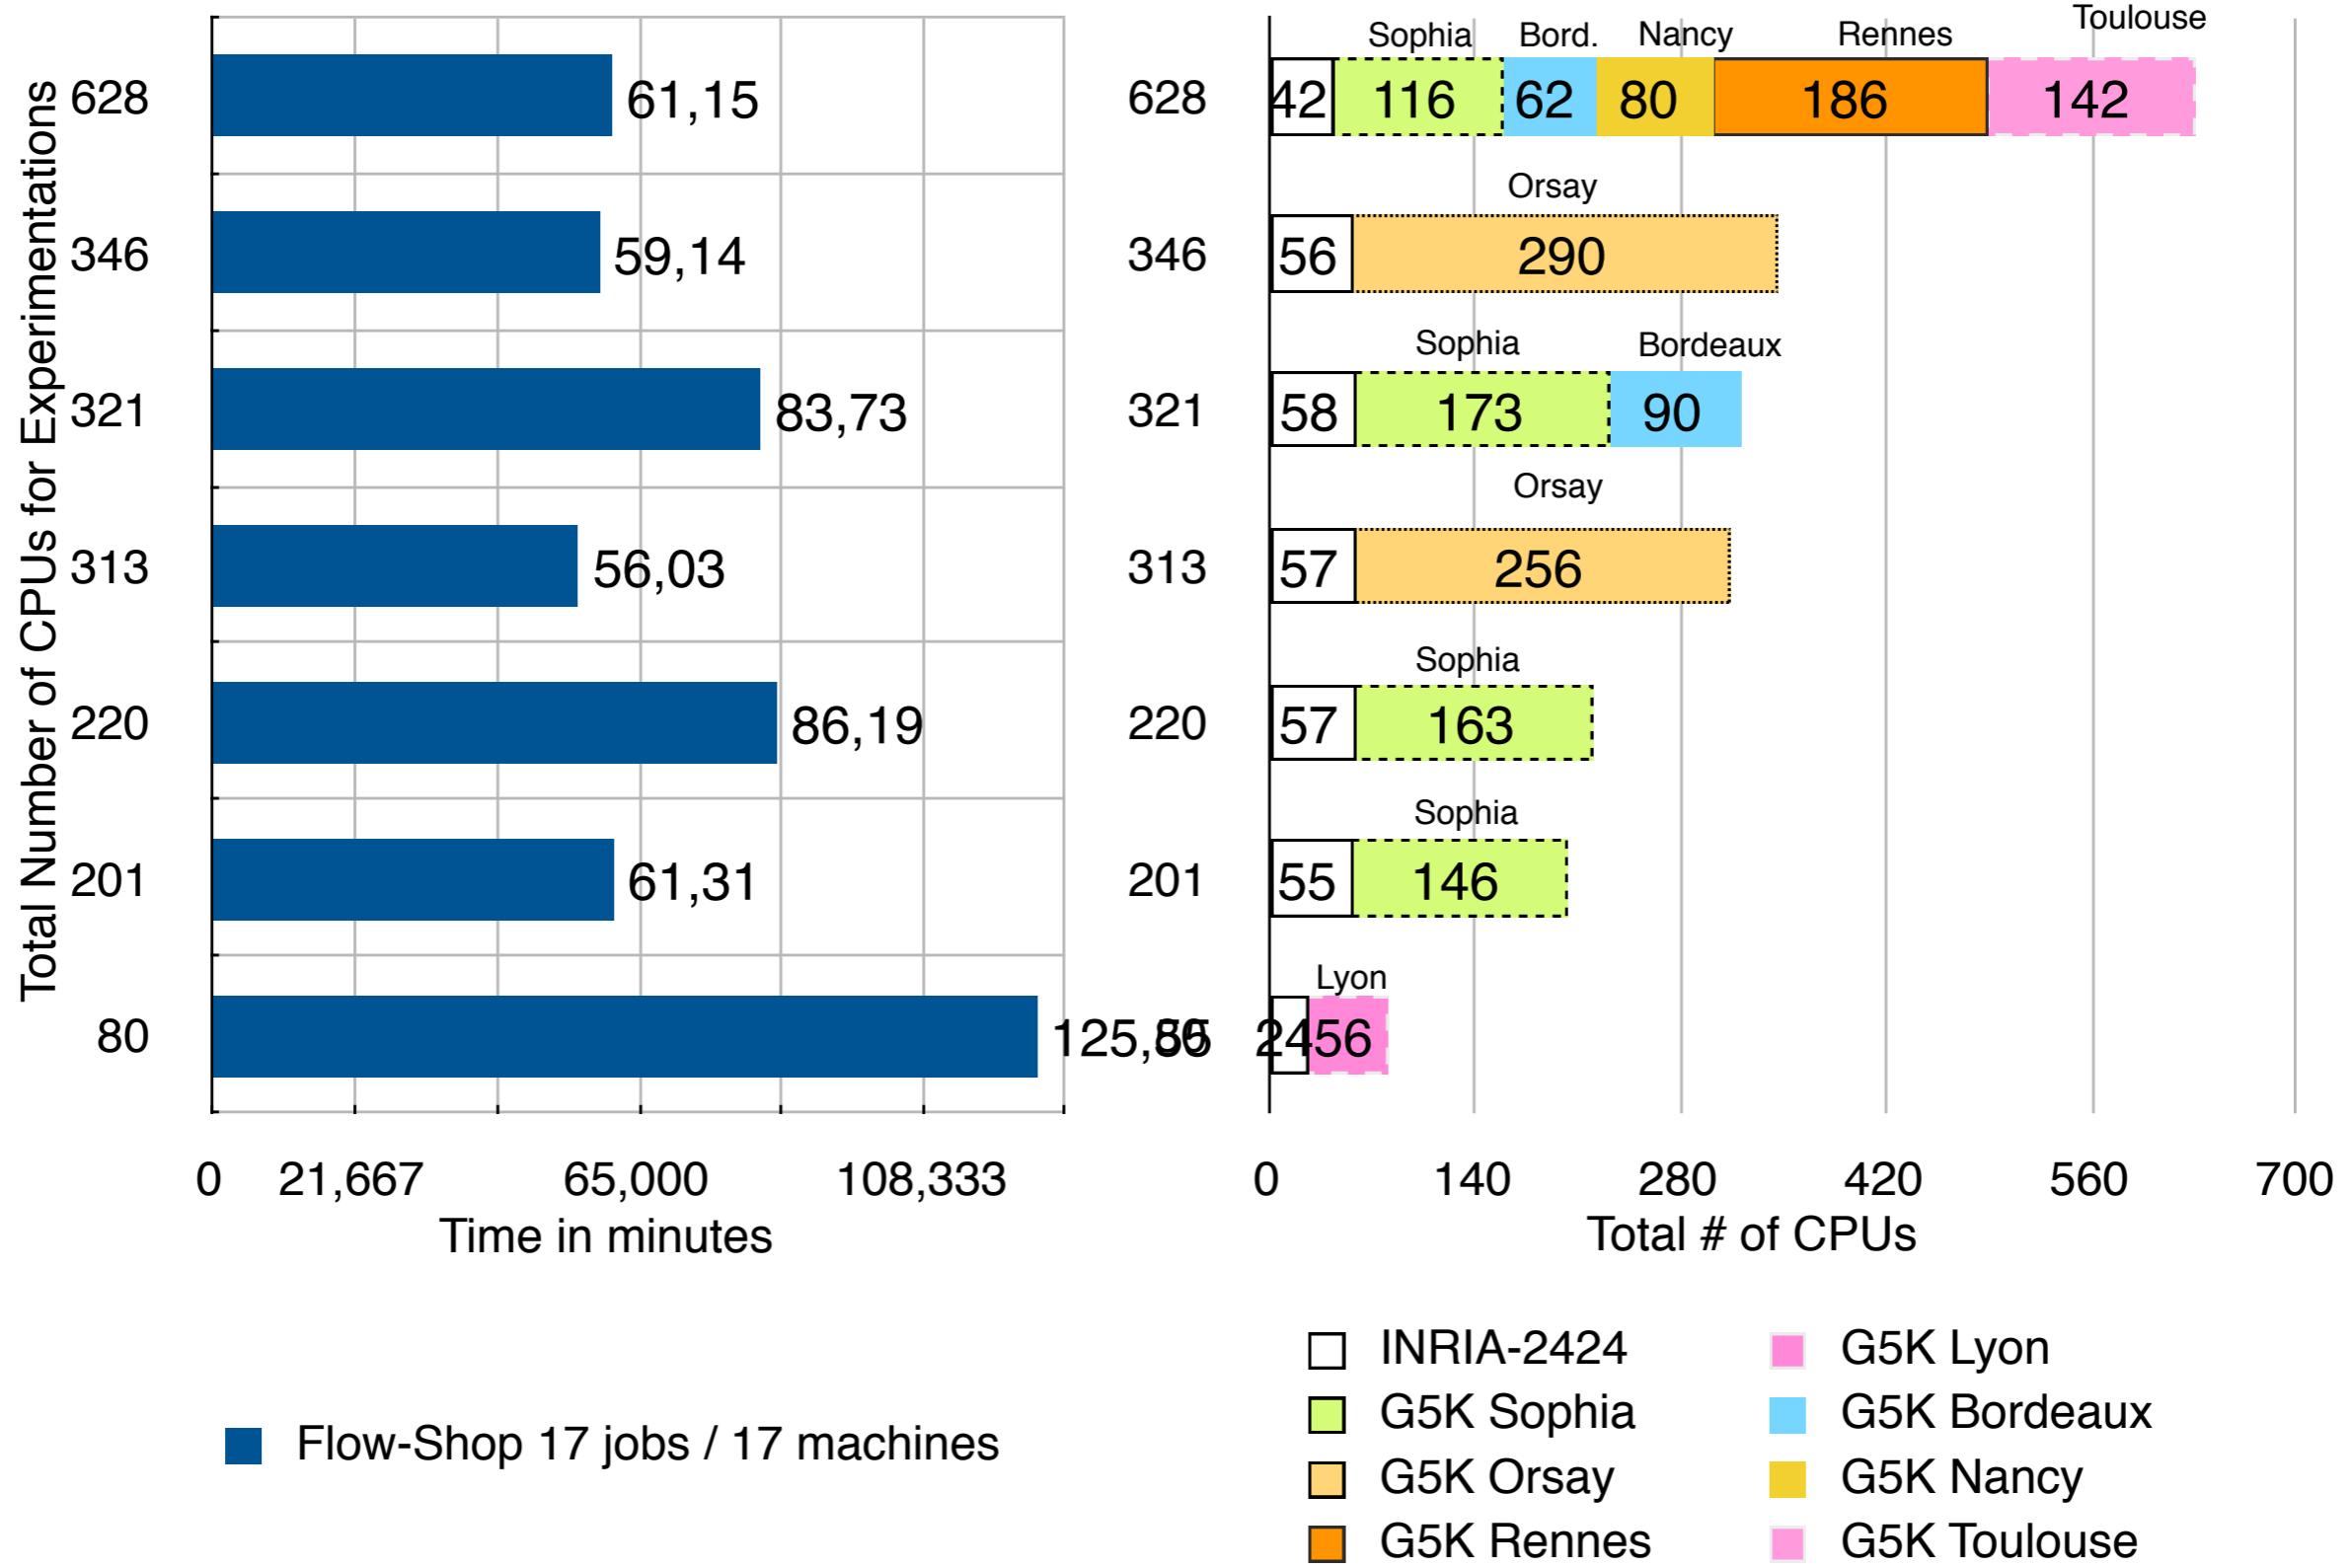
N-Queens Results



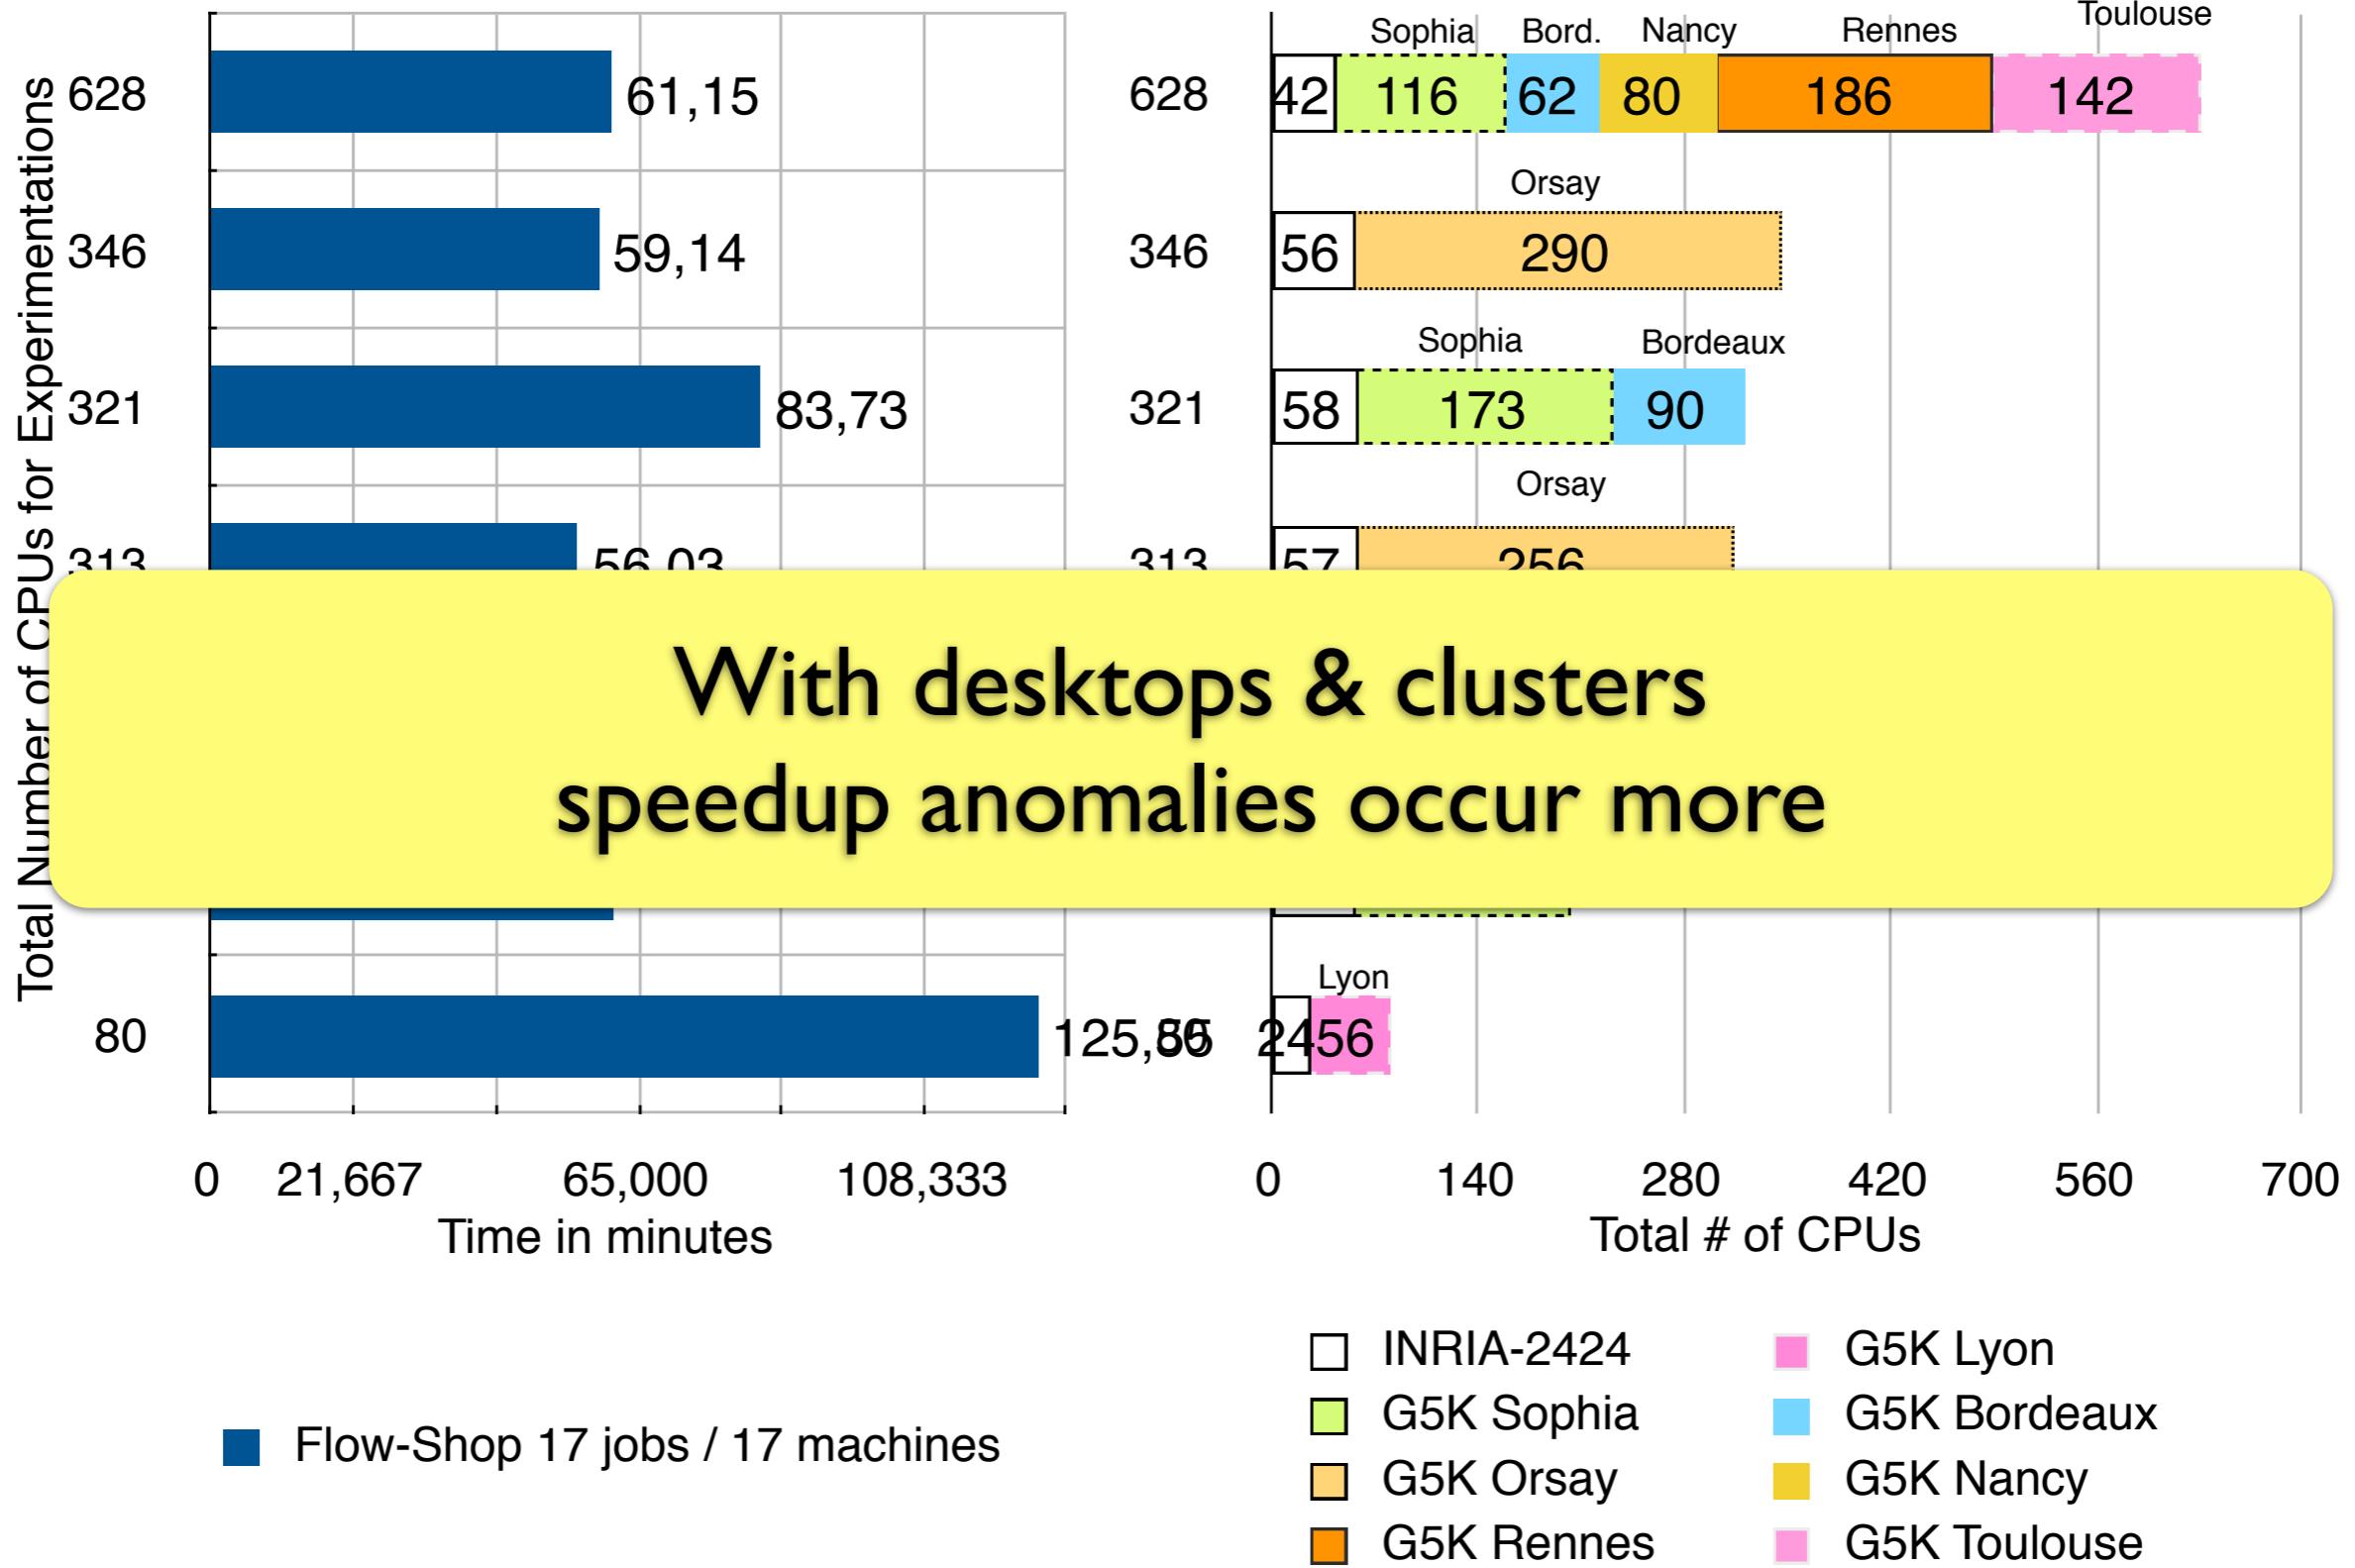
■ NQueens with n=22

□ INRIA-ALL	■ G5K Lyon
■ G5K Sophia	■ G5K Bordeaux
■ G5K Orsay	■ G5K Toulouse

Flow-Shop Results



Flow-Shop Results



Mixing - Analysis

- N-Queens problem scales well up to 1007 CPUs
- 349 Desktops + 5 Clusters
- Flow-Shop up to 628 CPUs
- 42 Desktops + 5 Clusters
- worse performances than using only clusters (anomalies are more frequent)
- Experimented in closed environments -- security
- Grid'5000 platform's success → hard for running long exp.

P2P as a Meta-Grid infrastructure

P2P as a Meta-Grid infrastructure

- **Observation** from Grid'5000:

- 300 nodes available for 2 minutes
- provide best-effort queue

P2P as a Meta-Grid infrastructure

- **Observation** from Grid'5000:

- 300 nodes available for 2 minutes
- provide best-effort queue
- **Idea:** take benefit from these nodes [Condor]
 - by hand: not easy
 - with the P2P infrastructure:
 - deploying peers in best-effort queue

P2P as a Meta-Grid infrastructure

- **Observation** from Grid'5000:

- 300 nodes available for 2 minutes
- provide best-effort queue
- **Idea:** take benefit from these nodes [Condor]
- by hand: not easy
- with the P2P infrastructure:
 - deploying peers in best-effort queue
- **Result:** a permanent P2P infrastructure over Grid'5000

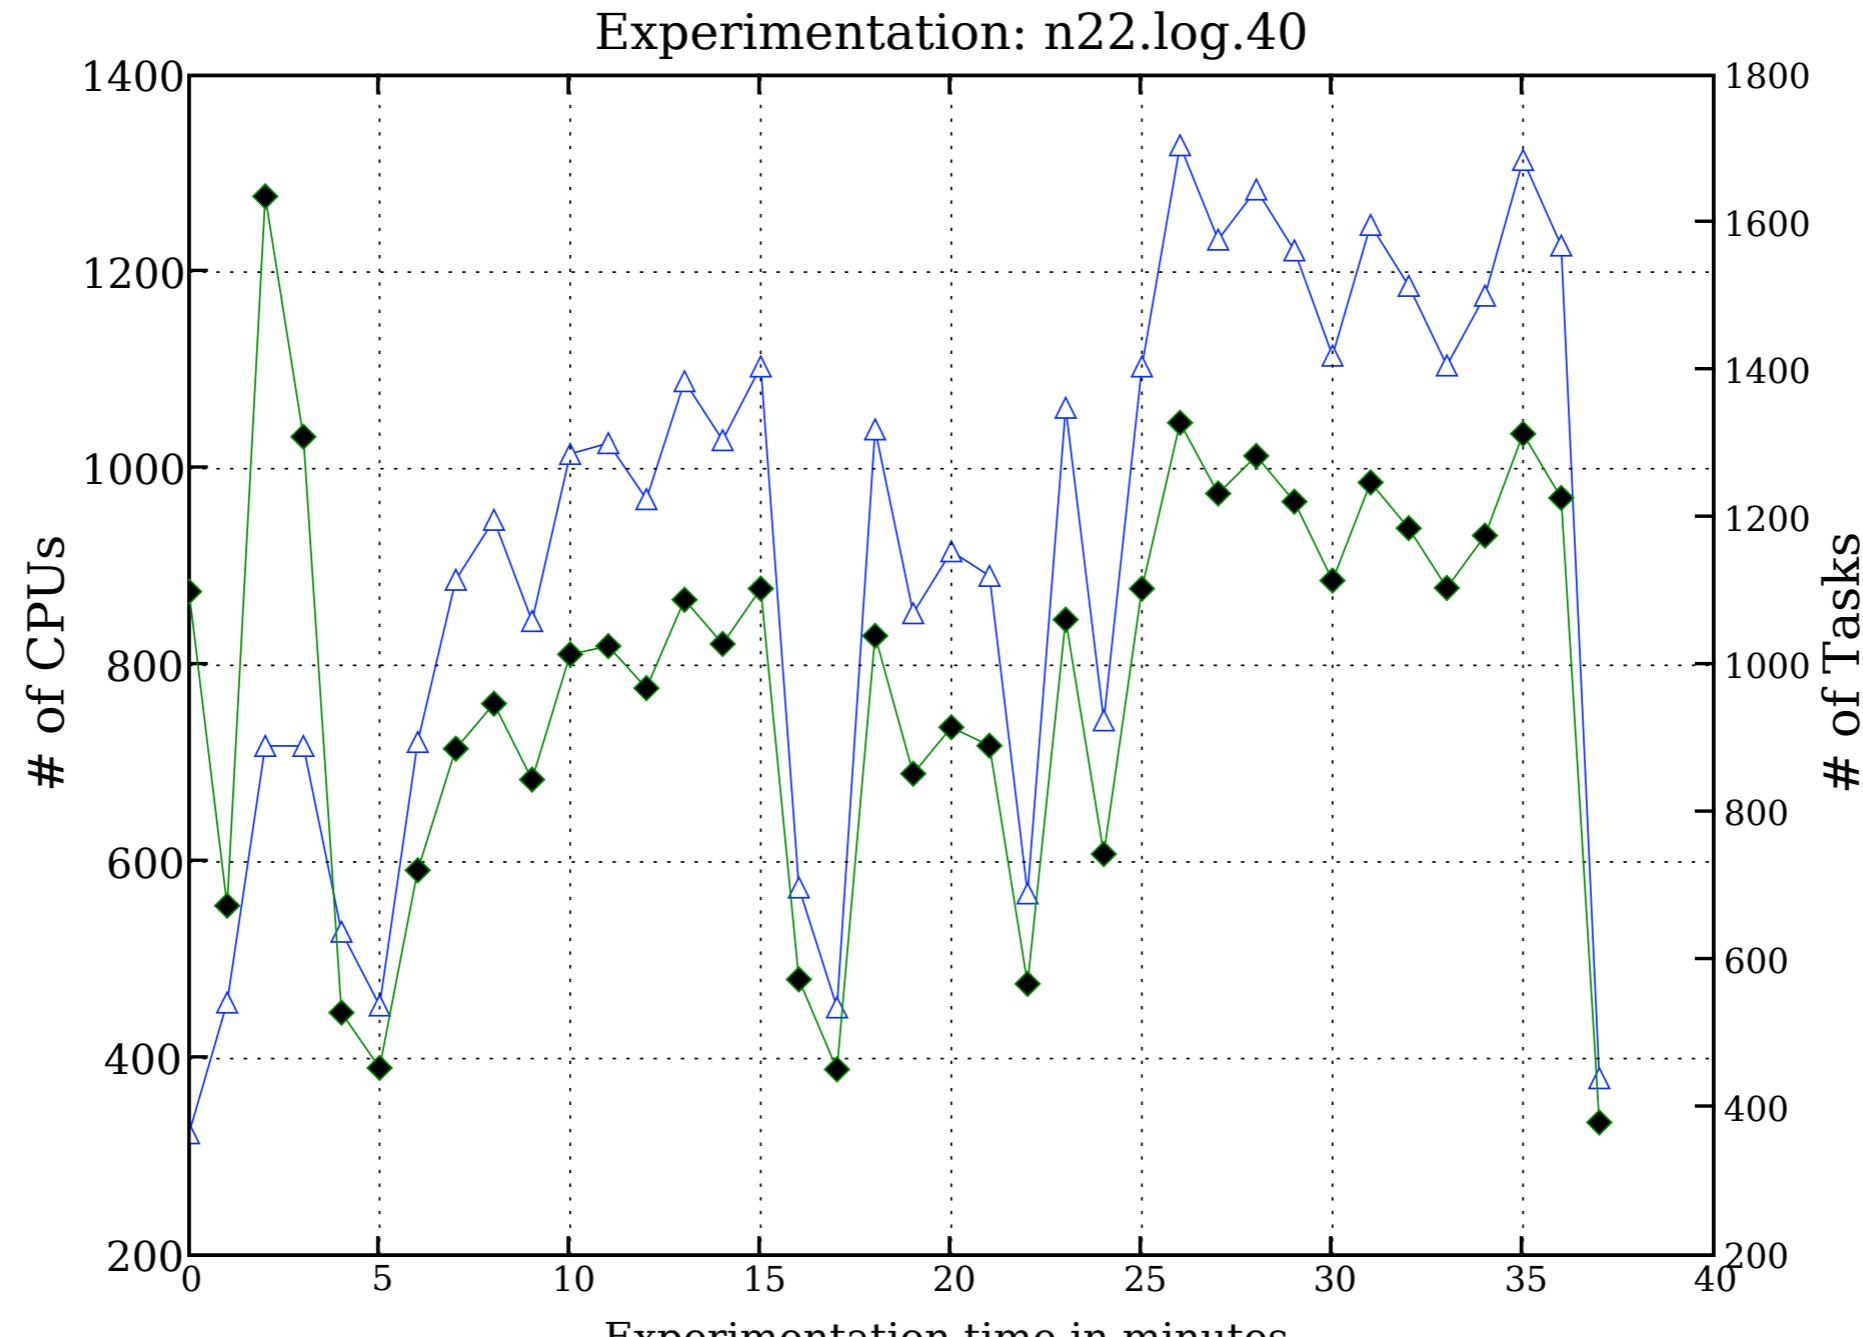
P2P as a Meta-Grid infrastructure

- **Observation** from Grid'5000:

- 300 nodes available for 2 minutes
- provide best-effort queue
- **Idea:** take benefit from these nodes [Condor]
- by hand: not easy
- with the P2P infrastructure:
 - deploying peers in best-effort queue
- **Result:** a permanent P2P infrastructure over Grid'5000

P2P infrastructure as a dynamic Grid middleware

P2P with N-Queens on Grid'5000

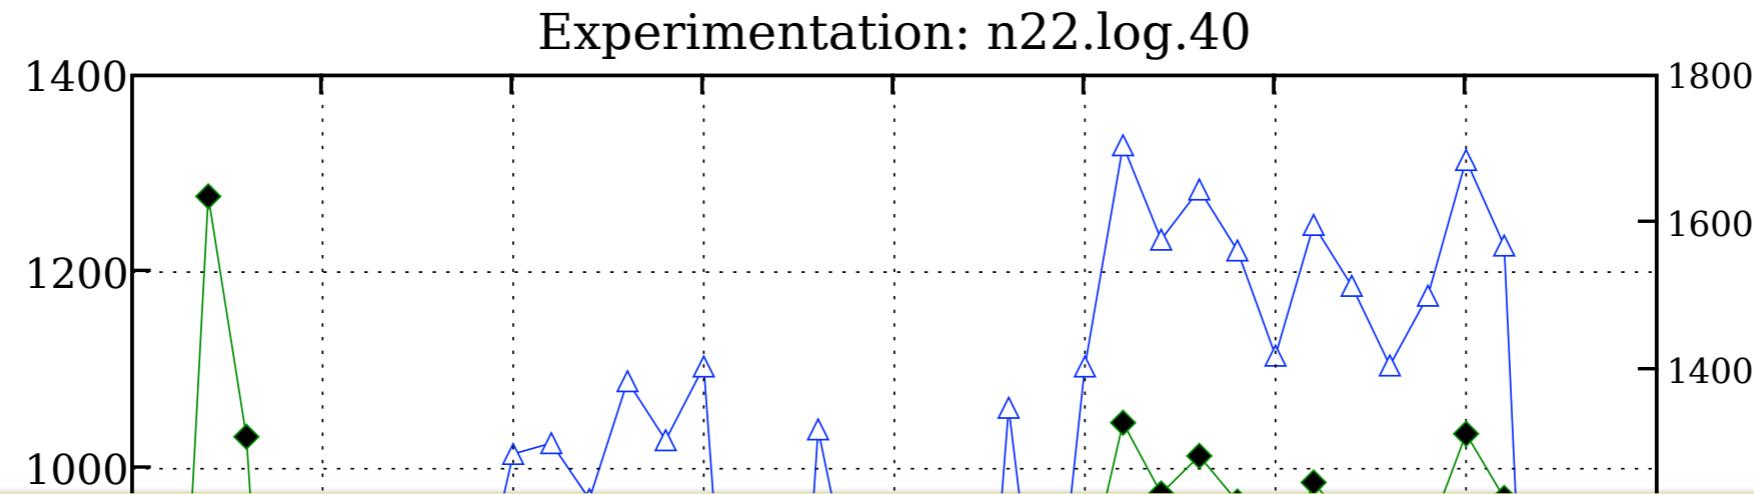


I384 CPUs - 9 sites

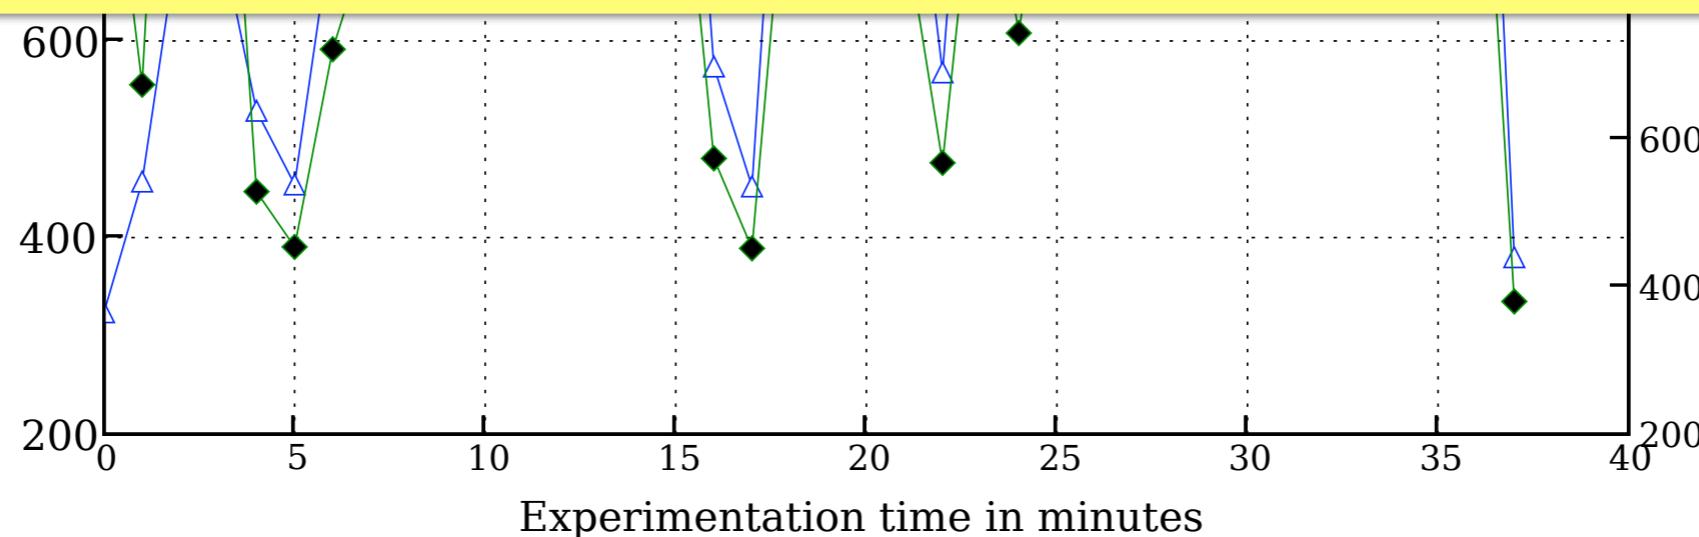
of workers by minutes

tasks computed by minutes

P2P with N-Queens on Grid'5000



Embarrassingly applications can take benefit from
Meta-Grid Infrastructure



of workers by minutes

tasks computed by minutes

I384 CPUs - 9 sites

Agenda

- About Me / Previous Experience
- PhD Main Contributions
 - Branch-and-Bound Framework for Grids
 - Desktop Grid with Peer-to-Peer
 - Mixing Desktops & Clusters
- **Perspectives**
- Possible Support & Conclusion

Perspectives & Ongoing Work

- Peer-to-Peer Infrastructure:
 - Job Scheduler
 - Resource Localization (PhD)
- Large-Scale Experiments:
 - International Grid: France, Japan, and Netherlands
 - Grid Pugtests
- Deployment:
 - Contracts in Grids (GCM Standard)
- Industrialization:
 - P2P → Desktop Resource Virtualization
 - CPER - P2P
 - 1M€/4years to professionalize the INRIA Grid

Conclusion

Conclusion

- Branch-and-Bound for Grids
 - solve COPs
 - hide Grid difficulties
 - communication between workers

Conclusion

- Branch-and-Bound for Grids
 - solve COPs
 - hide Grid difficulties
 - communication between workers

Conclusion

- Branch-and-Bound for Grids
 - solve COPs
 - hide Grid difficulties
 - communication between workers
- Peer-to-Peer as Grid infrastructure
 - mixing desktops and clusters
 - deployed at INRIA Sophia

Conclusion

- Branch-and-Bound for Grids
 - solve COPs
 - hide Grid difficulties
 - communication between workers
- Peer-to-Peer as Grid infrastructure
 - mixing desktops and clusters
 - deployed at INRIA Sophia
 - Tested and experimented

Conclusion

- Branch-and-Bound for Grids
 - solve COPs
 - hide Grid difficulties
 - communication between workers
- Peer-to-Peer as Grid infrastructure
 - mixing desktops and clusters
 - deployed at INRIA Sophia
- Tested and experimented
- Available in open source

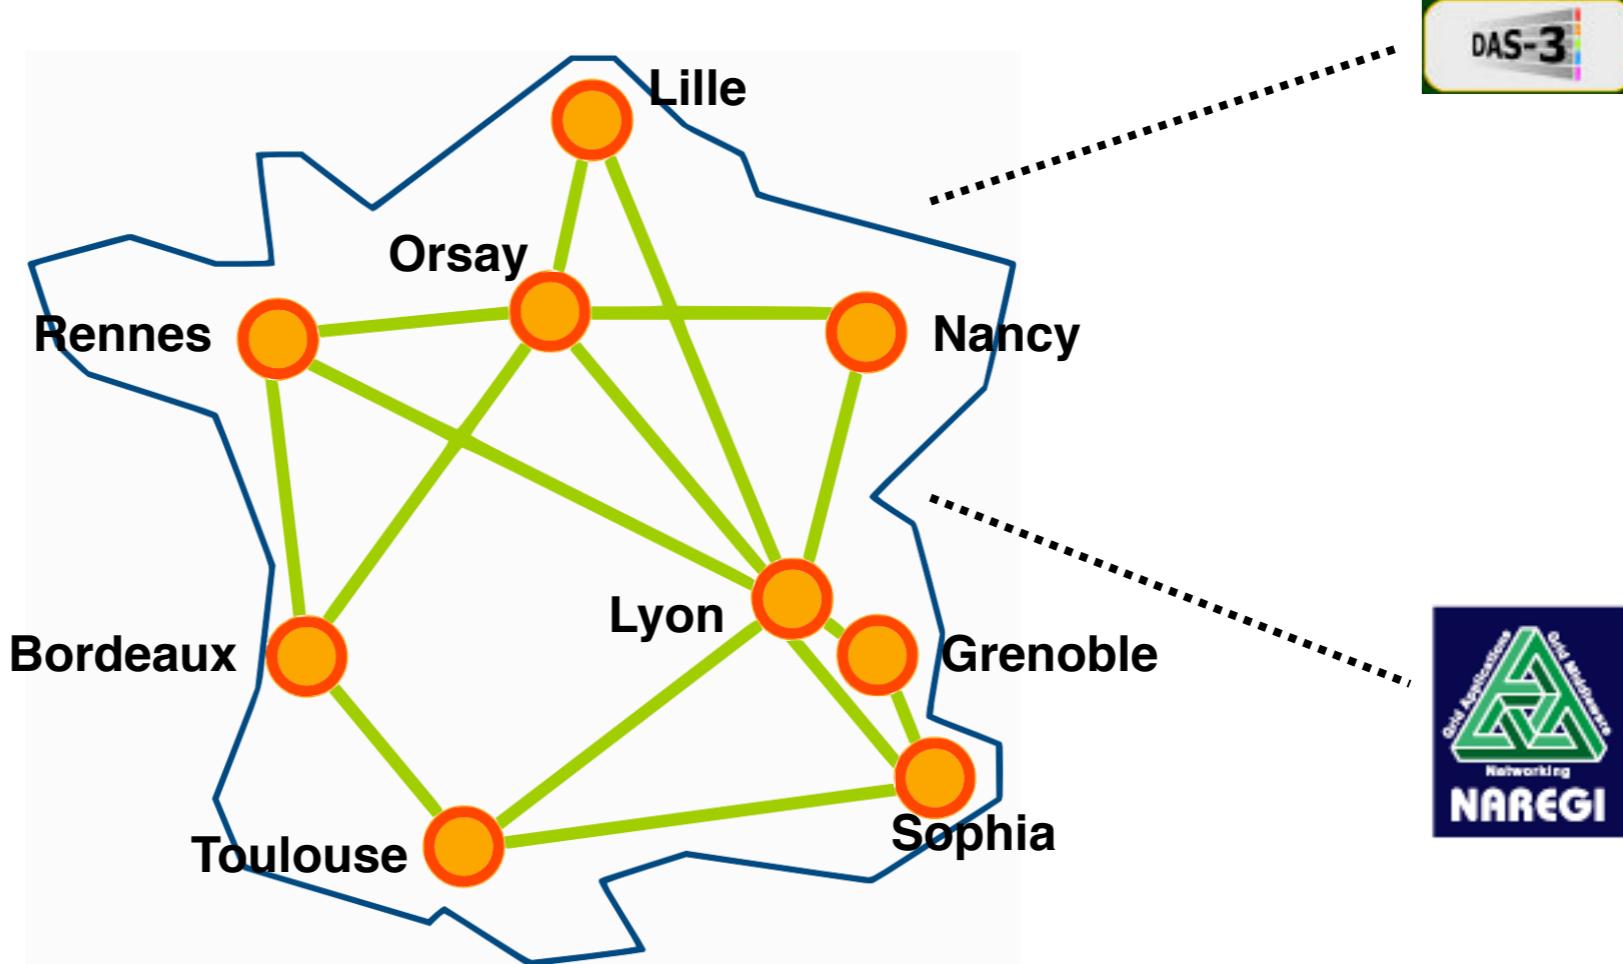
Conclusion

- Branch-and-Bound for Grids
 - solve COPs
 - hide Grid difficulties
 - communication between workers
 - Peer-to-Peer as Grid infrastructure
 - mixing desktops and clusters
 - deployed at INRIA Sophia
 - Tested and experimented
 - Available in open source
- **Provides framework and infrastructure
to hide Grid difficulties**

Agenda

- About Me / Previous Experience
- PhD Main Contributions
 - Branch-and-Bound Framework for Grids
 - Desktop Grid with Peer-to-Peer
 - Mixing Desktops & Clusters
 - Perspectives
- **Possible Support & Conclusion**

Grid'5000: The French Grid!

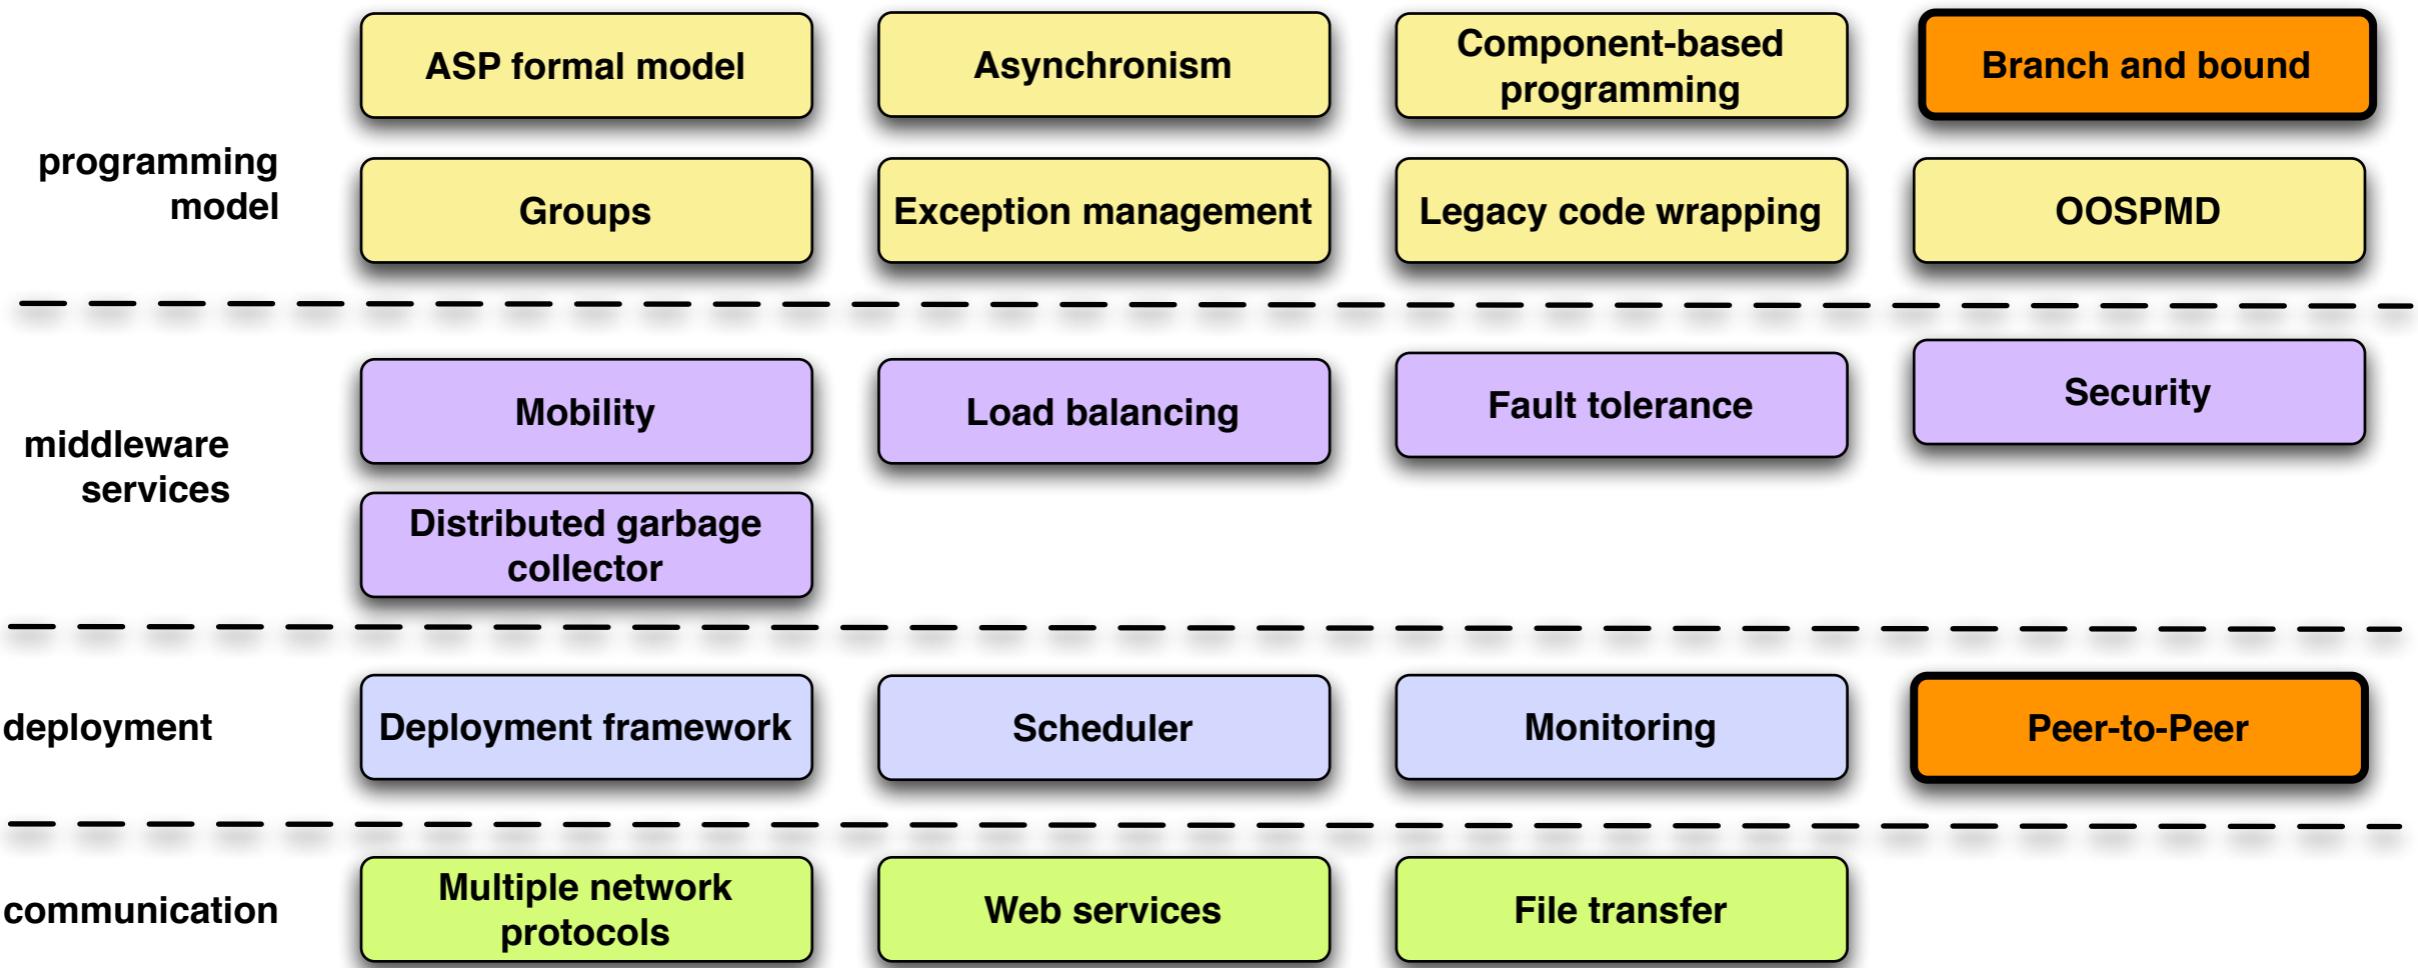


9 sites - 14 clusters - 4,790 Cores

Only large-scale experiments - aiming 3 sites / 1,000 cores

ProActive Grid Middleware

<http://proactive.ow2.org>





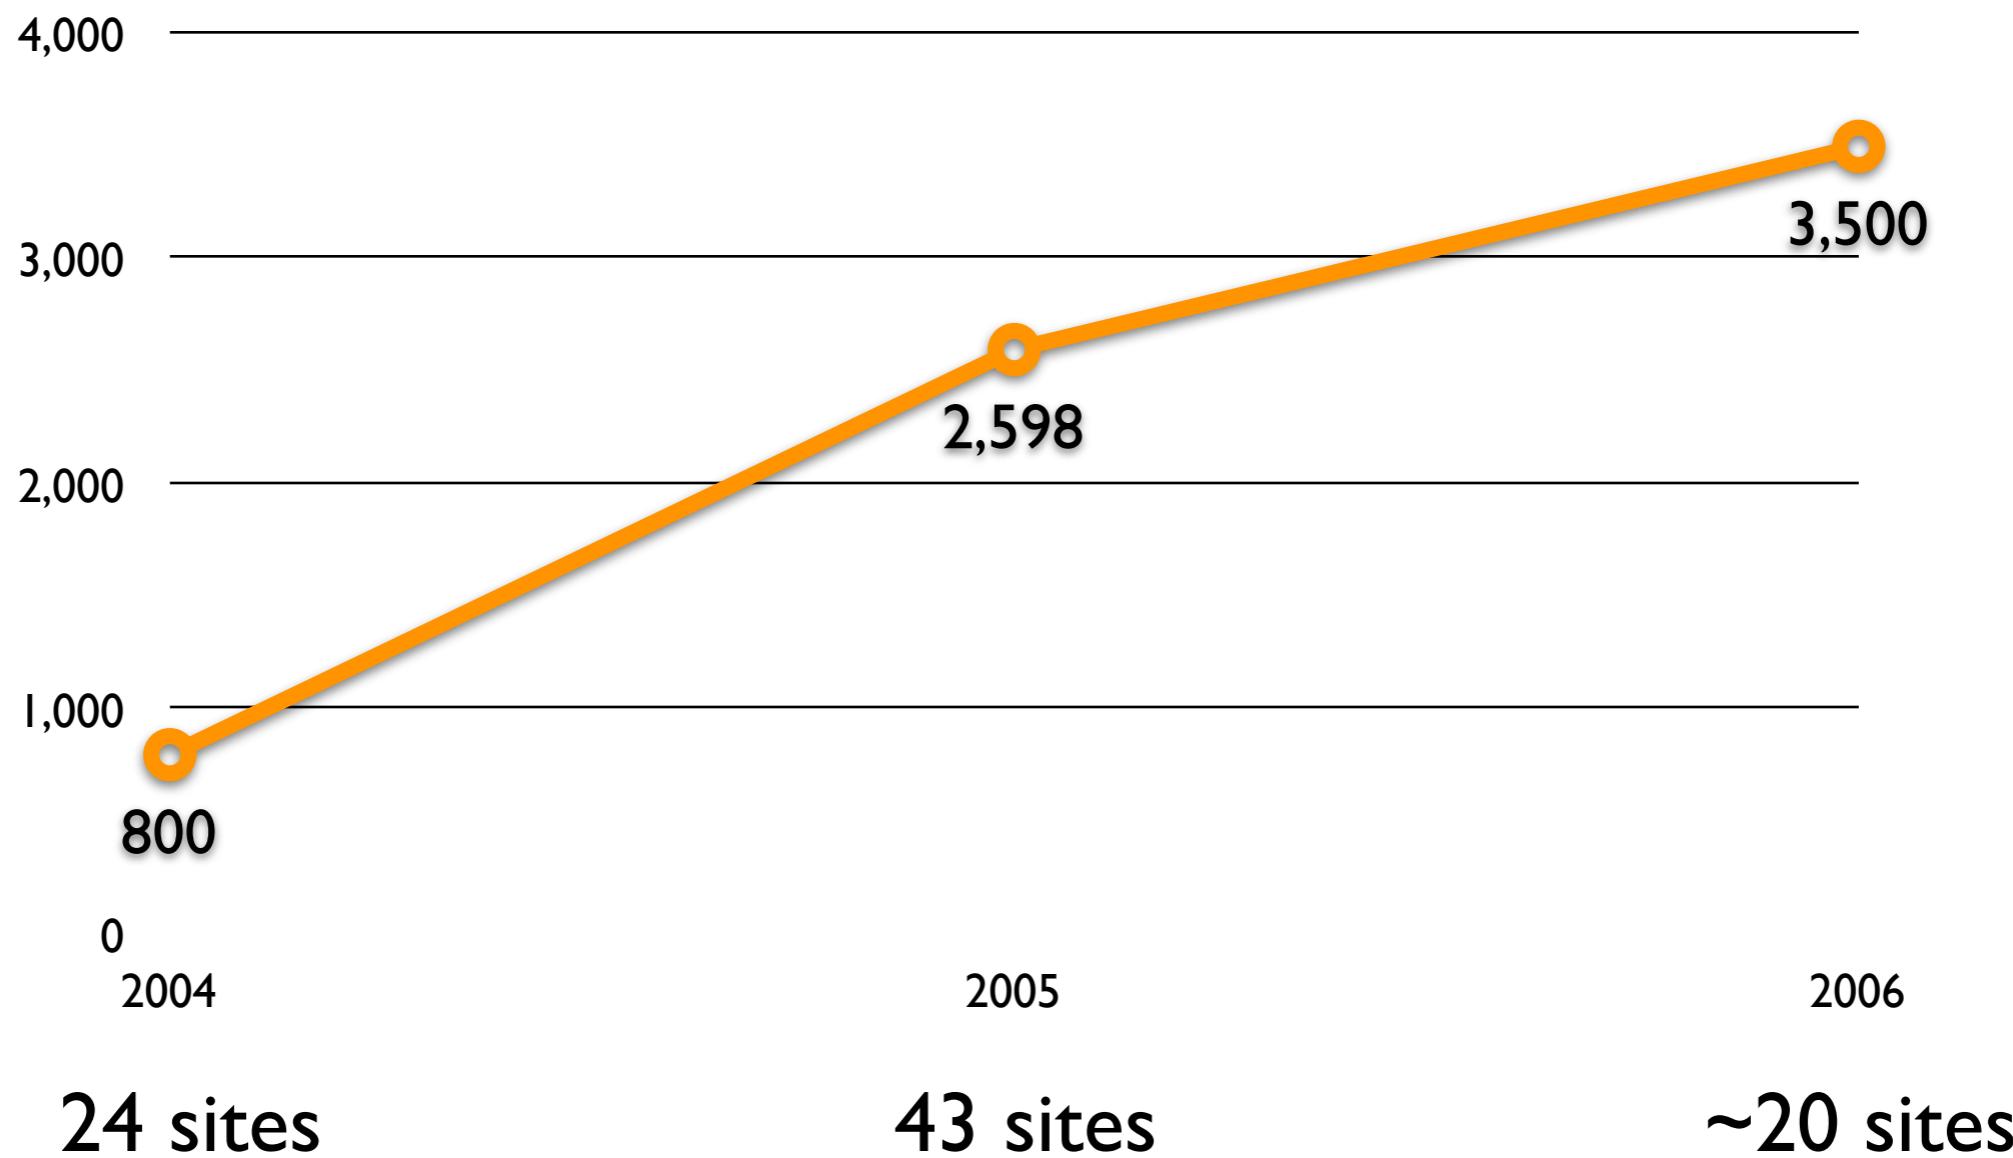
GRID@WORK Plugtest

- Taking place in Sophia Antipolis, France
- 20-24 October 2008
- **V Edition**
- CoreGrid meetings
- Conferences / Tutorials
- Grid Programming Challenge: Barrier Option Pricing

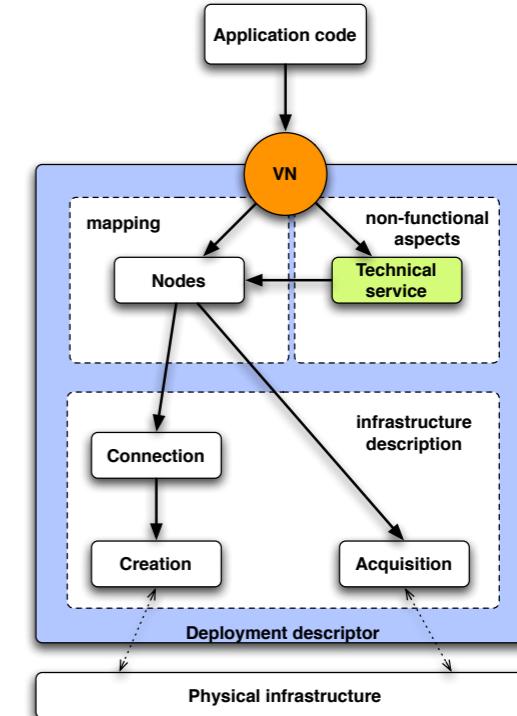
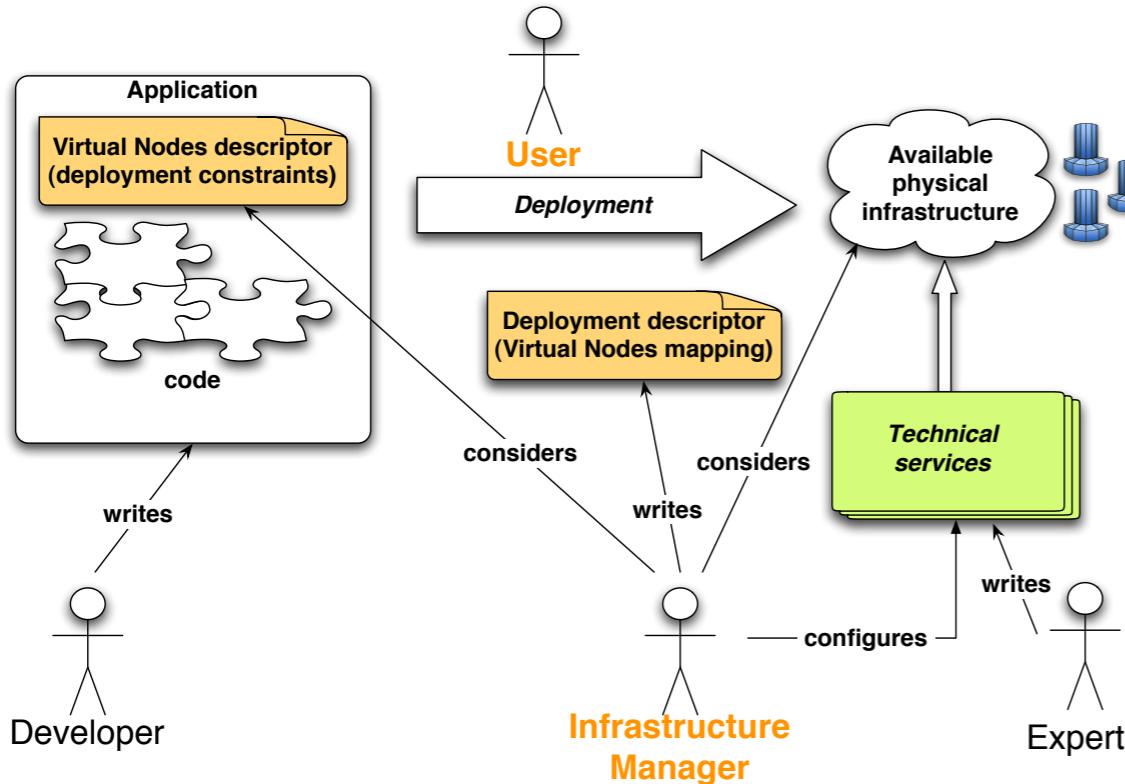
http://www.etsi.org/plugtests/GRID2008/About_GRID.htm



GRID@WORK Plugtest



GCM Deployment



- Part of GridComp Project in collaboration with ETSI
- Aim to standardize the Grid in complement of GGF
- Standard draft based on parts of my PhD



Support

- **Grid'5000:** Account creation & Experiment help
- **ProActive Java Grid Middleware**
- **Grid Plugtest and Challenge**
- **GCM Deployment**
- Java / Eclipse / Python

Last Things...

- **Office Room 5.21**
- **adc@csse.unimelb.edu.au**
- **<http://www.csse.unimelb.edu.au/~adc>**

[SCCC05] Balancing Active Objects on a Peer to Peer Infrastructure

Javier Bustos-Jimenez, Denis Caromel, Alexandre di Costanzo, Mario Leyton and Jose M. Piquer. Proceedings of the XXV International Conference of the Chilean Computer Science Society (SCCC 2005), Valdivia, Chile, November 2005.

[HIC06] Executing Hydrodynamic Simulation on Desktop Grid with ObjectWeb ProActive

Denis Caromel, Vincent Cavé, Alexandre di Costanzo, Céline Brignolles, Bruno Grawitz, and Yann Viala. HIC2006: Proceedings of the 7th International Conference on Hydroinformatics, Nice, France, September 2006.

[HiPC07] A Parallel Branch & Bound Framework for Grids

Denis Caromel, Alexandre di Costanzo, Laurent Baduel, and Satoshi Matsuoka. Grid'BnB: HiPC'07, Goa, India, December 2007.

[CMST06] ProActive: an Integrated platform for programming and running applications on grids and P2P systems

Denis Caromel, Christian Delb , Alexandre di Costanzo, and Mario Leyton. Journal on Computational Methods in Science and Technology, volume 12, 2006.

[PARCO07] Peer-to-Peer for Computational Grids: Mixing Clusters and Desktop Machines

Denis Caromel, Alexandre di Costanzo, and Cl  ment Mathieu. Parallel Computing Journal on Large Scale Grid, 2007.

[FGCS07] Peer-to-Peer and Fault-tolerance: Towards Deployment-based Technical Services

Denis Caromel, Alexandre di Costanzo, and Christian Delb . Journal Future Generation Computer Systems, to appear, 2007.

and Workshops, Master report, ...